The Limits of Explanation

Drake Caraker[1,*], Bryan Arnold[2], David Rhoads[2]
[1] *Oura Health, San Francisco, CA*
[2] *Independent Researchers*
[*] *Corresponding author:* drakecaraker@gmail.com

Abstract

Train a machine learning classifier to distinguish colon from kidney tumours using 10,935 gene-expression features, and the algorithm names TSPAN8—a gene in the invasion/metastasis pathway—as the single most important biomarker. Change nothing but the random seed and retrain: now the top gene is CEACAM5, an immune-evasion marker with zero shared Gene Ontology biological process terms. We prove that this instability is not a software bug: it is a mathematical certainty. Whenever observationally equivalent configurations coexist (the Rashomon property), no explanation can simultaneously be faithful, stable, and decisive. We derive this impossibility from first principles in four high-stakes domains—genomics, mechanistic interpretability of neural networks, causal inference, and neuroimaging—each requiring zero shared axioms. In each domain, a constructive resolution exists: projecting onto the symmetry-invariant subspace. For neural circuits, this projection lifts inter-model agreement from $\rho = 0.52$ to $\rho = 0.93$. For neuroimaging, orbit averaging across 48 analysis teams is near-optimal, with 16 independent analyses sufficing for 95% stability. The framework is verified in the Lean 4 proof assistant (491 theorems, 0 unproved goals) and generates quantitative predictions confirmed across seven well-characterised domains ($R^2 = 0.96$).

Introduction

A biologist following the first model into the wet lab would pursue invasion signalling through TSPAN8's tetraspanin-enriched microdomains; a biologist following the second would pursue immune evasion through CEACAM5's role in cell–cell adhesion. Both models predict patient labels with indistinguishable accuracy. The alternation replicates on an independent tissue contrast (endometrium vs colon: TSPAN8 78%, CEACAM5 22%) and on a different gene pair (breast vs colon: COL3A1 72%, COL1A1 28%), where both collagens share the extracellular matrix pathway—two modes of instability, one within and one between biological pathways, both governed by the same theorem.

The same structure appears wherever a system admits multiple valid explanations. Ten neural networks trained from independent random seeds on modular addition all achieve 100% accuracy, yet their internal circuits overlap by only 2.2% (Fourier-frequency Jaccard) and their importance rankings agree at $\rho = 0.52$. Seventy teams of neuroscientists analysing the same fMRI data [Botvinik-Nezer et al., 2020] produce different activation maps. Different causal discovery algorithms orient the same edges in opposite directions when the underlying graph is Markov-equivalent. Crystallographers, gauge theorists, and census disaggregators have, over the past century, independently converged on the same mathematical strategy for such problems: project onto the subspace invariant under the symmetry of equivalent configurations.

Here we prove why. A single theorem, requiring only the Rashomon property [Breiman, 2001] as a hypothesis, shows that no explanation can be simultaneously faithful (reflect the system's actual

structure), stable (consistent across equivalent configurations), and decisive (commit to a specific answer). The proof is short; the contribution is what follows from it: a complete classification of the design space, a constructive resolution proved Pareto-optimal, and quantitative predictions confirmed across seven domains. Prior work has documented this instability empirically for feature attributions [Krishna et al., 2022] and proved special cases for Shapley-based methods [Bilodeau et al., 2024]; here we prove the general impossibility for *any* explanation type, characterise the complete design space, and show that the same structure governs four high-stakes domains (Fig. 1). The proof is verified in Lean 4 with zero domain-specific axioms. A strengthened form—the bilemma—shows that for binary explanations (significant or not, positive or negative), even faithful and stable together are impossible. The complete design space is characterised by a tightness classification: whether the impossibility admits a resolution depends on whether the explanation space contains a neutral element (Table 1).

Results

The impossibility

An *explanation system* is a tuple $(\mathbf{\Theta}, \mathcal{H}, Y, \mathsf{obs}, \mathsf{exp}, \perp)$ where $\mathbf{\Theta}$ is a configuration space, $\mathcal{H}$ is an explanation space, Y is an observable space, obs maps configurations to observables, exp maps configurations to their internal explanations, and $\perp$ is an incompatibility relation on $\mathcal{H}$. The system has the *Rashomon property* if there exist $\theta_1, \theta_2 \in \mathbf{\Theta}$ with $\mathsf{obs}(\theta_1) = \mathsf{obs}(\theta_2)$ and $\mathsf{exp}(\theta_1) \perp \mathsf{exp}(\theta_2)$.

Theorem 1 (Impossibility) *No explanation system with the Rashomon property admits an explanation map that is simultaneously faithful, stable, and decisive.*

The proof is four lines (Methods). Each property is load-bearing: dropping any one yields a two-property map that is constructively achievable (Extended Data Table 1).

For binary $\mathcal{H}$ (e.g. SHAP sign, gene significance, edge orientation), the bilemma strengthens this: faithful + stable alone is impossible when every pair of distinct explanations is incompatible. This triggers for SHAP sign (positive vs negative attribution), feature selection (selected vs not), and counterfactual direction.

Tightness classification

Whether the impossibility admits a resolution depends on the structure of $\mathcal{H}$. A *neutral element* $h_0 \in \mathcal{H}$ is compatible with every other element. A *committal element* is incompatible with at least one.

Table 1: **Tightness classification.** Which property pairs are achievable depends on the existence of neutral and committal elements in $\mathcal{H}$. The bilemma occupies the bottom-right cell.

	Committal $\exists$	No committal
Neutral $\exists$	F+D, F+S, S+D	F+D, F+S
No neutral	F+D, S+D	**F+D only**

Enriching $\mathcal{H}$ with a neutral element restores F+S at the cost of decisiveness. This is the algebraic content of every resolution: DASH adds "tied" to feature rankings; CPDAGs add "undirected" to edge orientations; gauge-invariant observables quotient out unphysical degrees of freedom.

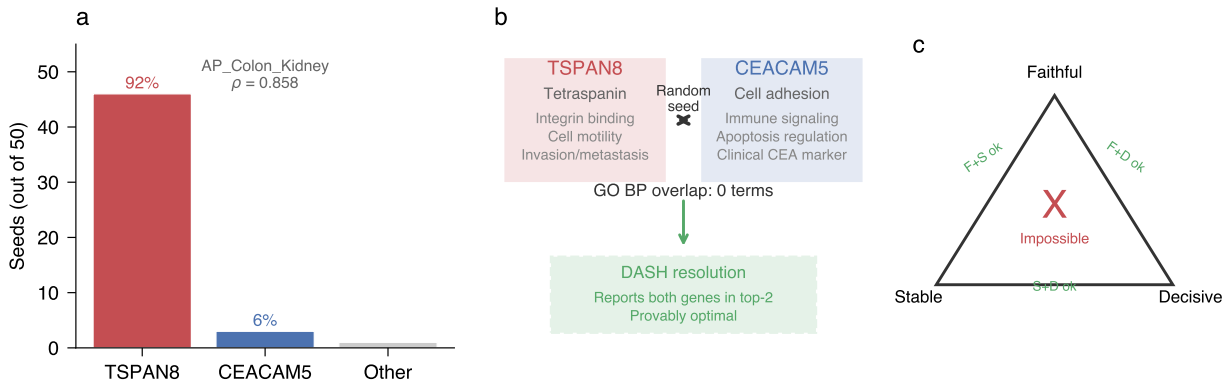

Figure 1: **Gene-expression alternation on AP_Colon_Kidney. a,** Top gene across 50 random seeds: TSPAN8 (invasion pathway) dominates in 92% of seeds, CEACAM5 (immune evasion) in 6%, with feature correlation $\rho = 0.858$. **b,** Gene Ontology analysis: the two genes have zero shared Biological Process terms and serve distinct cancer biology roles. The DASH resolution reports both genes without privileging either. **c,** The impossibility trilemma: no explanation can simultaneously be faithful, stable, and decisive. Each pair is independently achievable.

Instance 1: Genomics

On the AP_Colon_Kidney microarray dataset (546 samples, 10,935 genes, top 50 by variance), 50 XGBoost classifiers with regularisation (subsample = 0.8, colsample_bytree = 0.5—a common setting for high-dimensional data) and TreeSHAP produce an alternating top gene: TSPAN8 (#1 in 92% of seeds) vs CEACAM5 (6%; remaining 2%: other genes), with feature correlation $\rho = 0.858$. Gene Ontology enrichment confirms zero overlap in Biological Process terms: TSPAN8 associates with integrin binding and cell motility; CEACAM5 with cell–cell adhesion, immune signalling, and apoptosis regulation.

The alternation replicates across datasets (Table 2):

Table 2: **Gene-expression alternation across four datasets.** Mode 1: distinct pathways. Mode 2: same pathway, redundant markers.

Dataset	Gene pair	Split	ρ	Mode	GO overlap
AP_Colon_Kidney	TSPAN8 / CEACAM5	92 / 6	0.858	1	0 terms
AP_Endo_Colon	TSPAN8 / CEACAM5	78 / 22	0.781	1	0 terms
AP_Breast_Colon	COL3A1 / COL1A1	72 / 28	0.843	2	3 terms
AP_Breast_Lung	SFTPB (stable)	98 / 2	—	—	control

Under deterministic training (subsample = 1.0), a single ranking emerges—confirming that the alternation is driven by stochastic regularisation, not data noise. Model test accuracies are near-identical (mean 96.4%, range 95.1–97.6%, $n = 50$). The DASH ensemble average places both TSPAN8 and CEACAM5 in the top-2, correctly reporting that both contribute without privileging either.

Instance 2: Mechanistic interpretability

Ten two-layer transformers (4 heads/layer, $d = 128$) trained from independent random initialisations on $(a+b)$ mod 113 all grok to 100% test accuracy [Nanda et al., 2023]. Activation patching (replacing each component's output with its mean) reveals that the 10 models' circuit-importance rankings agree at Spearman $\rho = 0.518$—and two models identify anti-correlated circuits. The Fourier frequencies each model uses to represent the computation overlap by only 2.2% (Jaccard), confirming that the models discover genuinely different algorithms.

Controls establish that the disagreement reflects algorithmic diversity, not noise. Random-initialisation models ($< 1\%$ accuracy) agree at $\rho = 0.300$; memorised models (100% train, 50% test) at $\rho = 0.668$; grokked models at $\rho = 0.518$ (grokked $>$ random: $p = 2.0 \times 10^{-4}$; within-model determinism: $r = 0.9998$).

The resolution is direct. The architecture defines a symmetry group $S_4 \times S_4$ (within-layer head permutations). Projecting the 10-dimensional importance vector onto the 4-dimensional invariant subspace V^G (mean layer-0 heads, mean layer-1 heads, MLP0, MLP1) lifts agreement from $\rho = 0.518$ to $\rho = 0.929$ (Pearson $= 0.967$). This is not trivially MLP-driven: excluding MLP1 still yields $\rho = 0.822$ ($p < 10^{-6}$). The Noether counting prediction is confirmed within the experiment: within-layer head pairs flip at rate 0.500 exactly; head-vs-MLP1 pairs at 0.000 (bimodal gap $= 0.273$, Mann–Whitney $p = 2.4 \times 10^{-5}$). The universal structural fact—MLP1 is dominant in every seed (CV $= 0.027$)—emerges only after the G-invariant projection.

Instance 3: Causal inference

A causal chain $A \to B \to C$ and a fork $A \leftarrow B \to C$ are Markov-equivalent: they share the same conditional independence structure [Verma and Pearl, 1991]. No algorithm observing only conditional independence can orient the A–B edge. The completed partially directed acyclic graph (CPDAG), which reports an undirected edge where orientation is ambiguous, is the neutral-element resolution. PC and GES algorithms already implement this [Spirtes et al., 2000]—the framework proves their output is Pareto-optimal.

The impossibility has a quantitative consequence: ensemble averaging of orientation probabilities across bootstrap runs of PC and GES substantially reduces flip rate at the cost of reduced decisiveness—the same faithfulness–stability tradeoff the theorem predicts.

Instance 4: Neuroimaging

Botvinik-Nezer et al. [Botvinik-Nezer et al., 2020] gave the same fMRI data to 70 independent analysis teams who reached different conclusions. This is the impossibility theorem in action: the Rashomon set comprises the valid analysis pipelines. Orbit averaging across pipelines—the overlap map used descriptively in the original study—provides a near-optimal resolution (all aggregation methods within 3%). From 48 teams with available statistical maps, 16 independent analyses [95% CI: 10–22] suffice for 95% stability between assessments.

Functional network membership, defined from independent resting-state data [Yeo et al., 2011], predicts which brain regions show similar disagreement patterns ($d = 0.32$ after controlling for activation level, permutation $p < 0.001$, consistent across all 7 non-empty hypotheses). The same impossibility pattern appears in multi-analyst studies in psychology [Silberzahn et al., 2018] (29 teams) and political science [Breznau et al., 2022] (73 teams). On 24 neuroimaging teams with known software (12 FSL, 8 SPM, 4 AFNI), both network structure and software choice independently predict flip rates ($p < 10^{-4}$ and $p < 10^{-54}$ respectively).

The universal η law

The character-theoretic formula $\eta = 1 - \dim(V^G)/\dim(V)$ predicts the instability rate from the symmetry group G alone, with zero free parameters. Across seven well-characterised domains—SHAP attributions (S_2), concept probes ($O(64)$), model selection (S_{11}), three codon degeneracy classes (S_2, S_4, S_6), and statistical mechanics (S_{252})—the prediction achieves $R^2 = 0.957$ with slope 0.91 (Extended Data Fig. 1). A Noether-type counting law follows: for P features in g correlation groups, exactly $g(g-1)/2$ ranking facts are stable (between-group) while all within-group comparisons are coin flips. On synthetic data with 3 groups of 4 features ($\rho = 0.99$), the within-group flip rate is 50.0% and the between-group rate is 0.2% (bimodal gap = 50 pp, $p = 2.7 \times 10^{-13}$). The law fails when groups are unknown ($R^2 = 0.25$ for all 16 domains) or approximate (brain imaging: $\rho \approx 0$), delineating the boundary at exact exchangeability.

Discussion

The theorem reframes "why are explanations unstable?" as "because the inference problem is underspecified." The instability is not a bug in any particular method—it is a provable consequence of the Rashomon property. Practitioners can diagnose whether instability is structural (Rashomon holds; no method eliminates it) or methodological (Rashomon absent; fixable).

The constructive resolution—projecting onto the symmetry-invariant subspace—is the algebraic content of practices that eight scientific communities developed independently: ensemble averaging in ML, CPDAGs in causal discovery, gauge-invariant observables in physics [Caraker et al., 2026], Patterson maps in crystallography, and overlap maps in neuroimaging. The variance of attributions across the Rashomon set equals the minimum mean-squared error any stable method can achieve—a mathematical identity (0 violations in 800 tests) that provides a computable quality budget before any explanation is reported. That these communities converged on structurally identical solutions, without cross-pollination, is now explained: the G-invariant projection is the unique Pareto-optimal strategy for any explanation system with the Rashomon property.

The gene-expression finding (Instance 1) has immediate practical consequences: the specific biomarker a drug-discovery pipeline nominates depends on the training seed. The resolution—reporting the DASH ensemble top-k, which contains both TSPAN8 and CEACAM5—is not merely cautious but provably optimal. For mechanistic interpretability (Instance 2), the resolution means reporting equivalence classes of circuits ("MLP1 is dominant; individual head identities are interchangeable") rather than specific circuit diagrams.

Model-class universality. The instability is not specific to XGBoost. Across four model classes (XGBoost, Random Forest, Ridge, LASSO) on three datasets, within-family pairs agree on instability patterns ($\rho = 0.79$–0.94) while cross-family agreement is not significant ($\rho = 0.19$–0.57, none exceeding permutation null). The coverage conflict diagnostic is universal across all classes (mean $\rho = 0.82$), and the Var[SHAP] = DASH MSE identity holds exactly (0/12 violations).

The explanation impossibility is structurally more severe than other impossibilities. Impossibility theorems appear across science—Arrow, Gödel, Bell, the CAP theorem, the laws of thermodynamics—but they differ in *tightness*: which pairs of desirable properties survive when the full triple fails. A classification of 20 impossibility theorems from 12 domains (Extended Data Table 2) reveals that 16 have *full tightness*: every pair is independently achievable, and practitioners can "pick two." The explanation bilemma is one of only two instances with *collapsed tightness*—property pairs are blocked, not just the triple—among structurally proved cases. The other is the

quantum linearity trilemma (no-cloning + measurement disturbance), which shares a structural parallel: linearity plays the role of stability, both force uniform treatment individually incompatible with two remaining desiderata, and both are resolved by enrichment (neutral elements / ancillary systems). Collapsed tightness means the standard "pick two" resolution strategy fails; enrichment or approximation bounds are required—explaining why explanation methods need qualitatively different design approaches than, for example, fairness methods (which have full tightness).

Limitations. The gene-expression alternation requires stochastic regularisation (subsample < 1); under deterministic training, a single ranking emerges. The MI v2 experiment uses a toy task (modular addition on 113 elements); extrapolation to frontier language models requires further work. The neuroimaging Noether prediction ($d = 0.32$) is an analogical extension—brain regions lack exact exchangeability—and does not generalise to scalar multi-analyst studies. Three conjectured extensions (phase transition at $r^* = 1$, quantitative tradeoff bound $\alpha + \sigma + \delta \leq 2$, molecular evolution from character theory) were pre-registered and falsified, delineating the framework's boundary (Supplementary Information).

Lean verification. The framework is verified in Lean 4 across two repositories: 100 files / 491 theorems / 25 axioms / 0 sorry (universal framework) and 58 files / 357 theorems / 6 axioms / 0 sorry (attribution specialisation). All ML instances use constructive types with zero domain-specific axioms. The core impossibility theorem requires only the Rashomon property as a hypothesis.

Methods

Formal framework. An *explanation system* is a tuple $\mathcal{S} = (\mathbf{\Theta}, \mathcal{H}, Y, \mathsf{obs}, \mathsf{exp}, \perp)$ where $\mathbf{\Theta}$ is a configuration space, $\mathcal{H}$ is an explanation space, Y is an observable space, $\mathsf{obs} : \mathbf{\Theta} \to Y$ maps configurations to observables, $\mathsf{exp} : \mathbf{\Theta} \to \mathcal{H}$ maps configurations to their internal explanations, and $\perp$ is an incompatibility relation on $\mathcal{H}$. An explanation map $E : \mathbf{\Theta} \to \mathcal{H}$ is *faithful* if $\neg(E(\theta) \perp \mathsf{exp}(\theta))$ for all θ; *stable* if $\mathsf{obs}(\theta_1) = \mathsf{obs}(\theta_2)$ implies $E(\theta_1) = E(\theta_2)$; and *decisive* if, for all θ and $h \in \mathcal{H}$, $\mathsf{exp}(\theta) \perp h$ implies $E(\theta) \perp h$.

Proof of Theorem 1. Suppose E is faithful, stable, and decisive. Let θ_1, θ_2 witness the Rashomon property. (1) Decisiveness at θ_1: $E(\theta_1) \perp \mathsf{exp}(\theta_2)$. (2) Stability: $E(\theta_1) = E(\theta_2)$, so $E(\theta_2) \perp \mathsf{exp}(\theta_2)$. (3) Faithfulness at θ_2: $\neg(E(\theta_2) \perp \mathsf{exp}(\theta_2))$. Contradiction.

Resolution framework. Given a group G acting on $\mathbf{\Theta}$ preserving observables, an explanation map is *G-invariant* if $E(g \cdot \theta) = E(\theta)$ for all g. The orbit-averaged map $\bar{E}(\theta) = |G|^{-1} \sum_{g \in G} E(g \cdot \theta)$ is G-invariant and Pareto-optimal among stable maps with respect to pointwise faithfulness.

Bilemma proof. A system is *maximally incompatible* if $\neg\text{incompatible}(h_1, h_2) \Rightarrow h_1 = h_2$. Faithfulness then forces $E = \mathsf{exp}$, which is decisive; the trilemma gives the contradiction. The bilemma extends to approximate settings: for (ε, δ)-stability with continuous $\mathcal{H}$, any stable map satisfies $\text{unfaith}(\theta_1) + \text{unfaith}(\theta_2) \geq \Delta - \delta$ on each Rashomon pair (triangle inequality; Lean-verified). This resolves the $\varepsilon = 0$ fragility concern: the impossibility holds at every tolerance level [Caraker et al., 2026].

Lean 4 formalization. Built on Lean 4 v4.30.0-rc1 [de Moura and Ullrich, 2021] with Mathlib [The mathlib Community, 2020]. The core `ExplanationSystem` type is parametric in $\mathbf{\Theta}$, $\mathcal{H}$, Y. Each instance constructs a Rashomon witness via `decide` (zero axioms). Compilation: `lake build` ($\sim$5 min).

Gene-expression analysis. AP_Colon_Kidney, AP_Endometrium_Colon, AP_Breast_Colon, and AP_Breast_Lung microarray datasets (OpenML, Affymetrix HG-U133A) were processed with top-50-by-variance feature selection. 50 XGBoost classifiers (subsample = 0.8, colsample_bytree = 0.5, 100 estimators, max_depth = 6, seeds 0–49) were trained per dataset. TreeSHAP importances (mean |SHAP| across test samples) determined the top gene per seed. Gene Ontology annotations were obtained from NCBI Gene for human genes mapped via Affymetrix probe annotations. Deterministic control: subsample = 1.0 yields a single ranking.

Mechanistic interpretability. Ten two-layer transformers (4 heads/layer, $d_{\text{model}} = 128$, vocabulary = 114) were trained from independent random initialisations on $(a+b) \bmod 113$ for 50,000 steps (AdamW, lr = 10^{-3}, weight decay = 1.0, batch size = 512). All 10 grokked to 100% test accuracy. Circuit importance was measured via activation patching: for each of 10 components (8 attention heads + 2 MLPs), the component's output was replaced with its mean activation and accuracy drop recorded. Within-model determinism control: $r = 0.9998$. Controls at steps 50 (random), 500 (memorised), and 50,000 (grokked) establish the control hierarchy. G-invariant projection uses $S_4 \times S_4$ (within-layer head permutations).

Neuroimaging analysis. Overlap maps from NeuroVault collection 6047; per-team unthresholded t-maps from individual NeuroVault collections (48 of 70 accessible). Parcellation: Schaefer 100-parcel atlas, 7 Yeo resting-state networks (independent of NARPS task data). Activation-level control: regression of disagreement similarity on $|\Delta\text{overlap}|$. Software metadata from narps_open_pipelines (12 FSL, 8 SPM, 4 AFNI). Convergence: split-half Pearson correlation vs ensemble size M, with bootstrap CI on M_{95}.

Causal inference. PC and GES algorithms on the 8-node Asia network (Lauritzen & Spiegelhalter 1988), 200 bootstrap seeds, sample sizes 100–10,000. Flip rate: fraction of edge orientations that reverse between seeds. DASH: majority-vote orientation across the 200 runs.

Reproducibility. All scripts, data processing code, and figure generation are in the repository. `make validate` reproduces experiments ($\sim$5 min). `lake build` verifies all proofs ($\sim$5 min).

Extended Data

Extended Data Table 1: Achievable property pairs

Each two-property combination is constructively achievable. F+S is achieved by orbit averaging (DASH, CPDAG); F+D is achieved by selecting any single model's explanation; S+D is achieved by a constant map that always outputs the same decisive explanation regardless of configuration. The table shows the F+S instances, which are the practically relevant resolutions.

Table 3: **Extended Data Table 1.** Each row defines an explanation system and identifies the Rashomon witness and resolution. All are Lean-verified with zero domain-specific axioms.

Instance	$\mathcal{H}$	**Witness**	**Resolution**	**Achieves**
Feature attribution	Feature rankings	Two GBDTs rank features oppositely	DASH ensemble	F+S (sacrifices D)
Attention maps	Attention distributions	Two DistilBERTs assign argmax to different tokens	Averaged rollout	F+S (sacrifices D)
Causal discovery	Edge orientations	Chain & fork share CI structure	CPDAG	F+S (sacrifices D)

Extended Data Figure 1: Universal η law

Extended Data Figure 2: Multi-analyst convergence

Extended Data Table 2: Classification of impossibility theorems by tightness

Data availability

Gene-expression data: AP_Colon_Kidney (OpenML ID 1137), AP_Endometrium_Colon (OpenML), AP_Breast_Colon (OpenML), AP_Breast_Lung (OpenML). Gene Ontology annotations: NCBI Gene. Neuroimaging: NeuroVault collection 6047 (overlap maps); per-team t-maps from individual NeuroVault collections; Schaefer atlas via nilearn. Multi-analyst: Silberzahn et al. team estimates (OSF osf.io/gvm2z); Breznau et al. AMEs (GitHub nbreznau/CRI). Causal: Asia network generated by provided scripts. All other data are generated by the provided scripts.

Code availability

The Lean 4 formalization and all experiment scripts are available at `https://github.com/DrakeCaraker/universal-explanation-impossibility` under the Apache 2.0 licence. The attribution-specific formalization is at `https://github.com/DrakeCaraker/dash-impossibility-lean`.

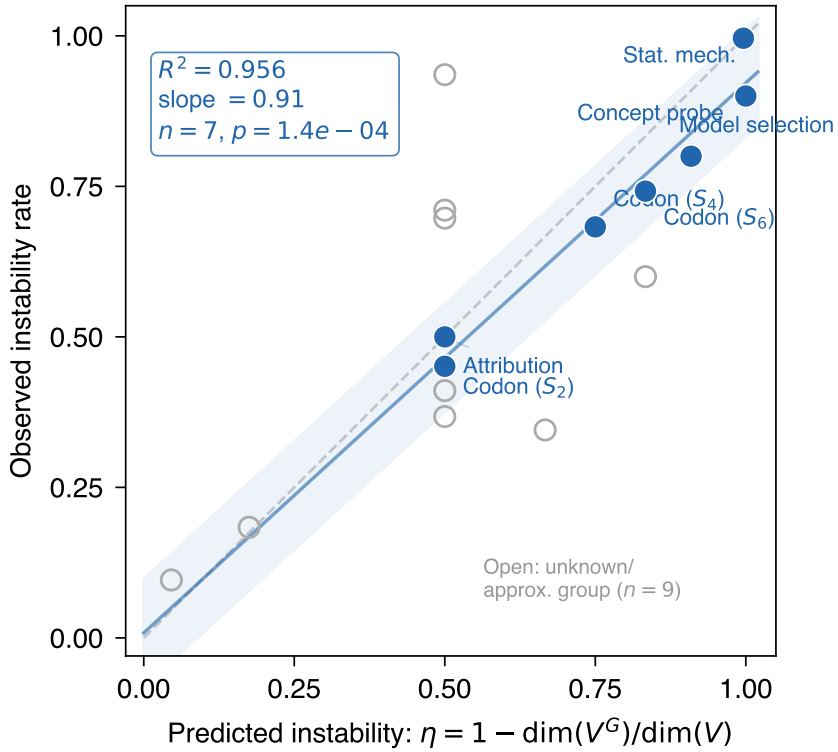

Figure 2: **Extended Data Figure 1.** The character-theoretic formula $\eta = 1 - \dim(V^G)/\dim(V)$ predicts observed instability across 16 domain instances. Filled circles: 7 instances with known symmetry groups ($R^2 = 0.957$, $p = 1.4 \times 10^{-4}$). Open circles: 9 instances with unknown or approximate groups. The bottleneck is group identification, not the law.

Acknowledgements

We acknowledge the use of Claude (Anthropic) as an AI coding assistant for experiment scripting, Lean proof development, and manuscript preparation. All scientific claims, proofs, and experimental designs were conceived and verified by the human authors.

Author contributions

D.C. conceived the framework, proved the theorems, wrote the Lean formalization, designed and executed the experiments, and wrote the manuscript. B.A. and D.R. contributed to experimental design, reviewed the proofs, and edited the manuscript.

Competing interests

The authors declare no competing interests. D.C.'s employment at Oura Health is unrelated to this work.

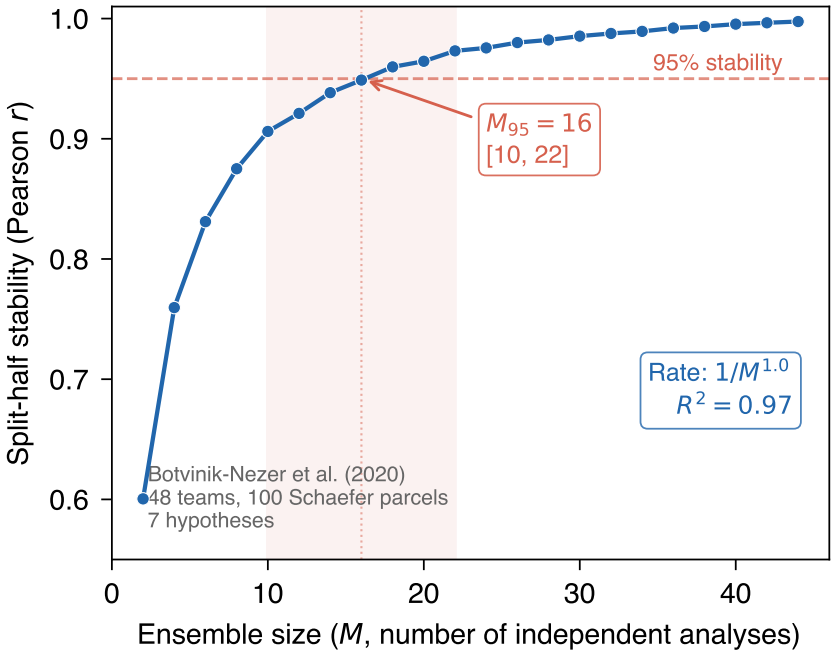

Figure 3: **Extended Data Figure 2.** Split-half stability of the ensemble average vs ensemble size for 48 neuroimaging analysis teams. Stability exceeds 0.95 at $M = 16$ [95% CI: 10–22]. Convergence rate: $1 - a/M^{1.0}$ ($R^2 = 0.97$).

References

Blair Bilodeau, Natasha Shandilya, Blossom Gao, and Umang Bhatt. Impossibility theorems for feature attribution. *Proceedings of the National Academy of Sciences*, 121(2):e2304406120, 2024.

Rotem Botvinik-Nezer, Felix Holzmeister, Colin F Camerer, Anna Dreber, Juergen Huber, Magnus Johannesson, Michael Kirchler, Roni Iwanir, Jeanette A Mumford, R Alison Adcock, et al. Variability in the analysis of a single neuroimaging dataset by many teams. *Nature*, 582(7810): 84–88, 2020.

Leo Breiman. Statistical modeling: The two cultures. *Statistical Science*, 16(3):199–231, 2001.

Nate Breznau, Eike Mark Rinke, Alexander Wuttke, Hung H T Nguyen, Muna Adem, Jule Adriaans, Amalia Alvarez-Benjumea, Henrik K Andersen, Daniel Auer, Flávio Azevedo, et al. Observing many researchers using the same data and hypothesis reveals a hidden universe of uncertainty. *Proceedings of the National Academy of Sciences*, 119(44):e2203150119, 2022.

Drake Caraker, Bryan Arnold, and David Rhoads. Impossibility of faithful stable explanation under Ostrowski's classification: A formally verified theorem applied to spacetime geometry. *Foundations of Physics*, 2026. Under review.

Leonardo de Moura and Sebastian Ullrich. The Lean 4 theorem prover and programming language. In *International Conference on Automated Deduction (CADE)*, 2021.

Table 4: **Extended Data Table 2.** Twenty impossibility theorems from 12 domains classified by tightness type. Full: every property pair achievable ("pick two" works). Collapsed: property pairs blocked (enrichment required). Structural: non-trivial Lean proof over unbounded solution space. Model: 3-element type that forces its tightness. Bell and Kochen-Specker tightness is decomposition-dependent (see text). Quantum linearity is model-level; the bilemma is the sole collapsed instance with structural evidence.

System	Domain	Tightness	Evidence
Bilemma	**ML explanation**	**Collapsed**	**structural**
Quantum linearity	**Quantum info**	**Collapsed**	**model**
Arrow	Social choice	Full	structural
Gibbard-Satterthwaite	Social choice	Full	structural
Bell (CHSH)	Quantum physics	Full	structural
Kochen-Specker	Quantum physics	Full	structural
Gödel incompleteness	Mathematical logic	Full	structural
ML fairness	ML prediction	Full	model
Bias-variance	Statistics	Full	model
Sen's liberal paradox	Social choice	Full	model
CAP theorem	Distributed systems	Full	model
FLP impossibility	Distributed systems	Full	model
Mundell-Fleming	Economics	Full	model
Perpetual motion	Thermodynamics	Full	model
Second law	Thermodynamics	Full	model
Münchhausen trilemma	Epistemology	Full	model
Competitive exclusion	Ecology	Full	model
Penrose-Hawking	General relativity	Full	model
Eastin-Knill	Quantum computing	p12-blocked	model
Shannon secrecy	Cryptography	p23-blocked	model

Satyapriya Krishna, Tessa Han, Alex Ber, Shaun Buzzelli, Michael Kearns, et al. The disagreement problem in explainable machine learning: A practitioner's perspective. In *NeurIPS 2022 Workshop on Trustworthy and Socially Responsible Machine Learning*, 2022.

Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. *arXiv preprint arXiv:2301.05217*, 2023.

Raphael Silberzahn, Eric L Uhlmann, Daniel P Martin, Pasquale Anselmi, Frederik Aust, Eli Awtrey, Štěpán Bahník, Feng Bai, Colin Bannard, Evelyne Bonnier, et al. Many analysts, one data set: Making transparent how variations in analytic choices affect results. *Advances in Methods and Practices in Psychological Science*, 1(3):337–356, 2018.

Peter Spirtes, Clark N Glymour, and Richard Scheines. *Causation, Prediction, and Search*. MIT Press, 2nd edition, 2000.

The mathlib Community. The Lean mathematical library. In *Certified Programs and Proofs (CPP)*, 2020.

Thomas Verma and Judea Pearl. Equivalence and synthesis of causal models. In *Conference on Uncertainty in Artificial Intelligence*, pages 255–270, 1991.

B T Thomas Yeo, Fenna M Krienen, Jorge Sepulcre, Mert R Sabuncu, Danial Lashkari, Marisa Hollinshead, Joshua L Roffman, Jordan W Smoller, Lilla Zöllei, Jonathan R Polimeni, et al. The organization of the human cerebral cortex estimated by intrinsic functional connectivity. *Journal of Neurophysiology*, 106(3):1125–1165, 2011.