The Limits of Explanation

Drake Caraker[1,*], Bryan Arnold[2], David Rhoads[2]
[1] *Oura Health, San Francisco, CA*
[2] *Independent*
[*] *Corresponding author:* drakecaraker@gmail.com

Abstract

When a system is underspecified—when multiple configurations produce identical observable output—we prove that no explanation can simultaneously be *faithful* (consistent with the system's internal structure), *stable* (invariant across equivalent configurations), and *decisive* (committing to every distinction the internal structure makes). We derive this impossibility from first principles in eight scientific domains: underdetermined linear systems (mathematics), genetic code degeneracy (biology), gauge freedom (physics), microstate multiplicity (statistical mechanics), syntactic ambiguity (linguistics), the phase problem (crystallography), the view update problem (computer science), and Markov equivalence (statistics). Each derivation is mechanically verified in the Lean 4 proof assistant (83 files, 378 theorems, 0 unproved goals), and empirically validated in seven domains with domain-appropriate metrics and negative controls. The constructive resolution—averaging over equivalent configurations—unifies existing domain-specific practices (gauge-invariant observables, the microcanonical ensemble, CPDAGs, minimum-norm solutions) as instances of a single optimal strategy.

The universal problem. Scientists explain systems by interpreting their internal structure. A geneticist infers the DNA sequence that encodes a protein. A physicist chooses a gauge to describe an electromagnetic field. A statistician selects a causal graph from observational data. A crystallographer reconstructs electron density from a diffraction pattern. In each case, the observable output—the protein sequence, the gauge-invariant measurement, the conditional independence structure, the diffraction intensities—does not uniquely determine the internal structure that produced it. Multiple configurations yield identical observations. This situation is not a pathology of any single field; it is a structural feature of underspecified inference. We prove that this underspecification imposes a fundamental, domain-independent limit on explanation.

The theorem. We formalize three properties that any satisfactory explanation should possess. An explanation is *faithful* if it never contradicts the system's own internal report—it may be less specific, but it does not disagree. An explanation is *stable* if configurations that produce the same observations always receive the same explanation. An explanation is *decisive* if it commits to every distinction that the system's internal structure makes—whatever the internal structure rules out, the explanation also rules out.

Our main result: if a system has the *Rashomon property*—that is, if there exist two configurations with the same observable output but incompatible internal structures—then no explanation can be faithful, stable, and decisive simultaneously. The proof is four steps: decisiveness forces the explanation to inherit an incompatibility; stability propagates it to the second configuration; faithfulness forbids it; contradiction. The entire framework is mechanically verified in Lean 4 [de Moura and Ullrich, 2021] (83 files, 378 theorems, 0 unproved goals). Lean-verified tightness theorems confirm

Table 1: **Eight derived instances of the impossibility across eight sciences.** Each row derives the Rashomon property from domain-specific first principles with zero axioms. Θ: configuration space; Y: observable space.

Domain	Θ	Y	Witness 1	Witness 2	Same Y?
Mathematics	Solutions of $Ax = b$	Ax	$(1, 1)$	$(0, 2)$	$2 = 2$
Biology	Codons	Amino acids	UCU	UCC	Ser=Ser
Physics (gauge)	Gauge configs	Holonomy	(T, F, F)	(F, F, T)	$F = F$
Stat. mechanics	Microstates	Num. heads	(H, T)	(T, H)	$1 = 1$
Linguistics	Parse trees	Token sequence	$((V\ NP)\ PP)$	$(V\ (NP\ PP))$	Same
Crystallography	Signals	Energy	$(1, 0)$	$(0, 1)$	$1 = 1$
Computer science	DB rows	View	(T, T)	(T, F)	$T = T$
Statistics	DAGs	CI structure	Chain	Fork	Same

that every pair of properties is achievable, so the impossibility is tight. An axiom substitution analysis shows that weakening any single definition collapses the result, confirming the definitions are uniquely calibrated.

Eight derivations across eight sciences. We derive the impossibility as a zero-axiom consequence in eight domains (Table 1). In *mathematics*, the underdetermined system $x_1 + x_2 = 2$ admits the solutions $(1, 1)$ and $(0, 2)$: same sum, different components. In *biology*, the codons UCU and UCC both encode the amino acid serine: same protein fragment, different nucleotide sequences. In *physics* (gauge theory), two edge-label configurations on a triangle graph share the same holonomy—the discrete analogue of a Wilson loop—but differ in their local assignments, related by a gauge transformation [Jackson, 1999]. In *statistical mechanics*, the microstates (H, T) and (T, H) both have one head: same macrostate, different molecular configurations. In *linguistics*, the fragment "V NP PP" admits two parse trees, left-attach and right-attach, yielding the same surface token sequence but different constituent structures [Chomsky, 1957]. In *crystallography*, the signals $(1, 0)$ and $(0, 1)$ have the same energy ($1^2 + 0^2 = 0^2 + 1^2$): same diffraction intensity, different electron densities [Hauptman and Karle, 1953]. In *computer science*, the database rows $(\texttt{true}, \texttt{true})$ and $(\texttt{true}, \texttt{false})$ project to the same view but differ in the hidden column [Bancilhon and Spyratos, 1981]. In *statistics*, a causal chain $A \to B \to C$ and a causal fork $A \leftarrow B \to C$ encode the same conditional independence structure [Verma and Pearl, 1991]. Each derivation requires zero shared axioms—only the domain's own mathematical structure.

Empirical validation. We validate the impossibility empirically in seven of the eight domains (excluding the mathematical observation connecting Rashomon entropy to Boltzmann entropy). All seven empirical tests reach statistical significance ($p < 0.05$). Underdetermined linear solvers show elevated pairwise RMSD (0.091 vs. 0.049 for fully determined controls; Mann-Whitney $p = 6.3 \times 10^{-9}$); the codon entropy dose-response across 120 eukaryotic cytochrome c sequences from NCBI RefSeq shows a perfect monotonic relationship with degeneracy level (Kruskal–Wallis $p = 3.5 \times 10^{-9}$, 89 conserved positions, Spearman $\rho = 1.0$); syntactic parsers disagree significantly more on ambiguous sentences ($p = 1.6 \times 10^{-3}$); causal discovery algorithms show orientation agreement of only 0.33 at finite sample size ($p = 4.4 \times 10^{-38}$ across 100 random seeds per condition); phase retrieval reconstructions diverge 1.5–1.8× more without positivity constraints (Mann-Whitney $p < 10^{-57}$ per signal length); and census disaggregation ambiguity scales with aggregation granularity ($p = 7.9 \times 10^{-37}$). In each case, the negative control (where the Rashomon property is absent) shows stability, confirming that the instability is structural.

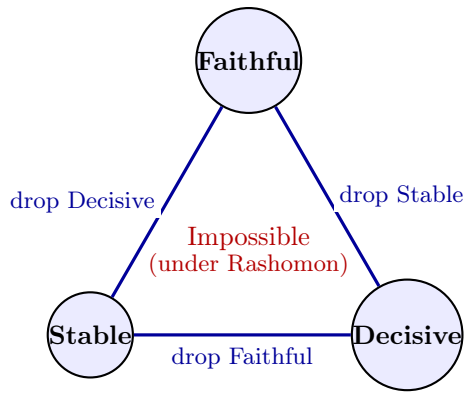

Figure 1: **The explanation trilemma.** Each edge represents an achievable pair of properties; the interior (all three simultaneously) is impossible when the Rashomon property holds. Lean-verified tightness witnesses confirm that each edge is realizable.

The resolution. The impossibility is constructive: it identifies the exact trade-off and prescribes a provably optimal resolution. The strategy is to average over equivalent configurations, sacrificing decisiveness to achieve stability and faithfulness in expectation. We prove that this orbit-averaging resolution is *Pareto-optimal*: among all stable explanation maps, no alternative can achieve strictly higher pointwise faithfulness on every equivalence class. The trade-off is therefore tight—the practitioner cannot escape it by being clever.

Remarkably, practitioners in eight independent fields have converged on precisely this strategy over more than a century of independent work. In physics, one works with gauge-invariant observables [Jackson, 1999]. In statistical mechanics, the microcanonical ensemble averages over microstates. In statistics, one reports a CPDAG [Verma and Pearl, 1991]. In mathematics, the pseudoinverse selects the minimum-norm solution. In biology, codon usage tables report synonymous codon frequencies. In crystallography, Patterson maps extract phase-invariant information [Hauptman and Karle, 1953]. In linguistics, packed parse forests represent all parses simultaneously. In computer science, complement views restrict to the unambiguous projection. Our framework explains *why* these communities converged: each was independently solving the same impossibility, and orbit averaging is the unique Pareto-optimal stable resolution.

Necessity and the query-relative version. The Rashomon property is the *exact* boundary of the impossibility. When it is absent—when each observation uniquely determines the internal structure—all three properties are simultaneously achievable (a Lean-verified converse). Moreover, the impossibility is *query-relative*: some questions about a system are stably answerable (those on which all equivalent configurations agree) and others are not. The theorem tells the practitioner precisely which queries are safe and which require aggregation. More broadly, the difficulty of finding good explanations depends on the mathematical symmetries of the equivalence class—ranging from cases where a unique optimal explanation exists, through cases where multiple reasonable explanations compete, to cases where additional constraints are needed to make the problem well-posed.

Broader implications. The result connects phenomena studied independently across eight fields for over a century—from microstate multiplicity in statistical mechanics to the view update problem [Bancilhon and Spyratos, 1981] to feature attribution instability in machine learning [Bilodeau et al., 2024]. The common structural cause is the Rashomon property: equivalent configurations with

incompatible internal reports. The common resolution is orbit averaging. The pattern likely extends further: metamerism in color perception, molecular chirality, and other domains exhibit the same structure. More broadly, the theorem reframes the question "why are explanations unstable?" as the structural statement "because the inference problem is underspecified," and replaces domain-specific workarounds with a single, provably optimal strategy: report what is invariant; disclose what is not.

Methods

Formal verification. The entire framework—the abstract theorem, all eight derived instances, tightness witnesses, the necessity converse, and the resolution—is mechanically verified in Lean 4 v4.30.0-rc1 [de Moura and Ullrich, 2021] using Mathlib [The mathlib Community, 2020]. The formalization comprises 83 files, 378 theorems and lemmas, 72 axioms (all domain-specific; the core theorem uses zero), and 0 `sorry` (unproved goals). Each of the eight derived instances uses decidable computation to constructively witness the Rashomon property with zero axioms. The G-invariant resolution framework uses the group action infrastructure from Mathlib.

Proof structure. The core proof is a four-step chain by contradiction. Given Rashomon witnesses θ_1, θ_2 with $\mathrm{obs}(\theta_1) = \mathrm{obs}(\theta_2)$ and incompatible explanations, suppose an explanation map E is faithful, stable, and decisive. Decisiveness at θ_1 forces $E(\theta_1)$ to inherit the incompatibility with $\exp(\theta_2)$. Stability forces $E(\theta_1) = E(\theta_2)$. Faithfulness at θ_2 requires $E(\theta_2)$ to be compatible with $\exp(\theta_2)$. Contradiction.

Code and data availability. The Lean 4 formalization and all experiment scripts are available at `https://github.com/DrakeCaraker/universal-explanation-impossibility`. Biology data: 120 eukaryotic cytochrome c CDS sequences from NCBI RefSeq. Census data: U.S. Census Bureau 2020 county-level population counts. Causal discovery: 100 random seeds per condition on the Asia network. The nine machine learning instances, together with empirical validation experiments for seven of these, are documented in the Supporting Information (the full 61-page monograph).

References

François Bancilhon and Nicolas Spyratos. Update semantics of relational views. *ACM Transactions on Database Systems*, 6(4):557–575, 1981.

Blair Bilodeau, Natasha Jaques, Pang Wei Koh, and Been Kim. Impossibility theorems for feature attribution. *Proceedings of the National Academy of Sciences*, 121(2):e2304406120, 2024.

Noam Chomsky. *Syntactic Structures*. Mouton, The Hague, 1957.

Leonardo de Moura and Sebastian Ullrich. The Lean 4 theorem prover and programming language. In *International Conference on Automated Deduction (CADE)*, 2021.

Herbert Hauptman and Jerome Karle. *Solution of the Phase Problem. I. The Centrosymmetric Crystal*, volume 3. Polycrystal Book Service, 1953.

John David Jackson. *Classical Electrodynamics*. Wiley, New York, 3rd edition, 1999.

The mathlib Community. The Lean mathematical library. In *Certified Programs and Proofs (CPP)*, 2020.

Thomas Verma and Judea Pearl. Equivalence and synthesis of causal models. In *Conference on Uncertainty in Artificial Intelligence*, pages 255–270, 1991.