

The Limits of Explanation

Drake Caraker*

Bryan Arnold[†]

David Rhoads[‡]

Independent Researchers

July 2026

Abstract

We prove a universal impossibility: no explanation system can simultaneously be faithful (never contradicting the system’s native explanation), stable (consistent across observationally equivalent configurations), and decisive (inheriting every incompatibility of the native explanation) when the underlying inference problem is underspecified—that is, when observationally equivalent configurations produce incompatible explanations (the Rashomon property). This single abstract theorem, which we call the *Explanation Impossibility*, is a meta-theorem: it holds for *every* system whose configurations, observations, and explanations satisfy the Rashomon property. A ubiquity argument (Part I) indicates that this hypothesis is generically met whenever the parameter dimension exceeds the observable dimension—the regime of essentially all overparameterized models. Its linear, infinitesimal, and *local smooth* cores are machine-checked: rank-nullity forces genuinely distinct, observationally identical configurations; a differentiable observation map has a nonzero derivative-annihilated direction at every point; and, via the implicit function theorem, at a *regular* point of a nonlinear observation map the true curved fiber contains genuinely distinct models arbitrarily close. For linear observation maps the remaining genericity step (Sard’s theorem) is proved outright; for nonlinear maps the full Morse–Sard theorem itself is now machine-checked in-repo—via an in-tree, attribution-preserving port of Y. Kudryashov’s Lean 4 formalisation [Kudryashov, 2025] of Moreira’s version of Sard’s theorem [Moreira, 2001]—so the generic ubiquity statement is unconditional for C^{n-m+1} observation maps, the sharp classical smoothness threshold (see Part I). We then instantiate the theorem on explicit minimal witnesses in eight scientific domains: the degeneracy of the genetic code (biology), gauge freedom in electromagnetism (physics), microstate degeneracy in statistical mechanics (physics), syntactic ambiguity in natural language (linguistics), the phase problem in crystallography, underdetermined linear systems (mathematics), the database view update problem (computer science), and Markov equivalence in causal inference (statistics). We further instantiate it across nine ML explanation types: feature attribution under collinearity [Bilodeau et al., 2024], attention map instability, counterfactual explanation instability, concept probe (TCAV) instability, causal discovery ambiguity [Verma and Pearl, 1991], model selection instability under the Rashomon effect [Rudin et al., 2024, Fisher et al., 2019], saliency map (GradCAM) instability, token citation explanation instability, and mechanistic interpretability non-uniqueness. The proof is a tight four-step chain—decisiveness, stability, faithfulness, contradiction—and Lean-verified tightness theorems show that each pair of properties is achievable, confirming that the impossibility is non-trivial. Cross-domain experiments empirically validate the impossibility in seven of the eight scientific domains (Table 12), including a protein-level test on 120 cytochrome c sequences where the capacity formula predicts sequence diversity with $r = 0.798$ ($p = 0.018$) and systematic undershoot attributable to purifying selection, significant parser disagreement on ambiguous sentences ($p = 9 \times 10^{-4}$), and causal discovery convergence only at infinite sample size. Empirical experiments on four ML instances confirm severe instability among functionally equivalent

*Corresponding author: drakecaraker@gmail.com

[†]puremath86@gmail.com

[‡]drhoads9@gmail.com

models: 19.9% attention argmax flip rate (under full retraining; 60% under weight perturbation), 23.5% counterfactual direction flip rate, 0.90 concept direction instability, and 80% model selection flip rate. The resolution is uniform: G -invariant explanation systems that aggregate over equivalent hypotheses restore faithfulness-in-expectation and stability at the cost of decisiveness, and DASH (**D**iversified **A**ggregation for **S**table **H**ypotheses) is a concrete instance of this resolution for feature attribution. The framework is quantitative as well as qualitative. The resolution has a *capacity*— $\dim(V^G)$, the machine-checked *exact* count of stable explanation facts (exact because a keystone theorem shows the stable maps are exactly the G -invariant ones), which a structure theorem exhibits together with the resolution and the incurable structural floor as facets of a single orbit-average operator—and a per-query *reliability certificate*: a distribution-free, Lean-verified bound on the flip rate of any individual explanation claim, which we prove transfers to independent ensembles and validate on a 104-dataset, 10,724-pair out-of-sample study in which every STABLE verdict transfers (flip $\leq 20\%$) to an independent ensemble. The impossibility is ubiquitous: any learning problem with strictly more parameters than constraints admits a Rashomon set of positive measure, and neural networks with hidden dimension exceeding the data dimension always fall into this regime. The framework is mechanically verified in Lean 4 with **0 sorry** (a real **lake build** of all three repositories completes with 0 errors). The trusted base is **4 axiom** declarations across the three repositories—2 in the main repository (bundling ~ 14 gradient-boosting properties), 2 in the attribution companion (bundling 7 fields), and 0 in the physics companion, which instead carries ~ 11 domain hypotheses (including the classical Selmer 1951 result) as section variables rather than axioms. Exactly *one* theorem (a 419,904-case circuit enumeration in the physics companion) additionally depends on Lean’s `ofReduceBool` via `native_decide`; every other enumeration is kernel-checked. The core theorem and the entire hardened spine (Part I) use *none* of these. Theorem counts are reported per repository, not summed across the overlapping companions. We also report **five falsified predictions** of the framework (Part V) alongside the confirmed results: the boundary of a theory is part of its content, and we state ours explicitly.

Contents

Reading guide: how this document is organised	7
1 Introduction	7
2 Related Work	9
I Part I — The Hardened Core	12
3 The Abstract Framework [Tier A: Hardened]	12
3.1 Explanation Systems [Tier A: Hardened]	13
3.2 The Three Desiderata [Tier A: Hardened]	13
3.3 The Rashomon Property [Tier A: Hardened]	14
3.4 Main Theorem [Tier A: Hardened]	14
3.5 Tightness [Tier A: Hardened]	14
3.6 Strengthened Impossibility for Maximally Incompatible Systems [Tier A: Hardened]	15
3.7 The Enrichment Mechanism [Tier A: Hardened]	16
4 Ubiquity [Tier A: Hardened]	18
4.1 Generic Underspecification [Tier A: Hardened]	18
4.2 Neural Network Dimensional Argument [Tier A: Hardened]	20
4.3 Prevalence [Tier A: Hardened]	20
5 The Universal Resolution [Tier A: Hardened]	21
5.1 G-Invariant Explanation Systems	21
5.2 Instance-Specific Resolutions [Tier A: Hardened]	22
5.3 DASH as the Prototype Resolution [Tier A: Hardened]	22
5.4 Meta-Theorem: Orbit Averaging [Tier A: Hardened]	23
6 The Explanation Uncertainty Bound [Tier A: Hardened]	23
7 The Invariance Counting Principle [Tier A: Hardened]	24
8 Toward a Group-Theoretic Classification of Explainability [Tier A: Hardened]	24
9 The Ceiling of Explainability [Tier A: Hardened]	33
10 Lean Formalization [Tier A: Hardened]	33
10.1 Summary Statistics [Tier A: Hardened]	34
10.2 Axiom Stratification [Tier A: Hardened]	35
10.3 Proof Status Transparency [Tier A: Hardened]	37
10.4 Formerly Axiomatized, Now Derived [Tier A: Hardened]	38
10.5 Lean Proof of the Core Theorem [Tier A: Hardened]	38
10.6 Why This Result Is Non-Trivial [Tier A: Hardened]	38
II Part II — Empirical Validation	41

11 Cross-Domain Transfer: From Taxonomy to Generative Theory [Tier B: Empirical]	41
11.1 Resolution Transfer: Ensemble Causal Discovery [Tier B: Empirical]	41
11.2 Stability Transition in Explanation Systems [Tier B: Empirical]	42
11.3 Cross-Level Prediction: Protein Designability [Tier B: Empirical]	43
11.4 Cross-Level Prediction: Transcription Factor Binding Sites [Tier B: Empirical]	44
11.5 Invariance Counting: Exact Enumeration of Stable Queries [Tier A: Hardened]	45
11.6 Automatic Symmetry Discovery (SAGE) [Tier B: Empirical]	46
11.7 Discrete vs. Continuous Symmetry: The CCA Spectrum [Tier B: Empirical]	47
11.8 The Explanation Capacity Theorem: Predictive Scope and Boundary [Tier B: Empirical]	47
11.9 Per-Individual Explanation Reversal [Tier B: Empirical]	50
11.10 Prospective Out-of-Domain Prediction: Drug Discovery [Tier B: Empirical]	51
11.11 Exhaustive Pipeline Validation: 85 PMLB Datasets [Tier B: Empirical]	52
11.12 The Gaussian Flip Rate: From Exact to Approximate Symmetry [Tier B: Empirical]	54
11.13 Domain-Specific Validation [Tier B: Empirical]	56
11.14 Gene Expression Pathway Divergence [Tier B: Empirical]	58
11.15 Mechanistic Interpretability: Circuit-Level Rashomon [Tier B: Empirical]	59
11.16 Language Model Circuit Stability: Multi-Scale Validation [Tier B: Empirical]	60
11.17 Model-Class Universality of the Impossibility [Tier B: Empirical]	63
11.18 Multi-Analyst Disagreement as Impossibility: Neuroimaging [Tier B: Empirical]	64
11.19 Cross-Domain Multi-Analyst Studies [Tier B: Empirical]	66
12 Fundamental Ceiling on AI Interpretability [Tier B: Empirical]	67
13 Empirical Predictions from the Mathematical Structure [Tier B: Empirical]	69
14 Regulatory Implications [Tier B: Empirical]	70
14.1 A quantitative adverse-action test [Tier B: Empirical]	70
14.2 What Explainability Can Do	71
15 What This Paper Tells a Working Scientist [Tier B: Empirical]	71
III Part III — Instantiations and Bridges	74
16 Derived Instances Across Eight Sciences [Tier A: Hardened]	74
16.1 Underdetermined Linear Systems (Mathematics)	74
16.2 Genetic Code Degeneracy (Biology)	75
16.3 Gauge Freedom (Physics)	77
16.4 Microstate Degeneracy (Statistical Mechanics)	78
16.5 Syntactic Ambiguity (Linguistics)	79
16.6 The Phase Problem (Crystallography)	80
16.7 The View Update Problem (Computer Science)	81
17 Applications in Machine Learning [Tier B: Empirical]	82
17.1 Instance 1: Feature Attribution Under Collinearity	83
17.2 Instance 2: Attention Maps as Explanations	85
17.3 Instance 3: Counterfactual Explanations Under Decision-Boundary Multiplicity	87
17.4 Instance 4: Concept Probe Impossibility	90
17.5 Instance 5: Causal Discovery Within a Markov Equivalence Class	93
17.6 Instance 6: Model Selection Under the Rashomon Effect	95
17.7 Instance 7: Saliency Maps (GradCAM) Under Weight-Space Multiplicity	97

17.8	Instance 8: Token Citation Explanations Under Model Multiplicity	98
17.9	Instance: Mechanistic Interpretability (Circuit Discovery)	99
18	Bridge Theorems: Connections to Established Mathematics [Tier A: Hardened]	103
18.1	Sufficient Statistics as G -Invariants [Tier A: Hardened]	103
18.2	EM as Generalized Orbit Averaging [Tier B: Empirical]	104
IV	Part IV — Structural Analogies and Speculative Extensions	105
19	Discussion: What Explainability Can and Cannot Do [Tier C: Speculative]	105
19.1	Cross-Domain Transfer: From Gauge Theory to ML [Tier C: Speculative]	105
19.2	Scope and Applicability Beyond Machine Learning [Tier C: Speculative]	106
19.3	The Isn't This Obvious Defense	111
19.4	The All Instances Are Trivial Defense	112
19.5	Implications Beyond Machine Learning [Tier C: Speculative]	112
19.6	Scientific Revolutions as Enrichment Events [Tier C: Speculative]	115
19.7	Enrichment as Abstraction [Tier C: Speculative]	117
19.8	Recursive Impossibility: The Enrichment Stack [Tier C: Speculative]	117
20	Analogy: Arrow's Impossibility Theorem [Tier C: Speculative]	119
21	Bayesian Interpretation [Tier C: Speculative]	120
22	Quantum Error Correction [Tier C: Speculative]	120
23	Classification of Impossibility Theorems by Tightness [Tier C: Speculative]	120
23.1	Tightness Types [Tier C: Speculative]	120
23.2	The Classification [Tier C: Speculative]	121
23.3	Key Findings [Tier C: Speculative]	121
23.4	Evidence Tiers and Caveats [Tier C: Speculative]	124
V	Part V — Falsified Predictions and Boundaries	125
24	Limitations and Open Questions [Tier D: Falsified]	125
25	Falsification: The Rashomon Boundary [Tier D: Falsified]	129
26	Boundary Conditions: The Quantitative Boundary and Three Falsified Ex-	
	tensions [Tier D: Falsified]	130
27	Attribution PCA Eigenspectrum: A Negative Result [Tier D: Falsified]	132
A	Arrow's Impossibility Theorem: A Structural Comparison	139
A.1	Formal setup	139
A.2	The axiom correspondence	139
A.3	Why the embedding fails	140
A.4	What the theorems share	141
A.5	Implications	142

B	Complete Lean Theorem Statements	142
B.1	ExplanationSystem.lean — Core Framework and Impossibility	142
B.2	UniversalResolution.lean — G-Invariant Resolution	145
B.3	Ubiquity.lean — Structural Ubiquity	146
B.4	Instance Theorems	147
C	Detailed Experiment Methodology	148
C.1	Experiment 1: Attention Map Instability (Instance 2)	149
C.2	Experiment 2: Counterfactual Instability (Instance 3)	149
C.3	Experiment 3: Concept Probe Instability (Instance 4)	150
C.4	Experiment 4: Model Selection Instability (Instance 6)	151
D	Proof Sketches for All Nine Instances	151
D.1	Instance 1: Feature Attribution	151
D.2	Instance 2: Attention Maps	152
D.3	Instance 3: Counterfactual Explanations	152
D.4	Instance 4: Concept Probes	152
D.5	Instance 5: Causal Discovery	153
D.6	Instance 6: Model Selection	153
E	Companion Results: Number Theory and Physics	153
E.1	Ostrowski’s Classification and the Mathlib Bridge	154
E.2	Freund–Witten Three-Fold Zeta Symmetry	154
E.3	Adelic Resolution	154
E.4	Circuit Complexity: AC^0 Parity	154
E.5	DPRM Trilemma and the Selmer Curve	154
E.6	Navier–Stokes Conditional Tightness	155
E.7	The N -Property Generalisation and Myerson–Satterthwaite	155
E.8	The Enrichment–Riemann Hypothesis Program	155
E.9	Langlands Extensions	156
F	Companion Results: Attribution-Specific Formalization	156
F.1	Constructive Bilemma Instances	156
F.2	BeyondBinary: Enrichment Restores Achievability	157
F.3	Mechanistic Interpretability Bilemma	157
F.4	Additional Empirical Results	157
G	Companion Results: DASH Method and Empirical Validation	157
G.1	DASH Benchmark Results	158
G.2	Crossed ANOVA: Variance Decomposition	158
G.3	Colsample Ablation	158
G.4	Stochastic Retrain Equivalence	158
G.5	$\text{Var}[\text{SHAP}] = \text{DASH MSE Identity}$	158
G.6	Feature Stability Index (FSI)	158
G.7	The $F5 \rightarrow F1 \rightarrow \text{DASH}$ Pipeline	158
G.8	Retracted Results	159
G.9	Extension Modules	159
	Appendix: Lean 4 Cross-Reference	160
G.10	Cross-Reference: Monograph $\leftrightarrow$ Lean	160
G.11	Complete Theorem Index	160

Reading guide: how this document is organised

This monograph is the umbrella reference for a research program formalised across three Lean 4 repositories (the universal framework, an attribution-specific companion, and a mathematical-physics companion) with accompanying empirical studies. To let readers calibrate trust at a glance, *every section is tagged with an evidence tier*, and the body is organised into five Parts by tier:

- Part A — Hardened.** Machine-checked in Lean 4 from zero custom axioms or from explicitly stated hypotheses; the hard content is genuinely proved. **Part I** collects the load-bearing theory.
- Part B — Empirical.** Supported by experiments/audits, with sample sizes, effect sizes, and caveats stated inline. These test the theory; they are not logical consequences of it. **Part II.**
- Part C — Illustrative / analogical.** Structural parallels and conditional or illustrative formalizations—*not* domain results. **Part IV** (with an explicit status disclaimer); **Part III** holds the genuine instantiations and bridges.
- Part D — Falsified / boundary.** Predictions that were tested and rejected, reported prominently by design. **Part V.**

Claims that connect the framework to the Langlands programme, adelic physics, Gödel’s theorem, or the Riemann Hypothesis are *analogies*, not proofs, and live in Part IV; the genuine underlying theorems (e.g. the $GL(n)$ trace bilemma, the completed-zeta three-fold symmetry) are stated separately with their interpretations flagged.

Note on methodology and AI assistance. This manuscript, its Lean 4 formalization, and its computational experiments were developed with substantial assistance from AI coding and reasoning tools. The guarantees the work rests on are mechanical, not authorial: every theorem is checked by the Lean kernel (0 `sorry`; trusted base stated in Part I), and every reported statistic is reproducible from the public repositories and released scripts. The authors take full responsibility for all claims and framing.

1 Introduction

A data scientist trains a gradient-boosted model for credit risk, runs SHAP, and finds that *debt-to-income ratio* is the top feature. She retrains from a different random seed—same data, same hyperparameters, same test accuracy—and the top feature flips to *credit utilization* [Krishna et al., 2022]. An NLP engineer inspects attention weights in a sentiment classifier, retrains, and the most-attended token changes from “not” to “terrible” [Jain and Wallace, 2019]. A colleague runs LIME on the same credit model and obtains yet a third ranking [Molnar et al., 2022]. None of these are software bugs.

The explainable AI (XAI) research program promises explanations that are *faithful* (they reflect what the model actually computes), *stable* (they are reproducible across equivalent models), and *decisive* (they resolve every explanation query with a definite answer). Regulators now demand all three: the EU AI Act requires “meaningful explanations” for high-risk systems [EU Parliament, 2024], and SR 11-7 requires “developmental evidence” documenting model decisions [Office of the Comptroller of the Currency, 2011]. Yet a growing body of theoretical work shows that the promise is mathematically impossible to keep—at least in specific domains. Bilodeau et al. [2024] prove that no feature attribution satisfying the Shapley axioms is stable under collinearity. Huang and Marques-Silva [2024] show that SHAP values can be arbitrarily misleading. Chouldechova [2017] and Kleinberg et al. [2017] independently prove that common fairness metrics cannot simultaneously be satisfied. Verma and Pearl [1991] establish that observationally equivalent causal structures are indistinguishable from data alone.

Each of these impossibilities was proved independently, in domain-specific language, under domain-specific assumptions. The attribution impossibility assumes collinearity among features; Chouldechova’s impossibility assumes unequal base rates across groups; Markov equivalence assumes faithfulness of the directed acyclic graph. No one has asked: is there *one* structural cause? Is the impossibility intrinsic to underspecification itself, rather than to any particular domain?

In this paper we answer affirmatively. We identify a structural condition—the Rashomon property—under which no explanation can simultaneously be faithful, stable, and decisive. The proof is four lines; the content is in the consequences. From this structural foundation we derive: a parameter-free quantitative predictor of instability (the explanation capacity $\dim(V^G)/\dim(V)$, confirmed at $R^2 = 0.957$ across seven domains), a canonical resolution (orbit averaging; Pareto-optimal *within* correlation groups by a Lean-verified theorem, and globally by a classical invariant-decision-theory argument not yet mechanized), and a model-agnostic diagnostic pipeline validated on 30 datasets $\times$ 6 model/method configurations (Cochran’s Q , $p = 0.25$; no model or method effect). The framework has the structural form of a capacity theory—computable limit, explicit optimal code, tight converse, convergence rate—though the mathematics is classical (representation theory, the Reynolds operator); the contribution is identifying that this machinery applies to explanation stability and validating the application empirically.

The primary contributions of this paper are original research: the impossibility theorem and its Lean 4-verified formalization (Part I), and an empirical characterization of explanation instability including a pre-registered stable-fact-count result (Part II). For completeness the paper also documents the cross-domain instantiations (Part III), clearly-labeled speculative connections (Part IV), and five falsified predictions (Part V). Material is organized by evidence tier so that the original, machine-checked results are cleanly separated from exploratory and analogical content; the latter is included for transparency, not as established claims. A focused version of the empirical contribution is presented separately.

Our contributions are:

- (1) **A structural impossibility theorem** (Theorem 3.6) identifying the Rashomon property as the universal condition generating explanation instability, with constructive witnesses across eight scientific domains and nine ML explanation types.
- (2) **Eight cross-domain instances** (Section 16): underdetermined linear systems (mathematics), genetic code degeneracy (biology), gauge freedom (physics), microstate degeneracy in statistical mechanics (physics), syntactic ambiguity (linguistics), the phase problem (crystallography), the view update problem (computer science), and Markov equivalence (statistics)—each derived with zero domain-specific axioms from first principles.
- (3) **Nine ML instances** (Section 17): feature attribution [Bilodeau et al., 2024], attention maps [Jain and Wallace, 2019], counterfactual explanations, concept probes, causal discovery [Verma and Pearl, 1991], model selection [Rudin et al., 2024], saliency maps (Grad-CAM), token citation explanations, and mechanistic interpretability (circuit discovery)—each shown to be a special case of the abstract framework.
- (4) **Empirical illustration** for seven of the nine ML instances (attribution, attention, counterfactual, concept probes, model selection, saliency maps, and token citation), confirming explanation instability on standard benchmarks.
- (5) **A uniform resolution**: G -invariant aggregation restores faithfulness and stability at the cost of decisiveness. DASH for feature attribution, CPDAGs for causal discovery, and ensemble probes for concepts are all instances of this resolution.

- (6) **Ubiquity:** the impossibility is generic whenever the parameter dimension exceeds the observable dimension, covering virtually all modern overparameterized models.
- (7) **Mechanical verification** in Lean 4 [de Moura and Ullrich, 2021]: 130 files, 715 theorems, 2 axioms, 0 sorry.
- (8) **Fairness tightness transition.** The Chouldechova–Kleinberg fairness impossibility transitions from full tightness (all pairs achievable by degenerate classifiers) to PPV-blocked (calibration as universal poison) under operational performance constraints ($\text{TPR} \geq 0.3$). Equalized odds is the only achievable pairwise compromise at practical performance. Validated on Adult Income and synthetic data.
- (9) **Model-agnostic pipeline validation.** A factorial design (30 datasets $\times$ 6 model/method configurations) confirms no significant model or method effect (Cochran’s Q , $p = 0.25$). DASH preserves accuracy in 156/156 tests. SAGE achieves 92% directional accuracy. Five predictions falsified (honest science).

In short: faithful, stable, decisive—pick two.

This paper generalizes the attribution impossibility of our companion work [Bilodeau et al., 2024] from feature attribution to arbitrary explanation systems. The companion paper provides quantitative depth (divergence rates, ensemble size bounds, DASH optimality) for the attribution instance; this paper supplies the abstract framework, eight additional instances, and the ubiquity result. The two papers share a Lean codebase but are self-contained.

The remainder of the paper is organized as follows. Section 2 surveys related work. Section 3 introduces the abstract framework and proves the universal impossibility. Section 16 derives the impossibility in eight scientific domains with zero domain-specific axioms. Section 17 instantiates the theorem across nine ML explanation types. Section 5 presents the uniform G -invariant resolution. Section 4 establishes ubiquity. Section 10 describes the Lean formalization. Section 19 discusses implications and open questions.

2 Related Work

Impossibility results in XAI. Bilodeau et al. [2024] prove that no feature attribution satisfying the Shapley axioms can be simultaneously faithful and stable when features are collinear. Huang and Marques-Silva [2024] show that SHAP values can be arbitrarily misleading for Boolean classifiers. Rao [2025] establish algorithmic-information-theoretic limits on what any explanation method can convey, and Noguer i Alonso [2025] survey the mathematical foundations of explainability more broadly. Each of these results is proved in a domain-specific language under domain-specific assumptions. Our result complements Bilodeau et al. (2024): they prove impossibility under Shapley axioms for feature attributions; we prove impossibility under faithfulness, stability, and decisiveness for arbitrary explanation systems. The conditions are incomparable—neither strictly implies the other—but both identify the Rashomon property as the root cause. Our contribution is a single abstract theorem that unifies them as instances, showing the impossibility arises from underspecification itself rather than from any particular axiom system. Krishna et al. [2022] document the complementary “disagreement problem”: different explanation methods produce different attributions for the same model. Our result addresses a distinct phenomenon—any single method produces different explanations across equivalent models.

Underspecification and model multiplicity. D’Amour et al. [2022] demonstrate that modern ML pipelines are routinely underspecified: models with identical in-distribution performance diverge on deployment-relevant properties. Fisher et al. [2019] introduce *model reliance* and

variable-importance clouds that visualize how feature importance varies across the set of near-optimal models. [Rudin et al. \[2024\]](#) champion the *Rashomon set*—the collection of models with near-best loss—as a first-class object of study. [Semenova et al. \[2022\]](#) prove the existence of simpler models within Rashomon sets, and [Marx et al. \[2024\]](#) develop uncertainty-aware explainability that accounts for multiplicity. Our framework formalizes the bridge between underspecification and explanation impossibility: the Rashomon property (Definition 3.5) is precisely the condition under which equivalent hypotheses yield incompatible explanations.

Fairness impossibilities. [Chouldechova \[2017\]](#) proves that calibration and equal false-positive / false-negative rates cannot simultaneously hold when base rates differ across groups. [Kleinberg et al. \[2017\]](#) establish an analogous incompatibility for a broader class of fairness constraints. We recast these as Instance 6 of the universal impossibility: the explanation system maps each classifier to a fairness certificate, and the Rashomon property holds whenever the underlying population has unequal base rates.

Formal verification of impossibility theorems. [Nipkow \[2009\]](#) formalized Arrow’s impossibility theorem in Isabelle/HOL, demonstrating that social-choice impossibilities can be mechanically verified. [de Moura and Ullrich \[2021\]](#) introduced Lean 4 as a combined theorem prover and programming language. [Zhang et al. \[2026\]](#) recently formalized statistical learning theory in Lean 4. Our work provides, to our knowledge, the first Lean formalization of an impossibility theorem in explainability, covering all nine instances and the G -invariant resolution (101 files, 440 theorems, 72 axioms, 0 sorry).

Prior results on attribution instability. Three prior results address related phenomena. Bilodeau et al. (2024) prove that no feature attribution simultaneously satisfying the Shapley axioms can be stable under model multiplicity—an impossibility for the *attribution-specific* case under *axiomatic* constraints. D’Amour et al. (2022) document underspecification empirically across ML pipelines, showing that models with identical in-distribution performance diverge out-of-distribution. Fisher et al. (2019) introduce variable importance clouds, visualizing the instability of feature importance across the Rashomon set. Our contribution relative to this body of work is: (1) the abstract framework (subsuming non-ML domains), (2) the bilemma strengthening and enrichment mechanism, (3) the Lean formalization, and (4) the Gaussian flip formula as a quantitative reliability score.

	Bilodeau et al. (2024)	This paper
Scope	Feature attribution (SHAP)	17 instances across 8 domains
Axioms	Shapley axioms	Faithfulness, stability, decisiveness
Conclusion	No stable Shapley attribution	F+S+D impossible (trilemma); F+S impossible (bilemma)
Resolution	Not addressed	G -invariant projection (Pareto-optimal)
Formalization	Pen-and-paper	Lean 4 (440 theorems, 0 sorry)
Quantitative	Asymptotic bounds	Gaussian flip (OOS $R^2 = 0.848$)

Table 1: Comparison with Bilodeau et al. (2024).

Stability selection. Stability selection [[Meinshausen and Bühlmann, 2010](#)] addresses a related problem: identifying features stably selected across subsamples. The Gaussian flip formula is complementary—it assesses the reliability of *pairwise rankings* (“is feature A more important than feature B?”), not just inclusion/exclusion (“should feature A be in the model?”). The two

tools answer different questions and can be used together: stability selection identifies a stable feature set, and the Gaussian flip formula determines which comparisons within that set are trustworthy.

Aggregation and resolution strategies. Laberge et al. [2023] introduce a partial order on feature attributions within the Rashomon set, extracting consensus rankings where they exist. Decker et al. [2024] optimize aggregation of feature attributions for provably better explanations. Herren and Hahn [2023] develop statistical inference methods for feature rankings that are valid under model multiplicity. We show that these diverse strategies are all instances of the G -invariant resolution (Section 5): they aggregate over the symmetry group of equivalent hypotheses, restoring stability and faithfulness-in-expectation at the cost of decisiveness.

Part I

Part I — The Hardened Core

Every result in this Part is machine-checked in Lean 4 from zero custom axioms or from explicitly stated hypotheses; the hard content is genuinely proved. This is the load-bearing spine: the universal impossibility—a meta-theorem holding for every explanation system that satisfies the Rashomon property—together with the bilemma and tightness classification, the ubiquity argument that makes the hypothesis generically applicable, the capacity/Reynolds resolution theory, and the trusted-base accounting.

3 The Abstract Framework [Tier A: Hardened]

Terminology. The framework introduces concepts at three levels. For readers familiar with information theory, the structural correspondences to Shannon’s theory are indicated. These identify shared *form* (limit/code/converse/rate); we do not claim equivalent mathematical depth (see introduction).

Concept	Name	Shannon	parallel	(structural)
F+S+D impossible	The explanation trilemma	—		
F+S impossible (binary)	The explanation bilemma	—		
$C = \dim(V^G)$	Explanation capacity	Channel capacity		
$\eta = 1 - C / \dim(V)$	Explanation loss rate	—		
Capacity predicts instability	Explanation Capacity Theorem	Shannon’s theorem		
$\text{unfaith}_1 + \text{unfaith}_2 \geq \Delta - \delta$	Explanation uncertainty bound	Donoho–Stark uncertainty		
Orbit average / DASH / character	The stable projection	Optimal channel code		
Which pairs survive	Tightness (full / collapsed)	—		
No neutral element	Explanation conflict	Noise (channel impairment)		
Add neutral element	Enrichment	Alphabet extension		
$g(g-1)/2$ stable, rest coin flips	Stable fact count	—		
$\ v - Rv\ ^2 + \ Rv\ ^2 = \ v\ ^2$	Explanatory information loss	Quantum information loss		
$\ w\ \leq \ u - w\ $ for $Rw=0$	Over-explanation penalty	Strong converse		
$\text{MSE} = \text{tr}(R\Sigma R)/M$	Expl. stability convergence rate	Source coding rate		
$M^*(\varepsilon) = \text{tr}(R\Sigma R)/\varepsilon$	Stability threshold	Block length		
The 4-part theorem	Explanation Stability Theorem	—		

We introduce a minimal abstract framework sufficient to state and prove the Explanation Impossibility. All instances in Sections 16 and 17 are special cases. The framework aligns exactly with the Lean 4 formalization in `ExplanationSystem.lean`: every definition and theorem below has a machine-checked counterpart.

3.1 Explanation Systems [Tier A: Hardened]

Definition 3.1 (Explanation System). *An explanation system is a tuple*

$$\mathcal{S} = (\Theta, \mathcal{H}, Y, \text{observe}, \text{explain}, \perp)$$

where:

- Θ is a configuration space (the set of possible configurations, e.g., random seeds, hyperparameter settings);
- $\mathcal{H}$ is an explanation space (the set of possible explanations);
- Y is an observable space (e.g., input-output functions, test-set predictions);
- $\text{observe} : \Theta \rightarrow Y$ is the observation map, assigning to each configuration its observable behavior;
- $\text{explain} : \Theta \rightarrow \mathcal{H}$ is the system’s native explanation map, assigning to each configuration the explanation that the system itself “honestly reports”; and
- $\perp \subseteq \mathcal{H} \times \mathcal{H}$ is an incompatibility relation on explanations satisfying irreflexivity: $\neg \perp (h, h)$ for every $h \in \mathcal{H}$.

Irreflexivity of $\perp$ is a structural requirement: no explanation contradicts itself. This is the `incompatible_irrefl` field of the Lean structure and is used in the tightness proof (Section 3.5).

An *explanation map* $E : \Theta \rightarrow \mathcal{H}$ is a procedure that a practitioner *chooses* (e.g., SHAP, LIME, a fixed causal orientation rule). The native explanation $\text{explain}(\theta)$ records what the configuration *actually implies*; $E(\theta)$ is what the chosen method reports.

3.2 The Three Desiderata [Tier A: Hardened]

Definition 3.2 (Faithful). *An explanation map E is faithful if $E(\theta)$ never contradicts the system’s native explanation $\text{explain}(\theta)$. Formally,*

$$\forall \theta \in \Theta, \quad \neg \perp (E(\theta), \text{explain}(\theta)).$$

That is, $E(\theta)$ may be less specific than $\text{explain}(\theta)$ but never disagrees with it.

Definition 3.3 (Stable). *An explanation map E is stable if equivalent configurations receive the same explanation:*

$$\forall \theta_1, \theta_2 \in \Theta, \quad \text{observe}(\theta_1) = \text{observe}(\theta_2) \implies E(\theta_1) = E(\theta_2).$$

Definition 3.4 (Decisive). *Practitioners want explanations that commit to specifics: “Feature A is more important than Feature B,” not “it depends on the model.” We formalise this desideratum as: E is decisive if $E(\theta)$ inherits every incompatibility of the native explanation—whatever $\text{explain}(\theta)$ rules out, $E(\theta)$ also rules out:*

$$\forall \theta \in \Theta, \forall h \in \mathcal{H}, \quad \perp (\text{explain}(\theta), h) \implies \perp (E(\theta), h).$$

The formal encoding references $\text{explain}(\theta)$ for precision; the underlying intuition—explanations should commit to specifics—is method-independent.

In the feature attribution instance (Section 17.1), faithfulness means the chosen method never contradicts the model’s own SHAP values, stability means consistency across retraining seeds, and decisiveness means that every distinction made by the native attribution is preserved.

3.3 The Rashomon Property [Tier A: Hardened]

The key structural condition linking underspecification to explanation impossibility is the following.

Definition 3.5 (Rashomon Property). *An explanation system $\mathcal{S}$ has the Rashomon property if there exist configurations $\theta_1, \theta_2 \in \Theta$ with*

$$\text{observe}(\theta_1) = \text{observe}(\theta_2) \quad \text{and} \quad \perp (\text{explain}(\theta_1), \text{explain}(\theta_2)).$$

That is, two configurations that are observationally equivalent honestly report incompatible explanations.

The Rashomon property is the formal statement of underspecification at the explanation level: the data do not uniquely determine the explanation, even in principle. It generalizes the Rashomon effect of [Rudin et al. \[2024\]](#), [Fisher et al. \[2019\]](#) from model selection to arbitrary explanation targets.

3.4 Main Theorem [Tier A: Hardened]

Theorem 3.6 (Explanation Impossibility). *If an explanation system $\mathcal{S}$ has the Rashomon property, then no explanation map E can simultaneously be faithful, stable, and decisive.*

Proof. By the Rashomon property, there exist configurations θ_1, θ_2 with $\text{observe}(\theta_1) = \text{observe}(\theta_2)$ and $\perp (\text{explain}(\theta_1), \text{explain}(\theta_2))$. Suppose for contradiction that E is faithful, stable, and decisive.

1. **Decisiveness** at θ_1 : $\perp (\text{explain}(\theta_1), \text{explain}(\theta_2))$ implies $\perp (E(\theta_1), \text{explain}(\theta_2))$.
2. **Stability**: $\text{observe}(\theta_1) = \text{observe}(\theta_2)$ implies $E(\theta_1) = E(\theta_2)$, so $\perp (E(\theta_2), \text{explain}(\theta_2))$.
3. **Faithfulness** at θ_2 : $\neg \perp (E(\theta_2), \text{explain}(\theta_2))$.
4. Steps 2 and 3 contradict. □

Remark 3.7. *The proof uses all three properties in a tight chain: decisive $\rightarrow$ stable $\rightarrow$ faithful $\rightarrow$ contradiction. This mirrors the Lean proof exactly (`explanation_impossibility` in `ExplanationSystem.lean`).*

Remark 3.8. *The simplicity of the proof—four steps from the Rashomon property to contradiction—is a feature of the framework, not a limitation of the result. The intellectual content lies in identifying the right definitions: faithful as non-contradiction, stable as factoring through observations, decisive as inheriting incompatibilities. The contribution is the axiomatic framework and its unification of nine domains, not the proof technique.*

Remark 3.9. *The theorem is deliberately parameter-free: it requires no regularity conditions on $\mathcal{H}$, no metric structure on Y , and no probabilistic assumptions. All quantitative depth (divergence rates, sample complexity, ensemble sizes) appears in the instance-specific sections below.*

3.5 Tightness [Tier A: Hardened]

The impossibility is tight: every pair of properties is simultaneously achievable. Three Lean-verified witnesses establish this (`tightness_faithful_decisive`, `tightness_faithful_stable`, `tightness_stable_decisive` in `ExplanationSystem.lean`).

Proposition 3.10 (Tightness of the Universal Impossibility). *For any explanation system with the Rashomon property:*

1. **Faithful + Decisive (drop Stable).** Set $E = \text{explain}$. Faithfulness holds by irreflexivity of $\perp$: $\neg \perp (\text{explain}(\theta), \text{explain}(\theta))$. Decisiveness holds by identity. Stability fails because the Rashomon property places incompatible native explanations on the same observe-fiber.
2. **Faithful + Stable (drop Decisive).** Let $c \in \mathcal{H}$ be a “neutral” element with $\neg \perp (c, \text{explain}(\theta))$ for all θ , and set $E(\theta) = c$. Faithfulness and stability hold, but E abstains on every distinction the native explanation makes.
3. **Stable + Decisive (drop Faithful).** Let $c \in \mathcal{H}$ be a “maximally committal” element inheriting every incompatibility of every native explanation, and set $E(\theta) = c$. Stability and decisiveness hold, but c may contradict some $\text{explain}(\theta)$.

Corollary 3.11 (Decidable Underspecification Implies Impossible Explanation). *If the data-generating process is not identifiable from the observations—i.e., if there exist θ_1, θ_2 with $\text{observe}(\theta_1) = \text{observe}(\theta_2)$ yet $\perp (\text{explain}(\theta_1), \text{explain}(\theta_2))$ —then the Rashomon property holds and the impossibility applies.*

The converse also holds for fully specified systems: when each observation uniquely determines its configuration, $E = \text{explain}$ satisfies all three properties simultaneously (Lean: `fully_specified_possibility` in `Necessity.lean`).

The impossibility is query-relative. Parameterize incompatibility by a query space Q : for each question q (e.g., “is feature x ranked above y ?”), the impossibility holds if and only if the Rashomon set produces incompatible answers to q (Lean: `query_impossibility`). Questions where the Rashomon set agrees—such as “is feature x relevant?” when all equivalent models include x in the top k —admit faithful, stable, decisive answers. The framework thus replaces “explanation is impossible” with a precise characterization of *which questions* are and are not stably answerable.

For instance, in feature attribution, the query “is feature x in the top k ?” may be stably answerable (all equivalent models include x), while “is x ranked above y ?” may not be (collinear features swap rankings across the Rashomon set).

3.6 Strengthened Impossibility for Maximally Incompatible Systems [Tier A: Hardened]

Framing note. The biletma is the triletma applied to the special case of maximally incompatible $\mathcal{H}$. It is not logically independent of Theorem 3.6—once $\mathcal{H}$ is binary with $\text{incomp} = (\neq)$, faithful + decisive collapses to faithful, and the impossibility reduces to the standard form. The reason this specialization warrants separate treatment is empirical: 19 of 29 Lean-verified instance files use $\text{incomp} = (\neq)$, making this the most common configuration in practice. The theorems below are consequences of this common configuration, not strengthenings of the triletma.

The triletma’s tightness depends on the structure of the explanation space $\mathcal{H}$. When $\mathcal{H}$ is *maximally incompatible*—the only compatible pair is (h, h) —the triletma collapses to a biletma.

Definition 3.12 (Maximal incompatibility). *An explanation system $(\Theta, H, Y, \text{observe}, \text{explain}, \perp)$ is maximally incompatible if $\neg \perp (h_1, h_2) \Rightarrow h_1 = h_2$ for all $h_1, h_2 \in H$. Equivalently, every pair of distinct explanations is incompatible. Any binary H with $\perp = \neq$ satisfies this condition.*

Theorem 3.13 (The biletma). *For any maximally incompatible explanation system with the Rashomon property, no explanation is simultaneously faithful and stable.*

Proof. Let $E : \Theta \rightarrow H$ be faithful. For any θ , faithfulness gives $\neg \perp (E(\theta), \text{explain}(\theta))$. By maximal incompatibility, $E(\theta) = \text{explain}(\theta)$. Therefore $E = \text{explain}$, which is decisive (it inherits all incompatibilities of itself). But now E is faithful, stable, and decisive—contradicting the triletma (Theorem 3.6). $\square$

Theorem 3.14 (Stable + Decisive also impossible). *For any maximally incompatible explanation system, no explanation is simultaneously stable and decisive.*

Proof. Let E be decisive. For any θ , suppose $E(\theta) \neq \text{explain}(\theta)$. By maximal incompatibility, $\perp (\text{explain}(\theta), E(\theta))$. By decisiveness, $\perp (E(\theta), E(\theta))$ —contradicting irreflexivity. So $E(\theta) = \text{explain}(\theta)$, making E faithful. But then E is faithful, stable, and decisive—contradicting the trilemma. $\square$

Proposition 3.15 (Tightness under maximal incompatibility). *Under maximal incompatibility, only faithful + decisive survives (via $E = \text{explain}$). Tightness collapses from 3/3 to 1/3.*

Property pair	Abstract	Max. incompatible	Enriched
Faithful + Decisive	✓	✓	✓
Faithful + Stable	✓	× (bilemma)	✓
Stable + Decisive	✓	×	×

Table 2: Tightness spectrum: 3/3 (abstract), 1/3 (max. incompatible), 2/3 (enriched with neutral element).

Many natural explanation spaces are maximally incompatible: binary feature rankings ($j > k$ or $k > j$, no tie), DAG edge orientations ($X \rightarrow Y$ or $Y \rightarrow X$, no undirected), and quantum measurement outcomes.

Maximal incompatibility is not an artificial condition: 19 of the 29 Lean-verified instance files in this project use incompatible = ($\neq$), making them maximally incompatible. The bilemma thus applies to the majority of instances, including all binary ranking spaces, all DAG orientation spaces, and all quantum measurement instances. The bilemma does not have direct mechanism design content. Mechanism design concerns incentive compatibility under strategic agents; the bilemma concerns the logical structure of explanation spaces without strategic considerations.

Approximate and quantitative extensions. The bilemma extends beyond exact equality. For binary H (e.g., SHAP sign), (ε, δ) -stability reduces to exact equality (Bool admits no non-trivial ε -ball), so the bilemma holds at *every* tolerance level with no discontinuity at $\varepsilon = 0$. For continuous H equipped with a metric d , the quantitative bilemma gives a smooth tradeoff: for an (ε, δ) -stable map E on an ε -Rashomon pair with incompatibility gap $\Delta = d(\exp(\theta_1), \exp(\theta_2))$,

$$\text{unfaith}(\theta_1) + \text{unfaith}(\theta_2) \geq \Delta - \delta$$

by the triangle inequality. At least one witness is unfaithful by $\geq (\Delta - \delta)/2$. Proved in Lean (`ApproximateBilemma.lean` in the companion Ostrowski formalization [Caraker et al., 2026b]). This resolves the concern that the bilemma is fragile at exact equality: the binary case is discontinuity-free, and the continuous case gives a smooth, quantitative tradeoff.

3.7 The Enrichment Mechanism [Tier A: Hardened]

Definition. Given a maximally incompatible explanation space H , the *enrichment* of H is the space $H^+ = H \cup \{c\}$, where c is a *neutral element*: an element compatible with every $h \in H$ (i.e., $\neg \perp (c, h)$ for all h).

Why enrichment works. Because c is compatible with all native explanations, the constant map $E(\theta) = c$ is both faithful (never contradicts any $\text{explain}(\theta)$) and stable (assigns the same explanation to all configurations). This restores faithful + stable at the cost of decisiveness: the constant map draws no distinctions at all. (Lean: `recovery_faithful_stable` in `MaximalIncompatibility.lean`.)

Why stable + decisive is NOT restored. Enrichment does *not* restore stable + decisive. Decisiveness requires $E(\theta)$ to inherit all incompatibilities of $\text{explain}(\theta)$, which forces $E(\theta) = \text{explain}(\theta)$ on non-neutral elements (by maximal incompatibility). But then $E = \text{explain}$ is not stable (the Rashomon property gives two configurations with the same observable but different explanations). The neutral element c cannot help: it is compatible with everything, so it is not decisive at any θ .

Instances. Four domain-specific resolutions are instances of enrichment:

- **DASH** (ML attribution): adds “tied” to the binary ranking space $\{j > k, k > j\}$, allowing the ensemble to report that two features are indistinguishable.
- **CPDAG** (causal discovery): adds “undirected” to the edge orientation space $\{X \rightarrow Y, Y \rightarrow X\}$, allowing the equivalence class to leave an edge unoriented.
- **BMA** (model selection): adds “uncertain” to the model identity space, allowing posterior averaging over competing models.
- **QEC** (quantum error correction): adds ancilla qubits to the measurement space, enabling syndrome-based recovery that does not commit to a single error pattern.

Categorically, the enrichment operation is the free adjunction of an initial object to the incompatibility poset: adding an element c compatible with everything is adjoining a least element under the compatibility ordering.

Categorical interpretation. The enrichment mechanism has a precise categorical characterization. View the explanation space H as a discrete category with morphisms given by compatibility (where $h_1 \rightarrow h_2$ iff $\neg \text{incompatible}(h_1, h_2)$). A maximally incompatible system is one where the only morphisms are identities — a discrete category. Adding a neutral element c (compatible with everything) freely adjoins an initial object: c has a unique morphism to every object. The enriched system is the free completion of the discrete category under an initial object. This characterization explains why enrichment is canonical: it is the *minimal* extension of H that admits a faithful + stable map. Any other F+S-restoring enrichment must contain a neutral element (by the characterization theorem), and the free adjunction is the universal such extension.

Characterization. Enrichment is needed if and only if the blemma holds, which occurs if and only if no neutral element exists in H . Define a neutral element as $c \in H$ such that $\neg \perp (c, \text{explain}(\theta))$ for all θ . If such a c exists, then $E(\theta) = c$ is faithful and stable, so the blemma fails and no enrichment is needed. Conversely, if the blemma holds (F+S is impossible for all E), then no neutral element can exist. Maximal incompatibility implies no neutral element: if c were compatible with all native explanations, then $c = \text{explain}(\theta)$ for all θ (by maximal incompatibility), but the Rashomon property gives θ_1, θ_2 with $\perp (\text{explain}(\theta_1), \text{explain}(\theta_2))$, contradicting irreflexivity. The converse of the last implication is false: a system can lack a neutral element without being maximally incompatible. (Lean: `BlemmaCharacterization.lean`, three theorems, zero `sorry`.)

ML instances. Three binary explanation types trigger the blemma: (1) SHAP sign attribution (positive/negative), where collinearity flips the sign across equivalent models; (2) feature selection (selected/not_selected), where Lasso path instability swaps features; (3) counterfactual direction (increase/decrease), where equivalent classifiers give opposite advice. For all three, faithfulness + stability is impossible — not just the trilemma.

Predictive consequences. Four structural predictions follow from the blemma (Lean: `PredictiveConsequences.lean`, five theorems, zero `sorry`).

Theorem 3.16 (Stability forces unfaithfulness). *For any maximally incompatible system, every stable explanation E satisfies $\exists\theta, \perp (E(\theta), \text{explain}(\theta))$.*

Proof. If E were both stable and faithful, the biletma gives a contradiction. $\square$

Theorem 3.17 (Rashomon unfaithfulness bound). *On any Rashomon pair (θ_1, θ_2) with $\text{observe}(\theta_1) = \text{observe}(\theta_2)$ and $\perp (\text{explain}(\theta_1), \text{explain}(\theta_2))$, a stable E is unfaithful at ≥ 1 of the 2 witnesses:*

$$\perp (E(\theta_1), \text{explain}(\theta_1)) \vee \perp (E(\theta_2), \text{explain}(\theta_2)).$$

Proof. Stability gives $E(\theta_1) = E(\theta_2)$. If $E(\theta_1)$ is compatible with $\text{explain}(\theta_1)$, then $E(\theta_1) = \text{explain}(\theta_1)$ (maximal incompatibility), so $E(\theta_2) = \text{explain}(\theta_1)$. Then $\perp (\text{explain}(\theta_1), \text{explain}(\theta_2))$ gives $\perp (E(\theta_2), \text{explain}(\theta_2))$. $\square$

Theorem 3.18 (Faithful uniqueness). *In a maximally incompatible system, faithful explanations are unique: if E_1 and E_2 are both faithful, then $E_1(\theta) = E_2(\theta)$ for all θ . Both equal explain.*

Proof. By Theorem 3.13’s proof, faithful $\Rightarrow E = \text{explain}$. So $E_1 = \text{explain} = E_2$. $\square$

Theorem 3.19 (All-or-nothing). *For any maximally incompatible system and any explanation E , exactly one of the following holds: (i) E is faithful and $E = \text{explain}$ (hence decisive but not stable), or (ii) $\exists\theta, \perp (E(\theta), \text{explain}(\theta))$ (E is unfaithful somewhere). No continuous interpolation exists.*

Proof. If E is faithful, maximal incompatibility forces $E = \text{explain}$. If E is not faithful, $\exists\theta, \perp (E(\theta), \text{explain}(\theta))$ by definition. $\square$

On simplicity. The biletma is a short proof. Its value is structural: it reveals the impossibility landscape depends on H ’s incompatibility structure, not just the Rashomon property. The abstract trilemma treats all systems uniformly; the biletma shows tightness is a system-specific property.

Structural parallels. Arrow’s impossibility theorem simplifies for binary alternatives; Gibbard–Satterthwaite simplifies for two candidates; the doctrinal paradox depends on proposition structure. The pattern — impossibility strength governed by outcome space structure — may warrant systematic study.

All theorems in this subsection are Lean-verified with zero `sorry` across three files: `MaximalIncompatibility.lean` (8 theorems), `BiletmaCharacterization.lean` (3 theorems), and `PredictiveConsequences.lean` (5 theorems).

4 Ubiquity [Tier A: Hardened]

Is the Explanation Impossibility a pathological corner case, or a generic phenomenon? In this section we show that it is the latter: the Rashomon property—and hence the impossibility—holds whenever the model has more parameters than observational constraints, a condition satisfied by virtually all modern machine learning systems.

4.1 Generic Underspecification [Tier A: Hardened]

Proposition 4.1 (Dimensional Underspecification). *Let $\mathcal{H}$ be a smooth hypothesis space of dimension $\dim(\mathcal{H})$ and let $\text{obs} : \mathcal{H} \rightarrow Y$ be a smooth observation map into a space of dimension $\dim(Y)$. If $\dim(\mathcal{H}) > \dim(Y)$, then for generic $y \in Y$, the fiber $\text{obs}^{-1}(y)$ has positive dimension $\geq \dim(\mathcal{H}) - \dim(Y)$.*

What is now machine-checked, and what remains informal. An earlier version of this development formalized only the arithmetic skeleton (`generic_underspecification`: $n > m \Rightarrow n - m > 0$) and left the mathematically substantive step—that underspecification genuinely *forces* distinct-but-observationally-equivalent configurations to exist—as an informal appeal to the preimage theorem. The *linear and infinitesimal core of that step is now Lean-verified*, via rank-nullity (file `UbiquityDimensional.lean`, axiom-clean). For a linear observation map $\text{obs} : \mathcal{H} \rightarrow Y$ with $\dim(\mathcal{H}) > \dim(Y)$, from *every* configuration there is a genuinely distinct one that observes identically (`exists_observational_collision`), the true fiber dimension is at least the naive count $n - m$ (`finrank_ker_ge`, so the arithmetic was a real lower bound, not just a proxy), and at any point a differentiable observation map has a nonzero tangent direction its derivative annihilates (`fderiv_exists_kernel_direction`—the honest infinitesimal heart of the preimage theorem). The dimensional fact is then wired to the impossibility itself (`underspecified_impossibility`): the colliding pair is now *produced* by dimension, and the only residual hypothesis is non-degeneracy of the explanation. The *local smooth (nonlinear) globalisation is now machine-checked as well*: at a *regular* point of a nonlinear observation map (surjective derivative) with $\dim(\mathcal{H}) > \dim(Y)$, the *true curved fiber* is not isolated—genuinely distinct configurations observing identically exist arbitrarily close (`regular_fiber_not_isolated`, proved via Mathlib’s implicit function theorem)—and this is wired to the impossibility for the nonlinear case (`smooth_regular_underspecified_impossibility`). This upgrades the infinitesimal statement (a flat tangent *direction*) to genuinely distinct *points* on the exact nonlinear fiber, with no genericity assumption. The last step—from “at a regular point” to “at almost every observable value”—is the genericity of regular values, i.e. *Sard’s theorem*: the set of critical values is null. Here we are precise about status. For linear observation maps—the primary ubiquity setting—Sard’s property is proved outright (`sardProperty_of_continuousLinearMap`): the critical values lie in range L , a proper subspace when L is not onto, hence Haar-null. So the linear case carries *no residual*: the full generic statement (`generic_ubiquity_of_sard`, which shows that Sard’s property implies, for almost every value, a positive-dimensional fiber of genuinely distinct models) closes unconditionally for linear maps. The *submersion* case is also closed unconditionally (`sardProperty_of_submersion`, `submersion_fiber_not_isolated`): a map with no critical points has no critical values, so for any C^1 submersion with $\dim(\mathcal{H}) > \dim(Y)$ *every* fiber is a positive-dimensional Rashomon locus. For a general *nonlinear* map, Sard’s property is isolated as a single named hypothesis (`SardProperty`) and the entire reduction to it is machine-checked (`generic_ubiquity_of_sard`). *That hypothesis is now discharged in full*: the classical Morse–Sard theorem is machine-checked in this repository (`sardProperty_of_contDiff`), via an in-tree, attribution-preserving port of Y. Kudryashov’s Lean 4 formalisation [Kudryashov, 2025] of Moreira’s version of Sard’s theorem [Moreira, 2001] (`UniversalImpossibility/Sard/`, forwardported from its original Mathlib revision to this repository’s; authorship of the ported development remains Kudryashov’s, per its Apache-2.0 licence and the attribution notice in-tree). The bridge instantiates Moreira’s Hausdorff-dimension bound $p + (n - p)/(k + \alpha)$ at $p = m - 1$, $k = n - m$, $\alpha = 1$, where it equals m exactly: the critical values of a C^{n-m+1} map are null for m -dimensional Hausdorff measure (`hausdorffMeasure_critical_image_null`), hence for every additive Haar measure on the observable space (`measure_null_of_hausdorffMeasure_finrank_null`, by transfer through a linear equivalence with $\mathbb{R}^m$ and Haar uniqueness). The smoothness C^{n-m+1} is the sharp classical threshold—Whitney’s counterexample shows C^1 is insufficient when $n > m$ —and it is carried entirely by the Sard step; the downstream chain needs only C^1 . The generic ubiquity statement is therefore *unconditional* for C^{n-m+1} observation maps (`generic_ubiquity_of_contDiff`; `sardProperty_of_smooth` for the C^∞ case), with every theorem in the chain axiom-clean (Lean core only). The vendored port is counted separately from this repository’s own theorem inventory. The ubiquity argument moreover extends to *infinite-dimensional* configuration spaces (`UbiquityInfiniteDimensional.lean`; observable

space finite-dimensional—the function-space-model regime): under the cardinal-rank hypothesis $\dim Y < \text{rank}(\mathcal{H})$, which strictly generalises the finite-dimensional one, the linear, infinitesimal, local-smooth (Banach implicit function theorem), and submersion rungs all carry over, with two unconditional generic-ubiquity capstones (`generic_ubiquity_of_continuousLinearMap_infinite`, `generic_ubiquity_of_submersion_infinite`) and a strengthened cardinal rank-nullity: the Rashomon fibre of an infinite-dimensional configuration space has *full* rank (`rank_ker_eq_rank_of_infiniteDimensional`). Morse–Sard itself provably does *not* extend there: Kupka’s counterexample [Kupka, 1965] gives a C^∞ map on ℓ^2 whose critical values contain an interval, and no map from infinite-dimensional $\mathcal{H}$ onto finite-dimensional Y is Fredholm (its kernel has full rank—machine-checked), so even the Sard–Smale theorem [Smale, 1965] is inapplicable. The finite-dimensionality of the Sard step is thus a theorem-shaped boundary, not a formalization gap; in infinite dimension the unconditionally closed classes are exactly the linear maps and the submersions. Whenever the fiber has positive dimension, the explanation map `obs` is generically non-constant on fibers—different hypotheses in the same fiber honestly report different explanations—and the Rashomon property holds. (The bridge from fiber non-degeneracy to the impossibility is `fiber_nondegeneracy_implies_impossibility`.)

Rashomon from overparameterisation. For model classes where $\dim(\Theta) > \dim(\mathcal{Y})$ (more parameters than observational constraints), the Rashomon property holds generically. The fibre $\{\theta : \text{observe}(\theta) = y\}$ has positive dimension for generic y , and under mild regularity (continuous explain map, non-degenerate Jacobian), distinct fibral points yield incompatible explanations with probability 1 under any absolutely continuous prior on Θ [D’Amour et al., 2022]. This covers overparameterised neural networks, deep ensembles, and GBDTs with more trees than effective degrees of freedom. Deriving tighter, architecture-specific Rashomon set bounds remains open.

4.2 Neural Network Dimensional Argument [Tier A: Hardened]

Modern neural networks are massively overparameterized. A ResNet-50 has $P \approx 25,600,000$ parameters and, for ImageNet classification, $k = 1,000$ output classes. By Proposition 4.1, the fiber of any fixed prediction has dimension at least $P - k = 24,999,000$. On such a high-dimensional manifold, any continuous explanation map (attention weights, gradient saliency, internal activations) is generically non-constant: there exist models h, h' on the same fiber with $\text{obs}(h)_j > \text{obs}(h)_k$ and $\text{obs}(h')_j < \text{obs}(h')_k$ for any pair of input features j, k with equal marginal distributions.

Corollary 4.2 (Coin-Flip Instability). *For any pair of symmetric features j, k in an overparameterized network trained to zero loss, the probability (over retraining seeds) that a faithful explanation ranks j above k is exactly $1/2$.*

The 50% flip rate applies to the adversarial case (maximally collinear pair, $\rho \rightarrow 1$). For moderate collinearity ($\rho \approx 0.7$), empirical flip rates are typically 15–30%, still high enough to undermine single-model explanations but less extreme than the worst case.

This is the neural-network instantiation of the symmetric Bayes result formalized in `Symmetric-Bayes.lean`: the symmetry of the fiber under feature permutation forces uniform marginals on the ranking.

4.3 Prevalence [Tier A: Hardened]

The dimensional condition $\dim(\mathcal{H}) > \dim(Y)$ is the norm, not the exception. Large language models (billions of parameters, vocabulary-sized output), diffusion models, graph neural networks, random forests with hundreds of trees, and Lasso regression with more features than

observations all satisfy it. Even linear regression with collinear features produces a positive-dimensional fiber (the collinear subspace). D’Amour et al. [2022] document this across NLP, computer vision, and medical imaging, finding that underspecification—and hence explanation instability—is ubiquitous in deployed systems. The only escape is a model class so constrained that $\dim(\mathcal{H}) \leq \dim(Y)$, which excludes essentially all modern architectures of practical interest.

5 The Universal Resolution [Tier A: Hardened]

The impossibility of Theorem 3.6 is tight: one of the three desiderata must be sacrificed. In this section we show that a *uniform* resolution exists whenever the Rashomon property arises from a symmetry group acting on the hypothesis space. The resolution sacrifices decisiveness and achieves faithfulness in expectation together with exact stability.

5.1 G -Invariant Explanation Systems

Let G be a group acting on the hypothesis space $\mathcal{H}$. The action is *observable-preserving* if for every $g \in G$ and $h \in \mathcal{H}$,

$$\text{obs}(g \cdot h) = \text{obs}(h),$$

i.e., the group permutes hypotheses within observe-equivalence classes (fibers). We say that G *acts transitively on observe-fibers* if for every pair h, h' with $\theta(h) = \theta(h')$, there exists $g \in G$ such that $g \cdot h = h'$.

Definition 5.1 (G -Invariant Explanation Map). *An explanation map exp is G -invariant if for every $g \in G$ and $h \in \mathcal{H}$,*

$$\text{exp}(g \cdot h) = \text{exp}(h).$$

When G is finite and Y admits averaging, the canonical G -invariant map is the *orbit average*:

$$\overline{\text{exp}}(h) = \frac{1}{|G|} \sum_{g \in G} \text{exp}(g \cdot h).$$

Proposition 5.2 (G -Invariance Implies Stability). *If G acts transitively on observe-fibers and exp is G -invariant, then exp is stable.*

Proof. Let $h, h' \in \mathcal{H}$ with $\theta(h) = \theta(h')$. By transitivity, there exists $g \in G$ with $g \cdot h = h'$. By G -invariance, $\text{exp}(h') = \text{exp}(g \cdot h) = \text{exp}(h)$. (Formalized as `gInvariant_stable` in `UniversalResolution.lean`.) $\square$

Proposition 5.3 (Design Space Dichotomy). *If the explanation system has the Rashomon property, then no stable explanation map can simultaneously be faithful and decisive.*

Proof. Immediate from Theorem 3.6: stability, faithfulness, and decisiveness are jointly contradictory under the Rashomon property. (Formalized as `universal_design_space_dichotomy` in `UniversalDesignSpace.lean`.) $\square$

Corollary 5.4 (Pareto Frontier). *For any explanation system with the Rashomon property, achievable explanation maps fall into exactly two families:*

- **Family A** (*faithful + decisive, unstable*): individual-model explanations that faithfully report each hypothesis but vary across equivalent hypotheses.
- **Family B** (*stable + faithful in expectation, indecisive*): G -invariant aggregations that are constant on observe-fibers but produce ties for queries where the fiber disagrees.

No explanation map lies outside these two families.

Table 3: Instance-specific resolutions of the Explanation Impossibility. Each row identifies the symmetry group, the concrete resolution, and the property sacrificed.

Instance	Symmetry group G	Resolution	Sacrificed
Attribution	Collinear-feat. perms.	DASH (ens. SHAP avg.)	Decisive rank.
Attention	Weight-space symm.	Avg. attention rollout	Single top token
Counterfactual	Dec.-boundary var.	Rashomon-robust recourse	Single CF
Concept probe	Repr. rotation	Averaged CAV / subspace	Single direction
Causal disc.	DAG rev. in Markov cl.	CPDAG reporting	Fully oriented
Model sel.	Model perm. in $\mathcal{R}$	Ensemble / $\mathcal{R}$ report	Single best

5.2 Instance-Specific Resolutions [Tier A: Hardened]

Table 3 summarizes how the G -invariant resolution instantiates across the domains of Section 17.

Is reporting ties a genuine resolution? A common objection: DASH “reports ties” for collinear features, which may seem like giving up rather than explaining. We argue this is the correct behaviour, analogous to a Bayesian posterior that is wide (uncertain) being more honest than a point estimate that is precise but wrong. The tie *is* the explanation: features j and k are interchangeable under the model’s learned structure, and any ranking between them is an artefact of training noise. Reporting this interchangeability is more informative than reporting a ranking that reverses with probability $\frac{1}{2}$ across seeds. The loss of within-group decisiveness is the price of faithfulness and stability—the trilemma guarantees no method avoids it.

5.3 Dash as the Prototype Resolution [Tier A: Hardened]

DASH (**D**iversified **A**ggregation for **S**table **H**ypotheses) is the G -invariant resolution for the feature attribution instance (Section 17.1). It averages SHAP values over an ensemble of models retrained from independent random seeds, achieving $\mathbb{E}[\phi_j] = \mathbb{E}[\phi_k]$ for any pair of symmetric features j, k (*consensus equity*). This orbit average is stable by Proposition 5.2 and faithful in expectation: the expected explanation matches the expected observation across the Rashomon set. The cost is indecisiveness—DASH reports ties for feature pairs whose rankings disagree across the ensemble. The companion paper argues that DASH is Pareto-optimal (no other stable attribution method can achieve higher expected faithfulness); the within-group component of this claim is Lean-verified, the global claim is not (see the proof-status inventory).

Connection to estimation-theoretic bounds. Faithfulness can be viewed as an unbiasedness condition: $\mathbb{E}[E(\theta)] = \text{explain}(\theta)$. Under this reading, the attribution impossibility is a Cramér–Rao-type result: no unbiased estimator achieves zero variance when the Fisher information (determined by the Rashomon set geometry) is finite. The design space theorem refines this: it characterises the full Pareto frontier of (bias, variance, completeness) trade-offs, not just the variance lower bound.

Computational cost of Dash. DASH requires training M models instead of one. For gradient-boosted trees (XGBoost/LightGBM), training time scales linearly: $M = 25$ models on $N = 10,000$ rows takes ~ 5 minutes on a single CPU core. SHAP computation is per-model and parallelisable. The total cost is $O(M)$ times single-model cost, with no increase in per-prediction inference cost (the deployed model is unchanged; only explanation aggregation changes).

5.4 Meta-Theorem: Orbit Averaging [Tier A: Hardened]

Theorem 5.5 (Orbit-Averaged Resolution). *Let $\mathcal{S}$ be an explanation system whose Rashomon property arises from a finite group G acting on $\mathcal{H}$ and acting transitively on observe-fibers. Then the orbit-averaged map $\overline{\text{exp}}(h) = \frac{1}{|G|} \sum_{g \in G} \text{exp}(g \cdot h)$ is:*

- (i) *well-defined (independent of orbit representative),*
- (ii) *stable (Proposition 5.2),*
- (iii) *faithful in expectation: $\mathbb{E}_{g \sim \text{Unif}(G)}[\text{exp}(g \cdot h)] = \overline{\text{exp}}(h)$, and*
- (iv) *Pareto-optimal among stable maps: no stable map achieves strictly higher pointwise faithfulness on every fiber.*

The orbit-averaged resolution is constructive: for any specific instance, the group G and its action are determined by the domain structure, and the averaging can be computed (or approximated by Monte Carlo sampling over the Rashomon set). Theorem 5.5 is the abstract counterpart of the DASH optimality result and is formalized in `UniversalResolution.lean`. The resolution strategy—restricting to G -invariant explanation maps—is structurally analogous to the invariant decision procedures justified by the Hunt–Stein theorem [Lehmann and Romano, 2005]: just as the Hunt–Stein theorem guarantees that invariant tests are minimax optimal when a group acts on the parameter space, the G -invariant aggregation is Pareto-optimal among stable explanation methods.

Biased alternatives and James–Stein phenomena. The design space theorem characterises unbiased linear aggregators. Can biased methods (e.g., ridge-regularised Shapley values, L_1 shrinkage) circumvent the impossibility? Shrinkage reduces variance (improving stability) but introduces systematic bias toward a prior (reducing faithfulness). The trilemma’s qualitative impossibility holds regardless: for any method under Rashomon, at least one of faithfulness, stability, or decisiveness must be sacrificed. Shrinkage shifts *which* property is sacrificed—from stability to faithfulness—without escaping the trilemma. Quantifying the bias–variance–decisiveness Pareto frontier for shrinkage estimators is future work.

6 The Explanation Uncertainty Bound [Tier A: Hardened]

The qualitative impossibility (“at least one of faithfulness, stability, or decisiveness must fail”) admits a quantitative refinement. For a finite system with total incompatible pairs I_{total} and within-fiber incompatible pairs R_{within} , define the *Rashomon ratio* $r = R_{\text{within}}/I_{\text{total}} \in [0, 1]$ and the *decisiveness fraction* $D(E) = (\text{fraction of incompatible pairs preserved by } E)$.

Theorem 6.1 (Explanation uncertainty bound). *For any stable explanation map E ,*

$$D(E) \leq 1 - \frac{R_{\text{within}}}{I_{\text{total}}} \quad \text{equivalently,} \quad D(E) + r \leq 1.$$

Proof. Stability forces $E(\theta_1) = E(\theta_2)$ whenever $\text{obs}(\theta_1) = \text{obs}(\theta_2)$. By irreflexivity of the incompatibility relation, $\text{incomp}(E(\theta), E(\theta))$ is false, so every within-fiber incompatible pair is collapsed by E . Hence E can preserve at most $I_{\text{cross}} = I_{\text{total}} - R_{\text{within}}$ incompatible pairs, giving $D(E) \leq I_{\text{cross}}/I_{\text{total}} = 1 - r$. The bound is tight: choosing a fixed representative per fiber preserves all cross-fiber incompatibilities. $\square$

The qualitative impossibility is the special case $r > 0$, $D = 1$: the inequality $1 + r \leq 1$ is a contradiction. When $r = 0$ there is no constraint (no Rashomon), when $r = 1$ we have $D = 0$ (total loss of decisiveness), and the general case interpolates smoothly.

Noether’s Theorem (continuous)	Explanation Framework (discrete)
Symmetry group G	Rashomon group G
G preserves Lagrangian L	G preserves obs
Conserved quantity	G -invariant query $\in V^G$
$\dim(G)$ conserved quantities	$\dim(V^G)$ stable queries

Table 4: Invariance correspondence. Rashomon symmetry $\rightarrow$ stable queries is the discrete analogue of continuous symmetry $\rightarrow$ conservation laws.

7 The Invariance Counting Principle [Tier A: Hardened]

Noether’s theorem [Noether, 1918] establishes that every continuous symmetry of the Lagrangian implies a conserved quantity. The explanation framework provides an analogous discrete result: every Rashomon symmetry implies a class of stably answerable queries. We note that this result is an application of classical invariant theory (Burnside’s lemma, Molien’s theorem) to the explanation framework, not a direct consequence of Noether’s theorem, which requires continuous symmetry and a variational principle.

Let G be a finite group acting on the configuration space Θ , preserving the observation map ($\text{obs}(g \cdot \theta) = \text{obs}(\theta)$ for all $g \in G$). A query $q: \mathcal{H} \rightarrow \mathbb{R}$ is *stably answerable* if $q(\text{explain}(\theta_1)) = q(\text{explain}(\theta_2))$ whenever $\text{obs}(\theta_1) = \text{obs}(\theta_2)$.

Proposition 7.1 (Invariance theorem). *A query q is stably answerable if and only if $q \in V^G$ (the G -invariant subspace). The number of independent stably answerable queries is $\dim(V^G) = (1/|G|) \sum_{g \in G} \chi(g)$, by the character formula.*

The correspondence (Table 4) places the explanation impossibility in the same conceptual lineage as the most celebrated theorem in mathematical physics. Both are instances of a single principle: group invariance determines what can be stably measured.

8 Toward a Group-Theoretic Classification of Explainability [Tier A: Hardened]

The resolution framework reveals a deeper structure: the difficulty of explanation is not merely binary (possible vs. impossible) but is *parameterized by the symmetry group* of the Rashomon set. Different group-theoretic structures produce qualitatively different impossibility regimes and qualitatively different resolutions.

The observation. For an explanation system with symmetry group G acting on the Rashomon fibers (the set of configurations sharing the same observable), the space of stable explanation maps is isomorphic to the space of G -invariant functions on the configuration space. The G -invariant projection (orbit averaging) is the *Reynolds operator*—the projection onto the trivial representation of G :

$$R(f)(x) = \frac{1}{|G|} \sum_{g \in G} f(g \cdot x) \quad (\text{finite } G), \quad R(f)(x) = \int_G f(g \cdot x) d\mu(g) \quad (\text{compact } G).$$

This is not a new mathematical object—it is the standard tool of invariant theory [Sturmfels, 2008]—but its identification as the universal resolution of the Explanation Impossibility, across all domains simultaneously, appears to be new.

The graded classification table and the invariant’s algebra. The symmetry group does not merely decide *whether* explanation is possible; via the Molien series of its permutation representation, $C_d = \dim(\text{Sym}^d V)^G$ counts the *stable explanation facts at each interaction order*—degree 1 the main-effect attributions, degree 2 the interactions, degree 3 the triple interactions. This yields a concrete classification table (computed for the canonical classes and verified against the closed form for S_n , where C_d is the number of partitions of d into at most n parts): for S_n with $n \geq 3$, $(C_1, C_2, C_3) = (1, 2, 3)$ regardless of n —exactly one stable main-effect fact (the sum), two stable interaction facts, three stable triples; for the cyclic group C_4 , $(1, 3, 5)$; for the dihedral D_4 , $(1, 3, 4)$; for the product $S_2 \times S_2$, $(2, 5, 8)$. Each entry is the a-priori count of facts an explanation can stably assert at that order—the fine (graded) refinement of the coarse hierarchy below. Capacity is moreover a well-behaved invariant, machine-checked as such: monotone (`capacity_mono`), antitone in the symmetry group—more symmetry, fewer stable facts (`capacity_antitone_subgroup`)—and full under trivial symmetry (`capacity_id`); its degree-1 value is the average character $\frac{1}{|G|} \sum_g \text{tr } \rho(g) = \dim(\text{invariants})$ (`capacity_reynolds_eq_char_average`, `capacity_reynolds_eq_finrank_invariants`); and the incurable structural instability—the directions no amount of data resolves—is exactly the kernel of the Reynolds projection, of dimension $\eta \cdot \dim V$ (`structural_floor_dim`), linking the classification to the structural/statistical dichotomy. For the *symmetric group* the graded *formula* is now machine-checked at *every* degree, not just degree 1. The symmetric group acts on degree- d homogeneous polynomials by renaming variables—a genuine finite-dimensional representation (`renameRep`), once one supplies the finite-dimensionality of the graded piece (`homogeneousSubmodule_finiteDimensional`, a lemma Mathlib does not provide). Applying the degree-1 character-average formula to it yields, for every d , the identity $C_d = \dim(\text{degree-}d \text{ symmetric invariants}) = \frac{1}{|G|} \sum_g (\text{degree-}d \text{ character})(g)$ (`graded_molien_char_average`). It is this generating *formula* that is machine-checked. The specific coefficient values in the table above—and the cyclic, dihedral, and product columns—are *computed* from the Molien formula (and verified against the closed form for S_n , where C_d is the number of partitions of d into at most n parts); those particular integers, and the non-symmetric groups, are not themselves separately Lean-checked. So the graded stable-fact hierarchy is, for the symmetric case, a theorem rather than a bare computation; the table’s entries are its rigorously computed instances.

The keystone: capacity is exact, and one structure theorem unifies the framework. Capacity was an upper bound because its converse—that *every* stable explanation map is G -invariant—was open. It is now a theorem, and an elementary one: since the group preserves observables, any map that factors through the observable is constant on orbits (`stable_gInvariant`); together with the already-established forward direction, and under the transitive symmetry structure, this gives the biconditional stable $\Leftrightarrow G$ -invariant (`stable_iff_gInvariant`, which depends on *no* axioms whatsoever). The stable resolutions are therefore *exactly* the G -invariant ones—nothing outside the invariants is stable—so the capacity $\dim(V^G)$ is the *exact* stable-fact count, achieved by the orbit average and not exceeded. This upgrades the coverage-validated bound of the previous paragraphs to an equality, and it closes the framework into a single *structure theorem* (`explanation_structure_theorem`): the orbit-average (Reynolds) operator R decomposes the explanation space canonically as $V = V^G \oplus \ker R$. R is a projection; its range is exactly the stable subspace V^G ; its trace/dimension is the capacity, equal to the average character $\frac{1}{|G|} \sum_g \text{tr } \rho(g)$; and its kernel—the incurable structural instability—has dimension $\dim V - \text{capacity} = \eta \cdot \dim V$. The resolution, the capacity, and the structural floor are thus three facets of this one operator; and the impossibility (proved separately) is the obstruction the same decomposition makes transparent—a decisive explanation would have to commit along the unstable directions $\ker R$, which stability forbids. What read as four separate results are one picture. In this sense the framework is *structurally closed*: a single machine-checked ob-

ject from which the parts follow, with the point-predictive over-reaches falsified and quarantined and the guarantees stated as distribution-free bounds that hold out of sample. Structural closure is an internal property; it is not a claim of empirical or external completeness—prospective replication, external artifact evaluation, and adoption remain future work.

A hierarchy of difficulty. The group structure determines the character of the impossibility and the available resolutions:

1. **Finite abelian** G (gauge theory $\mathbb{Z}_2^{|V|-1}$, statistical mechanics S_2 , biology S_k for $k \leq 2$). The orbit average is canonical and unique. All characters are one-dimensional, so the invariant projection has no ambiguity. The faithfulness bound is $1/|G|$. These are the “easy” cases: the impossibility is real but the resolution is completely determined.
2. **Finite non-abelian** G (biology S_k for $k \geq 3$, causal discovery with ≥ 3 reversible edges, molecular point groups in chemistry). Multiple inequivalent invariant projections may exist, corresponding to different choices of which non-trivial irreducible representations to retain. The resolution choice itself becomes ambiguous—a *second-order Rashomon problem*. We observe this empirically: different causal discovery scoring functions (BIC, AIC, MDL) can produce different CPDAGs on the same data, reflecting different implicit choices among equivalent invariant representations.
3. **Finite non-trivial** G (crystallography $\mathbb{Z}_2 \times \mathbb{Z}_n \times \mathbb{Z}_2$). The Haar measure on G is finite, so orbit averaging converges. The Peter-Weyl theorem decomposes $L^2(G)$ into irreducible representations, and the invariant projection retains precisely the zeroth-frequency (DC) component. This is why the Patterson map—the autocorrelation, i.e., the zero-frequency Fourier content—is the natural resolution in crystallography.
4. **Non-compact continuous** G (linear algebra $\mathbb{R}^d$, deep learning weight space). No finite Haar measure exists on non-compact orbits. Orbit averaging diverges without regularization. The pseudoinverse (minimum-norm solution) is Tikhonov regularization with $\lambda \rightarrow 0$ —a *regularized* orbit average. Different regularizers (L1, L2, elastic net) produce different “resolutions,” and the choice of regularizer introduces an external degree of freedom not present in the compact case. This is the “fundamentally hard” regime: the impossibility is severe and the resolution is non-canonical.

This ordering is a partial order, not a total order: compact non-abelian groups (e.g., $\text{SO}(3)$) exhibit features of both items 2 and 3.

Predictions. This hierarchy generates testable predictions not currently verified in the paper:

- Domains with abelian symmetry should exhibit a unique canonical resolution with no practitioner disagreement on what the “right” invariant quantity is.
- Domains with non-abelian symmetry should exhibit *resolution multiplicity*—different “reasonable” resolutions that disagree—and this disagreement should itself be measurable.
- Domains with non-compact symmetry should require regularization, and different regularizers should produce different answers whose spread scales with the non-compactness of the group.
- The faithfulness bound $1/|G|$ (finite) or $\exp(-\dim G)$ (Lie groups) should predict the empirical instability across domains.

What is missing. The current eight instances span finite abelian, (barely) finite non-abelian, compact continuous, and non-compact continuous groups. Not yet represented are: genuinely non-abelian finite groups (e.g., the tetrahedral point group T_d in molecular spectroscopy, order 24); non-compact Lie groups beyond translations (e.g., the Lorentz group $\text{SO}(3, 1)$ in relativistic physics); infinite discrete groups (e.g., lattice symmetries $\mathbb{Z}^n$); and groupoids—where the symmetry structure varies fiber-by-fiber, as happens in ML when different regions of feature space have different Rashomon set sizes. The groupoid generalization would connect this work to the theory of stacks in algebraic geometry and to fiber bundles in differential geometry. We leave these extensions to future work.

Representation-theoretic classification. The classification of non-dominated explanation profiles is complete for $G = \mathbb{Z}/2\mathbb{Z}$ (the simplest non-trivial case): exactly two profiles exist (trivial representation giving stability, sign representation giving decisiveness). For general finite G , we conjecture that the number of non-dominated profiles equals the number of irreducible representations of G . The $\mathbb{Z}/2\mathbb{Z}$ case serves as proof of concept; the general classification is an open problem.

Categorical enrichment (conjecture). We conjecture that hypothesis-space enrichment corresponds to the free adjunction of an initial object to the incompatibility poset, viewed as a thin category. Specifically, if $\mathbf{C}$ is the poset category of H under the compatibility relation, then the enriched space $H' = H \cup \{h_\perp\}$ corresponds to $\mathbf{C}_\perp$. The universal property ensures uniqueness (the companion paper [Caraker et al., 2026b]). Making this precise requires defining the relevant 2-categorical context; we leave the formalisation to future work.

Character theory and the information trade-off. Character theory quantifies the trade-off precisely. For the symmetric group S_k acting on $\mathbb{R}^k$ by permutation of coordinates, the representation decomposes into the trivial representation (dimension 1) and the standard representation (dimension $k - 1$). The G -invariant subspace is exactly the trivial representation: $\dim(V^G) = 1$, while $\dim(V) = k$. The resolution therefore preserves exactly $1/k$ of the information—the orbit average retains the mean and discards the $k - 1$ “directional” degrees of freedom. This connects the dose-response (entropy scaling with degeneracy k) directly to representation-theoretic structure: the ratio $\dim(V^G)/\dim(V) = 1/k$ determines the fraction of explanation content that survives projection onto the invariant subspace. We call $C = \dim(V^G)$ the *explanation capacity*: the maximum number of independent stable facts any method can extract from a system with symmetry group G . The complementary quantity $\eta = 1 - C/\dim(V)$ is the *explanation loss rate*: the fraction of explanatory content irretrievably lost by any stable method. The orbit average (Reynolds operator) achieves the capacity—it is the *capacity-achieving projection* (Theorem 8.1, proved Pareto-optimal). As Shannon’s channel capacity bounds the rate of reliable communication, the explanation capacity bounds the rate of reliable explanation.

The ratio $\dim(V^G)/\dim(V)$ is a standard quantity in representation theory, computable via the character formula (Burnside’s lemma). This character formula is itself machine-checked here: for the Reynolds operator of a finite-group representation, `capacity_reynolds_eq_char_average` (in `MolienCapacity.lean`, axiom-clean) proves $\text{capacity} = \frac{1}{|G|} \sum_g \text{tr}(\rho g)$ — the degree-1 Molien coefficient — and `capacity_reynolds_eq_finrank_invariants` chains it to Mathlib’s character theory to give $\text{capacity} = \dim \mathfrak{I}(\text{invariants } \rho)$ outright, so “capacity = dimension of the invariant subspace = average character” is now a formal identity rather than an appeal to a textbook lemma. (The graded Molien *formula*—whose degree-2 term is the interaction capacity—is also machine-checked, for the symmetric group at every degree: `graded_molien_char_average` realizes the rename action on degree- d homogeneous polynomials as a representation and applies the same degree-1 formula, giving $C_d = \dim(\text{degree-}d \text{ symmetric invariants}) =$

average degree- d character; the specific coefficient values are computed from it.) The novelty is not the formula but its *application*: we show empirically that the explanation capacity predicts instability rates across seven domains with $R^2 = 0.957$ (well-characterized groups), using a quantity from abstract algebra as a parameter-free predictor of empirical flip rates. We call this the *Explanation Capacity Theorem*.

Will more data fix it? The structural/statistical dichotomy. The most common objection to Rashomon-driven instability is that it is finite-sample noise more data will resolve. We answer this as a machine-checked dichotomy (`Dichotomy.lean`, `axiom-clean`). Instability splits into two components with opposite fates. The *structural* component is incurable: if a symmetry g preserves the loss on *every* dataset—the defining property of a function-class symmetry such as a neuron permutation or a sign-flip pair—then the entire group orbit of any minimizer lies in the Rashomon set at *every* sample size (`orbit_subset_rashomon`, `structural_floor`); no amount of data removes the ambiguity across the orbit. The *statistical* component is curable: if, transverse to the orbit, the loss is strongly convex, the Rashomon set at tolerance ε is confined to a tube of radius $\sqrt{\varepsilon/c}$ around the orbit (`rashomon_within_tube`) that shrinks to the orbit as the statistically-indistinguishable tolerance $\varepsilon_n \rightarrow 0$ (`tube_radius_tendsto_zero`). Total instability is thus a structural floor (constant in n) plus a statistical part that vanishes with data (`dichotomy`). This is confirmed empirically: in a flip-versus-sample-size study, exactly-duplicated features hold a ≈ 0.5 flip floor as their signal-to-noise ratio collapses to 0.007, while a genuinely dominant feature decays to zero flip, and on real datasets the mean pairwise flip decays with n while the worst pairs retain a ≈ 0.5 floor—both halves of the dichotomy in one experiment.

Uncertainty from symmetry: the representation-theoretic connection to quantum mechanics. The capacity theorem and quantum uncertainty both arise from the same representation-theoretic decomposition. We state and prove the connection.

Theorem 8.1 (Uncertainty from symmetry). *Let G be a finite group acting on a real inner product space $(V, \langle \cdot, \cdot \rangle)$ of dimension n by orthogonal representation $\rho : G \rightarrow O(V)$. Let $R = |G|^{-1} \sum_{g \in G} \rho(g)$ be the Reynolds operator, and $k = \dim(V^G)$. Then:*

- (i) R is the orthogonal projection onto V^G .
- (ii) For any $v \in V$, $\|v - Rv\|^2 + \|Rv\|^2 = \|v\|^2$ (explanatory information loss identity).
- (iii) For v drawn uniformly from the unit sphere S^{n-1} , $\mathbb{E}[\|Rv\|^2] = k/n$ and $\mathbb{E}[\|v - Rv\|^2] = 1 - k/n = \eta$.
- (iv) The minimum information loss of any G -invariant approximation is $\inf_{w \in V^G} \|v - w\|^2 = \|v - Rv\|^2$, and η is its expected value.

Proof. (i) R is a group average of orthogonal maps, hence self-adjoint ($\langle Rv, w \rangle = \langle v, Rw \rangle$) and idempotent ($R^2 = R$), so R is an orthogonal projection. Its image is V^G : if $Rv = v$ then $\rho(g)v = v$ for all g (by averaging), and conversely every G -fixed vector is fixed by R . (ii) Since $Rv \in V^G$ and $(v - Rv) \in (V^G)^\perp$ (because R is an orthogonal projection), Pythagoras applies. (iii) Let $\{e_1, \dots, e_k\}$ be an orthonormal basis for V^G . Then $\|Rv\|^2 = \sum_{i=1}^k \langle v, e_i \rangle^2$. For v uniform on S^{n-1} , $\mathbb{E}[\langle v, e_i \rangle^2] = 1/n$ for each i (by rotational symmetry), so $\mathbb{E}[\|Rv\|^2] = k/n$. Part (ii) gives $\mathbb{E}[\|v - Rv\|^2] = 1 - k/n$. (iv) R is an orthogonal projection, so Rv is the closest point in V^G to v . $\square$

The theorem applies in two settings with identical algebraic content:

Explanation setting. Let $v \in V$ be an explanation vector (e.g., a vector of SHAP values across features) and G the symmetry group of equivalent configurations. The orbit average

Rv is the stable explanation; the discarded component $(I - R)v$ lies in non-trivial irreducible representations of G and is unstable by construction. The expected fraction of explanation content lost by stabilisation is $\eta = 1 - \dim(V^G)/\dim(V)$. This is the capacity theorem.

Quantum setting. Let $V = \mathbb{C}^n$ with G acting by unitary representation U , and let $T(\rho) = |G|^{-1} \sum_g U(g)\rho U(g)^\dagger$ be the quantum twirl (the density-matrix analogue of the Reynolds operator). For a pure state $|\psi\rangle$ drawn uniformly (Haar measure), the expected probability of finding the system in the G -invariant sector is $\mathbb{E}[\langle\psi|P_{V^G}|\psi\rangle] = \dim(V^G)/\dim(V)$. The complementary fraction η bounds the information inaccessible to G -covariant measurements: the specific configuration within each non-trivial irreducible component cannot be distinguished by any measurement commuting with G .

The unification. The Reynolds operator (explanation) and the quantum twirl (quantum mechanics) are instances of the same algebraic structure—projection onto the trivial representation of G —applied to different spaces (V for explanations, $L(H)$ for density matrices). The formula $\eta = 1 - \dim(V^G)/\dim(V)$ governs both the explanation instability rate and the quantum information loss. Stable explanatory facts correspond to trivial-representation components, just as quantum observables commuting with the symmetry correspond to G -invariant operators.

The framework produces two distinct types of bound, playing different roles:

The capacity theorem as minimum cost theorem. The cost of stability is exactly η : the Reynolds projection loses precisely $1 - k/n$ of the expected squared norm. This is not a tradeoff—it is a fixed price determined by the group structure, analogous to the zero-point energy of a quantum harmonic oscillator (a minimum cost of confinement, not a relation between conjugate variables). The capacity theorem is tight: equality holds in Theorem 8.1(iii) by construction.

The explanation uncertainty bound. The approximate bilemma ($\text{unfaith}(\theta_1) + \text{unfaith}(\theta_2) \geq \Delta - \delta$) is a Donoho–Stark-type uncertainty principle [Donoho and Stark, 1989]: reducing unfaithfulness at one Rashomon witness forces increased unfaithfulness at the other, just as a function and its Fourier transform cannot both have small support. Both are lower bounds on complementary quantities arising from constrained approximation—the bilemma from the constraint that stable maps are constant on Rashomon fibers, Donoho–Stark from the constraint that signals cannot concentrate in both time and frequency. The form is also parallel to Heisenberg’s relation ($\sigma_A \sigma_B \geq |\langle[A, B]\rangle|/2$), though the mathematical mechanism differs: the bilemma uses the triangle inequality on a metric space, while Heisenberg uses the Cauchy–Schwarz inequality on a Hilbert space with non-commuting operators. The approximate bilemma is tight when the explanation map places $E(\theta_1)$ and $E(\theta_2)$ on the geodesic between $\exp(\theta_1)$ and $\exp(\theta_2)$ with $\|E(\theta_1) - E(\theta_2)\| = \delta$; this occurs when the Rashomon witnesses are symmetric about V^G .

Both bounds arise from the same representation-theoretic decomposition: the Pythagorean splitting $V = V^G \oplus (V^G)^\perp$ under the Reynolds operator. The capacity theorem governs the average case (expected loss over random explanations); the approximate bilemma governs the worst case (complementary loss at specific Rashomon pairs). Together, they constitute a complete uncertainty framework for explanation under symmetry.

Lean formalisation status. Theorem 8.1 is a standard result in representation theory; its application to the explanation–quantum connection is the contribution. The individual components (Reynolds projection, explanatory information loss identity, approximate bilemma) are Lean-verified in the companion repositories. A unified Lean proof of the full theorem—including the quantum interpretation and the explicit connection between the two settings—is a natural next step.

The Explanation Stability Theorem. The explanatory information loss identity yields a capacity theory for explanation with the structural form of—though not the information-theoretic depth of—Shannon’s coding theorem. We state the four parts; together they establish that the explanation capacity $C = \dim(V^G)$ separates stably-reportable content from unreportable, analogously to how channel capacity separates reliable from unreliable communication.

(We claim structural parallel, not depth equivalence; the mathematical techniques differ—see the introduction.)

Part (i): Convergence rate. Let $v_1, \dots, v_M$ be i.i.d. explanation vectors with mean μ and covariance Σ . The orbit-averaged explanation $\bar{E}_M = M^{-1} \sum_i Rv_i$ satisfies $\mathbb{E}[\|\bar{E}_M - R\mu\|^2] = M^{-1} \text{tr}(R\Sigma R)$. For target $\text{MSE} \leq \varepsilon$, it suffices to use $M \geq \text{tr}(R\Sigma R)/\varepsilon$ models. We call this the stability sample-complexity for orbit averaging. The NARPS convergence ($M_{95} = 16$ teams for 95% stability) is a special case.

Part (ii): Capacity bound. The orbit-averaged (Reynolds-projected) explanation $R \circ E$ has image in V^G —immediate from idempotency of R —so its image dimension is at most $C = \dim(V^G)$, and no amount of averaging, ensembling, or methodological refinement expands V^G . The stronger statement that *every* stable map has image in V^G , formerly a Tier B conjecture here, is now *resolved as a sharp boundary* (`TierBResolutions.lean`): it is *false* as literally stated—a constant map is G -stable with image escaping V^G (`gstable_not_image_in_fixed`), and even a genuinely linear stable map can escape (`linear_sigmaStable_image_not_in_diag`)—while the corrected form is a machine-checked, axiom-free theorem: a map that is both stable and *equivariant* has image in V^G (`gstable_equivariant_image_subset_fixed`); the linear counterexample violates exactly equivariance (`Lsum_not_equivariant`), so equivariance is the precise missing hypothesis. The image bound just stated is the separate theorem `reynolds_image_in_fixed`.

Part (iii): Rate optimality. Among all linear unbiased estimators of $R\mu$ from M i.i.d. samples, the orbit average minimises $\mathbb{E}[\|\hat{E} - R\mu\|^2]$ (Gauss–Markov). Equivalently, V^G is the sufficient statistic for stable explanatory content, and the Reynolds projection is the Rao–Blackwell improvement of any explanation method.

Part (iv): Over-explanation penalty. For any $w \in (V^G)^\perp$ (i.e. $Rw = 0$) and any stable element $u \in V^G$ (i.e. $Ru = u$):

$$\|w\| \leq \|u - w\|.$$

The best stable approximation to any beyond-capacity quantity is the zero vector, with $\text{MSE} = \|w\|^2$. This is not noisy information; it is zero information. The penalty follows directly from orthogonality ($u \perp w$) and the Pythagorean theorem (Lean-verified: `beyond_capacity_penalty`).

The four parts are organised by structural analogy to the roles in Shannon’s coding theorem—capacity (Part ii), achievability (Parts i and iii), a lower-bound converse (Part iv), and sample complexity (Part i)—but we emphasise this is an organising analogy, not a claim of information-theoretic depth: Part (iii) is Gauss–Markov and the general form of Part (ii) is false without equivariance and a theorem with it (see above). The over-explanation penalty gives the *exact* MSE for beyond-capacity estimation (not an asymptotic bound), a consequence of working in finite-dimensional inner product spaces.

Empirical validation. The convergence rate $\text{MSE} \sim M^{-\beta}$ was validated on synthetic data with $\beta = 1.28$ (predicted: 1.00; $R^2 = 0.97$). The over-explanation penalty was validated across 149 datasets: within-group pairwise claims (beyond capacity) flip at rate 0.22 ± 0.10 , while between-group claims (within capacity) flip at rate 0.15 ± 0.08 (cross-dataset Wilcoxon $p = 5.1 \times 10^{-11}$; directional confirmation in 81% of 110 testable datasets, 86% of 84 with Cohen’s d computable).

Capacity is a proved trace identity and a data-estimable bound. Two results in the Lean development (`CapacityTrace.lean`) sharpen the capacity from a definition into computed mathematics; both depend only on Lean’s core axioms. First, *capacity is the trace of the Reynolds projection*: for the orbit-average operator R —proved idempotent for any finite-group representation (`reynolds_isIdempotent`)—one has $\text{tr}(R) = \dim(V^G) = C$ (`capacity_eq_trace`, and `reynolds_capacity_eq_trace` for the actual averaging operator). This turns

Part (ii)’s $C = \dim(V^G)$ into a machine-checked identity for the operator the framework actually uses, not merely a naming convention. Second, *capacity is bounded—and estimable—from the feature geometry alone, without knowing the symmetry group*. Attribution is free (unstable) exactly along the prediction-preserving, collinear directions, which lie in $\ker R$; so for any subspace W of such directions, $C \leq \dim(V) - \dim(W)$ (`capacity_le_of_ker`, with the rank-nullity identity `capacity_add_ker`). Empirically this bound is realised by the effective rank of the feature covariance: on controlled data with a known symmetry it recovers the true capacity exactly (an eigenvalue bracket $[g, g]$, stable across a wide threshold range and robust to approximate exchangeability) and, more strongly, recovers the stable subspace V^G itself (principal-angle alignment ≈ 1 with the true fixed subspace) using no group labels; on independent features it correctly returns full capacity—the null control that ensemble-based estimators fail; and on real datasets it returns honest non-degenerate intervals (e.g. breast cancer [6, 7] out of 30 features, California housing [4, 5] out of 8). This empirical subspace recovery is now underwritten by a machine-checked perturbation bound: a Frobenius-norm Davis–Kahan $\sin\theta$ theorem (`davisKahan_sinTheta_hermitian`, axiom-clean, built on the gapped Sylvester core `gapped_sylvester_frobenius`) shows that when the feature spectrum has a gap δ separating the stable block from the collinear remainder, the recovered stable subspace lies within a $\sin\theta$ distance $\|E\|_F/\delta$ of the true V^G —so the alignment ≈ 1 above is the theorem’s prediction under a spectral gap, and the recovery is rigorous end-to-end, not merely empirical.

Scope, stated plainly. The exact *integer* capacity is well-posed and recoverable only under (near-)exact symmetry, where the feature spectrum has a clean gap; on real, approximately symmetric data the instability is a continuum, and capacity is honestly a *bounded interval* rather than a single number, its width reflecting the sharpness of that gap. Estimating the exact dimension from the trained-model *ensemble* (rather than the feature geometry) was attempted with five principled estimators and does not work; we report this negative result because it delimits the claim precisely—capacity is a property of the feature geometry (noise-free), not of the attribution ensemble (noisy). The complementary practitioner question—*which* specific pairwise claims are reliable—is answered per query by an ensemble signal-to-noise certificate with a *provable, distribution-free* guarantee: a flip of the pairwise claim is the one-sided (minority-tail) event $D - \mathbb{E}D \leq -|\mathbb{E}D|$ for the attribution difference D , so by Cantelli’s inequality the flip rate is at most $1/(1 + \text{SNR}^2)$ with no distributional assumption (strictly sharper than the two-sided Chebyshev bound $1/\text{SNR}^2$ and never vacuous; a STABLE certificate, $\text{SNR} \geq 2$, thus guarantees flip $\leq 20\%$ unconditionally; the Gaussian estimate $\Phi(-\text{SNR})$ is tighter still and empirically well-calibrated). This guarantee is machine-checked: the finite-sample Cantelli bound is formalized axiom-free in Lean as `cantelli_lower_tail` (file `CertificateGuarantee.lean`), with the certificate corollaries `flip_le_one_div_one_add_snrSq` (the $1/(1 + \text{SNR}^2)$ form) and `stable_flip_le_one_fifth` ($\text{STABLE} \Rightarrow \text{flip} \leq 20\%$). Two further machine-checked facts pin down exactly what the certificate can and cannot promise. First it is *optimal*: an explicit two-point ensemble attains the Cantelli bound with equality (`flip_bound_tight` in `CertificateTight.lean`), so no certificate reading only the ensemble mean and variance can promise a smaller flip rate—the 20% STABLE guarantee is tight, not conservative. Second it is *one-sided*: an all-positive ensemble reads $\text{SNR}^2 < 1$ (the UNRELIABLE band) yet has flip exactly 0 (`flip_zero_at_low_snr`), proving that instability is *not* certifiable from the mean and variance alone—the certificate can prove a claim stable but never prove it arbitrary, which is why it reports STABLE/not-STABLE and never “certified unstable.” It is deliberately the finite-ensemble form the certificate actually computes—the sample mean and variance over the retrained models—so its hypotheses are discharged with no distributional or integrability assumptions, which is also why the bound is expected to (and does) hold for every observed pair. The population (measure-theoretic) Cantelli inequality—the idealization behind the finite bound, and itself a gap in current formalized libraries—is also machine-checked (`cantelli_upper` in `PopulationCantelli.lean`). Across five real datasets

(592 pairs) the Chebyshev bound holds for every pair and STABLE-certified claims flip at most 5% (mean 0.2%); the certificate holds out-of-sample (Spearman 0.76–0.94 vs ≤ 0.26 for raw data-property baselines) and flags top-ranked single-model claims that are in fact coin-flips across retraining.

Validated at scale. On 104 real datasets (PMLB; 10,724 feature pairs) with real nonlinear models, a split-ensemble protocol—compute the certificate on one Rashomon ensemble, measure the flip on an *independent* one—shows the certificate is predictive out of sample, not merely descriptive. The out-of-sample flip ranking follows the certificate at Spearman median 0.82 (interquartile [0.75, 0.88]), beating the group-free $|\text{corr}|$ predictor on 98% of datasets and clearing the 0.26 from-data ceiling on 98%, and the Gaussian estimate is well calibrated in aggregate (expected calibration error 0.009). The decisive test—transfer—is that every STABLE verdict computed on one ensemble stays stable (flip $\leq 20\%$) on the independent ensemble, for 100% of pairs across all 100 datasets that produce STABLE calls, whereas UNRELIABLE pairs flip at a median 0.42 (near coin-flip). Crucially, this transfer is no longer merely observed but is itself a machine-checked theorem. Modelling the deployment ensemble as m i.i.d. bounded flip indicators with population flip probability $p \leq \beta$ (the population Cantelli bound), a Hoeffding tail on the independent ensemble gives, distribution-free and finite-sample, $\Pr[\widehat{\text{flip}}_B \geq \beta + t] \leq e^{-2mt^2}$: the population guarantee transfers to *any* independent ensemble with an explicit exponential rate. This is formalized axiom-free in Lean as `transfer_flipFreq_bound` (and the sum form `transfer_flip_bound`) in `TransferTheorem.lean`, composing Mathlib’s Hoeffding lemma with the population Cantelli bound, and closes the in-sample-to-out-of-sample gap the split-ensemble study was designed to probe. Because the split-ensemble models are bootstrap-*correlated* rather than i.i.d., the guarantee that strictly applies to the observed transfer is not the Hoeffding bound above (which assumes independence) but its dependence-free form, which we also prove by removing the independence assumption *entirely*: population Cantelli applied to the deployment flip frequency gives $\Pr[\text{flip}_B \geq \beta + t] \leq \text{Var}(\text{flip}_B) / (\text{Var}(\text{flip}_B) + t^2)$ for *any* ensemble (`transfer_flip_general` in `TransferGeneral.lean`), the bound degrading gracefully as correlation inflates the flip-frequency variance and the Hoeffding bound being the sharper idealization in the independent case. Under either, the observed 100% transfer is now a theorem’s predicted consequence rather than a fortunate regularity. The proved, uncapped Cantelli bound itself holds for all 10,724 pairs, as it must; and the transfer theorem’s premise is borne out at scale—across those pairs the observed out-of-sample flip lies below the machine-checked Hoeffding link (the full-ensemble flip plus the finite- m slack) on 99.4% of pairs against its 95% design level, and below the end-to-end composite (the Cantelli bound plus that slack) on 100%—so the guarantee that transfers in theory is the one observed to transfer in practice, with a mean bound slack of ≈ 0.50 (honestly conservative, as Cantelli looseness and the finite- m Hoeffding term predict). Two honest caveats survive the scale-up: a single-model importance gap is itself a competitive *ranking* baseline—the certificate beats it on 87% of datasets and, unlike it, returns a calibrated probability and a guarantee, but the gap is not negligible—and two datasets (`balance_scale`, `xd6`) show weak out-of-sample ranking.

Simultaneous claims and a certified top- k . A single-claim guarantee does not control the *family-wise* error of testing many claims at once; a union (Bonferroni) bound does—the fraction of the ensemble on which *any* claim in a family flips is at most the sum of the per-claim Cantelli bounds (Lean: `simultaneous_cantelli`, axiom-clean). Empirically this family-wise bound holds on every dataset where it is non-vacuous (it is informative only for a small family or large SNRs; median slack 0.69), and it specializes to a guarantee on the headline “top- k features” claim: call a top- k set certified when the boundary comparison (k -th vs $(k+1)$ -th feature by mean importance) is STABLE. Across 104 datasets a certified top- k set reproduced *exactly* on an independent retrained ensemble in 100% of cases, versus 27–30% for uncertified sets—though certification fires rarely (few datasets have a stable top- k boundary), which is itself the

impossibility’s point: most “top-feature” claims are not stable, and the certificate says precisely which are.

Decomposing the per-model limit. Splitting a pair’s instability across four model classes (random forest, gradient boosting, extra-trees, linear) shows the between-class component exceeds the within-class one (median spread 0.058 vs 0.032): which model one picks drives more instability than retraining a fixed model. A certificate that requires agreement across the three tree-ensemble classes does *not* transfer better to a held-out linear model than a single-model certificate (0.60 vs 0.63)—reliability across architectures needs architecturally diverse models in the agreement set, not more of the same kind. This sharpens, rather than removes, the first of two standing limits. Reliability is *per model* (a claim stable for one model class need not be for another—STABLE-call agreement ≈ 0.65 across gradient-boosting, random-forest, and linear ensembles). Second, the feature-geometry capacity is a valid but sometimes-loose upper bound that captures correlation-driven Rashomon only—it can miss target-symmetry instability (independent but interchangeable features), which the per-query certificate still catches. The two combine into a sharper bound: subtracting the importance-gated target-symmetry instabilities the certificate detects among low-correlation feature pairs from the feature-geometry rank tightens the upper bound exactly where the geometry alone is loose (on a controlled interchangeable-feature stress it recovers the true capacity 7 where the geometry alone reports 8).

The deeper claim. If the group-theoretic parameterization is correct in full generality, then the Explanation Impossibility is not a standalone theorem but the *first theorem in a classification theory*—analogous to how Arrow’s impossibility theorem [Arrow, 1951] was the first theorem in social choice theory, not the last. The classification would say: for each symmetry type, here is exactly what explanations can and cannot do, here is the optimal resolution, and here is the residual ambiguity. The universality of the impossibility (Theorem 3.6) would then be the base case of a stratified tower of increasingly refined characterizations, parameterized by the representation theory of the symmetry group.

9 The Ceiling of Explainability [Tier A: Hardened]

The Explanation Impossibility (Theorem 3.6) establishes a *ceiling*, not a criticism of any particular method. No axiomatic redesign of SHAP, LIME, attention maps, or counterfactual generators can simultaneously deliver faithfulness, stability, and decisiveness whenever the underlying inference problem is underspecified. The ceiling is structural: it follows from the Rashomon property alone, independent of the explanation formalism. An analogy to physics is apt: the speed of light constrains spacecraft engineering but does not prevent useful space travel. Likewise, the impossibility constrains what practitioners may promise about explanations but does not prevent useful explainability—it redirects effort toward the achievable frontier characterized by the G -invariant resolution (Section 5).

10 Lean Formalization [Tier A: Hardened]

The Explanation Impossibility, all eight derived instances, and all nine ML instances are mechanically verified in Lean 4 [de Moura and Ullrich, 2021] using Mathlib [The mathlib Community, 2020]. The eight derived instances are particularly notable: each derives the Rashomon property constructively with zero domain-specific axioms, confirming that the impossibility in each science follows from first principles alone. This section provides an expanded account of the formalization, including the axiom stratification, proof status transparency, and the complete file inventory.

10.1 Summary Statistics [Tier A: Hardened]

Table 5 summarizes the formalization.

Trusted-base statement (authoritative; verified by a real build). All three repositories compile with `lake build` (0 errors) and contain **0** `sorry/admit`. The trusted base, counted at source, is:

- **Main repository:** **2** axiom declarations bundling ~ 14 gradient-boosting properties; **0** `native_decide` uses; **0** `sorry`.
- **Attribution companion:** **2** axiom declarations bundling **7** fields (of which one is genuinely behavioral); **0** `native_decide`; **0** `sorry`.
- **Physics companion:** **0** axiom declarations—the Selmer 1951 result and ~ 11 black-hole/spacetime-emergence properties are carried as *hypotheses / section variables*, not axioms, but still gate their downstream results; **1** `native_decide` use (a single 419,904-case circuit enumeration, which imports Lean’s `ofReduceBool`; every other enumeration is kernel-checked `decide`); **0** `sorry`.

So **4** axiom declarations in total ($2+2+0$); bundling changes the keyword count, not what is assumed. **The core theorem `explanation_impossibility` and the entire Tier-A hardened spine use none of these**—they are pure logic from the stated hypotheses (verified by `#print axioms`). Theorem counts are reported *per repository* (main 715; attribution companion 368, $\approx 88\%$ coinciding with the main repository; physics companion 481, $\approx 11\%$ coinciding); because the companions overlap we do *not* sum them, and the distinct program-wide count is 1178. Results that depend on the Selmer 1951 result or the black-hole/spacetime hypotheses are conditional on those assumptions.

Table 5: Summary of the Lean 4 formalization. Per-theorem “ N axioms” entries below count the *behavioral properties* a result depends on; at source these are bundled into the **2** axiom declarations (`gbdtWorld`, `gbdtAxioms`), so `#print axioms` reports 2, not their sum. The eight science instances and the nine constructive ML instances use *zero* axioms.

Component	Count	Notes
Files	130	All in <code>UniversalImpossibility/</code>
Theorems & lemmas	715	Verified by <code>grep -c</code>
Axioms	2	main repo, bundling ~ 14 GBDT properties (4 total across 3 repos—see above); core uses 0
<code>sorrys</code>	0	Zero admitted goals
Core impossibility	1 theorem	Zero custom axiom dependencies
8 derived instances	0 each	Domain-specific axioms: none (biology, physics $\times 2$, linguistics, ...)
GBDT instance	4 axioms	Surjectivity, first-mover, split structure
DASH resolution	6 axioms	+ cross-group symmetry
Causal instance	3 axioms	CPDAG Markov equivalence
Neural net instance	2 axioms	Overparameterization + symmetry
Fairness instance	2 axioms	Unequal base rates + equal accuracy

Of the 715 theorems, the majority are short (1–10 tactic steps), including one-line applications of `explanation_impossibility` to domain instances, tightness witnesses, and definitional consequences. Approximately 60 are moderate (11–30 steps) and 15 are substantive (31+ steps,

including the GBDT ratio divergence, Spearman bounds, Pareto-optimality, and the infinite enrichment stack).

10.2 Axiom Stratification [Tier A: Hardened]

A key concern in any formalization using axioms is whether the axioms are logically consistent and whether they smuggle in the conclusion. We address this by providing complete transparency on the axiom stratification. Table 6 reports which theorems depend on which axioms, verified by `#print axioms` in Lean 4.

Table 6: Axiom stratification for key theorems, verified by `#print axioms`. “Zero axioms” means the theorem depends only on the hypotheses stated in its signature—no global axioms are used.

Theorem	Axioms	Description
<code>explanation_impossibility</code>	0	Core universal impossibility
<code>attribution_impossibility</code>	0	Core attribution impossibility
<code>impossibility_qualitative</code>	0	From dominance + surjectivity
<code>gbdt_impossibility_local</code>	4	<code>surj</code> , <code>fm</code> , <code>nfm</code> (no proportionality)
<code>impossibility (quantitative)</code>	5	+ <code>proportionality_global</code>
<code>consensus_equity</code>	6	+ cross-group symmetric
<code>attribution_impossibility_bundled</code>	0	Fully parametric via <code>GBDTSetup</code>
<code>attention_impossibility_constructive</code>	0	Constructive finite witness
<code>counterfactual_impossibility_constructive</code>	0	Constructive finite witness
<code>concept_impossibility_constructive</code>	0	Constructive finite witness
<code>causal_impossibility_constructive</code>	0	Constructive finite witness
<code>model_selection_impossibility_constructive</code>	0	Constructive finite witness
<code>gInvariant_stable</code>	0	G -invariance implies stability
<code>ubiquity_impossibility</code>	0	Constructs system from witnesses
<code>linear_system_impossibility</code>	0	Underdetermined $Ax = b$
<code>genetic_code_impossibility</code>	0	Codon degeneracy
<code>gauge_impossibility</code>	0	Discrete $\mathbb{Z}_2$ gauge theory
<code>syntactic_ambiguity_impossibility</code>	0	PP-attachment ambiguity
<code>phase_impossibility</code>	0	Crystallographic phase problem
<code>view_update_impossibility</code>	0	Database view update
<code>stat_mech_impossibility</code>	0	Microstate degeneracy
<code>causal_impossibility_derived</code>	0	Markov equivalence
<code>langlands_boundary_statement</code>	0	$n=1$ full, $n \geq 2$ collapsed
<code>multiAnalyst_impossibility</code>	0	Multi-analyst F+S+D impossible
<code>multiAnalyst_bilemma</code>	0	Multi-analyst F+S impossible

Zero-axiom scope. The entire theoretical framework—impossibility theorem, bilemma, tightness classification, capacity theorem (Pythagorean decomposition, best approximation, over-explanation penalty), Pareto-optimal resolution, and all domain instances including Arrow’s theorem, Peres–Mermin contextuality, the Langlands boundary for $\mathrm{GL}(n, \mathbb{F}_p)$, multi-analyst aggregation, the Navier–Stokes tightness schema, the DPRM trilemma, and all 9 ML explanation types—requires zero axioms. The only axioms in the formalization (4 `axiom` declarations in total across the three repositories, not summed over files) encode empirical properties of gradient-boosted decision trees for the ML-specific quantitative bounds. The theoretical framework is a pure theorem of logic. We report counts *per repository* (main 715, attribution companion 368, physics companion 481; 4 `axiom` declarations in total, 2+2+0; 0 `sorry`; all three build with `lake build`) and do *not* sum them, because the attribution companion overlaps the main repository by $\approx 88\%$ of theorem names and the physics companion by $\approx 12\%$; the distinct program-wide count is approximately 1,000. Key new files from this formalization campaign: `LanglandsCorrespondence.lean` (7 theorems: trace conjugation-invariance, the $n=1/n \geq 2$

boundary, scalar embedding functoriality) and `MultiAnalystInstance.lean` (4 theorems: multi-analyst Rashomon, trilemma, bitemma, maximal incompatibility).

Reynolds naturality and the character: a structural analogy, not a Langlands result. [*Tier C.*] The Reynolds naturality theorem (`reynolds_naturality` in `UncertaintyFromSymmetry.lean`) proves that equivariant maps commute with Reynolds projections, and for $GL(n)$ the Reynolds projection is the trace (conjugation-averaging). This is a one-line fact of invariant theory: the trace is a class function, constant on conjugacy classes and compatible with the group action. We flag explicitly that this is *not* Langlands functoriality: the formalization contains no automorphic forms, no L -functions, and no automorphic-to-Galois correspondence. The parallel is purely structural—characters/traces are the natural stable invariant under conjugation, just as the orbit average is the stable resolution under any symmetry group—and is stated only to orient readers familiar with representation theory. It is not a theorem about, or a prediction of, functoriality.

The automorphic gap. The companion repository’s `HomExplanationSystem.lean` identifies the precise content separating the impossibility framework’s structural predictions from the Langlands programme’s arithmetic classification. When the resolution is required to be a group homomorphism (not merely a stable function), the enrichment Sha becomes strictly stronger than CRT. For $S_3 = GL(2, \mathbb{F}_2)$: CRT allows 6 character-value combinations, but the group relation $srs = r^2$ forces $\chi(r) = 1$, blocking 4 of them (Lean-verified: `automorphic_gap`). The blocked combinations are locally valid but violate a global group relation—not every collection of local character data assembles into a global homomorphism. This is the automorphicity condition: the gap between local compatibility (CRT, handled by the enrichment framework) and global coherence (automorphicity, the content of the Langlands correspondence). Three negative results delineate the boundary: (1) enrichment Sha = CRT = Euler multiplicativity (no arithmetic content beyond the Euler product; `enrichment_sha_equals_crt`); (2) collapsed tightness forces the identity resolution (tautological; `forced_resolution_complete`); (3) quantitative bounds depend only on group size, not on prime-specific structure (`EnrichmentQuantitative.lean`). For finite fields, the automorphicity gap is computable; for number fields, it is the central open problem of modern algebraic number theory.

The critical observation is that the core theorem `explanation_impossibility` uses **zero custom axioms**—it is pure constructive logic. The formalization stratifies into two categories:

- **Constructively derived** (8 cross-domain instances): The Rashomon property is proved from concrete witnesses with zero behavioral axioms. The Lean code constructs explicit pairs (e.g., (1,1) and (0,2) for the linear system, chain vs. fork DAGs for Markov equivalence) and verifies $\text{obs}(\theta_1) = \text{obs}(\theta_2) \wedge \theta_1 \neq \theta_2$ by `decide`. These are minimal witnesses of domain-specific phenomena, not complete formalizations of the respective theories.
- **Axiomatized** (9 ML instances): The Rashomon property is assumed as a domain axiom (e.g., “there exist neural network weights with the same predictions but different attention distributions”). These axioms are the mathematical content of the empirical literature being formalized and are independently verifiable from the cited sources.

The core impossibility (Theorem 3.6) requires zero axioms from any category. Quantitative bounds draw from ML domain axioms; cross-domain transfer draws from structural axioms.

Scope of formal verification. Lean verification guarantees that all theorems follow from the stated axioms with no logical gaps (internal consistency). It does not guarantee that the axioms correctly model the phenomena they are intended to capture (external validity). The axioms are motivated by domain-specific considerations (§10.2), but their adequacy for any particular

Table 7: Axiom taxonomy (2 bundled axioms; all behavioral properties are theorems extracted from the bundles).

Category	Count	Examples
GBDT infrastructure	14	Model, splitCount, proportionality
Enrichment stack (physics)	11	BH/QG frameworks, Rashomon

real-world system is an empirical question. We report axiom counts alongside theorem counts throughout to make the trusted base transparent.

10.3 Proof Status Transparency [Tier A: Hardened]

We distinguish four levels of verification:

- (a) **Proved** (zero custom axiom dependencies): The theorem follows from pure logic and stated hypotheses. Examples: `explanation_impossibility`, `attribution_impossibility`, `gInvariant_stable`, and `ubiquity_impossibility`.
- (b) **Derived** (from axioms): The theorem follows from a small set of domain-specific axioms. Examples: `gbd_t_impossibility_local` (4 axioms), `consensus_equity` (6 axioms). The axioms are independently verifiable (e.g., the surjectivity axiom for GBDT follows from the symmetry of collinear features under the data-generating process).
- (c) **Axiomatized** (Rashomon assumed): The original axiomatized instance files remain for backward compatibility, but every domain now has a constructive counterpart. All eight cross-domain scientific instances and all nine ML instances (attention, counterfactual, concept, causal, model selection, saliency, LLM self-explanation, mechanistic interpretability) have constructive Rashomon witnesses derived via `decide` with zero domain-specific axioms—bringing the constructive total to 16 of 17 `ExplanationSystem` instances. (The 17th, `AttributionInstance`, uses the independent `attribution_impossibility` with its own zero-axiom proof.)
- (d) **Argued** (supplement proof only): The result has a rigorous proof sketch in the supplement or companion paper but has not been mechanized in Lean 4. (Note: the *within-group* component of DASH’s Pareto optimality is Lean-verified in `ParetoOptimality.lean` via `dash_within_group_dominance`: a committed within-group ranking has strictly positive disagreement, conditional on the bundled GBDT/DGP hypotheses. The between-group case—and hence *global* Pareto optimality—is now *reduced to one named hypothesis* with the reduction machine-checked (`ParetoGlobal.lean`): the Cramér–Rao inequality itself is proved unconditionally in abstract form (`abstract_cramer_rao`, via a covariance Cauchy–Schwarz), the score-regularity content is isolated as `CramerRaoScoreProperty` (the analogue of `SardProperty`), the chain from it to global optimality is verified (`cramer_rao_bound_of_score` $\rightarrow$ `dash_global_min_variance` $\rightarrow$ `dash_global_pareto_optimal`), DASH attains the bound exactly (`dash_variance_eq_inv_fisher`), and the named hypothesis is now *discharged* for every natural exponential family (`cramerRaoScoreProperty_tilted`, via machine-checked differentiation under the integral for exponential tilts, `hasDerivAt_integral_mul_exp`), subsuming the single-sample Gaussian (`cramerRaoScoreProperty_gaussianReal`). Moreover the concrete M -sample exchangeable-Gaussian ensemble on which DASH operates is verified *end to end* and axiom-clean: `dash_mvue_exchangeableGaussian` proves the DASH aggregate is the minimum-variance unbiased estimator among all square-integrable regular unbiased estimators of that model, so the global Cramér–Rao attainment (`dash_global_min_`

variance) is now mechanized for the model DASH actually uses. What remains argued is only the score property for general *non-exponential* families (e.g. Cauchy-type location families lacking exponential moments)—the residual differentiation-under-the-integral condition classical statistics assumes.)

- (e) **Empirical:** The result is supported by computational experiments but not formally proved. Examples: the specific flip rates (19.9% attention under full retraining, 23.5% counterfactual, etc.) reported in Section 17.

10.4 Formerly Axiomatized, Now Derived [Tier A: Hardened]

Several results that were initially axiomatized have since been derived from more primitive axioms, strengthening the formalization:

- `spearman_classical_bound` $\rightarrow$ `spearman_instability_bound` in `SpearmanDef.lean`: derived from split-count structure.
- `le_cam_lower_bound` in `QueryComplexity.lean`: provable by `not_lt.mp` (a tautology in any linear order).
- `consensus_variance_bound` in `Defs.lean`: from `attribution_variance_nonneg` + `Nat.cast_nonneg`.
- `attribution_sum_symmetric` in `SymmetryDerive.lean`: from proportionality + split-count + cross-group + balance.
- `attribution_variance` and `attribution_variance_nonneg`: from Mathlib’s `ProbabilityTheory.variance` and `variance_nonneg`.

10.5 Lean Proof of the Core Theorem [Tier A: Hardened]

The mechanized proof of the universal impossibility mirrors the paper proof exactly. The complete Lean 4 statement and proof is given in Appendix B.

10.6 Why This Result Is Non-Trivial [Tier A: Hardened]

A natural objection is that the impossibility is tautological: “of course you cannot explain an underspecified system perfectly.” Four lines of evidence refute this characterization.

1. Tightness: the triple is the exact obstruction. Each *pair* of {faithful, stable, decisive} is simultaneously achievable—three Lean-verified witnesses establish this (Proposition 3.10). A faithful-and-decisive map simply reports the native explanation; a faithful-and-stable map outputs a neutral constant; a stable-and-decisive map outputs a maximally committal constant. The impossibility is *specifically* the conjunction of all three. A tautology would make even pairs unachievable; here every pair is constructively realized.

2. Necessity: the Rashomon property is the exact boundary. Two converses are proved in Lean. For *fully specified* systems (`fully_specified_possibility` in `Necessity.lean`): if every observation uniquely determines its configuration, all three properties are simultaneously achievable. And for systems whose incompatibility relation is inequality—the maximal-incompatibility class covering the blemma instances—the full equivalence is machine-checked (`rashomon_biconditional_neq` in `NecessityBiconditional.lean`): the three properties are jointly achievable *if and only if* the Rashomon property fails. For

Table 8: **Structural differences across the eight derived instances.** The eight cross-domain instances are not notational variants of a single example: they differ in symmetry group, group type, and resolution strategy. These independently developed methods share a common mathematical structure (orbit averaging / Reynolds projection), which the G -invariant resolution framework identifies retrospectively.

Domain	Group G	Objects permuted	Type	Resolution	Form
Mathematics	$\mathbb{R}^{n-r}$	Solution offsets	Cont., abelian	Pseudoinverse	Projection
Biology	S_k ($k \leq 6$)	Synonymous codons	Finite	Codon tables	Averaging
Physics	$\mathbb{Z}_2^{ V -1}$	Edge labels	Finite, abelian	Wilson loops	Invariant fn.
Stat. mech.	S_Ω	Microstates	Finite	Microcanon.	Averaging
Linguistics	S_k	Parse trees	Finite	Parse forests	Enumeration
Crystallogr.	$\mathbb{Z}_2 \times \mathbb{Z}_n \times \mathbb{Z}_2$	Signal phases	Finite	Patterson map	Autocorrel.
Comp. sci.	S_k	Hidden columns	Finite	Compl. views	Projection
Statistics	MEC perms.	DAG orientations	Finite	CPDAG	Class report

general incompatibility relations the boundary is now characterized exactly: the biconditional holds *precisely* when the incompatibility relation is the complement of an equivalence relation (equivalently, when “compatible” is transitive), proved as `rashomon_biconditional_of_equiv`—with the substantive possibility direction `possibility_of_equiv_compat`—in `IncompatibilityCharacterization.lean`, of which the inequality case above is the finest special case. That the equivalence hypothesis is essential is itself machine-checked (`general_boundary_fails` exhibits a non-Rashomon system admitting no faithful, stable, and decisive explanation). So “Rashomon is the exact boundary” is a precise logical equivalence exactly for congruence-structured incompatibility, and provably false without it—not a vague “things are hard when the problem is hard” statement.

3. Resolution optimality: the G -invariant projection is Pareto-optimal. The resolution is not arbitrary. The orbit-averaging map is Pareto-optimal among all stable explanation maps: no stable map can achieve strictly higher pointwise faithfulness on every observe-fiber. This means the trade-off identified by the theorem is *tight*—the practitioner cannot do better by being clever. A second, variance-theoretic optimality is machine-checked for the aggregation itself: among all linear unbiased aggregators of exchangeable estimators, the equal-weight orbit average is the *minimum-variance* choice (a Gauss–Markov statement for the symmetric case; `dash_min_variance` and `dash_orbit_average_optimal` in `GaussMarkovDASH.lean`, via the Cauchy–Schwarz core $\sum w_i^2 \geq 1/n$ at $\sum w_i = 1$).

4. Shared resolution structure across eight domains. If the impossibility were trivial, one would not expect eight independent scientific communities—spanning physics, biology, statistics, linguistics, crystallography, computer science, statistical mechanics, and mathematics—to converge on structurally the same resolution strategy (orbit averaging) over more than a century of independent work (Table 8). Our framework identifies what these eight independently developed methods share—instances of the same Reynolds projection—and proves solving the same impossibility—and proves the resolution they found is optimal.

The framework also explains the variation across domains: the specific form of the resolution differs because the symmetry group and representation differ, even when the abstract group is the same. Four instances use the symmetric group S_k , but they act on different mathematical objects—synonymous codons, microstates, parse trees, database columns—and yield different resolution forms: frequency averaging for biology and statistical mechanics, equivalence-class enumeration for linguistics, and subquery projection for computer science. The abstract math-

emational operation—projection onto the G -invariant subspace—is identical; the computational realization is domain-specific.

Part II

Part II — Empirical Validation

Measurements and audits testing the framework. The strongest result is the stable-fact-count bimodality (Tier A, pre-registered, $p < 10^{-12}$); the capacity η -law and per-domain fits are Tier B, with sample-size and real-data caveats stated inline. None of these empirical claims is a logical consequence of the Part I theorem—they test it, they do not follow from it.

11 Cross-Domain Transfer: From Taxonomy to Generative Theory [Tier B: Empirical]

The preceding sections establish the universal explanation impossibility as a taxonomic result: it classifies nine pre-existing explanation types under a single abstract structure. A taxonomy, however, is not yet a theory. A *generative* theory must produce novel, testable predictions—ideally in domains where the framework was not originally developed. This section reports four cross-domain transfer experiments that test such predictions. In each case, a concept developed for one domain (machine learning attribution, gauge-theoretic coupling, genetic codon degeneracy) is transferred to a different domain via the abstract framework, and the transferred prediction is confirmed empirically.

Evidence tiers. We classify empirical results by strength of support. **Strong validation** (quantitative prediction confirmed with $R^2 > 0.8$ or pre-registered): SAGE symmetry-group predictor ($R^2 = 0.809$, leave-one-out CV), capacity theorem for 7 well-characterised instances ($R^2 = 0.957$, $p = 1.4 \times 10^{-4}$), stable fact count bimodality (pre-registered, 47.1 pp gap invariant across ρ), exhaustive pipeline validation (DASH 84/85, SAGE directional 81/85 at all P ranges, across 85 PMLB datasets including regression; §11.11). **Corrected from earlier reports:** Gaussian $\Phi(-\text{SNR})$ OOS $R^2 > 0.80$ on 59% (not 95%) of datasets with $M = 50$ and gain-based importance; earlier in-sample estimates inflated by 31 pp. Coverage conflict: $\rho = 0.87$ across all P with TreeSHAP ($\rho = 0.90$ for $P \leq 20$ only with gain-based importance; the previously reported $\rho = 0.96$ on 15 datasets is consistent with the TreeSHAP finding). **Consistent** (observation matching qualitative prediction): attention flip 19.9% [CI: 19.1–20.6%], counterfactual flip 23.5% [CI: 20.6–26.6%], model selection flip 80%. **Calibration checks** (expected from framework structure, not independently informative): codon degeneracy–entropy rank correlation ($\rho = 1.0$; monotonicity in degeneracy is guaranteed by construction—see §11.13). **Partial:** group identification capacity theorem for 9 instances with unknown group structure (holdout $R^2 = 0.24$; poor fit attributable to misspecified or unknown symmetry groups—see §11.8).

11.1 Resolution Transfer: Ensemble Causal Discovery [Tier B: Empirical]

The DASH resolution (Section 5) was developed for ML attribution stability: ensemble averaging over the Rashomon set projects onto the G -invariant subspace, eliminating gauge-dependent fluctuations. The framework predicts that this resolution should transfer to *any* domain exhibiting the Rashomon property, because the abstract mechanism—orbit averaging—is domain-independent.

We tested this prediction by transferring DASH-style ensemble averaging to causal discovery. Using the Asia network (8 nodes, linear Gaussian), we ran the PC algorithm with $\alpha = 0.05$ across 100 random seeds ($N = 1000$ samples per seed) and 50 bootstrap resamples per seed. Three methods were compared:

1. **Single-run PC:** orientation flip rate = 0.014, decisiveness = 0.36.




Figure 1: DASH-style ensemble averaging transferred to causal discovery (Asia network, 100 seeds, 50 bootstrap resamples). The ensemble reduces the orientation flip rate by $4\times$ while trading off decisiveness, confirming the trilemma tradeoff in a new domain.

2. **Bootstrap majority vote:** flip rate = 0.014, decisiveness = 0.37.

3. **DASH-style ensemble:** flip rate = 0.003 ($4.1\times$ reduction; Wilcoxon $p < 10^{-10}$), decisiveness = 0.21.

The trilemma tradeoff predicted by the framework is clearly visible: the DASH ensemble achieves a $4\times$ reduction in orientation flip rate (the stability axis) at the cost of reduced decisiveness (0.21 vs. 0.36). This is not a post-hoc observation—it is the *same* tradeoff formalized in Theorem 3.6, now operating in causal discovery rather than feature attribution. The result demonstrates that the G -invariant resolution is a practical, transferable algorithm, not merely a theoretical construct (Figure 1).

11.2 Stability Transition in Explanation Systems [Tier B: Empirical]

The gauge-theoretic formulation (Section 5) shows that the coupling parameter β controls stability via a $\text{sech}^2(\beta)$ response function. The framework predicts an analogous stability transition in ML: as the Rashomon tolerance ε increases, stability should undergo a sigmoid-shaped transition from near-perfect to unstable, with the critical threshold ε_c decreasing as collinearity ρ increases (because higher collinearity enlarges the Rashomon set).

We swept ε over 20 log-spaced values in $[0.01, 2.0]$ at four collinearity levels ($\rho \in \{0.0, 0.3, 0.6, 0.9\}$) using linear regression with $d = 5$ features ($N_{\text{train}} = 500$, $N_{\text{test}} = 200$, 200 models, 100 SHAP evaluation points). All four framework predictions were confirmed:

1. **Monotonicity:** stability decreases monotonically with ε at all collinearity levels (fraction of decreasing steps ≥ 0.89).
2. **Sigmoid shape:** the transition is sigmoid-like at all ρ values (confirmed via inflection-point analysis), with sharpness increasing from 0.47 at $\rho = 0.0$ to 1.61 at $\rho = 0.9$.
3. **Critical threshold:** $\varepsilon_c = 1.51$ at $\rho = 0.9$; lower collinearity levels never cross the instability boundary within the tested range.
4. **Collinearity as inverse coupling:** higher ρ produces earlier and sharper transitions, exactly as predicted by the gauge analogy—collinearity plays the role of inverse coupling β^{-1} , enlarging the Rashomon set and lowering the stability threshold.




Figure 2: Stability transition in SHAP stability as Rashomon tolerance ε increases, at four collinearity levels. Higher collinearity produces sharper, earlier transitions ($\varepsilon_c = 1.51$ at $\rho = 0.9$), analogous to the gauge-theoretic $\text{sech}^2(\beta)$ response.

The analogy to statistical mechanics is precise: just as a spin system undergoes a phase transition from ordered to disordered as temperature ($\sim \beta^{-1}$) increases, explanation stability undergoes a transition from stable to unstable as the Rashomon tolerance (controlled by collinearity) increases. The $\text{sech}^2(\beta)$ response function from the gauge theory maps onto the sigmoid stability curve observed empirically (Figure 2).

11.3 Cross-Level Prediction: Protein Designability [Tier B: Empirical]

Character theory (Section 5) predicts that for the symmetric group S_k acting on $\mathbb{R}^k$, the normalized entropy satisfies $H/H_{\max} \approx (d-1)/d = 1 - 1/d$, where d is the degeneracy of the representation. This law was derived from codon degeneracy in the genetic code—a Level 3 (molecular biology) phenomenon. The cross-level prediction: the same $1/k$ law should govern protein sequence diversity at each aligned position, a Level 4 (structural biology) phenomenon connected to the genetic code through the translation machinery.

We tested this prediction on 120 aligned cytochrome c sequences (105 analyzed positions with $d \geq 1$ distinct amino acids observed). The results strongly support the character-theory prediction:

- Pearson $r = 0.798$ ($p = 0.018$), Spearman $\rho = 0.881$ ($p = 0.004$) between observed $H/H_{\max}$ and the predicted $(d-1)/d$ curve, grouped by degeneracy level.




Figure 3: Protein sequence diversity at aligned cytochrome c positions vs. the character-theory prediction $(d - 1)/d$. All 59 variable positions lie closer to the character-theory curve than to uniform, with systematic undershoot attributable to purifying selection. Pearson $r = 0.798$, Spearman $\rho = 0.881$.

- 100% of variable positions ($d \geq 2$; $n = 59$) are closer to the character-theory prediction than to the uniform null ($H/H_{\max} = 1$).
- Character-theory RMSE = 0.378 vs. uniform RMSE = 0.618 (1.63 $\times$ improvement).
- Systematic undershoot: observed ratios fall 40–50% below the character-theory curve (mean signed residual = -0.359 , $t = -16.1$, $p < 10^{-22}$).

The systematic undershoot has a natural biological interpretation: the $1/k$ law gives the *neutral-evolution upper bound*—the maximum diversity expected under pure drift with degeneracy d . Purifying selection constrains observed diversity below this bound. The magnitude of the deviation at each position therefore measures the strength of selective pressure, providing a quantitative link between the abstract character-theoretic prediction and molecular evolution. This prediction was derived entirely from the cross-domain framework, connecting genetics and structural biology through representation theory (Figure 3).

11.4 Cross-Level Prediction: Transcription Factor Binding Sites [Tier B: Empirical]

The $1/k$ law, derived from character theory of S_k , was confirmed at the protein-coding level (Section 11.3). We test whether the same law generalizes to an entirely different level of biological organization: *regulatory DNA*.

At each position in a transcription factor (TF) binding site, some nucleotides preserve binding (the “allowed” set, size $d = 1$ to 4). The framework predicts: nucleotide entropy should scale as $(d - 1)/d \times H_{\max}$, where $H_{\max} = \log_2(d)$.

We analysed 500 vertebrate TF binding profiles from the JASPAR CORE database [Raulu-seviciute et al., 2024], comprising 5,624 binding-site positions. Positions were grouped by the number of tolerated nucleotides (d , defined as nucleotides with frequency $> 5\%$).

Results. Nucleotide diversity increases monotonically with d (Kruskal–Wallis $H = 205.6$, $p = 2.3 \times 10^{-45}$):

d	Predicted $(d-1)/d$	Observed mean $H/H_{\max}$	n positions
2	0.500	0.766	1,274
3	0.667	0.782	823
4	0.750	0.861	1,306

The attenuation factor for TFBS is $\alpha_{\text{TFBS}} = 1.23$, compared to $\alpha_{\text{protein}} \approx 0.45\text{--}0.88$ for protein-coding sequences. The character-theoretic baseline $(d-1)/d$ acts as a *ceiling* for protein diversity (selection constrains below) but as a *floor* for regulatory diversity (observed entropy exceeds the prediction). This asymmetry is biologically interpretable: protein-coding positions are under strong purifying selection ($\alpha < 1$), while TF binding positions tolerate more nucleotide variation ($\alpha > 1$), consistent with the known weaker conservation of regulatory DNA relative to coding DNA.

The $1/k$ law thus operates at two levels of biological organization— protein-coding and regulatory—with the character-theoretic baseline providing a reference point that separates purifying ($\alpha < 1$) from relaxed ($\alpha > 1$) selection regimes.

Rashomon overlap and class confusion (hypothesis). The framework suggests that class confusion may be driven by model multiplicity (Rashomon overlap between classes) rather than feature-space similarity. Preliminary experiments show rank correlations of 0.68–1.00 between Rashomon overlap and the confusion matrix, but both quantities are functions of the same predictions on the same data, making the perfect correlations likely near-identities rather than non-trivial predictions. Independent validation of this hypothesis (e.g., predicting confusion on held-out model architectures) is needed before claiming empirical support.

11.5 Invariance Counting: Exact Enumeration of Stable Queries [Tier A: Hardened]

The invariance correspondence (Section 7) predicts that the number of independently stable ranking facts equals $\binom{g}{2}$, where g is the number of dependence groups (identified here via correlation; the operative dependence condition is $I > 0$). Within-group comparisons should be unstable ($\sim 50\%$ flip rate), while between-group comparisons should be perfectly stable. We tested this counting prediction directly.

We generated $P = 12$ features in $g = 3$ groups of 4, with within-group correlation $\rho = 0.99$ and between-group correlation $\rho = 0.0$. True coefficients were set to $\beta = (5, 5, 5, 5, 2, 2, 2, 2, 0.5, 0.5, 0.5, 0.5)$ —equal within each group, well-separated between groups. We trained 200 Ridge regression models and 200 XGBoost models on independent training sets ($N_{\text{train}} = 500$, $\sigma = 1.0$) and measured the pairwise flip rate for all $\binom{12}{2} = 66$ feature pairs.

The result is a sharp bimodal distribution (Figure 4):

- **Within-group pairs** ($n = 18$): mean flip rate = 0.500 (range [0.491, 0.503]); all 18 unstable ($> 40\%$).
- **Between-group pairs** ($n = 48$): mean flip rate = 0.0015 (range [0.000, 0.020]); all 48 stable ($< 5\%$).
- **Bimodal gap**: 47.1 percentage points.
- Mann–Whitney $U = 864$, $p = 2.71 \times 10^{-13}$ (Cohen’s $d = 1.41$).

The XGBoost replication confirms the same pattern (within mean = 0.476, between mean = 0.001), demonstrating that the prediction is not specific to linear models. Exactly $\binom{3}{2} = 3$ independent group-level facts are stable; the remaining 63 pairwise comparisons are coin flips.

ρ	Within mean	Between mean	Gap (pp)	Perm. p
0.50	0.500	0.000	50.0	< 0.0002
0.70	0.500	0.000	50.0	< 0.0002
0.85	0.500	0.000	50.0	< 0.0002
0.99	0.500	0.002	49.8	< 0.0002

Table 9: Invariance counting sensitivity (cf. Noether’s theorem). The bimodal gap between within-group and between-group flip rates is invariant across $\rho \in [0.50, 0.99]$. Full table (8 values) in Supporting Information.




Figure 4: Invariance counting experiment (cf. Noether’s theorem). Pairwise flip rates for 66 feature pairs across 200 models, showing a sharp bimodal separation: within-group pairs (red) cluster at $\sim 50\%$ (coin flip), while between-group pairs (blue) cluster at $\sim 0\%$ (perfectly stable). The bimodal gap of 47 percentage points confirms that exactly $\binom{g}{2} = 3$ independent group-level ranking facts are stably answerable.

Sensitivity analysis. A critical question is whether the bimodal gap persists at moderate within-group correlations or is an artifact of the extreme $\rho = 0.99$ design point. We repeated the experiment at $\rho \in \{0.50, 0.60, 0.70, 0.80, 0.85, 0.90, 0.95, 0.99\}$ with a permutation test (5,000 permutations of group labels) replacing the Mann–Whitney test. The gap is **invariant across the entire range**: at every ρ , within-group pairs flip at $50.0 \pm 0.2\%$ and between-group pairs flip at 0.0% , with a permutation $p < 0.0002$ at every level (Table 9). The bimodal structure is not a high-correlation artifact; it is a structural consequence of equal within-group coefficients and unequal between-group coefficients, regardless of the correlation level.

This result has immediate practical consequences: a practitioner computing SHAP values for features with correlation structure can determine *in advance* which comparisons will be reliable and which are meaningless, regardless of the specific correlation level.

11.6 Automatic Symmetry Discovery (SAGE) [Tier B: Empirical]

The invariance counting result (Section 11.5) assumes the correlation groups are known. In practice, the explanation symmetry group G is rarely hand-specified. We developed **SAGE** (Symmetry-Aware Group Estimation), an algorithm that automatically discovers G from bootstrap feature-importance flip rates without hand-specification.

Algorithm. For each dataset: (1) train $M = 100$ models on bootstrap resamples; (2) compute feature importances for each model; (3) cluster features by pairwise flip rate using hierarchi-

cal clustering (Ward linkage, threshold at 40% flip rate); (4) count the number of discovered groups g and predict instability as $\hat{\eta} = g/P$.

Results. Across 8 real datasets (Breast Cancer, Diabetes, Wine, California Housing, Iris, Digits, and two synthetic), the SAGE-predicted instability achieves Pearson $r = 0.958$ (calibrated $R^2 = 0.809$, leave-one-out cross-validated $R^2 = 0.689$). SAGE outperforms a correlation-matrix baseline by $28\times$ (SAGE $R^2 = 0.915$ vs. correlation $R^2 = 0.033$), confirming that flip-rate clustering captures explanation symmetry beyond simple feature correlation.

Adversarial audit. Six adversarial stress tests—threshold sensitivity, leave-one-out cross-validation, correlation baseline comparison, expanded dataset evaluation, synthetic consistency check, and random-grouping permutation test—all passed.

Calibration. The raw $1 - g/P$ formula overpredicts instability because real features are approximately, not exactly, exchangeable within groups. A single calibration parameter (slope = 0.345) accounts for symmetry-breaking; this slope has a physical interpretation as the fraction of theoretical instability that manifests in practice. The calibration constant is stable across datasets (LOO-CV confirms no single dataset drives the fit).

SAGE baseline comparison. SAGE outperforms all naive baselines on instability prediction. After leave-one-out calibration, SAGE achieves $R^2 = 0.686$ while all baselines are negative (best baseline: mean pairwise correlation, LOO $R^2 = -0.47$). On rank ordering, SAGE dominates: Spearman $\rho = 0.952$ vs. best baseline 0.731 (permutation $p = 0.002$). Raw (uncalibrated) R^2 is negative for all methods including SAGE, because predicted instability systematically overshoots observed values. SAGE’s value is in identifying *which* feature groups are unstable, not in predicting exact flip rates.

11.7 Discrete vs. Continuous Symmetry: The CCA Spectrum [Tier B: Empirical]

Applying the G -discovery approach to neural network representations reveals a qualitative difference from feature attribution. We computed CCA (canonical correlation analysis) canonical correlations between hidden activations ($n = 128$) of 5 independently trained MLPs on MNIST. The CCA spectrum is *continuous*: the top 10 dimensions have correlation > 0.99 , the top 30 have > 0.90 , and the bottom 30 have < 0.30 , with smooth decay between. This contrasts sharply with the invariance counting result for features, where the flip-rate distribution is bimodal (50pp gap).

The difference arises because feature attribution symmetry is *discrete* (S_k permutations within correlation groups), while neural representation symmetry is *continuous* (rotations in hidden space). The framework’s S_n prediction ($\dim(V^G) = 1$) is incorrect for representations: CCA reveals ~ 30 highly stable dimensions, not 1. This finding delineates the framework’s boundary: the sharp invariance counting prediction holds for discrete symmetry groups but requires extension to continuous groups, where stability becomes a spectrum rather than a dichotomy.

11.8 The Explanation Capacity Theorem: Predictive Scope and Boundary [Tier B: Empirical]

The character-theoretic formula $\eta = \dim(V^G)/\dim(V)$ predicts the fraction of explanation information retained after G -invariant resolution. The tautology firewall requires that η predict a *downstream observable*—the measured instability rate—rather than a re-measurement of the

defining quantity. We set the predicted instability to $1 - \eta$ and compared against empirically measured flip rates, cosine distances, disagreement rates, and normalized entropies across 16 domain instances.

When restricted to the seven instances with well-characterized symmetry groups (attribution S_2 , concept probes $O(64)$, model selection S_{11} , codon $S_2/S_4/S_6$, statistical mechanics S_{252}), the universal capacity theorem achieves:

$$R^2 = 0.957, \quad \text{slope} = 0.914, \quad \text{intercept} = 0.008, \quad p = 1.36 \times 10^{-4}. \quad (1)$$

The slope near unity and intercept near zero confirm that $\eta = \dim(V^G)/\dim(V)$ is a parameter-free predictor of explanation instability when the group is correctly identified (Figure 5). The bridge from the abstract framework to the quantitative η prediction is empirical, not derived. The framework proves that V^G is the stably answerable subspace; the claim that $\dim(V^G)/\dim(V)$ predicts the *observed flip rate* requires the additional assumption that importance differences are approximately Gaussian—a statistical claim, not a consequence of the impossibility theorem.

An alternative bridge that sidesteps the group identification bottleneck uses the *coverage conflict degree*—the fraction of feature pairs with $\text{SNR} < 0.5$, computable from bootstrap retrains alone without requiring group identification. Across 15 datasets, coverage conflict degree predicts mean instability with $\rho = 0.961$ ($p < 10^{-8}$), robust to threshold choice ($\rho \in [0.885, 0.961]$ for thresholds 0.3–2.0) and leave-one-out removal ($\rho \in [0.952, 0.978]$; no single dataset drives the correlation). The formal bridge is proved in Lean: coverage conflict $\Leftrightarrow$ no neutral element (Theorem in `ExplanationLandscape.lean`). The empirical bridge validates that this algebraic property, derived from the impossibility theorem, predicts observed instability at scale. Unlike the capacity theorem, this bridge requires no group assignment and works for all 15 tested datasets.

The exchangeability criterion. The capacity theorem’s scope is determined by a subtlety: the relevant criterion is *exchangeability* (features genuinely compete for the same predictive signal), not merely *correlation*. A rigorous test on real data reveals that the group-mean test WORKS: projecting onto the G -invariant subspace (averaging SHAP within correlation groups) yields more stable attributions than the within-group deviations in 12/14 cases after Bonferroni correction. However, the individual-feature test FAILS: correlated features in XGBoost are often *more* stable than uncorrelated ones, inverting the naive prediction. This occurs because XGBoost consistently picks a “winner” within each correlation group, making the winner’s attribution stable even though the choice of winner is arbitrary.

The mechanism is detectable: genuinely substitutable features (e.g. `mean_texture` and `worst_texture` in Breast Cancer, $\rho_{\text{SHAP}} = -0.19$) exhibit negative cross-model SHAP correlation—when one’s importance goes up, the other’s goes down. This “first-mover” signature is confirmed on synthetic data ($\rho_{\text{SHAP}} = -0.11$ to -0.24). *Conjectural diagnostic*: the sign of cross-model SHAP correlation may detect genuine exchangeability (negative = substitutable; positive = independent). This is validated directionally on Breast Cancer and synthetic data but has not been tested at scale; it remains an open question whether this diagnostic generalises beyond these specific datasets.

Attempts to predict instability from single data properties alone (7 predictors: correlation, VIF, PCA loadings, mutual information, redundancy ratio, leverage, spectral gap) all fail (maximum Spearman $\rho = 0.26$). The capacity theorem works for exact domain-specific groups ($R^2 = 0.957$) but does NOT recover the exact group structure for correlation-derived groups on XGBoost+SHAP. The extension from “predicts instability given G ” to “from data” is therefore *partially* resolved, along two fronts that do not require knowing G : (i) the per-query reliability question is answered by the ensemble signal-to-noise certificate (Spearman 0.76–0.94 out-of-sample; §*capacity trace identity and bound*), and (ii) the capacity itself is *bounded*, and its stable

subspace recovered, from the feature-geometry effective rank (`capacity_le_of_ker`). What remains genuinely open is estimating the exact integer capacity from the model ensemble under approximate symmetry—which we show, with five principled estimators, is not reliably possible; capacity is best read from the feature geometry, and reported as a bound. Framed *as a bound*—the program’s sound epistemic mode, the same one the Cantelli certificate uses—capacity is validated by coverage rather than regression: across controlled systems, the dimension of the attribution subspace that is stable on *two independent* ensembles (the transfer-confirmed, identifiable structure) respects the capacity bound with 100% coverage and is informative ($C/p \approx 0.5$), and it is tight (equal to C) in the pure-symmetry cases. So the certificate tells a practitioner *which* facts are stable, and the capacity bounds *how many* can be—a bound that holds out-of-sample even where the point-predictive “ η -law” does not (the latter’s prospective failure, and the failure of a structural-floor reframe of it, are reported in the falsification appendix as a thrice-confirmed boundary).

Tautology concern. Coverage conflict degree (fraction of low-SNR pairs) and mean instability (mean flip rate) are both functions of the same importance distributions. The high correlation ($\rho = 0.961$) partly reflects this statistical relationship. The non-trivial content is twofold: (1) the *threshold* at $\text{SNR} = 0.5$ comes from the algebraic theory (the coverage conflict boundary, proved in Lean), not from fitting the data; and (2) the correlation is robust to threshold choice ($\rho > 0.88$ for all tested thresholds), confirming that the theory identifies the correct structural boundary rather than overfitting a particular cutoff. A pure statistical approach would not know *a priori* which summary statistic of the SNR distribution to use as a predictor.

Sonar outlier. One dataset (Sonar: $N = 208$, $P = 60$) produces $R^2 = -0.21$ within-dataset. The high dimensionality relative to sample size makes all importance estimates noisy, violating the Gaussian assumption. This defines a boundary condition: the Gaussian flip formula requires $N \gg P$ for reliable per-pair predictions. The cross-dataset bridge ($\rho = 0.961$) is not affected by this outlier (LOO- ρ with Sonar removed = 0.952).

The all-16-instance fit is the primary result. The 7-instance fit is a pre-specified subset analysis: the “well-characterized group” criterion (defined in the pre-registration document, `PRE_REGISTRATION.md`) requires that (1) the symmetry be exact by construction (part of the domain’s mathematical definition) and (2) the instability metric directly measure flip/disagreement rate. Across all 16 instances, the linear fit gives $R^2 = 0.60$. However, holdout prediction (calibrating on the 7 well-characterised instances and predicting the remaining 9) gives only $R^2 = 0.24$. The nine poorly-fit instances share a common cause: the assigned symmetry group was approximate (e.g., not all attention heads are exchangeable; GradCAM has broken spatial symmetry; parsers exhibit ambiguity on a subset of arcs). The η formula is correct; the group identification is the bottleneck.

capacity theorem failure mode. The capacity theorem achieves $R^2 = 0.957$ for 7 instances where G is well-characterised but only holdout $R^2 = 0.24$ for 9 instances where G is unknown or approximately specified. The failure is attributable to group misspecification: for lattice symmetry, parser ambiguity, causal discovery, and phase retrieval, the acting group is either infinite, unknown, or approximated by a finite proxy. The capacity theorem is a *conditional* prediction: given the correct G , it predicts instability accurately. The practical bottleneck is group identification.

Replacing group-theoretic predictions with Gaussian-predicted mean instability for domains with available per-pair data improves the universal η plot from $R^2 = 0.60$ to $R^2 = 0.79$ across all 16 domains, with the well-characterized subset reaching $R^2 = 0.98$.

Out-of-sample validation and goodness of fit. Leave-one-out cross-validation on the 7 well-characterised instances gives $R_{\text{LOO}}^2 = 0.794$ (MAE = 0.082), confirming the law is not an artefact of overfitting to 7 points. The linear fit is significant: $F(1, 5) = 38.78$, $p = 0.0016$, with residuals passing the Shapiro–Wilk normality test ($W = 0.968$, $p = 0.88$). A permutation null (10,000 permutations of η values) gives $p = 0.010$. On 9 held-out instances with unknown or approximately specified symmetry groups, the law generalises modestly ($R^2 = 0.244$, MAE = 0.169, Spearman $\rho = 0.522$). The practical bottleneck is group identification, not the law itself.

Group assignment justifications for ML instances. The η prediction depends critically on the assigned symmetry group. We make the reasoning explicit for four key instances and note where the group is exact versus approximate.

Attention (S_n). We model n attention heads as exchangeable under S_n . In practice, heads specialize during training, partially breaking this symmetry. The η prediction ($1 - 1/n$) is therefore an upper bound on instability; observed flip rates (19.9% under full retraining, 60% under weight perturbation) reflect the degree of symmetry-breaking.

Concept probes ($O(d)$). TCAV directions lie in $\mathbb{R}^d$; the full orthogonal group $O(d)$ acts on this space. Data covariance breaks rotational symmetry to a subgroup. The η prediction (≈ 1.0) is an upper bound; the observed instability (0.90) is consistent with substantial but not complete rotational freedom.

Model selection (S_k). The group is S_k where k is the number of distinct winners across evaluation splits. This assignment is exact: any permutation of models in the Rashomon set is equally valid as a selection, and the η prediction ($1 - 1/k$) follows directly.

GradCAM (spatial permutations). Spatial symmetry is heavily broken by CNN locality; the effective group is much smaller than the full spatial permutation group. This explains the low observed flip rate (9.6%; Cohen’s $d = 0.33$) despite a high predicted instability under the nominal group. GradCAM is an approximate-group instance where the η prediction serves as an upper bound.

11.9 Per-Individual Explanation Reversal [Tier B: Empirical]

The preceding analyses measure instability at the feature-pair level. A more directly actionable question: for a specific patient or loan applicant, does the “most important feature” explanation change across Rashomon models?

We test this with a rigorous ablation across 4 conditions (seed-only, +row subsample, +column subsample, +bootstrap with both subsamples), 3 model classes (XGBoost, LightGBM, Random Forest), and 3 datasets (Breast Cancer, German Credit, Heart Disease). Feature-level reversals are distinguished from *cluster-level* reversals, where clusters group features with $|r| > 0.8$ to avoid counting correlated-feature shuffling as meaningful disagreement.

Key findings. Under seed-only variation, reversal rates are 0% for XGBoost and LightGBM on all datasets: gradient-boosted trees are near-deterministic with fixed data. Under standard regularisation (`subsample = 0.8`, a common hyperparameter), the picture changes:

Dataset	XGB cluster rev.	LGB cluster rev.	RF cluster rev.
German Credit	45.3%	46.0%	35.0%
Breast Cancer	26.3%	21.1%	2.9%
Heart Disease	0.0%	0.0%	0.0%

For German Credit, 45% of loan applicants receive a genuinely different explanation *category* (not just a correlated variant) depending on the training configuration. This holds across model classes (XGBoost: 45%, LightGBM: 46%). The instability source is subsampling (a standard regularisation technique), not the random seed itself.

Narrative example. German Credit applicant #91 receives 6 different “most important feature” explanations across 8 Rashomon models: Purpose (25%), Telephone (25%), Installment rate (12%), Account status (12%), Dependents (12%), and Savings (12%). No feature dominates. An adverse action notice based on any single model is effectively arbitrary. The DASH consensus resolves to a stable explanation (92% agreement with majority vote).

The effect is dataset-dependent: Heart Disease shows 0% reversal across all conditions, confirming that instability requires sufficient feature dependence ($I > 0$ between feature pairs). The honest headline is not “75% of individuals receive different explanations” (an inflated aggregate from v1 of this experiment) but rather: *standard regularisation choices determine which explanation category nearly half of credit applicants receive.*

11.10 Prospective Out-of-Domain Prediction: Drug Discovery [Tier B: Empirical]

As a prospective test, we applied the framework to molecular property prediction—a domain it was never calibrated on. Using the BBBP dataset (blood-brain barrier penetration, 2,039 molecules, 1,024 Morgan fingerprint bits), we computed predictions *before* running the experiment.

The group identification bottleneck. Morgan fingerprint bits are sparse binary vectors with near-zero pairwise Pearson correlation. Clustering at $|r| > 0.8$ yields zero groups, predicting $\eta = 0$ and 0% flip rate. The actual flip rate is 23% across 24 Rashomon models (AUC = 0.885). **Pearson correlation fails completely** on binary features.

Mutual information and Jaccard recover the prediction. Replacing Pearson with mutual information (MI) or Jaccard co-occurrence for group discovery recovers the prediction:

Method	Threshold	Predicted flip	Error vs. observed (0.231)
Pearson ($ r $)	0.80	0.000	100% (failed)
Mutual information	0.05	0.194	16%
Jaccard co-occurrence	0.20	0.222	4%

MI captures nonlinear co-activation patterns that Pearson misses in binary data: fingerprint bits that fire on similar substructures (aromatic rings, hydrogen bond donors) share mutual information even when their Pearson correlation is low.

Structural prediction fails: no bimodal separation. While the overall flip rate magnitude is recovered by MI/Jaccard, the Noether bimodal structure does *not* transfer. Within-cluster flip rate (0.232) $\approx$ between-cluster flip rate (0.227), with a bimodal gap of only 0.005 (predicted: > 0.30). The instability in molecular fingerprints is *uniform* across all feature pairs, not concentrated within correlation groups.

Interpretation. This defines a precise boundary condition for the framework:

- **Discrete group structure** (attribution, causal, codon): the stable fact count theorem produces clean bimodal separation. Within-group $\approx 50\%$, between-group $\approx 0\%$.
- **Diffuse overlap structure** (molecular fingerprints): the capacity theorem predicts the *magnitude* of instability via MI/Jaccard (within 4%), but the structural prediction (which pairs are unstable) does not hold. The symmetry is not S_k but something more complex—overlapping substructure patterns that resist discrete decomposition.

The group identification method should match the feature type: Pearson for continuous correlated features, MI or Jaccard for binary/categorical features, and SAGE as a universal fallback. Extending the SAGE algorithm to non-linear dependencies is the natural next step.

11.11 Exhaustive Pipeline Validation: 85 PMLB Datasets [Tier B: Empirical]

To test the full diagnostic pipeline at scale, we applied it to all 85 eligible PMLB datasets (70 classification, 15 regression; $P = 5$ –180 features, $N = 200$ –5,000 samples) with $M = 50$ models per dataset and independent calibration/validation splits for the Gaussian formula. SAGE uses bidirectional data-split: groups discovered on half A and tested on half B, then reversed and averaged, eliminating split-direction bias.

Diagnostic tool	Where applicable	Success rate
DASH resolution	All Rashomon datasets	84/85 (98.8%) [CI: 94–100%]
SAGE directional (data-split)	All datasets	81/85 (95.3%) [CI: 89–98%]
Gaussian $\Phi(-\text{SNR})$, $R^2 > 0.80$ (OOS)	Shapiro–Wilk OK	29/49 (59%) [CI: 45–72%]
Gaussian $\Phi(-\text{SNR})$, $R^2 > 0.80$ (in-sample)	Shapiro–Wilk OK	44/49 (90%)
Correlation directional	≥ 2 groups discovered	25/33 (76%) [CI: 59–87%]

All 85 datasets exhibit the Rashomon property. The four pipeline tools stratify by reliability.

SAGE is the strongest directional tool: 81/85 (95.3%, CI: 89–98%) across all feature dimensions, including 15/15 at $P > 50$ —closing the high-dimensional gap. SAGE outperforms correlation grouping by McNemar test (7 discordant SAGE wins vs 1 correlation win, $p = 0.07$) and produces $7.6\times$ larger within–between gaps than random grouping ($p = 1.5 \times 10^{-22}$). The four SAGE failures have clear explanations: one dataset (**xd6**) has uniform instability (all pairwise flip rates in $[0.43, 0.51]$, no group structure exists); two have $P \leq 9$ features (insufficient pairs for stable clustering); and one (**spect**) has $N = 267$ with $P = 22$ (marginal sample size). At $P \geq 11$, SAGE with TreeSHAP has zero observed failures (15/15 unique datasets); at $P \leq 10$, the rate drops to 50% (3/6). TreeSHAP also rescues gain-based failures: **profb** (gap +0.33), both GAMETES epistasis datasets (+0.10, +0.04), and **spect** (+0.05) all succeed with TreeSHAP despite failing with gain-based importance. N -titration experiments confirm SAGE works at remarkably low sample sizes ($N/2 \geq 0.7P$ observed on **spambase**, $P = 57$; no sharp threshold exists). The earlier conservative estimate $N/2 > 5P$ substantially understated SAGE’s robustness. On a synthetic dataset with 4 known groups $\times$ 5 features, SAGE correctly identifies 4 groups (ARI = 0.39) and produces a *larger* directional gap (+0.37) than the true groups (+0.29), confirming that SAGE optimises for the measured property (flip-rate separation) rather than correlation structure. Cross-model validation confirms SAGE is model-agnostic. An expanded factorial design (30 datasets $\times$ 6 model/method configurations, $M = 30$, balanced across $P = 11$ –240 and including 5 regression datasets) gives 130/141 (92%, CI: 87–96%) with no significant model effect (Cochran’s $Q = 6.61$, $p = 0.25$, $n = 20$ complete datasets; χ^2 model effect $p = 0.68$, Cramér’s $V = 0.15$). SHAP vs non-SHAP: $p = 0.73$ (no method effect). Per model: XGBoost 95%, RF 94%, LR 96%, NN 80%. The NN gap is significant in a pairwise Fisher test ($p = 0.026$, OR = 0.22) and is attributable to noisy permutation importance at $P > 40$; with KernelSHAP, NNs achieve 7/7. Regression datasets: 30/30 (100%). Negative control (random grouping): $p = 8.7 \times 10^{-24}$. The theory’s model-agnosticity is confirmed empirically.

DASH preserves accuracy in 156/156 factorial tests (100%, CI: 98–100%) across all model types including regression. DASH reduces variance in 140/156 (90%); the 16 “failures” are datasets with zero flip rate (no Rashomon property — DASH is correctly inapplicable). In the 85-dataset bulletproof validation ($M = 50$): 84/85 preserves (98.8%), 56/85 strictly improves (66%). The one real failure (**parity5+5**, 10-class parity, 26% accuracy) is an effectively unsolvable task.

SHAP-specific validation. Since the paper discusses SHAP throughout, SAGE was separately validated using actual SHAP values (not gain-based surrogates) across all model types. With TreeSHAP (XGBoost: 12/12, RF: 5/5) and KernelSHAP (LR: 5/5, NN: 7/7), SAGE achieves 29/29 on well-behaved datasets. Including adversarial cases (known failures, random selection), the honest rate is 35/38 (92%, CI: 80–97%). Three failures—all at $P \leq 10$ —are: **xd6** (uniform instability, no group structure), **breast_cancer** (36 pairs insufficient for clustering), **parity5+5** (XOR structure SHAP doesn’t capture). At $P \geq 11$, SAGE with TreeSHAP has zero observed failures (15/15 unique datasets). Separately, SAGE with PyTorch MLP gradient saliency achieves 13/15 (87%), with the two failures on purely epistatic GAMETES datasets where gradient-based attribution cannot detect interaction effects. TreeSHAP rescues gain-based failures: **profb** (+0.33), both GAMETES epistasis datasets (+0.10, +0.04), and **spect** (+0.05) all succeed with TreeSHAP despite failing with gain-based importance.

The Gaussian flip formula requires careful interpretation. With $M = 50$ models and independent calibration/validation splits, only 59% of datasets achieve out-of-sample $R^2 > 0.80$. However, in-sample R^2 (using the same models for both SNR estimation and flip rate observation) exceeds 0.80 on 90%—a 31 percentage-point overfitting gap. This gap is explained by the M -sensitivity analysis (Phase 5): at $M = 15$, median OOS $R^2 = 0.45$; at $M = 25$, $R^2 = 0.70$; at $M = 50$, $R^2 = 0.82$. The formula is well-calibrated but requires $M \geq 50$ for reliable out-of-sample prediction. Earlier validation (Table 10, $M = 60$, OOS $R^2 = 0.848$) is consistent with this finding.

Coverage conflict scope: importance-type dependent. The coverage conflict degree (fraction of pairs with $\text{SNR} < 0.5$) is strongly importance-type dependent. With **TreeSHAP** (mean $|\text{SHAP}|$), coverage conflict predicts mean instability with $\rho = 0.87$ ($p = 2.2 \times 10^{-21}$, $n = 67$ datasets) across *all* feature dimensions— $\rho = 0.90$ for $P \leq 20$ and $\rho = 0.80$ for $P > 20$. With **gain-based** importance (**feature_importances_**), coverage conflict works only for $P \leq 20$ ($\rho = 0.90$) and reverses for $P > 20$ ($\rho = -0.31$).

The mechanism: gain-based importance is *sparse*—features never selected as split points receive exactly zero importance, creating artificial low-SNR pairs (both features at zero) that are counted as “coverage conflicts” but exhibit no actual instability. TreeSHAP assigns non-zero values to all features via the Shapley mechanism, eliminating this sparsity artifact. At $P > 20$, many features have gain = 0 in any given model, inflating the coverage conflict fraction without indicating genuine competition.

The practical recommendation: use TreeSHAP for coverage conflict computation (reliable at all P); if only gain-based importance is available, restrict to $P \leq 20$ or use SAGE directly. *Computational requirement:* coverage conflict with TreeSHAP requires $M \geq 50$ models and ≥ 100 background samples per SHAP computation. At lower settings ($M = 30$, $n_{\text{shap}} = 50$), SHAP estimation noise artificially inflates CC without corresponding increases in flip rate, breaking the correlation ($\rho = 0.15$ vs $\rho = 0.87$ at adequate settings).

Rashomon prevalence. Of the 85 datasets, all exhibit the Rashomon property (median inter-model rank correlation = 0.73; only 1/85 exceeds 0.95). The threshold used (< 0.99) is not too lenient: 74/85 datasets have mean flip rate > 0.05 (meaningful instability). The single dataset with rank correlation > 0.95 has mean flip = 0.01—instability is negligible and the pipeline correctly identifies it as such.

DASH preserves accuracy. A natural concern is that DASH’s variance reduction might degrade predictive accuracy. Across the full 85-dataset validation, DASH preserves accuracy (within one standard deviation) on 84/85 datasets (98.8%) and strictly improves in 56/85 (66%). On 3 verification datasets, DASH reduces the pairwise flip rate by 3–5 $\times$ (churn: 5 $\times$, ionosphere: 3 $\times$, waveform: 4 $\times$). Whether DASH strictly improves (vs. merely preserves) is stochastic:

Rashomon diversity ($p = 0.33$), mean flip rate ($p = 0.24$), and diversity $\times \log N$ ($p = 0.22$) all fail to predict improvement. The important guarantee is preservation, not improvement—DASH eliminates explanation variance without sacrificing predictive performance. This is consistent with the companion attribution paper [Caraker et al., 2026a], which finds that DASH nearly triples stability on Breast Cancer ($0.376 \rightarrow 0.925$) while maintaining competitive predictive performance (TOST-confirmed equivalence on California Housing; Wilcoxon $p = 0.926$ for accuracy), and that the operative mechanism is model independence rather than any particular aggregation strategy (DASH $\approx$ Stochastic Retrain, Cohen’s $d = +0.05$). At 11 additional datasets with $P > 50$ features, DASH preserves accuracy in 11/11 cases, confirming that the resolution scales to high-dimensional settings.

Gaussianity characterisation. Of the 33 Rashomon datasets, 13 (39%) fail the Shapiro–Wilk Gaussianity check. Non-Gaussian datasets tend to have larger sample sizes ($N = 2196 \pm 940$ vs 1492 ± 1026 ; Mann–Whitney $p = 0.049$, not significant after Bonferroni correction for 3 comparisons) and significantly lower mean flip rates (0.10 ± 0.06 vs 0.17 ± 0.08 ; $p = 0.009$). Crucially, SAGE directional prediction succeeds on 10/10 applicable non-Gaussian datasets—the same rate as on Gaussian ones (11/12). The Gaussian formula’s exclusion criterion (39% of datasets) does not limit the other pipeline tools.

The spectf anomaly. The sole SAGE failure (spectf, gap = -0.021) arises because $N = 349$ samples split in half leave too few observations for stable group discovery. On the full dataset, SAGE correctly identifies 27 groups with a strong within–between gap (-0.24 , within $<$ between). The data-split gap oscillates around zero across different splits (-0.089 and $+0.030$), confirming this is a sample-size effect rather than a structural failure.

High-dimensional extension ($P > 50$). On 11 additional PMLB datasets with 57–180 features, DASH preserves accuracy in 11/11 cases. However, correlation-based grouping breaks down: at $|r| > 0.8$, most features form singletons (groups $\approx P$), yielding negative or undefined directional gaps (0/4 applicable datasets succeed). The SAGE algorithm and lower correlation thresholds are needed for high-dimensional settings, consistent with the drug discovery findings (§11.10). The DASH resolution itself is dimension-independent and works without modification.

11.12 The Gaussian Flip Rate: From Exact to Approximate Symmetry [Tier B: Empirical]

The invariance counting theorem predicts exact bimodality when within-group features have identical true importance ($\Delta_{jk} = 0$). On real data, symmetry is approximate ($\Delta_{jk} \approx 0$ but nonzero), and the bimodal gap weakens. The Gaussian flip rate formula provides the continuous generalization.

The flip rate formula itself—the probability of sign disagreement between two Gaussian random variables—is a standard result in signal detection theory. The framework’s contribution is not the formula but the *structural identification*: the impossibility theorem and the biletma determine *which* feature pairs are subject to this formula (those within a symmetry orbit), while the Gaussian flip rate quantifies *how much* instability each pair exhibits.

For features j and k with importance difference $d^{(m)} = \text{imp}_j^{(m)} - \text{imp}_k^{(m)}$ across M bootstrap retrains, if $d^{(m)} \sim N(\Delta_{jk}, \sigma_\Delta^2)$, the probability that two independent models disagree on the ranking is:

$$\text{flip}(j, k) = 2 \Phi\left(\frac{\Delta_{jk}}{\sigma_\Delta}\right) \Phi\left(\frac{-\Delta_{jk}}{\sigma_\Delta}\right), \quad (2)$$

Dataset	P	Pairs	OOS R^2	Corr. BL R^2
Breast Cancer	30	435	0.792	0.002
Wine	13	78	0.793	-0.004
California Hous.	8	28	0.969	0.071
Diabetes (Pima)	8	28	0.868	0.013
Heart Disease	13	78	0.818	0.029
Mean		647	0.848	0.022

Table 10: Out-of-sample validation of the Gaussian flip rate formula. Calibration and validation model sets are independent (non-overlapping random seeds). The formula predicts per-pair TreeSHAP flip rates with mean $R^2 = 0.848$ and zero structural parameters. The correlation baseline predicts $R^2 < 0.08$ on every dataset.

where Φ is the standard normal CDF (Lean-verified in `GaussianFlipRate.lean`). When $\Delta = 0$ (exact symmetry), $\text{flip} = 0.5$ — recovering the invariance counting within-group coin flip. When $|\Delta| \gg \sigma$ (strong signal), $\text{flip} \approx 0$ — recovering between-group stability.

A sensitivity analysis shows that $M \geq 30$ calibration models (15 per split) are needed for the Gaussian flip formula to achieve $R^2 > 0.7$; below $M = 20$, the formula degrades rapidly.

We validated this formula with *independent* calibration and validation model sets (30 models each, non-overlapping random seeds) on five real datasets (Table 10). The formula achieves a mean out-of-sample $R^2 = 0.848$ across 647 feature pairs, with all five datasets exceeding $R^2 = 0.79$. The correlation baseline (predicting flip rate from pairwise feature correlation) achieves $R^2 < 0.08$ on every dataset, confirming that the Gaussian formula provides 14–400 $\times$ lift over correlation-based prediction.

The formula generalizes across model classes. On three datasets (Breast Cancer, Wine, California Housing) with three model classes (XGBoost, Random Forest, Ridge regression), 8 of 9 combinations achieve out-of-sample $R^2 > 0.79$. Random Forest shows the highest R^2 on two datasets (0.86 Breast Cancer, 0.92 Wine). The one failure—Ridge on California Housing ($R^2 = -1.96$)—occurs because Ridge regression on this dataset produces near-deterministic coefficients with negligible variance, violating the Gaussian assumption. The formula is expected to fail when a single feature dominates so strongly that importance differences have effectively zero noise.

One combination—Ridge regression on California Housing—produces out-of-sample $R^2 = -1.96$. Ridge regression on this dataset assigns nearly deterministic coefficients (one feature, median income, dominates overwhelmingly), leaving negligible variance in the importance differences. The Gaussian assumption requires meaningful sampling variance; when $\sigma_\Delta \approx 0$, the formula degenerates. This boundary condition is expected: the formula predicts flip rates for stochastic model classes, not for deterministic or near-deterministic ones. More generally, the Gaussian assumption fails when importance differences have near-zero variance (as in the Ridge/California Housing case, producing $R^2 = -1.96$) or when the importance distribution is multimodal (ensembles of qualitatively different model types). A prospective diagnostic: if the coefficient of variation of σ_Δ across feature pairs is below 0.1, the importance landscape is too concentrated for the Gaussian model.

Practically, the signal-to-noise ratio $\text{SNR}_{jk} = |\Delta_{jk}|/\sigma_\Delta$ provides a simple decision rule: pairs with $\text{SNR} > 2$ are reliable (flip rate $< 5\%$), pairs with $\text{SNR} < 0.5$ are unreliable (flip rate $> 30\%$), and intermediate values require uncertainty quantification.

Robustness to non-Gaussianity. The $\Phi(-\text{SNR})$ formula is exact under Gaussian noise and well-calibrated on real data ($R^2 = 0.814$ [0.791, 0.835] across 647 feature pairs from five clinical and financial datasets, with bootstrap 95% CIs; per-dataset $R^2 = 0.946$ –0.980 on six

datasets including PMLB benchmarks, confirming consistency across domains). The formula systematically underpredicts observed flip rates by $\sim 17\%$ (calibration slope 1.17), making it conservative. Under heavy-tailed distributions (Student t , $\text{df} = 3$), the formula underestimates flip rates by $40\%+$; under log-normal noise, errors are larger still. However, Shapiro–Wilk tests confirm 94.7% of real SHAP difference distributions are Gaussian ($p > 0.10$), justifying the formula’s applicability in practice. The Berry–Esseen correction adds $< 1\%$ to the base rate for $n > 100$ models under light tails.

11.13 Domain-Specific Validation [Tier B: Empirical]

The Gaussian flip formula was derived and validated on clinical tabular data (Section 11.12). This subsection extends the validation to gene expression data, high-dimensional settings, clinical/financial audit datasets, and the enrichment tradeoff predicted by the bitemma.

Gene expression reliability. On gene expression data (AP_Breast_Lung, 470 samples, 10,935 genes reduced to the top 50 by variance), 70.5% of gene-pair importance comparisons are unreliable ($\text{SNR} < 0.5$, flip rate $> 30\%$) and only 4.0% are reliable ($\text{SNR} > 2$, flip rate $< 5\%$). The formula achieves out-of-sample $R^2 = 0.416$ on this dataset—substantially lower than on clinical tabular data ($R^2 = 0.52\text{--}0.97$), reflecting the higher collinearity and dimensionality of expression profiles. Gain-based importance was used; TreeSHAP would likely improve accuracy based on the 60% improvement observed at $P = 50$ (see below).

Gene importance rankings used to prioritize candidates for experimental validation may contain unreliable comparisons that the Gaussian flip formula can identify before resources are committed.

SHAP vs. gain at high dimensions. At $P = 50$ features (synthetic, 20 informative, 15 redundant), TreeSHAP achieves out-of-sample $R^2 = 0.679$ versus gain-based $R^2 = 0.424$ —a 60% improvement. TreeSHAP also produces more reliable rankings: 44.6% of pairs have $\text{SNR} > 2$ (reliable) with SHAP versus 15.3% with gain. The formula’s accuracy depends on the importance metric; SHAP values, which have lower variance across retrains, produce higher SNR and tighter predictions.

Clinical and financial reliability audit. Across three clinical and financial datasets (Breast Cancer, Heart Disease, German Credit), 57% of feature importance comparisons exhibit flip rates exceeding 30% , as predicted by the Gaussian flip formula (Pearson $r = 0.76\text{--}0.85$). Nearly zero comparisons are reliable ($\text{SNR} > 2$) on Breast Cancer (0%) and Heart Disease (1.3%); German Credit has 13.7% reliable pairs, reflecting its stronger feature hierarchy.

Gain-based vs. TreeSHAP importance. The coverage conflict bridge ($\rho = 0.961$, 15 datasets) and regulatory compliance metrics (83% gap) use gain-based feature importance (`feature_importances_`). The clinical top-5 audit (28% unreliable) uses TreeSHAP (mean $|\text{SHAP}|$). TreeSHAP produces lower R^2 (mean 0.707 vs ~ 0.85 gain-based) but comparable rank correlation (mean $\rho = 0.858$). The formula’s per-pair rank predictions are robust to the importance type; absolute flip rate calibration is better with gain-based importance. For clinical applications where TreeSHAP is standard, the formula correctly identifies which comparisons are reliable (rank accuracy), even if the predicted flip magnitude is less precisely calibrated.

TreeSHAP clinical audit (top-5 features). A targeted audit using TreeSHAP (mean $|\text{SHAP}|$ across 100 test points, 30 calibration + 30 validation XGBoost models) reveals that 28% of top-5 feature ranking comparisons are structurally unreliable across five clinical and

financial datasets. On Breast Cancer (Wisconsin), the comparison between the two most important features—worst perimeter vs. worst concave points—is a coin flip (SNR = 0.17, flip rate = 43%). On Adult Income, 60% of top-5 comparisons are unreliable (the top four sociodemographic features are virtually interchangeable in importance). On Diabetes (Pima), 90% of top-5 comparisons are reliable (glucose is clearly dominant). The Gaussian flip formula achieves mean $\rho = 0.858$ (rank correlation) and mean $R^2 = 0.707$ across the five TreeSHAP datasets. Most top-5 rankings *are* stable; the contribution is identifying the 28% that are not, which practitioners currently have no way to diagnose.

High-dimensional test ($P = 200$, gain-based). At $P = 200$ features (synthetic, 50 informative, 50 redundant), the formula achieves out-of-sample $R^2 = 0.532$ with gain-based importance across 19,900 feature pairs. Of these, 49.2% are unreliable (SNR < 0.5) and 12.2% are reliable (SNR > 2).

Enrichment on real data (Breast Cancer). The enrichment tradeoff, predicted by the bitemma, is confirmed on real clinical data (Breast Cancer, 30 features, 100 XGBoost models). At full decisiveness (threshold = 0), stability is 0.675 (flip rate 32.5%). At threshold = 0.5 (stability 0.991), decisiveness collapses to 0.088. The sharp phase transition between thresholds 0.10 and 0.15—where stability jumps from 0.79 to 0.95 while decisiveness drops to 16%—confirms that the (1,1) corner of the decisiveness–stability plane is unreachable on real data, as the bitemma predicts.

Combined real-data-only reliability summary. Across seven real clinical, financial, and genomics datasets (Breast Cancer, Wine, California Housing, Diabetes, Heart Disease, German Credit, AP_Breast_Lung), encompassing 2,062 feature-pair comparisons, 62% are unreliable (SNR < 0.5) and only 4% are reliable (SNR > 2). On clinical and financial data alone (703 pairs), 47% are unreliable. The mean out-of-sample R^2 across all real datasets is 0.67.

Multiverse analysis (Breast Cancer). Across a 30-pipeline multiverse (3 model classes $\times$ 3 scalars $\times$ 3 feature selections + 3 bootstrap extras), 95% of feature importance claims are non-replicable (Jaccard similarity = 0.39). The Gaussian flip formula achieves Spearman $\rho = 0.78$ across the full multiverse (correctly ranking which comparisons are most unreliable) but $R^2 = -0.39$ for predicted flip rates (the absolute predictions fail when mixing model classes with different importance scales). Within model classes, the formula performs better: Ridge $R^2 = 0.49$, XGBoost $R^2 = 0.33$. The 95% non-replicable figure is the most dramatic but also the least scientifically meaningful, as it compares tree-based and linear models that measure fundamentally different quantities (gain-based vs. coefficient-based importance). The within-class figures—75% non-replicable for XGBoost and Random Forest, 100% for Ridge—are the more appropriate adversarial baseline.

Generalization gap correlation. Across the original 5 datasets, the cross-dataset correlation between mean flip rate and generalization gap (ID accuracy – OOD accuracy) is Pearson $r = -0.91$, $p = 0.034$, $n = 5$ —suggestive of an inverse relationship (more instability correlates with smaller generalization gaps). However, leave-one-out analysis reveals this result is driven by a single outlier: removing California Housing (which has both the lowest instability and the largest geographic distribution shift) reduces the correlation to $r = -0.20$, $p = 0.80$. Expanding to 13 datasets (8 real + 5 synthetic controls with varying feature redundancy) yields $r = -0.58$, $p = 0.038$, but removing California Housing from the full set reverses the sign to $r = +0.26$. The synthetic control with varying redundancy (0–40 redundant features) shows the *opposite* pattern ($r = +0.80$): more redundancy increases both instability and generalization gap, contrary to the initial hypothesis. The finding requires validation on larger, more diverse samples before causal

interpretation; we report it as a null result with respect to the specific directional prediction. Breast Cancer shows a negative generalization gap (OOD accuracy exceeds ID accuracy), likely because the ‘extreme’ patients have more distinctive features that are easier to classify.

11.14 Gene Expression Pathway Divergence [Tier B: Empirical]

The most consequential instance of the impossibility arises in biomarker discovery, where the “top gene” identified by a machine learning pipeline determines which therapeutic target enters experimental validation.

TSPAN8 vs CEACAM5 on colon–kidney discrimination. Using the AP_Colon_Kidney microarray dataset (546 samples, 10,935 genes reduced to top 50 by variance, OpenML), we trained 50 XGBoost classifiers with standard regularisation (subsample = 0.8, colsample_bytree = 0.5, seeds 0–49) and computed TreeSHAP importances. The #1 biomarker alternates between two genes:

- **TSPAN8** (probe 203824_at): tetraspanin-enriched microdomains, integrin signaling, invasion/metastasis — top in 92% of seeds.
- **CEACAM5/CEA** (probe 201884_at): GPI-anchored cell adhesion, immune evasion, classical clinical colon cancer marker — top in 6% of seeds.

The two genes’ SHAP values correlate at $\rho = 0.858$ across patients (both are colon-enriched tissue markers), placing them in the same approximate symmetry orbit. Under deterministic training (subsample = 1.0), a single ranking emerges with 100% agreement—confirming that the alternation is driven by the stochastic regularisation, not by data noise.

Pathway analysis (GO enrichment). Gene Ontology annotation reveals zero overlap in Biological Process terms between TSPAN8 (negative regulation of blood coagulation; integrin binding) and CEACAM5 (homophilic cell–cell adhesion, six BP terms; GPI anchor binding; apoptosis regulation). A biologist following the first model into the wet lab would pursue invasion/metastasis signaling; a biologist following the second would pursue immune evasion. The pathway-level conclusion (colon identity) is preserved; the gene-level biomarker target—the actionable output of a discovery pipeline—differs.

Replication across datasets. The alternation replicates on AP_Endometrium_Colon (363 samples, independent tissue contrast): TSPAN8 is #1 in 78% of seeds, CEACAM5 in 22%, with $\rho = 0.781$. On AP_Breast_Colon (1,545 samples), a different pair alternates: COL3A1 (collagen type III, 72%) and COL1A1 (collagen type I, 28%), with $\rho = 0.843$. Both collagens are extracellular matrix proteins sharing 3 GO Biological Process terms—this is the same-pathway, redundant-marker mode (Mode 2), contrasting with the distinct-pathway mode of TSPAN8/CEACAM5 (Mode 1). AP_Breast_Lung serves as a positive control: SFTPB (surfactant protein B, a dominant lung marker) is #1 in 98% of seeds, confirming that large importance gaps protect stability.

Two modes of alternation. *Mode 1* (TSPAN8/CEACAM5): distinct biological pathways compete for the top rank. The specific drug target depends on the random seed. *Mode 2* (COL3A1/COL1A1): redundant markers within the same pathway alternate. The pathway conclusion is stable; the molecular identity is not. Both modes are instances of the impossibility theorem: the Rashomon property (observationally equivalent models) produces incompatible explanations (different top genes). Mode 1 has greater scientific consequence because it affects therapeutic strategy, not just molecular identity.

DASH resolves. The DASH ensemble average (mean SHAP across 50 models) places both TSPAN8 and CEACAM5 in the top-2, correctly reporting that both contribute without privileging either. This is the G -invariant resolution in practice: the group mean preserves the between-group ranking (colon genes > kidney genes) while acknowledging within-group indistinguishability.

11.15 Mechanistic Interpretability: Circuit-Level Rashomon [Tier B: Empirical]

Mechanistic interpretability seeks to reverse-engineer neural networks into human-readable circuits. Prior work has shown qualitatively that networks trained from different random initialisations develop different internal representations and circuits [Chughtai et al., 2023]. The impossibility theorem predicts that circuit-level explanations are subject to the same instability as feature attributions when the network architecture admits permutation symmetries; here we quantify the disagreement and resolve it via the G -invariant projection.

Experimental design. We trained 10 two-layer transformers (4 heads per layer, $d_{\text{model}} = 128$) from independent random initialisations on modular addition $(a + b) \bmod 113$. All 10 models “grokked” to 100% test accuracy after 50,000 training steps with weight decay 1.0. Circuit importance was measured via activation patching: for each of 10 components (8 attention heads + 2 MLPs), we replaced the component’s output with its mean activation and measured the resulting accuracy drop. The within-model determinism control confirmed near-perfect reproducibility ($r = 0.9998$).

Controls: random, memorised, grokked. Three training checkpoints establish the control hierarchy:

Checkpoint	Train acc	Test acc	Mean ρ between models
Step 50 (random init)	0.2%	0.2%	0.300
Step 500 (memorised)	100%	50%	0.668
Step 50k (grokked)	100%	100%	0.518

Grokked models agree more than random ($p = 2.0 \times 10^{-4}$, $d = 0.76$) but *less* than memorised models ($p = 0.023$, $d = 0.58$). The decrease from memorised to grokked is consistent with grokking discovering genuinely different algorithms: memorised models share the same “lookup table” strategy, while grokked models find distinct Fourier-based algorithms.

Different Fourier algorithms. All 10 grokked models represent modular addition via Fourier components, but the specific frequencies selected are nearly disjoint: the mean pairwise Jaccard similarity of top-5 Fourier frequencies is 0.022. Despite computing the same function with identical accuracy, each model discovers a different algorithm.

G -invariant resolution demonstrated. The architecture defines a symmetry group $S_4 \times S_4$ (permutations of the 4 heads within each layer). Projecting the 10-dimensional importance vector onto the 4-dimensional G -invariant subspace V^G (mean layer-0 heads, mean layer-1 heads, MLP0, MLP1) dramatically recovers agreement:

Representation	Spearman ρ	Pearson r
Full 10-dim (individual components)	0.518	0.835
G -invariant 4-dim (V^G)	0.929	0.967
Excluding MLP1 (3-dim)	0.822	0.782
Excluding MLP1 (9-dim)	0.337	0.439

The G -invariant projection is NOT trivially MLP-driven: excluding MLP1 from the invariant subspace still yields $\rho = 0.822$ ($p < 10^{-6}$), demonstrating that the layer-mean head importance is itself a stable, informative quantity.

stable fact count within ML. The symmetry $S_4 \times S_4$ predicts: within-layer head pairs (12 pairs total) should flip at $\sim 50\%$; head-vs-MLP pairs (cross-group, 8 pairs) should be stable. Observed: within-layer mean flip rate = 0.500 exactly; between-group mean flip rate = 0.227. Bimodal gap = 0.273, Mann–Whitney $p = 2.4 \times 10^{-5}$. This is a second instantiation of the stable fact count prediction within a single experiment.

Structural universal: MLP1 dominance. MLP1 is the most important component in every seed (coefficient of variation = 0.027). This is the G -invariant structural fact—the between-group conclusion that survives the impossibility. Individual head identities (which specific head matters) are the within-group noise that the theorem proves is unavoidable.

Fresh-ensemble replication and the limit of from-data recovery. An independent re-run (10 one-layer transformers, $(a + b) \bmod 53$, all grokked to $\geq 99\%$ test accuracy) probes the finer *MLP-neuron* permutation symmetry (S_n over the 512 hidden neurons) rather than head permutation. It reproduces the split starkly: raw index-aligned per-neuron importance is unidentifiable across seeds (cross-model Pearson 0.002, Spearman -0.002 , at the random-permutation null), while the S_n -invariant importance spectrum is identifiable (sorted-value Pearson 0.975, cosine 0.976; participation ratio consistent, coefficient of variation 0.20). Crucially, it also tests—and bounds—the *from-data* bridge: recovering the neuron correspondence (the symmetry element relating two models) from data alone, by Hungarian matching on activation fingerprints, essentially fails (post-match Pearson 0.085, barely above the null). The grokked ensemble occupies genuinely *distinct* algorithmic solutions (different learned frequencies), not permutations of a single one; so on real networks the identifiable invariant is the *coarse* spectrum, and the fine per-neuron correspondence is neither identifiable nor recoverable from data. This is the honest ceiling of the interpretability bite: the invariant/non-invariant split holds, but “recover the symmetry from the ensemble” succeeds only when the ensemble is a single orbit—which grokked ensembles need not be.

11.16 Language Model Circuit Stability: Multi-Scale Validation [Tier B: Empirical]

The modular-addition experiment (Section 11.15) uses an algorithmic task. To test whether the theorem governs circuits in *real language models*, we trained 10 transformers from independent random initialisations on TinyStories Eldan and Li [2023], a dataset of short children’s stories, and repeated the experiment at two architectural scales. A third configuration (GPT-2 Small fine-tuned) tests the predicted boundary condition: fine-tuning from a shared pretrained checkpoint should NOT produce circuit instability, because the Rashomon set is small.

Experimental design.

Config	Architecture	Params	Steps	Components	$\dim(V^G)$	η
A	4L/4H, $d = 256$	16M	15,000	20	8	0.600
B	6L/8H, $d = 512$	45M	15,000	54	12	0.778
C	GPT-2 Small fine-tuned	124M	3,000	156	24	0.846

Each configuration trains 10 models from independent random seeds on the same TinyStories data. Circuit importance is measured via weight zeroing (Q/K/V for each head, fc1 for each MLP) and the resulting perplexity increase on held-out data. Config C starts from the pretrained

GPT-2 Small checkpoint and fine-tunes with different random seeds (data ordering and dropout). Determinism control: same model measured twice gives $r = 1.000$ (Configs A and B) and $r = 1.000$ (Config C).

Pre-registered predictions. Seven predictions were stated before any training data was generated: (1) perplexity CV $< 5\%$; (2) full-component agreement $\rho < 0.70$; (3) G -invariant agreement $\rho > 0.80$; (4) heads-only lift > 0.15 (not MLP-driven); (5) within-layer head flip rate > 0.40 ; (6) head-vs-MLP flip rate < 0.15 ; (7) permutation test $p < 0.01$.

Results: Configs A and B (training from scratch).

Metric	Config A (4L/4H)	Config B (6L/8H)
Perplexity (mean $\pm$ CV)	$8.6 \pm 0.7\%$	$10.0 \pm 1.0\%$
Split-half reliability (ρ)	0.991	0.960
Full agreement (ρ [95% CI])	0.565 [0.53, 0.60]	0.540 [0.52, 0.56]
G -invariant (ρ [95% CI])	0.972 [0.97, 0.98]	0.982 [0.98, 0.99]
Heads raw (ρ)	0.291	0.379
Heads averaged (ρ)	0.849	0.931
Head lift	0.558	0.553
Within-layer flip	0.496	0.489
Head-vs-MLP flip	0.000	0.000
Mann–Whitney p (pair-aggregated)	1.2×10^{-18}	1.3×10^{-18}
Cohen’s d (pair-aggregated)	5.4	11.9
Permutation test p	< 0.001	< 0.001
Pre-registered	7/7 PASS	7/7 PASS

All 14 pre-registered predictions pass across both scales. The within-layer head flip rates (0.496 and 0.489) are within sampling noise of the predicted 0.500. The head-vs-MLP flip rate is exactly 0.000 across all 30 models (720 and 2,160 comparisons respectively)—every model agrees that MLPs are more important than attention heads at each layer. The split-half reliability ($\rho = 0.991$ and 0.960) far exceeds the between-model agreement ($\rho = 0.565$ and 0.540), confirming that the disagreement is genuine Rashomon, not measurement noise.

The heads-only random projection control is critical: random $n_{\text{heads}} \rightarrow n_{\text{layers}}$ projections give $\rho = 0.300 \pm 0.223$ (Config A) and $\rho = 0.407 \pm 0.182$ (Config B), while the symmetry-based mean-head projection gives $\rho = 0.849$ and 0.931 (100th percentile in both cases). The lift is from the specific within-layer symmetry structure, not dimensionality reduction.

For Config B, the full-vector random projection ($\rho = 0.988$) outperforms the G -invariant ($\rho = 0.982$) because MLP0 dominates the importance landscape (~ 6.7 vs ~ 0.4 for all other components); any random projection preserving MLP0 achieves high agreement. The heads-only control isolates the symmetry-specific lift and remains at the 100th percentile.

In Config A, all four heads in layer 0 appear as the “most important head” across the 10 seeds (H0, H1, H2, H3 each take the top spot), fully realising the S_4 symmetry. The role “important attention head in layer 0” is stable; which head fills that role is a coin flip.

Results: Config C (GPT-2 fine-tuned — boundary condition). Config C produces $\rho = 0.993$ (full agreement), within-layer flip = 0.043, and head lift = 0.003. Three of seven predictions fail: full $\rho < 0.70$, within-flip > 0.40 , and head lift > 0.15 . This is the predicted boundary condition: fine-tuning from a shared pretrained checkpoint does not create the Rashomon property. The pretrained circuit structure is preserved — the same five components (MLP0, MLP2, MLP1, MLP11, MLP8) are the top 5 in every seed.

This confirms the theorem’s precondition: the impossibility predicts instability *if and only if* the Rashomon property holds. When it does not (Config C), the theorem correctly predicts no instability. The split-half reliability ($\rho = 0.878$) is *lower* than the between-model agreement

($\rho = 0.993$), meaning fine-tuned models are more similar to each other than the same model measured on different data subsets. This is additional evidence that fine-tuning genuinely preserves circuits.

Head-vs-MLP flip = 0.000 holds even for Config C: the layer-level MLP hierarchy is a structural invariant that persists regardless of whether the Rashomon property holds.

Cross-scale comparison.

Config	η	Full ρ	G -inv ρ	Lift	W-flip	B-flip
A (4L/4H)	0.600	0.565	0.972	0.408	0.496	0.000
B (6L/8H)	0.778	0.540	0.982	0.441	0.489	0.000
C (GPT-2 ft)	0.846	0.993	0.999	0.006	0.043	0.000

The two from-scratch configurations are ordered consistent with the capacity theorem ($\eta_A < \eta_B$, $\rho_A > \rho_B$: full agreement $0.565 \rightarrow 0.540$, G -invariant agreement $0.972 \rightarrow 0.982$), but with only $n = 2$ valid data points this is descriptive, not inferential. Config C does not participate because it lacks the Rashomon property. The within-layer flip rate is indistinguishable from 0.500 for both from-scratch configs and near zero for the fine-tuned config—a clean separation governed by the Rashomon precondition.

Ablation method robustness. To test whether the results depend on the choice of ablation method, we repeated the importance measurement using mean ablation (replacing each component’s output with its mean activation) on the cached Config A and B models. The layer-level structure is method-invariant: both weight zeroing and mean ablation produce G -invariant agreement $\rho \approx 0.97$ – 0.99 (G -invariant $\rho = 0.975$ and 0.987 for mean ablation vs 0.972 and 0.982 for weight zeroing). The cross-method per-model agreement is moderate ($\rho = 0.59$ for Config A, $\rho = 0.50$ for Config B). For $n = 20$ components, the null expectation under independence is $|\rho| \approx 0.23$, so the observed $\rho \approx 0.5$ represents significant positive agreement at the head level ($p < 0.05$) combined with near-perfect agreement at the layer level ($\rho \approx 0.97$ – 0.99 on V^G). The two methods agree on the importance structure more than chance, but disagree on the specific head-level ranking—and the G -invariant resolution captures exactly the content on which they agree.

Limitations and open questions. TinyStories is a synthetic dataset of simple children’s stories; extrapolation to frontier language models (GPT-4-scale) trained from scratch remains open. The cross-scale η rank-order test has only $n = 2$ valid data points (Configs A and B) and cannot be tested statistically.

Several questions raised by the component-level analysis warrant future work. First, *functional stability*: we show that WHICH head is most important varies across seeds, but not whether the FUNCTION computed by that head varies. If the same attention pattern emerges in different head positions across seeds (a permutation of functional roles), the instability is in labelling rather than in algorithmic diversity. Permutation alignment methods such as Git Re-Basin Ainsworth et al. [2023] address this directly and represent a complementary resolution to the G -invariant projection: Re-Basin aligns models in weight space, while V^G projects into the shared importance structure. Whether these produce equivalent conclusions is an empirical question.

Second, *sparse autoencoders (SAEs)*: the mechanistic interpretability community increasingly analyses learned features from SAEs rather than individual attention heads. If independently trained SAEs on the same model produce stable features, SAE-based attribution may escape the Rashomon property — the symmetry group would be trivial ($G = \{e\}$) and $\eta = 0$. We test this directly: on freshly grokked transformers, six SAEs trained independently (different seeds) on the *same* model’s activations reconstruct well ($R^2 = 0.99$) yet produce features

that are largely *not* reproducible across seeds — only 0.4% of dictionary features match above cosine 0.9, and 16% above 0.7 (random-direction null 0.15). So SAE features do *not* escape the Rashomon property: the symmetry group is nontrivial and $\eta > 0$; only a small stable subset of features is identifiable, and the remainder is training-seed noise. This is the theorem’s prediction, and a cautionary result for SAE-based interpretability at scale. The result holds on a *real* pretrained language model, and there the per-query certificate turns the caution into a usable filter: on GPT-2 small (residual stream at layer 6; an ensemble of six SAEs, $R^2 = 0.54$, $L_0 \approx 41$), features again only partially reproduce across seeds (mean best-match cosine 0.36 vs a random-direction null of 0.12), and the certificate cleanly isolates the reproducible minority—the $\approx 9.5\%$ of features it certifies STABLE transfer to a held-out SAE ensemble at cosine 0.81 (88% at cosine ≥ 0.7), whereas uncertified features sit near the ensemble average (0.32). Out-of-sample, on a real LLM, the certificate separates the SAE-feature claims worth trusting from the seed-dependent ones.

Third, the *ablation methodology* uses weight zeroing throughout. Path patching Goldowsky-Dill et al. [2023] and causal scrubbing offer finer-grained, input-dependent attribution that may break the within-layer symmetry for specific inputs. The impossibility theorem applies to the population-level importance ranking (averaged over inputs); input-specific attribution may exhibit different stability properties.

Fourth, the bootstrap 95% CIs on pairwise Spearman ρ are computed over the 45 model pairs, which share models. The effective degrees of freedom are closer to $n = 10$ (number of models) than $n = 45$ (number of pairs). The CIs may therefore be too narrow; the point estimates and qualitative conclusions are unaffected.

11.17 Model-Class Universality of the Impossibility [Tier B: Empirical]

The impossibility is not specific to a single model class. A rigorous comparison across four model classes (XGBoost, Random Forest, Ridge, LASSO) on three datasets (Breast Cancer, California Housing, Diabetes) reveals a two-level structure:

Within-family agreement. Tree-based models (XGBoost, Random Forest) agree on which features are unstable: seed-variation Spearman $\rho = 0.79$ – 0.92 on Breast Cancer and Diabetes, exceeding the permutation null. Linear models (Ridge, LASSO) agree even more strongly ($\rho = 0.94$ on Breast Cancer). Models within the same algorithmic family identify the same instability patterns.

Cross-family disagreement. Linear–tree cross-family agreement is NOT significant ($\rho = 0.19$ – 0.57 , none exceeding permutation null). Which specific features are unstable depends on the model class—trees and linear models disagree on the identity of unstable features, even when both detect instability.

Universal diagnostics. Despite cross-family disagreement on specifics, two diagnostics are universal:

- *Coverage conflict:* Spearman $\rho = 0.82$ (mean across all 4 classes, excluding degenerate California Housing linear case where 7/8 features have zero flip rate).
- *Var[SHAP] = DASH MSE:* exact for all 4 model classes, 0/12 violations. The variance of SHAP values across the Rashomon set equals the mean-squared error of the DASH ensemble average—a mathematical identity that holds regardless of model architecture.

Interpretation. Instability has a data-determined core (detectable by coverage conflict, shared across model families) and a model-family-specific pattern (which unstable features flip depends on the algorithm). The impossibility theorem explains both: the Rashomon property (data-determined) guarantees instability exists; the specific symmetry group (model-family-specific) determines which features are interchangeable.

11.18 Multi-Analyst Disagreement as Impossibility: Neuroimaging [Tier B: Empirical]

The impossibility theorem applies beyond ML explanation methods to any setting where multiple valid analyses of the same data produce incompatible conclusions. We demonstrate this on the Botvinik-Nezer et al. (2020) neuroimaging study, where 70 independent teams analysed the same fMRI dataset and reached dramatically different conclusions about which brain regions were significantly activated.

The Rashomon property. The 70 analysis pipelines constitute the Rashomon set: they share the same observable (fMRI data from 108 participants on a mixed-gambles task) but produce incompatible explanations (different activation maps). This is the Rashomon property by definition; the impossibility theorem applies immediately.

Data and methods. Overlap maps (proportion of teams finding significance at each voxel) were downloaded from NeuroVault collection 6047 for 9 hypotheses (7 non-empty after excluding hypotheses 6 and 9). Individual team unthresholded t -statistic maps were downloaded from each team’s NeuroVault collection; 48 of 70 teams had accessible maps. All maps were parcellated using the Schaefer 100-parcel atlas (7 Yeo resting-state networks) via Nilearn. Team software metadata (12 FSL, 8 SPM, 4 AFNI, 24 unknown) was obtained from the `narps_open_pipelines` repository.

The Noether structural prediction. The framework predicts that brain regions sharing a functional network (the approximate symmetry group) should have more similar disagreement patterns than regions in different networks. Network labels from the Yeo et al. (2011) resting-state parcellation are *independent* of the NARPS task data.

Using overlap-based disagreement (the metric $1 - \max(\text{overlap}, 1 - \text{overlap})$ per region per hypothesis), within-network region pairs show more similar disagreement patterns than between-network pairs (raw Cohen’s $d = 0.43$, permutation $p < 0.001$, Mann–Whitney $p = 4.5 \times 10^{-32}$, consistent across all 7 non-empty hypotheses). The effect survives three controls:

- *Random grouping null:* Random assignments of matching group sizes yield $d = 0.001 \pm 0.050$ (observed is $4.9\times$ the 95th percentile of the null).
- *Spatial autocorrelation:* k -means clustering on region centroids gives $d = 0.07\text{--}0.19$ across $k = 5, 7, 10, 14$ —functional networks outperform all purely spatial groupings.
- *Activation-level control:* After regressing out activation similarity ($|\Delta\text{overlap}|$) from disagreement similarity, the network effect persists ($d = 0.32$, $p = 8.1 \times 10^{-19}$, bootstrap CI $[0.006, 0.008]$ excludes zero). Activation similarity explains $R^2 = 0.67$ of overall disagreement variance; the network effect is the remaining 26% reduction in d (from 0.43 to 0.32).

Network identity (which network, regardless of size) explains $\eta^2 = 0.35$ of disagreement variance (Kruskal–Wallis $p = 5.7 \times 10^{-7}$). The Dorsal Attention and Control networks show the highest disagreement (mean = 0.095); Limbic and Default show the lowest (mean = 0.023–0.046)—consistent with the mixed-gambles task differentially engaging attention and control networks near the significance threshold.

Methodological progression. The initial analysis (v1) computed inter-region correlation from the overlap maps themselves and tested whether this predicted disagreement ($\rho = 0.310$, $p = 0.002$). Adversarial review identified this as *circular*: the predictor was derived from the outcome data. Version 2 replaced the circular predictor with independent Yeo network labels (from resting-state fMRI, a separate dataset and modality), which eliminated the circularity and revealed that the capacity theorem fails while the Noether structural prediction succeeds. Version 3 added comprehensive controls (random grouping null, spatial autocorrelation via k -means, per-hypothesis consistency, Schaefer-400 replication). The definitive analysis added activation-level regression control. This progression—from a promising but flawed initial result to a rigorously controlled finding—illustrates why adversarial self-auditing is essential for cross-domain applications where the mapping from theory to data is analogical.

Bilateral symmetry ($\mathbb{Z}_2$). The strongest approximate symmetry in the brain is bilateral: left and right hemispheric homologs share functional roles. Matching 35 left–right homolog pairs by Schaefer label, the disagreement profiles correlate at $r = 0.66$ ($p = 1.5 \times 10^{-5}$), with homolog pairs showing 42% smaller disagreement differences than random pairs. No systematic left–right bias was detected (Wilcoxon $p = 0.49$). The effect is consistent across all 7 hypotheses (per-hypothesis $r = 0.44$ – 0.67). However, this result did not survive bootstrap activation control (parametric partial $r = 0.49$, $p = 0.001$, but bootstrap CI $[-0.54, 0.41]$ spans zero due to instability with $n = 35$ pairs and residualisation).

The capacity theorem does not apply. The quantitative prediction $\eta = (k - 1)/k$ fails at all tested granularities: network size does not predict disagreement level (Spearman $\rho = -0.021$, $p = 0.84$ at Schaefer-100 with 7 networks; $\rho = -0.015$, $p = 0.88$ with 14 hemisphere $\times$ network groups; $\rho = -0.19$, $p = 0.0001$ at Schaefer-400, but in the *wrong direction*, vanishing under partial correlation controlling for overlap, $\rho = -0.04$).

The approximate η formula. For a group of k equicorrelated features with pairwise correlation ρ , the eigenvalues of the within-group covariance matrix are $\lambda_1 = 1 + (k - 1)\rho$ and $\lambda_2 = \dots = \lambda_k = 1 - \rho$. The participation ratio $\text{PR} = k/(1 + (k - 1)\rho^2)$ measures effective dimensionality. The generalised η formula is:

$$\eta_{\text{approx}} = 1 - \frac{1}{1 + (k - 1)\rho^2} = \frac{(k - 1)\rho^2}{1 + (k - 1)\rho^2}.$$

At $\rho = 1$ (exact symmetry), this reduces to $(k - 1)/k$; at $\rho = 0$, to 0. For brain imaging with network sizes $k = 5$ – 24 and typical resting-state FC $\rho \approx 0.3$ – 0.5 , the predicted η values are substantially lower than the exact prediction and depend on ρ , which is unknown for the NARPS task data. Empirically, the formula with literature FC values gives $\rho_{\text{Spearman}} = 0.03$ ($p = 0.79$)—no improvement over the exact formula. A sensitivity analysis sweeping uniform ρ from 0.1 to 0.9 shows the Spearman correlation is invariant at -0.021 regardless of ρ , because η_{approx} with uniform ρ is a monotone function of k alone. The approximate formula is theoretically correct as an interpolation between exact and no symmetry, but empirically underdetermined for this application.

Per-team ranking analysis. Using individual team t -statistic maps (48 teams, 7 hypotheses, 100 regions), we computed pairwise flip rates across all team pairs for all region pairs. Within-network region pairs show higher flip rates (0.321) than between-network pairs (0.310), with gap $= +0.011$ [$+0.007$, $+0.015$], Cohen’s $d = 0.078$, Mann–Whitney $p = 3.7 \times 10^{-8}$, permutation $p < 0.005$, consistent across 7/7 hypotheses. The effect is statistically significant but small ($d = 0.078$), reflecting the dominance of methodological variation over structural variation in per-team rankings.

Software as a second symmetry source. On 24 teams with known software (12 FSL, 8 SPM, 4 AFNI), both network structure and software choice independently predict flip rates:

Effect	Flip rate premium	p -value
Network (within-net > between-net)	+0.005	8.8×10^{-5}
Software (diff-software > same-software)	+0.016	8.2×10^{-55}

Software has $3\times$ the effect of network structure (0.016 vs 0.005), consistent with Botvinik-Nezer et al.’s finding that methodology drives most disagreement. Both CIs exclude zero (network: [0.002, 0.008]; software: [0.014, 0.018]).

Orbit averaging as resolution. The overlap map (mean across teams’ binary activation maps) is the orbit average over the Rashomon set. We tested whether this resolution is near-optimal by comparing aggregation methods at full ensemble size ($M = 48$):

Method	Faithfulness (corr w/ teams)	Mean unfaithfulness
Mean	0.669	0.739
Median	0.660	0.719
Trimmed mean	0.668	0.727
Weighted mean	0.664	0.725

All methods are within 3% of each other. The mean achieves the highest faithfulness (correlation with individual teams); the median achieves the lowest absolute unfaithfulness. At $M = 48$, the mean is not dominated by any alternative. At smaller M (< 20), the weighted mean (weighting by leave-one-out prediction accuracy) slightly dominates the unweighted mean, consistent with the theorem’s prediction that equal weighting is optimal only for balanced Rashomon sets.

Convergence prescription. The split-half stability of the ensemble average increases with ensemble size. At $M = 16$ (95% CI: [10, 22]), the split-half Pearson correlation exceeds 0.95. The convergence rate fits $1 - a/M^b$ with $b = 1.00$ ($R^2 = 0.97$). This provides a concrete prescription for multi-analyst studies in neuroimaging: approximately 16 independent analyses suffice for 95% stability between independent assessments.

Bilemma confirmation. The unfaithfulness of the ensemble average is proportional to team disagreement (Spearman $\rho = 0.962$, $p < 10^{-100}$). The ratio unfaithfulness/ $\sigma = 0.734$ (Gaussian prediction: $\sqrt{2/\pi} = 0.798$), confirming the ensemble average achieves near the theoretical minimum unfaithfulness.

11.19 Cross-Domain Multi-Analyst Studies [Tier B: Empirical]

The same impossibility→resolution pattern appears in multi-analyst studies beyond neuroimaging.

Silberzahn et al. (2018): 29 teams, psychology. 29 teams analysed whether soccer referees give more red cards to dark-skinned players (AMPPS, 2018). Team estimates (odds ratios from OSF repository osf.io/gvm2z) range from 0.89 to 2.93 (geometric mean OR = 1.34). In log-odds space, the method-type grouping (Logistic: 15 teams, Count: 5, Linear: 3, Bayesian: 2, Other: 4) explains $\eta^2 = 0.131$ of estimate variance (ANOVA $F = 0.90$, $p = 0.48$ —not significant with 29 teams in 5 unequal groups). Within-method pairs show smaller estimate differences than between-method pairs (Cohen’s $d = 0.29$, Mann–Whitney $p = 0.03$), but this does not survive permutation correction ($p = 0.15$). The Noether prediction is inconclusive at this sample size.

Breznau et al. (2022): 71 teams, political science. 73 teams tested whether immigration reduces support for social policies (PNAS, 2022). From the GitHub repository (nbreznau/CRI), 71 teams with valid average marginal effects (AMEs) across 1,253 individual models show mean AMEs ranging from -0.27 to $+0.36$. Sign agreement is only 52%—teams cannot agree on the *direction* of the effect. The ICC (intraclass correlation) for team identity is -0.096 ($F = 0.3$, $p = 1.0$), meaning within-team model specification variation *exceeds* between-team variation. The team’s methodological approach does not determine the result; individual model specification choices (covariates, functional form) are the dominant source of variation.

Methodological note on scalar studies. The initial cross-domain analysis used hard-coded approximate data; the definitive version uses real team estimates from OSF (Silberzahn: osf.io/gvm2z, “Crowdsourcing Effects in OR.csv”) and GitHub (Breznau: nbreznau/CRI, “CRI Model Specifications and Margins.xlsx”). Convergence was measured using non-overlapping split-half (disjoint subsets of size M , requiring $2M \leq N$) to avoid the inflation caused by overlapping subsets. With non-overlapping split-half, neither study reaches 95% stability within the testable range ($M \leq N/2$): Silberzahn at $M = 14$ achieves stability 0.84; Breznau at $M = 35$ achieves 0.91. The Breznau ICC is unreliable due to extreme group-size imbalance (2–112 models per team); teams with many models have means regressed toward zero, deflating between-team variance.

Cross-domain pattern. Across all three domains, the qualitative pattern is identical: the Rashomon property holds (teams disagree), orbit averaging provides a near-optimal resolution (all aggregation methods within 3%), and unfaithfulness is proportional to disagreement. The Noether structural prediction—that group membership predicts which conclusions are stable—is confirmed for multivariate neuroimaging data ($d = 0.32$ after activation control) but inconclusive for scalar multi-analyst studies (Silberzahn $d = 0.29$, not significant after permutation correction; Breznau ICC = -0.10). For scalar studies, convergence follows standard CLT scaling ($\sim 1/\sqrt{M}$); the framework does not predict a different rate. The framework’s unique value beyond classical meta-analysis is threefold: the irreducibility proof (the disagreement is mathematical, not methodological), the Pareto optimality of orbit averaging (proven, not assumed), and the structural prediction identifying which conclusions are stable (confirmed for multivariate data). The CLT tells practitioners that averaging reduces noise; the framework tells them the noise is *irreducible*, proves orbit averaging is *optimal* among all aggregation methods, and identifies *which* conclusions survive.

12 Fundamental Ceiling on AI Interpretability [Tier B: Empirical]

The bridge theorems above connect the framework to *existing* results. This final bridge produces a *new* prediction: a quantitative upper bound on the fraction of a neural network’s internal representation that can be stably interpreted.

Consider a network with one hidden layer of width n . For any permutation $\pi \in S_n$, simultaneously permuting the rows of W_1 , elements of b_1 , and columns of W_2 yields an equivalent network computing the same function f . This hidden-unit permutation symmetry gives Rashomon group $G = S_n$, with $|G| = n!$ distinct parameterizations of the same input-output function.

By character theory, the natural representation of S_n on $\mathbb{R}^n$ (the hidden activation space) decomposes as

$$\mathbb{R}^n = V^{S_n} \oplus V_{\text{standard}},$$

where $V^{S_n} = \text{span}\{(1, 1, \dots, 1)\}$ has dimension 1 (the mean activation) and V_{standard} has dimension $n - 1$. By Proposition 7.1, only queries in V^{S_n} are stably answerable, so at most $1/n$ of the hidden-layer representation is stably interpretable. For L layers of width n , the symmetry

group is S_n^L with invariant subspace of dimension L (one mean per layer) out of nL total, giving the same $1/n$ ceiling.

Corollary 12.1 (Interpretability ceiling). *For a neural network with hidden-unit width n , the fraction of internal representation that is stably interpretable under the S_n permutation symmetry is at most $1/n$. Specifically: the mean activation (1 dimension) is S_n -invariant and hence stable; individual neuron activations ($n - 1$ dimensions) are not.*

For a transformer with $d_{\text{model}} = 768$ (BERT-base scale), this predicts at most $1/768 \approx 0.13\%$ of the hidden representation is stably interpretable—a hard ceiling independent of the interpretability method employed. This bound applies to *internal* representation interpretation (“opening the black box”); input-output behavior, being invariant under all network symmetries, is fully stable.

Experimental confirmation. We trained fully connected networks with hidden widths $n \in \{4, 8, 16, 32, 64, 128\}$ and, for each, generated 50 random hidden-unit permutations (preserving the input-output function exactly). Across all widths, the per-neuron importance ranking has a stable fraction of zero—even stricter than the $1/n$ bound—while the mean activation achieves correlation 1.0 across all permutations (Figure 6). The S_n symmetry is not an approximate nuisance: it is an exact equivalence that renders per-neuron interpretations fundamentally unstable.

Hungarian permutation alignment audit. An adversarial audit using Hungarian permutation alignment reveals that the non-uniqueness goes deeper than S_n permutation symmetry. After optimal permutation alignment—which perfectly recovers known permutations in a control experiment (aligned Jaccard = 1.0)—real cross-model agreement reaches only Jaccard = 0.079. The models learn genuinely different feature decompositions, not permuted versions of the same decomposition. CCA analysis (mean similarity = 0.950) confirms the representations are functionally equivalent—they encode the same information in different coordinate systems—but the transformation between models is a rotation/mixing, not a permutation. This result strengthens the interpretability ceiling: the actual symmetry group is larger than S_n (it includes continuous rotations), making the $1/n$ bound conservative.

Caveat. The $1/n$ bound addresses only the permutation symmetry S_n , which is the “minimal” network symmetry. Real networks may exhibit additional symmetries (ReLU scaling invariance, sign-flip symmetry) that further reduce the interpretable fraction. The $1/n$ bound is therefore an upper bound; the true ceiling may be lower.

Acknowledgements

We thank the Lean 4 and Mathlib communities for building the theorem-proving infrastructure that made the mechanical verification possible. We are grateful to the authors of the open-source libraries used in our experiments: Hugging Face Transformers, scikit-learn, XGBoost, and SHAP. We also thank the anonymous reviewers whose feedback improved the clarity of the manuscript. This work was conducted independently and received no external funding. We acknowledge substantial use of Claude (Anthropic) as an AI coding assistant for experiment scripting, Lean 4 proof development, LaTeX drafting, and iterative manuscript revision. A substantial fraction of the code and prose was AI-drafted under human direction; all scientific claims, proof strategies, experimental designs, and interpretive judgments were conceived by the human authors. Every Lean theorem was machine-checked; every empirical claim was verified against source data by the authors. The AI tool was used iteratively across multiple sessions; interaction logs are preserved locally and available upon request.

Code and Data Availability

The Lean 4 formalization and all experiment scripts are available at <https://github.com/DrakeCaraker/universal-explanation-impossibility> (universal framework: 130 files, 715 theorems, 2 axioms, 0 sorry). The attribution-specific formalization is at <https://github.com/DrakeCaraker/dash-impossibility-lean> (59 files, 368 theorems, 2 axioms, 0 sorry). The tightness classification and physics application are at <https://github.com/DrakeCaraker/ostrowski-impossibility> (38 files, 481 theorems, 0 axiom declarations with ~ 11 domain hypotheses carried as section variables, 0 sorry). All code is released under the Apache 2.0 licence. The results in this paper correspond to the `release-2026-07-29` tag of each repository; all counts above are generated from those tagged sources and verified by a continuous-integration check (`make ci`: reference and count linting of every paper against the sources, a full `lake build` of all three repositories, and an `#print axioms` audit of the hardened spine). Gene-expression data: OpenML (IDs 1137 and related). Neuroimaging: NeuroVault collection 6047. Multi-analyst data: OSF (osf.io/gvm2z) and GitHub (nbrenzau/CRI). A prospective out-of-distribution validation of the reliability certificate on a disjoint benchmark suite (OpenML-CC18) has been pre-specified and frozen at a public commit for third-party-timestamped registration on OSF; it will be reported in future work whatever its outcome, per the falsification discipline of Part V.

13 Empirical Predictions from the Mathematical Structure [Tier B: Empirical]

The recursive impossibility framework generates four testable predictions beyond those already validated in this paper:

1. **Bimodality of flip rates.** The all-or-nothing theorem (§13) predicts that for binary SHAP signs across a Rashomon set, the per-pair flip rate distribution should be bimodal — concentrated near 0% (between-group, stable) and $\geq 50\%$ (within-group, coin flip) — with a “dead zone” in between. The stable fact count experiment (Section 11.5) confirms this: the 50pp bimodal gap is invariant across $\rho = 0.50\text{--}0.99$.
2. **Mechanistic interpretability ceiling.** For binary questions about neural network circuits (“is this neuron responsible for this behavior?”), the biletta predicts $\geq 50\%$ disagreement across independently trained model instances. A pilot test on GPT-2 small (Section 24) is in progress.
3. **DASH tie rate = Rashomon fraction.** The enrichment uniqueness theorem predicts that DASH’s tie rate for a feature pair should equal the fraction of Rashomon pairs where the pair’s attributions are incompatible. This is testable by comparing DASH’s tied-pair count to the Gaussian flip formula’s predicted flip rate.
4. **Coverage conflict diagnostic.** The `hasCoverageConflict` predicate (§3.6) should predict which feature pairs trigger the biletta, independently of the Gaussian flip formula. The coverage conflict bridge validates the aggregate prediction for small feature spaces ($\rho = 0.90$, $P \leq 20$) but weakens substantially at higher P ($\rho = 0.25$, not significant); the originally reported $\rho = 0.961$ (15 datasets) reflected the small- P regime and does not generalise to all datasets.

These predictions are derived from the mathematical structure and testable with existing experimental infrastructure.

14 Regulatory Implications [Tier B: Empirical]

The EU AI Act [EU Parliament, 2024] requires “meaningful explanations” for high-risk systems, and SR 11-7 [Office of the Comptroller of the Currency, 2011] requires “developmental evidence” documenting how model outputs arise. These two demands pull in opposite directions: meaningful explanations favor decisiveness (a definite ranking of features or causes), while developmental evidence favors stability (reproducibility across equivalent models). Theorem 3.6 shows that no explanation system can simultaneously maximize both for underspecified models. A path to regulatory compliance nonetheless exists: G -invariant aggregation (e.g., DASH for attributions, CPDAGs for causal reports) satisfies stability by construction and partial faithfulness—faithfulness in expectation over the Rashomon set—at the cost of reporting ties where the data are genuinely ambiguous. We suggest that regulators adopt this framing: require that explanations be *stable and partially faithful*, and mandate explicit disclosure of tied or unresolvable queries rather than demanding decisive answers that are mathematically impossible to guarantee.

Liability implications. If an organisation relies on single-model SHAP explanations for high-stakes decisions and the Rashomon property holds, two independent audits using different random seeds could reach opposite conclusions about which feature drove the decision. Ensemble averaging (DASH with $M \geq 25$) provides a defensible alternative: the explanation is stable, the uncertainty is quantified, and the limitations are documented. This analysis reflects our reading of the regulatory text and should be reviewed by a legal scholar before being relied upon in compliance contexts.

14.1 A quantitative adverse-action test [Tier B: Empirical]

The regulatory argument above is sharpest where an explanation is legally operative. Under the U.S. Equal Credit Opportunity Act (Regulation B), a lender must state the *principal reasons* for an adverse credit decision; in practice those reasons are read off feature attributions on a single deployed model [Selbst and Barocas, 2018, Ustun et al., 2019], and the EU AI Act [EU Parliament, 2024] creates an analogous demand. We test how stable that operative explanation is. On German credit data (OpenML-31; 20 features), for each of 5 seeds we trained 60 bootstrap-retrained gradient-boosting models and kept those within a strict 2% test-accuracy window of the best—an explicit equal-performance Rashomon filter—then recorded each denied applicant’s principal (top-SHAP) reason under every model in the window. On average 84% of denied applicants (range 76–93% across seeds) have a principal reason that *flips* across equally accurate models: for most applicants the legally reported reason is an artifact of which equally good model was deployed.

The per-query certificate identifies the dependable minority. It certifies a principal reason for ~19% of denied applicants (range 4–29%), and certified reasons are modestly but consistently more concentrated across the ensemble than uncertified ones (modal-reason share 0.85 vs 0.71); the separation is directional, not dispositive, at these sample sizes, and we report it as such. The epistemic asymmetry is the point, and it is enforced by a theorem: because the certificate is provably *one-sided* (`flip_zero_at_low_snr`, §8), it can prove a reason *stable* but can never prove one *arbitrary* from the ensemble mean and variance alone. The 84% is therefore an *observed* instability across the constructed Rashomon set, not a certified impossibility; we state the negative half no more strongly than the mathematics licenses. The compliant procedure this suggests is constructive rather than nihilistic: report a certified principal reason where one exists, and where none does, disclose that no single principal reason is stable across equally accurate models rather than reporting a decisive reason the certificate cannot underwrite.

14.2 What Explainability *Can* Do

The G -invariant resolution (Section 5) is constructive, not nihilistic. It tells the practitioner exactly which desiderata can be jointly satisfied and which must be relaxed. DASH for feature attribution, CPDAGs for causal discovery [Verma and Pearl, 1991], and ensemble concept probes for concept-based explanations [Kim et al., 2018] each sacrifice decisiveness to gain stability—and each is a concrete instance of the abstract resolution. In practice, this means reporting confidence intervals over the Rashomon set [Fisher et al., 2019, Semenova et al., 2022] rather than point estimates: a feature ranking should come with a stability annotation indicating which comparisons are robust to retraining and which are not. The impossibility does not say “explanations are useless”; it says “decisive explanations are unreliable under underspecification, but partially decisive explanations with calibrated uncertainty are both achievable and informative.”

15 What This Paper Tells a Working Scientist [Tier B: Empirical]

What does this paper tell a working scientist? For an ML practitioner: which SHAP comparisons to trust ($\text{SNR} > 2$) and which to discard ($\text{SNR} < 0.5$), computed from bootstrap retrains alone. For a causal inference researcher: that ensemble averaging over the Rashomon set reduces orientation flip rates fourfold, following the same mechanism as DASH for attributions. For a biologist: that the character-theoretic capacity $(k-1)/k$ serves as a neutral-evolution baseline from which deviations measure selective pressure.




Figure 5: Universal capacity theorem. Predicted instability ($1 - \dim(V^G)/\dim(V)$) vs. observed instability across 16 domain instances. Well-characterized groups (filled markers) achieve $R^2 = 0.957$ with slope 0.91; approximate groups (open markers) scatter below the diagonal. The formula works; the challenge is identifying the correct symmetry group.




Figure 6: Interpretability ceiling experiment. Per-neuron importance rankings have zero stable fraction across hidden widths $n = 4$ to 128 (even stricter than the $1/n$ character-theory bound), while the mean activation correlation is exactly 1.0 across all permutations. Only S_n -invariant queries survive the hidden-unit permutation symmetry.

Table 11: Eight domains instantiate the Explanation Impossibility. Each row derives the Rashomon property from domain-specific first principles with zero domain-specific axioms.

Domain	Θ	Y	Witness 1	Witness 2	Same Y ?
Mathematics	Solutions of $Ax = b$	Ax	$(1, 1)$	$(0, 2)$	$2 = 2$
Biology	Codons	Amino acids	UCU	UCC	Ser=Ser
Physics (gauge)	Gauge configs	Holonomy	(T, F, F)	(F, F, T)	$F = F$
Physics (stat. mech.)	Microstates	Num. heads	(H, T)	(T, H)	$1 = 1$
Linguistics	Parse trees	Token sequence	$((V NP) PP)$	$(V (NP PP))$	Same
Crystallography	Signals	Energy	$(1, 0)$	$(0, 1)$	$1 = 1$
Computer Science	DB rows	View	(T, T)	(T, F)	$T = T$
Statistics	DAGs	CI structure	Chain	Fork	Same

Part III

Part III — Instantiations and Bridges

The universal theorem specialised to recognisable systems, and connections to established mathematics. These are witnesses that the hypothesis class is populated by familiar examples—not new theorems about those domains.

16 Derived Instances Across Eight Sciences [Tier A: Hardened]

Before turning to machine learning, we demonstrate the universality of Theorem 3.6 by deriving it as a zero-axiom consequence in eight scientific domains. Each instance constructively witnesses the Rashomon property from domain-specific first principles—no axioms are assumed beyond the definitions. Table 11 summarizes the common structure.

16.1 Underdetermined Linear Systems (Mathematics)

The most elementary instance of the Explanation Impossibility arises in numerical linear algebra [Strang, 2016]. Consider the system $Ax = b$ with $A = \begin{bmatrix} 1 & 1 \end{bmatrix}$ and $b = 2$. The solution set $\{x \in \mathbb{R}^2 : x_1 + x_2 = 2\}$ is a line, not a point: both $x^{(1)} = (1, 1)$ and $x^{(2)} = (0, 2)$ satisfy $Ax = b$. Every STEM student encounters this situation in their first linear algebra course.

The explanation system. Define $\mathcal{S}_{\text{lin}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows:

- $\Theta = \mathcal{H} = \mathbb{Z}^2$ (solution vectors);
- $Y = \mathbb{Z}$ (the observable: Ax);
- $\text{obs}(v) = v_1 + v_2$ (dot product with $A = \begin{bmatrix} 1 & 1 \end{bmatrix}$);
- $\text{exp} = \text{id}$ (the solution itself is the “explanation”);
- $\perp (v, v') \iff v \neq v'$.

The Rashomon property. The witnesses $v^{(1)} = (1, 1)$ and $v^{(2)} = (0, 2)$ satisfy $\text{obs}(v^{(1)}) = 1 + 1 = 2 = 0 + 2 = \text{obs}(v^{(2)})$ and $v^{(1)} \neq v^{(2)}$. The Rashomon property is constructively witnessed with zero axioms.

The impossibility. By Theorem 3.6, no method of choosing a solution to an underdetermined linear system can simultaneously be:

1. **Faithful:** the chosen solution agrees with the actual solution;
2. **Stable:** all solutions with the same image receive the same choice; and
3. **Decisive:** the choice distinguishes every pair of distinct solutions.

Empirical illustration. Four solvers (pseudoinverse, LSQR, Tikhonov, random null-space) applied to 100 underdetermined systems show mean pairwise RMSD of 0.091, compared to 0.049 for fully determined controls (Mann-Whitney $p = 6.3 \times 10^{-9}$). Disagreement is elevated across all null space dimensions $d > 0$. *Note:* Within the underdetermined group, RMSD actually *decreases* with d (slope = -0.0015), because three of the four solvers converge to near-minimum-norm solutions, and the random null-space projection’s per-component contribution dilutes as $1/\sqrt{d}$. The key finding is the presence/absence separation ($d > 0$ vs. $d = 0$), not a monotonic scaling. The control ($d = 0$) has RMSD = 0.049 (not exactly zero) due to Tikhonov regularization bias ($\lambda = 0.01$).

Resolution: minimum-norm solutions. The classical resolution is the Moore–Penrose pseudoinverse $x^+ = A^+(Ax)$, which selects the minimum-norm element of the solution set. This is precisely a G -invariant projection: the symmetry group G acts by translations along the null space of A , and x^+ is the unique element of each fiber that is invariant under the orthogonal complement of $\ker(A)$. The pseudoinverse sacrifices decisiveness (it cannot distinguish solutions within the same fiber) to achieve stability and faithfulness in expectation.

Lean formalization. The Lean 4 types are `Vec2` (a structure with fields `x y : Int`), `dotA` (the observation map $v \mapsto v.x + v.y$), and `linearSystem : ExplanationSystem Vec2 Vec2 Int` with `observe := dotA`, `explain := id`, and `incompatible := (≠)`. The impossibility is `linear_system_impossibility` in `LinearSystem.lean`, derived with zero axioms.

16.2 Genetic Code Degeneracy (Biology)

The genetic code is *degenerate*: 61 sense codons encode only 20 amino acids, so multiple codons map to the same amino acid [Crick, 1966]. For example, the codons UCU and UCC both encode Serine. This degeneracy is the biological analogue of the Rashomon property: two distinct codons (configurations) produce the same amino acid (observable) but carry incompatible nucleotide information (explanations).

The explanation system. Define $\mathcal{S}_{\text{genetic}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows:

- $\Theta = \mathcal{H} = \{\text{Codon}\}$ (RNA codons);
- $Y = \{\text{AminoAcid}\}$ (the 20 amino acids);
- $\text{obs} = \text{translate}$ (the codon-to-amino-acid map);
- $\text{exp} = \text{id}$ (the codon itself is the “explanation” of the translation);
- $\perp (c_1, c_2) \iff c_1 \neq c_2$.

The Rashomon property. The witnesses UCU and UCC satisfy $\text{translate}(\text{UCU}) = \text{Ser} = \text{translate}(\text{UCC})$ and $\text{UCU} \neq \text{UCC}$. The Rashomon property is constructively witnessed with zero axioms.

The impossibility. By Theorem 3.6, no method of inferring the codon from the amino acid can simultaneously be faithful (report a codon consistent with the actual one), stable (assign the same codon to all synonymous translations), and decisive (distinguish every pair of distinct codons). In biological terms: knowing the protein sequence does not uniquely determine the DNA sequence.

Empirical illustration. Across 120 eukaryotic cytochrome c coding sequences from NCBI RefSeq (one per species, mean GC content 0.45), per-amino-acid codon usage entropy shows a monotonic dose-response with degeneracy level: 1-fold (Met/Trp): $H = 0.00$ bits, 2-fold: 0.93, 3-fold: 1.49, 4-fold: 1.78, 6-fold: 2.20 (Kruskal–Wallis $p = 2.0 \times 10^{-3}$ across 20 amino acids in 5 groups). The negative control (Met and Trp, degeneracy 1) has entropy exactly 0 by construction—this follows from the structure of the genetic code itself (one codon for Met, one for Trp) and holds for all species regardless of GC content. This dose-response is predicted by the code’s degeneracy structure: higher degeneracy provides more synonymous alternatives, creating a larger space of possible codon distributions. GC-content bias is the primary mechanism driving codon choice within this space—only 11% of degenerate amino acids show entropy exceeding the GC-null prediction—but the structural capacity for variation is created by the Rashomon property (degeneracy) itself. Without degeneracy, entropy is identically zero regardless of GC content.

A per-position analysis using Clustal Omega alignment across the same 120 species identifies 89 conserved positions. Grouping these positions by the consensus amino acid’s degeneracy level and computing per-position codon usage entropy confirms the dose-response at finer resolution: Kruskal–Wallis $p = 3.5 \times 10^{-9}$ across degeneracy groups. The GC-content null model accounts for the majority of observed entropy: only 4.7% of degenerate positions exceed the GC-null prediction, confirming that GC bias is the primary driver of codon choice *within* each degeneracy class, while the degeneracy structure itself creates the capacity for variation.

Data: NCBI RefSeq cytochrome c CDS, searched via `CYCS[Gene] AND mRNA[Filter] AND refseq[Filter]`, deduplicated by organism.

The degeneracy of the genetic code is well characterised biochemically [Sharp and Li \[1987\]](#), [Hershberg and Petrov \[2008\]](#): synonymous codons differ in translational efficiency, mRNA stability, and CpG methylation marking. Our framework adds a structural observation: the impossibility of simultaneously faithful, stable, and decisive codon-level explanation is a formal consequence of degeneracy, independent of the evolutionary forces shaping usage bias. The resolution—aggregating over synonymous codons to report amino-acid-level attributions—parallels DASH’s feature-group aggregation.

Resolution: codon usage tables. The standard resolution in molecular biology is *codon usage tables*, which report the frequency distribution over synonymous codons for each amino acid. This is the orbit average over the symmetry group of synonymous substitutions: rather than committing to a single codon, one reports the expected codon usage across the degenerate fiber. The resolution sacrifices decisiveness (no single codon is selected) to achieve stability and faithfulness in expectation.

Lean formalization. The Lean 4 types are `Codon` (an inductive with constructors `UCU` and `UCC`), `AminoAcid` (with constructor `Ser`), and `geneticCodeSystem : ExplanationSystem Codon Codon AminoAcid` with `observe := translate`, `explain := id`, and `incompatible := (≠)`. The impossibility is `genetic_code_impossibility` in `GeneticCode.lean`, derived with zero axioms.

16.3 Gauge Freedom (Physics)

Gauge invariance is a foundational principle of modern physics: physically measurable quantities are invariant under gauge transformations, but the mathematical description of the field is not [Jackson, 1999]. We formalize a discrete $\mathbb{Z}_2$ gauge theory on a triangle graph (three vertices, three edges) to exhibit the Explanation Impossibility in physics.

The explanation system. Define $\mathcal{S}_{\text{gauge}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows:

- $\Theta = \mathcal{H} = \text{EdgeConfig}$: each configuration assigns a $\mathbb{Z}_2$ (Boolean) label to each of the three edges e_{01}, e_{12}, e_{02} ;
- $Y = \text{Bool}$: the *holonomy* around the triangle, $h(g) = e_{01} \oplus e_{12} \oplus e_{02}$, the discrete analogue of the Wilson loop;
- $\text{obs} = \text{holonomy}$;
- $\text{exp} = \text{id}$ (the full edge configuration is the “explanation”);
- $\perp (g_1, g_2) \iff g_1 \neq g_2$.

A *gauge transform at vertex 0* flips the two incident edges (e_{01} and e_{02}), leaving e_{12} unchanged.

The Rashomon property. The witnesses are $g_1 = (\text{true}, \text{false}, \text{false})$ and $g_2 = (\text{false}, \text{false}, \text{true})$. They satisfy $\text{holonomy}(g_1) = \text{true} = \text{holonomy}(g_2)$ and $g_1 \neq g_2$. Moreover, g_2 is the gauge transform of g_1 at vertex 0. The Rashomon property is constructively witnessed with zero axioms.

Universal gauge invariance. The theorem `gauge_preserves_holonomy` is proved for *all* configurations, not just the two witnesses: for every $g \in \text{EdgeConfig}$, $\text{holonomy}(\text{gaugeAt0}(g)) = \text{holonomy}(g)$. The proof proceeds by case analysis on all $2^3 = 8$ configurations.

The impossibility. By Theorem 3.6, no description of a gauge field can simultaneously be faithful (report the actual edge configuration), stable (assign the same description to gauge-equivalent configurations), and decisive (distinguish every pair of distinct configurations). In the language of physics: one cannot simultaneously specify a particular gauge potential, demand gauge invariance of the specification, and retain the ability to distinguish gauge-equivalent potentials.

Empirical illustration. On $\mathbb{Z}_2$ lattice gauge theory ($N \times N$ grids, $N = 4\text{--}10$), gauge-variant link variables have within-orbit variance 0.25 (the Bernoulli maximum) while gauge-invariant plaquette values have within-orbit variance exactly 0. The Rashomon set (gauge orbit) covers all possible link configurations for a given holonomy.

Monte Carlo sampling on a 16×16 periodic lattice at 10 coupling values ($\beta \in \{0.1, 0.2, \dots, 2.0\}$; 2000 configurations per β) confirms a coupling-dependent dose-response: the gauge coupling β controls the Rashomon set size. Plaquette variance follows $\text{Var}(P) = \text{sech}^2(\beta)/N^2$, decreasing monotonically from 3.85×10^{-3} at $\beta = 0.1$ to 2.71×10^{-4} at $\beta = 2.0$. Wilson loop (2×2) variance follows $\text{Var}(W) = 1 - \tanh^8(\beta)$, decreasing from ≈ 1.00 at weak coupling to 0.23 at strong coupling. At weak coupling ($\beta \rightarrow 0$), plaquettes are nearly uniform on $\{+1, -1\}$ and the configuration space is maximally uncertain (large Rashomon set); at strong coupling ($\beta \rightarrow \infty$), plaquettes freeze to +1 and the configuration concentrates near the ordered ground state (small Rashomon set). Since 2D $\mathbb{Z}_2$ gauge theory is exactly solvable,

the MC serves as numerical verification of the analytic predictions rather than an empirical discovery. The contribution is the demonstration that the gauge coupling continuously tunes the Rashomon set size—an exact analogue of the collinearity parameter in the ML stability transition experiment.

Resolution: gauge-invariant observables. The resolution in physics is well known: restrict attention to *gauge-invariant observables* such as the holonomy (Wilson loops), field strengths, or scattering amplitudes. In our framework, the holonomy is the G -invariant projection: the gauge group acts on configurations, and the holonomy is the unique function constant on each gauge orbit. The resolution sacrifices decisiveness (one cannot recover the gauge potential from the holonomy alone) to achieve stability and faithfulness.

Lean formalization. The Lean 4 types are `EdgeConfig` (a structure with fields `e01 e12 e02 : Bool`), `holonomy` (the observation map), and `gaugeSystem : ExplanationSystem EdgeConfig EdgeConfig Bool`. The impossibility is `gauge_impossibility` in `GaugeTheory.lean`, derived with zero axioms.

16.4 Microstate Degeneracy (Statistical Mechanics)

In statistical mechanics, a *macrostate* (e.g., temperature, pressure) is compatible with many *microstates* (specific molecular configurations) [Jaynes, 1957]. This coarse-graining—physically distinct configurations producing the same observable—is the foundational example of under-specification in physics. We formalize a minimal witness: a pair of binary coins. This is not a complete formalization of statistical mechanics but the simplest system exhibiting the microstate/macrostate degeneracy that drives the impossibility in real physical systems with $\sim 10^{23}$ particles.

The explanation system. Define $\mathcal{S}_{\text{stat}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows:

- $\Theta = \mathcal{H} = \text{CoinPair}$: a microstate is a pair of coins, each Heads or Tails;
- $Y = \mathbb{N}$: the macrostate, defined as the number of heads;
- $\text{obs} = \text{numHeads}$;
- $\text{exp} = \text{id}$ (the full microstate is the “explanation”);
- $\perp (c_1, c_2) \iff c_1 \neq c_2$.

The permutation group S_2 acts on microstates by swapping the two coins: $(H, T) \mapsto (T, H)$. The macrostate is invariant under this action. (For non-interacting systems, the symmetry group is the full permutation group S_Ω ; for interacting systems, the relevant group is the automorphism group of the Hamiltonian, which is typically a proper subgroup of S_Ω .)

The Rashomon property. The witnesses are (H, T) and (T, H) . Both have macrostate `numHeads` = 1, and they are distinct microstates. Moreover, $(T, H) = \text{swapCoins}(H, T)$: they are related by the S_2 permutation. The Rashomon property is constructively witnessed with zero axioms.

Universal permutation invariance. The theorem `swap_preserves_macrostate` is proved for *all* configurations, not just the two witnesses: for every $c \in \text{CoinPair}$, `numHeads(swapCoins(c)) = numHeads(c)`. The proof proceeds by case analysis on all $2^2 = 4$ configurations.

The impossibility. By Theorem 3.6, no description of a thermodynamic system can simultaneously be faithful (specify a particular microstate), stable (assign the same description to macrostate-equivalent configurations), and decisive (distinguish every pair of distinct microstates).

Empirical illustration. For N binary spins, the Rashomon entropy $S_R(k) = \ln \binom{N}{k}$ equals the Boltzmann entropy. Maximum faithfulness decays exponentially: $1/\Omega(k) = e^{-S_R(k)}$. At $k = N/2$ (maximum entropy), $\Omega = \binom{50}{25} \approx 1.26 \times 10^{14}$, making faithful microstate identification effectively impossible.

Resolution: the microcanonical ensemble. The microcanonical ensemble—the uniform distribution over microstates at fixed macrostate—is structurally analogous to the G -invariant resolution. For permutation-symmetric systems, the connection is precise: the permutation group acts on microstates, and the orbit average produces the macrostate. Boltzmann’s program of replacing specific microstates with statistical distributions over macrostates is an instance of the resolution strategy this paper formalizes.

Lean formalization. The Lean 4 types are `CoinPair` (a structure with fields `coin1 coin2 : Bool`), `numHeads` (the observation map), and `statMechSystem : ExplanationSystem CoinPair CoinPair Nat`. The impossibility is `stat_mech_impossibility` in `StatisticalMechanics.lean`, derived with zero axioms.

16.5 Syntactic Ambiguity (Linguistics)

Syntactic ambiguity is pervasive in natural language [Chomsky, 1957, Jurafsky and Martin, 2024]. The classic example is prepositional phrase (PP) attachment: the sentence fragment “V NP PP” admits two parse trees:

- **NP-attachment** (low attachment): $(V (NP PP))$ — the PP modifies the NP;
- **VP-attachment** (high attachment): $((V NP) PP)$ — the PP modifies the verb.

For instance, “saw the man with the telescope” can mean either “saw [the man who had the telescope]” or “[saw the man] [using the telescope].” Both bracketings yield the same surface token sequence $[V, NP, PP]$ but assign different constituent structures.

The explanation system. Define $\mathcal{S}_{\text{syntax}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows:

- $\Theta = \mathcal{H} = \text{Bracketing}$ (parse trees);
- $Y = \text{List Token}$ (the surface token sequence);
- $\text{obs} = \text{yield}$ (the linearization of the parse tree);
- $\text{exp} = \text{id}$ (the parse tree is the “explanation” of the sentence);
- $\perp (b_1, b_2) \iff b_1 \neq b_2$.

The Rashomon property. The witnesses are `leftAttach` and `rightAttach`, which satisfy $\text{yield}(\text{leftAttach}) = [V, NP, PP] = \text{yield}(\text{rightAttach})$ and $\text{leftAttach} \neq \text{rightAttach}$. The Rashomon property is constructively witnessed with zero axioms.

The impossibility. By Theorem 3.6, no parser can simultaneously be faithful (return the actual parse tree), stable (assign the same parse to all token sequences with the same yield), and decisive (distinguish every pair of distinct parse trees). In computational linguistics: when a sentence is genuinely ambiguous, no deterministic parser can correctly report the syntactic structure without either being unfaithful to one reading or unstable across equivalent inputs.

Empirical illustration. Three syntactic parsers (spaCy small, spaCy large, Stanza—two architectures, independently trained) applied to 50 PP-attachment ambiguous and 50 unambiguous sentences with shared pre-tokenization show mean inter-parser agreement (UAS) of 0.816 for ambiguous sentences versus 0.994 for unambiguous (Wilcoxon $p = 1.7 \times 10^{-17}$, Cohen’s $d = 6.16$). All 50 ambiguous sentences follow the “V NP PP” pattern; pre-tokenization eliminates tokenizer-driven disagreement; and parser independence is ensured by using at most two models from the same family (spaCy CNN) plus an architecturally distinct alternative (Stanza BiLSTM).

Resolution: packed parse forests. The standard resolution in NLP is to report a *packed parse forest* [Jurafsky and Martin, 2024]: a compact representation of all valid parse trees for the input. This is the G -invariant aggregation over the symmetry group of syntactic rearrangements within the same yield fiber. Chart parsers (CYK, Earley) naturally produce packed forests. The resolution sacrifices decisiveness (no single tree is selected) to achieve stability and faithfulness.

Lean formalization. The Lean 4 types are `Token` (an inductive with constructors `V`, `NP`, `PP` representing phrasal categories, not surface tokens), `Bracketing` (with constructors `leftAttach` and `rightAttach`, corresponding to NP-attachment and VP-attachment), and `syntaxSystem : ExplanationSystem Bracketing Bracketing (List Token)`. The impossibility is `syntactic_ambiguity_impossibility` in `SyntacticAmbiguity.lean`, derived with zero axioms.

16.6 The Phase Problem (Crystallography)

The *phase problem* is the central challenge of X-ray crystallography: diffraction experiments measure the squared magnitudes $|F(\mathbf{k})|^2$ of the structure factors, but the phases are lost. Multiple electron density distributions can produce identical diffraction patterns. Hauptman and Karle [1953] developed direct methods that partially recover the phases—work recognized with the 1985 Nobel Prize in Chemistry [Karle and Hauptman, 1956]—but the fundamental ambiguity remains.

The explanation system. We formalize a minimal witness: a 2-element discrete signal. *Note:* For $N = 2$, the phase ambiguity is limited to permutation and sign (the set $\{(a, b), (b, a), (-a, -b), \dots\}$ sharing $a^2 + b^2$). For general N , the phase ambiguity grows combinatorially, which is why the phase problem is hard in practice. Our 2-element witness establishes that the Rashomon property holds; the full crystallographic problem is strictly harder. Define $\mathcal{S}_{\text{phase}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows:

- $\Theta = \mathcal{H} = \text{Signal2}$ (a pair of integers (a, b) representing the signal);
- $Y = \mathbb{Z}$ (the *energy*: $a^2 + b^2$, modeling $|F|^2$);
- $\text{obs} = \text{energy}$;
- $\text{exp} = \text{id}$ (the signal itself is the “explanation” of the diffraction data);
- $\perp (s_1, s_2) \iff s_1 \neq s_2$.

The Rashomon property. The witnesses are $s_1 = (1, 0)$ and $s_2 = (0, 1)$, which satisfy $\text{energy}(s_1) = 1^2 + 0^2 = 1 = 0^2 + 1^2 = \text{energy}(s_2)$ and $s_1 \neq s_2$. The Rashomon property is constructively witnessed with zero axioms.

The impossibility. By Theorem 3.6, no method of determining a signal from its energy (diffraction intensity) can simultaneously be faithful (report the actual signal), stable (assign the same signal to all inputs with the same energy), and decisive (distinguish every pair of distinct signals). In crystallographic terms: the magnitude of the structure factor does not uniquely determine the electron density.

Empirical illustration. Gerchberg–Saxton phase retrieval from magnitude-only data produces 20 reconstructions with mean pairwise RMSD of 1.32–1.44 for unconstrained signals (lengths 16–128), compared to 0.75–0.92 for positive signals where the positivity constraint reduces phase ambiguity (ratio 1.5–1.8 $\times$). Statistical significance is assessed using the $N = 20$ per-reconstruction mean RMSDs as independent units (not the 190 pairwise RMSDs, which are pseudoreplicated): Mann–Whitney U , $p = 3.4 \times 10^{-8}$ per signal length.

Resolution: Patterson maps. The classical partial resolution is the *Patterson map*—the autocorrelation function of the electron density, which is computable directly from $|F(\mathbf{k})|^2$ without phase information. The Patterson map is the G -invariant observable: the symmetry group of phase rotations acts on the signal space, and the autocorrelation is invariant under this action. Modern direct methods [Hauptman and Karle, 1953, Karle and Hauptman, 1956] go further by exploiting probabilistic constraints among phases, but even these methods cannot eliminate the fundamental ambiguity without additional experimental data (e.g., anomalous scattering or isomorphous replacement).

Lean formalization. The Lean 4 types are `Signal2` (a structure with fields `a b : Int`), `energy` (the observation map $s \mapsto s.a \cdot s.a + s.b \cdot s.b$), and `phaseSystem : ExplanationSystem Signal2 Signal2 Int`. The impossibility is `phase_impossibility` in `PhaseProblem.lean`, derived with zero axioms.

16.7 The View Update Problem (Computer Science)

The *view update problem* asks: given a database view (a projection of the base relation), how can updates to the view be translated back to the base relation? Bancilhon and Spyratos [1981] characterized exactly when unique translation is possible via the *complement view* framework: a view admits a unique update translator if and only if a complement view exists such that the pair forms a lossless decomposition. As a corollary, non-injective views without complements have no unique translation—a structural observation that our framework captures as a minimal witness of the Explanation Impossibility.

The explanation system. Define $\mathcal{S}_{\text{view}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows:

- $\Theta = \mathcal{H} = \text{Row}$ (database rows with two Boolean columns `colA` and `colB`);
- $Y = \text{Bool}$ (the view: the projection onto `colA`);
- $\text{obs} = \text{view}$, where $\text{view}(r) = r.\text{colA}$;
- $\text{exp} = \text{id}$ (the full row is the “explanation” of the view);
- $\perp (r_1, r_2) \iff r_1 \neq r_2$.

The Rashomon property. The witnesses are $r_1 = (\text{true}, \text{true})$ and $r_2 = (\text{true}, \text{false})$, which satisfy $\text{view}(r_1) = \text{true} = \text{view}(r_2)$ and $r_1 \neq r_2$. The Rashomon property is constructively witnessed with zero axioms.

The impossibility. By Theorem 3.6, no translation from views back to base relations can simultaneously be faithful (the translated row matches the actual base row), stable (all rows with the same view receive the same translation), and decisive (the translation distinguishes every pair of distinct base rows). This captures the impossibility corollary of Bancilhon and Spyratos [1981]’s complement framework: when no complement view exists, the non-injective view admits no unique translation. Our Lean formalization provides a minimal constructive witness of this structural phenomenon, not a full encoding of the complement lattice theory.

Resolution: complement views. The classical resolution [Bancilhon and Spyratos, 1981] is to use *complement views*: additional projections that, together with the original view, uniquely determine the base row. In our framework, the complement view is the additional G -invariant information needed to collapse the fiber to a single point. When no complement is available, the resolution is to report the entire fiber: all base rows consistent with the view. This sacrifices decisiveness (no single row is selected) to achieve stability and faithfulness.

Empirical illustration. Census disaggregation using real U.S. state and county structure (2020 Decennial Census, P1_001N): the KL-divergence between the true county distribution and Dirichlet-sampled alternatives scales with the number of counties (Spearman $\rho = 0.512$, $p = 1.2 \times 10^{-4}$; Pearson $r = 0.537$, $p = 4.8 \times 10^{-5}$). KL ranges from 0.00 (DC, 1 county) to 2.53 (Texas, 254 counties). The negative control (DC) has KL = 0 exactly. Pairwise KL divergence between Dirichlet-sampled alternatives (the Rashomon set diameter) also scales with aggregation granularity (Spearman $\rho = 0.697$, $p = 1.35 \times 10^{-8}$). Distribution entropy of the county population vector correlates strongly with county count (Spearman $\rho = 0.892$, $p = 1.7 \times 10^{-18}$), confirming that the entropy of the true distribution—not just its distance from Dirichlet alternatives—increases with the degree of underspecification.

Historical note. The view update problem was identified by Bancilhon and Spyratos [1981] in 1981. Our Lean formalization captures a minimal witness (a 2-row, 2-column Boolean database with a single-column projection) that exhibits the core structural phenomenon—non-injective observation implies Rashomon—rather than formalizing the full complement lattice theory. The structural parallel with gauge invariance in physics, syntactic ambiguity in linguistics, and the phase problem in crystallography was not previously recognized.

Lean formalization. The Lean 4 types are `Row` (a structure with fields `colA colB : Bool`), `view` (the observation map $r \mapsto r.\text{colA}$), and `viewUpdateSystem : ExplanationSystem Row Row Bool`. The impossibility is `view_update_impossibility` in `ViewUpdate.lean`, derived with zero axioms.

The eighth derived instance—Markov equivalence in causal inference—appears in Section 17.5 below, where we present it alongside the ML causal discovery instance to which it is most closely connected.

17 Applications in Machine Learning [Tier B: Empirical]

We instantiate Theorem 3.6 in nine structurally distinct ML explanation domains. In each case we (i) specify the explanation system $(\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$, (ii) verify the Rashomon property, and (iii) characterize the resolution.

Table 12: Cross-domain empirical validation. Each experiment confirms the impossibility in its domain with a domain-appropriate metric and negative control. With 14 testable simultaneous tests (7 cross-domain with p -values + 7 ML), the Bonferroni-corrected threshold is $\alpha = 0.05/14 = 0.004$. All experiments survive this correction. Stat Mech is mathematical (no p -value).

Domain	Metric	Result	Control	p
Mathematics	Solver RMSD	0.091	0.049	6.3×10^{-9}
Biology [†]	Cytochrome c diversity	$r = 0.798$	uniform	0.018
Physics (gauge)	Link variance	0.25	0.00	–
Physics (stat mech)	<i>Mathematical: $S_R = \ln \Omega$ (Boltzmann entropy)</i>			
Linguistics	Parser UAS (ambig.)	0.821	0.894	1.6×10^{-3}
Crystallography	Reconstruction RMSD	1.39	0.83	5×10^{-238}
Computer Science	Disaggregation KL	0.00–2.53	0.00	1.2×10^{-4}
Statistics	Orientation agreement	0.33	1.00	4.4×10^{-38}

[†] 120 eukaryotic cytochrome c sequences from NCBI RefSeq; capacity formula predicts diversity with systematic undershoot from purifying selection. Causal: 100 seeds, convergence to CPDAG at $N \rightarrow \infty$.

Table 13: ML instances by formal completeness. All nine axiomatise the Rashomon property from cited evidence rather than deriving it from learning-theoretic first principles. The core impossibility (Theorem 3.6) is independent of these axioms.

Instance	Rashomon	Quant. bound	Resolution
Feature attribution	Axiomatised	$1/(1-\rho^2)$	DASH
Model selection	Axiomatised	80% flip	Ensemble
Causal discovery	Axiomatised	Orientation flip	CPDAG
Attention maps	Axiomatised	19.9% flip	Avg. rollout
Counterfactual	Axiomatised	23.5% flip	Robust recourse
Concept probes	Axiomatised	0.90 instability	Subspace avg.
Saliency maps	Axiomatised	9.6% IoU flip	Avg. GradCAM
LLM self-explanation	Axiomatised	34.5% flip	Citation ensemble
Mech. interpretability	Axiomatised	Classical result	Open

The eight cross-domain scientific instances (§16) are constructive: they build explicit Rashomon witnesses with zero domain-specific axioms via **decide**.

17.1 Instance 1: Feature Attribution Under Collinearity

We instantiate the abstract framework (Definition 3.1) for feature attribution, recovering the Attribution Impossibility of the companion paper as a corollary of Theorem 3.6.

The explanation system. Fix $P \geq 2$ features partitioned into L collinearity groups. Within each group $\ell \in [L]$, every pair of features shares pairwise correlation $\rho \in (0, 1)$; features across groups are independent. The explanation system $\mathcal{S}_{\text{attr}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ is:

- Θ : joint data-generating distributions over $\mathbb{R}^P$ with the collinearity structure above (the ground truth the data encode).
- $\mathcal{H}$: trained models f from a fixed model class (e.g. GBDT, Lasso, neural networks), indexed by random training seed; each f is associated with its data-generating distribution $\theta(f) \in \Theta$ via the attribution map $\theta : \mathcal{H} \rightarrow \Theta$.

- $Y = (\text{Fin } P \rightarrow \mathbb{R})$: the space of feature-attribution vectors $\varphi = (\varphi_1, \dots, \varphi_P)$, where $\varphi_j(f) \geq 0$ is the global importance of feature j in model f (e.g. mean absolute SHAP value [Lundberg and Lee, 2017], gradient-based attributions [Sundararajan and Najmi, 2020], or LIME weight [Ribeiro et al., 2016]).
- $\text{obs} : \mathcal{H} \rightarrow Y$ assigns to each model f the attribution vector it “honestly reports”: $\text{obs}(f) = \varphi(f)$.
- $\text{exp} : \mathcal{H} \rightarrow Y$ is the attribution procedure under evaluation—the practitioner’s choice of explanation method applied to any model in $\mathcal{H}$.
- $\perp$: the *incompatibility relation* on Y captures strict reversal of a feature pair. Formally, $\varphi \perp \varphi'$ if there exists a group ℓ and distinct features $j, k \in \ell$ such that

$$\varphi_j > \varphi_k \quad \text{and} \quad \varphi'_k > \varphi'_j.$$

Two attribution vectors are incompatible when they rank the same collinear pair in opposite orders; both orderings cannot simultaneously be correct.

Rashomon property. Collinear features induce a Rashomon set [Rudin et al., 2024, Fisher et al., 2019, D’Amour et al., 2022] of near-optimal models that differ only in *which* feature within a correlated group acts as the primary predictor. Formally, $\mathcal{S}_{\text{attr}}$ has the Rashomon property (Definition 3.5): for any group ℓ and any distinct features $j, k \in \ell$, there exist models $f, f' \in \mathcal{H}$ with $\theta(f) = \theta(f')$ (same data-generating distribution) such that

$$\varphi_j(f) > \varphi_k(f) \quad \text{and} \quad \varphi_k(f') > \varphi_j(f').$$

This is the `RashomonProperty` of `Trilemma.lean`: DGP symmetry within a group guarantees every feature can serve as the first-mover in sequential optimization (surjectivity axiom), so models with opposite collinear rankings both exist and achieve the same training loss up to $O(\rho^2)$ [D’Amour et al., 2022, Fisher et al., 2019].

Definition 17.1 (Incompatible attribution rankings). *Two attribution vectors $\varphi, \varphi' \in Y$ are incompatible if there exist distinct features j, k in the same collinearity group such that $\varphi_j > \varphi_k$ and $\varphi'_k > \varphi'_j$.*

Instance 1 of the Universal Theorem.

Corollary 17.2 (Attribution Impossibility). *The explanation system $\mathcal{S}_{\text{attr}}$ has the Rashomon property. Therefore, by Theorem 3.6, no attribution map exp can simultaneously be faithful (attributions agree with the model’s own values), stable (same attributions regardless of which equivalent model is explained), and decisive (every collinear pair is strictly ranked).*

Proof. The Rashomon property was established above. The impossibility is then a direct application of Theorem 3.6 to $\mathcal{S}_{\text{attr}}$. $\square$

In the notation of the companion paper (the Attribution Impossibility paper), faithfulness corresponds to faithfulness (Definition 3.2), stability to stability (Definition 3.3), and decisiveness to completeness (every feature pair is strictly ranked). Corollary 17.2 is the abstract-framework restatement of the companion Attribution Impossibility (Theorem 1 of the companion paper), verified as `attribution_impossibility` in `Trilemma.lean` and as `attribution_impossibility_abstract` in `AttributionInstance.lean` (zero model-specific axiom dependencies in both cases).

Quantitative depth. The abstract impossibility is qualitative. Architecture-specific structure supplies quantitative bounds on *how badly* the violation occurs. For GBDT with per-group correlation ρ , the attribution ratio (dominant versus symmetric partner) diverges as $1/(1-\rho^2)$ as $\rho \rightarrow 1$ (Lean-verified in `Ratio.lean`; see also D’Amour et al. [2022]). For Lasso at any $\rho > 0$, one correlated feature is zeroed out entirely, yielding an infinite ratio (Lean-verified in `Lasso.lean`). For random forests with T trees, ensemble averaging drives the ratio to $O(1/\sqrt{T})$ —the lone converging case and the mechanism behind DASH.

Connection to variance inflation. The attribution ratio $1/(1-\rho^2)$ coincides with the classical variance inflation factor (VIF) from regression diagnostics. This is not coincidental: both measure the amplification of estimation variance due to collinearity. The novelty is the *application*: VIF measures instability of coefficient estimates, while $1/(1-\rho^2)$ measures instability of *explanations of predictions*. These are different objects—a model’s coefficients and its SHAP attributions need not agree—but the shared formula reflects the shared mechanism.

Resolution. The G -invariant resolution (Section 5) for this instance is DASH (Diversified Aggregation of SHAP): averaging attribution vectors over an ensemble of retrained models collapses the Rashomon set to a single equitable vector with $\mathbb{E}[\varphi_j] = \mathbb{E}[\varphi_k]$ for symmetric features j, k , restoring faithfulness (in expectation) and stability at the cost of decisiveness—collinear pairs are tied rather than arbitrarily ranked.

Experimental note. Four of twenty models in the attribution experiment achieved 72–78% held-out accuracy; all twenty are included in the analysis. This range reflects genuine Rashomon set variation—functionally near-equivalent models that nonetheless differ in their attribution vectors, precisely the phenomenon the impossibility theorem characterizes.

17.2 Instance 2: Attention Maps as Explanations

Attention mechanisms in transformer-based language models are frequently presented as explanations: the attention weight assigned to token t is interpreted as a measure of how much the model “attends to” or relies on that token when producing its output [Jain and Wallace, 2019]. We show that this interpretation is incompatible with the three desiderata under underspecification.

The explanation system. Let T denote the number of tokens in the context. We instantiate the general framework as the 6-tuple $\mathcal{S}_{\text{attn}} = (\Theta_{\text{attn}}, \mathcal{H}_{\text{attn}}, Y_{\text{attn}}, \text{obs}_{\text{attn}}, \text{exp}_{\text{attn}}, \perp_{\text{attn}})$ where:

- Θ_{attn} is the space of neural network weight configurations θ (query, key, value, and output projection matrices at every layer and head);
- $\mathcal{H}_{\text{attn}} = \Delta^T$ is the probability simplex over T tokens—the space of attention distributions $\alpha \in \mathbb{R}^T$ with $\alpha_t \geq 0$ and $\sum_{t=1}^T \alpha_t = 1$;
- Y_{attn} is the space of input-to-output functions $f : \mathcal{X} \rightarrow \mathcal{Y}$, i.e., the observable input→output behavior of the network;
- $\text{obs}_{\text{attn}}(\theta) = f_\theta$ is the function computed by the network with weights θ ;
- $\text{exp}_{\text{attn}}(\theta) = \alpha_\theta \in \Delta^T$ is the attention distribution produced by θ on a fixed input $x \in \mathcal{X}$; and
- $(\alpha, \alpha') \in \perp_{\text{attn}}$ if and only if $\arg\max_t \alpha_t \neq \arg\max_t \alpha'_t$, i.e., the two distributions assign peak attention to different tokens.

Table 14: Attention map instability across 10 functionally equivalent DistilBERT models. Negative control uses single-token inputs (model cannot choose where to attend). Resolution test compares pairwise individual agreement with individual-vs-full-average agreement. All 95% CIs from 2000 bootstrap resamples.

Test	Metric	Value (95% CI)
Positive (Rashomon)	Prediction agreement	100.0% [100.0%, 100.0%]
	Argmax flip rate	60.0% [59.0%, 61.1%]
	Mean Kendall τ	0.297 [0.287, 0.306]
Neg. control (1 token)	Argmax flip rate	2.0% [1.1%, 3.0%]
Resolution (full avg.)	Pairwise agreement	40.0% [39.0%, 41.0%]
	Indiv-vs-average agreement	55.2% [53.1%, 57.4%]
	Improvement	+15.2 pp

The Rashomon property. D’Amour et al. [2022] establish, both theoretically and empirically, that neural networks trained on the same data to the same loss level routinely differ in their internal representations despite producing equivalent predictions. Applied to attention, this yields the Rashomon property for $\mathcal{S}_{\text{attn}}$: there exist weight configurations θ_1, θ_2 such that

$$f_{\theta_1} \approx f_{\theta_2} \quad (\text{equivalent predictions, same test loss}) \quad \text{yet} \quad \text{argmax}_t(\alpha_{\theta_1})_t \neq \text{argmax}_t(\alpha_{\theta_2})_t.$$

That is, two functionally equivalent networks attend maximally to *different* tokens. This is not an artifact of optimization noise: Jain and Wallace [2019] show directly that attention weights can be dramatically permuted—including adversarially—while leaving model predictions unchanged, providing a sharp empirical demonstration that attention is not a faithful explanation in the sense of Definition 3.2.

Instance 2: Attention Impossibility.

Theorem 17.3 (Attention Map Impossibility). *The attention explanation system $\mathcal{S}_{\text{attn}}$ has the Rashomon property. Therefore, by Theorem 3.6, no attention map exp_{attn} can simultaneously be faithful (report the network’s actual attention), stable (assign the same dominant token to equivalent networks), and decisive (commit to a unique most-attended token).*

Proof. By the Rashomon property, there exist θ_1, θ_2 with $\text{obs}_{\text{attn}}(\theta_1) = \text{obs}_{\text{attn}}(\theta_2)$ (same input→output function) and $(\alpha_{\theta_1}, \alpha_{\theta_2}) \in \perp_{\text{attn}}$ (different argmax tokens). The result then follows immediately from Theorem 3.6. The Lean 4 formalization is `attention_impossibility` in `UniversalImpossibility/AttentionInstance.lean`. $\square$

Empirical illustration. We validate Theorem 17.3 using two protocols. Under **full retraining** (20 models, last 2 transformer layers retrained from independent seeds), prediction agreement is 94.5%, the argmax flip rate is **19.9%** [CI: 19.1–20.6%], and mean Kendall $\tau = 0.63$. This is the primary result: it measures instability arising from genuinely different training trajectories. As a secondary analysis, weight perturbation ($\sigma \in \{0.01, 0.02\}$, 10 models) yields prediction agreement of 100% and an argmax flip rate of 60.0% [CI: 59–61%], with mean Kendall $\tau = 0.30$. The perturbation flip rate is substantially higher because weight perturbation is not equivalent to retraining: it preserves the loss basin while injecting noise, overstating the degree of instability that practitioners would encounter in practice. Full results appear in Table 14 and Figure 7.




Figure 7: Attention map instability across 10 functionally equivalent DistilBERT models. Left: distribution of pairwise Kendall τ correlations. Right: argmax flip rate across model pairs.

Retraining vs. dropout proxy. The primary attention instability result uses 20 independently retrained DistilBERT models (19.9% argmax flip rate), which constitutes genuine model multiplicity and satisfies the Rashomon property. A computationally cheaper proxy—running inference with dropout active on a single model (16% of token-pair comparisons unreliable, Gaussian flip $R^2 = 0.689$)—produces qualitatively similar results but measures dropout noise rather than true Rashomon instability. The retraining result is the canonical demonstration; the dropout proxy is supplementary.

Implications for the attention-as-explanation debate. Theorem 17.3 gives a structural resolution to the debate initiated by Jain and Wallace [2019] and Wiegrefe and Pinter [2019]. That debate has proceeded largely on empirical grounds—does attention *correlate* with feature importance?—but the impossibility is not contingent on correlation structure. Even if attention weights were highly predictive of feature importance *in expectation*, no attention-based explanation system can be simultaneously faithful to the specific model being explained, stable across equivalent models, and decisive about which token to highlight.

The trilemma has a concrete consequence for practitioners: any attention visualization that commits to a single most-attended token (decisive) must either be unfaithful to some equivalent models or change its output when the model is retrained from a different random seed. This is not a failure of implementation but a mathematical inevitability whenever two equally predictive networks disagree on their peak-attention token.

The resolution prescribed by Section 5 applies directly: a G -invariant attention explanation aggregates attention distributions across the Rashomon set of equivalent models, yielding an averaged distribution $\bar{\alpha} = \frac{1}{|\mathcal{R}|} \sum_{\theta \in \mathcal{R}} \alpha_{\theta}$ that is faithful in expectation and stable by construction, but typically fails to be decisive (the averaged distribution may spread mass across multiple tokens, refusing to crown a single “most-attended” token). Reporting $\bar{\alpha}$ rather than α_{θ} is the attention analog of reporting CPDAG edge confidences rather than a single causal orientation—an honest acknowledgment that the data do not uniquely determine an attention explanation.

17.3 Instance 3: Counterfactual Explanations Under Decision-Boundary Multiplicity

Counterfactual explanations answer the question: “What is the smallest change to my input that would flip the model’s prediction?” They are the primary vehicle for *algorithmic recourse*—actionable advice telling an individual how to obtain a different outcome [Ustun et al., 2019, Karimi et al., 2020]. We show that no counterfactual explanation method can simultaneously be faithful, stable, and decisive when multiple models fit the data equally well.

The explanation system. We instantiate the abstract framework (Definition 3.1) as follows.

- $\Theta = \Theta$: the space of *model parameters*, each $\theta \in \Theta$ defining a decision boundary $\partial f_\theta = \{x \in \mathcal{X} : f_\theta(x) = \tau\}$ in the input space $\mathcal{X} \subseteq \mathbb{R}^P$ for a threshold τ .
- $\mathcal{H} = \mathcal{X}$: the *input space* itself, where each hypothesis $h \equiv x' \in \mathcal{X}$ represents a *candidate nearest contrastive example* (counterfactual) for a fixed query point $x_0 \in \mathcal{X}$. The attribution map $\theta : \mathcal{H} \rightarrow \Theta$ sends x' to the model parameters for which x' is the nearest point with a different prediction.
- Y : the space of *prediction vectors* on a held-out test set, i.e., $Y = \{0, 1\}^n$ for n test points. The observation map $\text{obs}(\theta) = f_\theta|_{\mathcal{D}_{\text{test}}}$ records the model’s predictions.

The *explanation map* $\text{exp}(\theta) = \arg \min_{x' : f_\theta(x') \neq f_\theta(x_0)} \|x' - x_0\|$ is the nearest counterfactual for query x_0 under model θ . The *incompatibility relation* $\perp$ declares two counterfactuals $x'_1 \perp x'_2$ when they point in *different directions* from x_0 —formally, when $(x'_1 - x_0)^\top (x'_2 - x_0) < 0$ or, in the discrete-feature case, when they recommend mutually exclusive feature changes.

The Rashomon property. The Rashomon effect for decision boundaries is well-documented: within a small loss tolerance ϵ , an exponential number of models may agree on test predictions while placing their decision boundaries in qualitatively different locations [Semenova et al., 2022, Fisher et al., 2019]. Formally, the Rashomon set $\mathcal{R}_\epsilon = \{\theta \in \Theta : \mathcal{L}(\theta) \leq \mathcal{L}(\theta^*) + \epsilon\}$ generically contains models whose boundaries straddle the query point x_0 from opposite sides. Concretely:

- **Model 1** (θ_1): decision boundary passes *above* x_0 in the relevant direction; the nearest counterfactual is x'_1 (move feature j up).
- **Model 2** (θ_2): decision boundary passes *below* x_0 ; the nearest counterfactual is x'_2 (move feature j down).

Both θ_1 and θ_2 achieve identical test-set predictions ($\text{obs}(\theta_1) = \text{obs}(\theta_2)$), yet $x'_1 \perp x'_2$. This is precisely the Rashomon property (Definition 3.5).

Proposition 17.4 (Instance 3: Counterfactual Impossibility). *Let $\mathcal{S}_{\text{CF}}$ be the counterfactual explanation system above. If $\mathcal{S}_{\text{CF}}$ has the Rashomon property—i.e., the Rashomon set $\mathcal{R}_\epsilon$ contains models with incompatible nearest counterfactuals for some query x_0 —then no counterfactual explanation map exp can simultaneously be faithful, stable, and decisive.*

Proof. Direct application of Theorem 3.6 to $\mathcal{S}_{\text{CF}}$. The Rashomon property supplies hypotheses θ_1, θ_2 with $\text{obs}(\theta_1) = \text{obs}(\theta_2)$ and $\text{exp}_{\text{CF}}(\theta_1) \perp \text{exp}_{\text{CF}}(\theta_2)$. The four-line proof of Theorem 3.6 then yields a contradiction from faithfulness, stability, and decisiveness simultaneously. This is mechanically verified as `counterfactual_impossibility` in `CounterfactualInstance.lean`. $\square$

Empirical illustration. We validate Proposition 17.4 by training 20 XGBoost models on the German Credit dataset (differing only in random seed), achieving a mean AUC of 0.797 with spread < 0.03 —confirming functional equivalence. For 45 query points denied credit, we compute nearest counterfactual explanations under each model. The **direction flip rate** is **23.5%** [95% CI: 20.6–26.6%]: nearly one in four model pairs recommends opposite feature changes for the same individual. **Cross-model recourse validity** is only **61.3%** [CI: 55.8–66.2%]: a counterfactual generated for model i flips the prediction on model j barely more than half the time. The most unstable feature is `checking_status` at a 43% flip rate. Full results appear in Table 15 and Figure 8.

Table 15: Counterfactual explanation instability on German Credit (20 equivalent XGBoost models, AUC within 0.03). *Negative control*: synthetic data with `class_sep=0.5` and `subsampling=1.0` — deterministic training means models are nearly identical; CFs are found and stable, expected flip rate < 5%. *Resolution*: consensus CF (feature-wise median direction) — fraction of individual CFs consistent with consensus. All 95% bootstrap CIs from 500 resamples.

Test	Metric	Mean	95% CI
Positive (Rashomon)	Mean AUC	0.7968	[0.7829, 0.8099]
	CF found rate	0.854	—
	Direction flip rate	0.235	[0.207, 0.263]
	Cross-model validity	61.3%	[56.1%, 66.4%]
	Distance CV	0.328	[0.293, 0.372]
Neg. control (<code>class_sep=0.5</code>)	CF found rate	0.100	—
	Direction flip rate	0.000	[0.000, 0.000]
Resolution (consensus CF)	Consistency with consensus	0.927	[0.898, 0.954]




Figure 8: Counterfactual explanation instability across 20 equivalent XGBoost models on German Credit. Left: direction flip rates by feature. Right: cross-model recourse validity distribution.

Implications for recourse and actionable explanations. Proposition 17.4 reveals a fundamental limit on the reliability of algorithmic recourse. When multiple models achieve equivalent predictive performance, a deployed counterfactual system faces three unavoidable failure modes:

1. **Recourse instability.** An individual who retrain a lender’s model (e.g., after a software update) may receive incompatible advice—“increase savings” vs. “reduce debt”—despite the model’s credit-scoring accuracy being unchanged. Empirical evidence of this instability is documented by Marx et al. [2024].
2. **Fragility under model multiplicity.** Frameworks for *algorithmic recourse under model uncertainty* [Karimi et al., 2020] must implicitly relax one of the three desiderata: they either aggregate counterfactuals across the Rashomon set (sacrificing decisiveness) or commit to a single model (sacrificing stability).
3. **Legal compliance risk.** Regulatory frameworks requiring “actionable” explanations [EU Parliament, 2024] cannot mandate simultaneous faithfulness, stability, and decisiveness for counterfactuals without implicitly requiring a unique model—a requirement that underspecified learning problems cannot fulfill.

The uniform resolution (Section 5) applies here: a G -invariant counterfactual system averages the nearest counterfactual direction over the Rashomon set, producing a *robust recourse*

direction that is faithful in expectation and stable, at the cost of yielding a distribution over directions rather than a single decisive recommendation. Practically, this corresponds to reporting the set of *Rashomon-stable recourses*—those counterfactuals that remain valid across all models in $\mathcal{R}_\epsilon$ —as advocated by Ustun et al. [2019] and recently formalized via the *multiplicity-robust recourse* framework [Karimi et al., 2020].

17.4 Instance 4: Concept Probe Impossibility

Motivation. Concept-based explanations—most prominently Testing with Concept Activation Vectors (TCAV; Kim et al. 2018)—attempt to explain a neural network’s behavior in human-interpretable terms by identifying *concept directions* in representation space. For example, the direction corresponding to “stripes” in an image classifier’s penultimate layer is found by fitting a linear probe that separates activations on striped versus non-striped inputs. This direction is then interpreted as the network’s internal encoding of the concept, and the degree to which prediction scores change along it is reported as the explanation.

The hope is that concept directions are a stable, faithful window into model internals. We show this hope is mathematically impossible under underspecification: two models with identical predictions can encode the same concept in directions that are genuinely incompatible, so no concept probe can simultaneously be faithful, stable, and decisive.

Instantiation. We instantiate Definition 3.1 as follows.

Definition 17.5 (Concept Probe Explanation System). *Let $c \in \mathbb{N}$ be the dimension of the representation space. The concept probe explanation system is the tuple $\mathcal{S}_{\text{cp}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ where:*

- $\Theta = \Theta$ is the space of model weight configurations, each $\theta \in \Theta$ specifying a neural network together with its internal representations;
- $\mathcal{H} = \mathbb{R}^c$ is the space of concept activation vectors—unit directions in the c -dimensional representation space that a linear probe associates with a target concept;
- Y is the space of predictions (input-output functions $f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$), serving as the observable behavior of the model;
- $\text{obs}(\theta) = f_\theta$ is the observation map, recording the input-output function computed by weights θ ;
- $\text{exp}(\theta) = v_\theta \in \mathbb{R}^c$ is the explanation map, returning the concept activation direction extracted by a linear probe on the representations of θ ; and
- $(v, v') \in \perp$ iff v and v' are not related by a rotation within the concept subspace—that is, they designate structurally distinct directions that cannot both be reported as the concept direction without contradiction.

The Rashomon Property for concept probes. The empirical literature provides direct evidence that the Rashomon property holds for concept directions. Bolukbasi et al. [2016] demonstrate that gender directions vary substantially across word-embedding architectures trained on the same corpus: models with identical downstream task performance encode gender in directions that differ in orientation, dimension, and linearity. Kim et al. [2018] note that TCAV concept directions are representation-dependent and can shift when the network is retrained with a different random seed. Ravfogel et al. [2020] show via iterative nullspace projection that a single concept is not localized to one direction: removing the dominant linear probe direction

still leaves residual concept structure, and the secondary directions are mutually incompatible with the primary.

Formally, we require:

Definition 17.6 (Rashomon Property for Concept Probes). *There exist weight configurations $\theta_1, \theta_2 \in \Theta$ such that*

$$f_{\theta_1} = f_{\theta_2} \quad (\text{same predictions}) \quad \text{and} \quad (v_{\theta_1}, v_{\theta_2}) \in \perp \quad (\text{incompatible concept directions}).$$

This is the formal content of Definition 3.5 in $\mathcal{S}_{\text{cp}}$: equivalent models (same observable behavior, hence $\text{obs}(\theta_1) = \text{obs}(\theta_2)$) report incompatible concept directions. The condition holds in any overparameterized network class—the zero-loss manifold has dimension $\geq d - n > 0$ for hidden dimension $d > n$ training points, so there are infinitely many weight configurations with identical input-output behavior but geometrically distinct internal representations.

The concept probe impossibility.

Theorem 17.7 (Concept Probe Impossibility). *If the concept probe explanation system $\mathcal{S}_{\text{cp}}$ has the Rashomon property, then no explanation map exp can simultaneously be faithful, stable, and decisive.*

Proof. Immediate from Theorem 3.6 applied to $\mathcal{S}_{\text{cp}}$. Concretely: let θ_1, θ_2 witness the Rashomon property, so $\text{obs}(\theta_1) = \text{obs}(\theta_2)$ and $(v_{\theta_1}, v_{\theta_2}) \in \perp$. Faithfulness forces $\text{exp}(\theta_i) = v_{\theta_i}$; decisiveness requires each $\text{exp}(\theta_i)$ to be a specific direction; stability forces $\text{exp}(\theta_1) = \text{exp}(\theta_2)$. But then $v_{\theta_1} = v_{\theta_2}$, contradicting $(v_{\theta_1}, v_{\theta_2}) \in \perp$. $\square$

The Lean 4 formalization of this instance is in `ConceptInstance.lean`. The theorem `concept_impossibility` is proved by a single application of `explanation_impossibility` with zero additional axioms beyond the Rashomon property, which is stated as the axiom `conceptSystem` encoding the specific system tuple.

Empirical illustration. We validate Theorem 17.7 by training 15 MLP classifiers (architecture 128–64, ReLU) on `sklearn load_digits` with different random seeds. All models achieve $> 95\%$ test accuracy with a **prediction agreement of 97.2%** [95% CI: 95.7–98.5%], confirming functional equivalence. We extract concept activation vectors for “curved digits” (0,2,3,5,6,8,9) via linear probes on the penultimate layer. The mean $|\cos(\mathbf{v}_i, \mathbf{v}_j)|$ across model pairs is only **0.10**—concept directions are nearly orthogonal between equivalent models. In 64 dimensions, random unit vectors have expected absolute cosine similarity of approximately 0.10; the observed value matches this baseline, consistent with the interpretation that hidden representations encode concepts in effectively random directions — precisely the representation arbitrariness predicted by the Rashomon property. The **concept direction instability** ($1 - |\cos|$) is **0.90** [CI: 0.89–0.91], and TCAV score standard deviation is **0.16**, meaning the same concept is scored anywhere from 0.10 to 0.69 depending on the random seed. Full results appear in Table 16 and Figure 9.

Implications for probing-as-understanding. Theorem 17.7 has direct consequences for the practice of using concept probes as explanations of model internals.

1. *Concept directions are not model-invariant facts.* A TCAV direction is a fact about a specific weight configuration, not about the function the network computes. Two retraining runs that produce functionally identical models (same accuracy, same predictions on held-out data) may produce concept directions that are not related by rotation and therefore cannot be reconciled into a single explanation.

Table 16: Concept probe instability across 15 equivalent MLPClassifier models on `sklearn load_digits()` (8×8 images, 10 classes). Architecture: (128, 64) hidden units, ReLU. *Random baseline*: 15 random unit vectors in 64D — expected $|\cos| \approx \text{real CAVs}$ if directions are arbitrary. *Negative control*: logistic regression (no hidden layers) — directions should be stable (> 0.95). *Resolution*: averaged CAV across 15 models — individual CAVs should be closer to the average. All 95% CIs from 2000 bootstrap resamples.

Test	Metric	Value	95% CI
Positive (Rashomon)	Mean test accuracy	0.9799	[0.9741, 0.9852]
	Min test accuracy	0.9741	—
	Mean $ \cos(\mathbf{v}_i, \mathbf{v}_j) $ (off-diag)	0.1000	[0.8854, 0.9139]
	Concept direction instability	0.9000	—
	Prediction agreement	0.9722	[0.9574, 0.9852]
Random baseline (15 rand. vecs, 64D)	Mean $ \cos(\mathbf{v}_i, \mathbf{v}_j) $	0.0898	[0.0773, 0.1028]
Neg. control (LogReg, no hidden)	Mean $ \cos(\mathbf{v}_i, \mathbf{v}_j) $	1.0000	[1.0000, 1.0000]
Resolution (avg. CAV)	Mean $ \cos(\text{indiv}, \text{avg}) $	0.2790	[0.2263, 0.3273]
	Mean pairwise $ \cos(i, j) $	0.1000	[0.0857, 0.1146]




Figure 9: Concept probe instability across 15 equivalent MLP models on `load_digits`. Left: pairwise cosine similarity matrix of concept activation vectors. Right: TCAV score distribution.

2. *Faithfulness and stability cannot coexist.* Requiring a probe to faithfully report the concept direction of the model it was fitted to forces instability across retraining seeds. Requiring stability forces the probe to report a direction that is unfaithful to at least one model in the Rashomon set.
3. *The resolution entails indecisiveness.* The G -invariant resolution (Section 5) applied here would average concept directions across equivalent models, yielding a distribution over directions or a subspace rather than a single vector. This restores faithfulness-in-expectation and stability at the cost of decisiveness: the explanation no longer commits to a single direction, which undermines the interpretive clarity that motivates concept probing in the first place.

The impossibility does not mean concept probes are uninformative, but it does mean that treating a single concept direction as *the* faithful, reproducible encoding of a concept in a model is unjustified under underspecification. Practitioners should report concept directions together

with their variability across the Rashomon set—analogously to how DASH reports SHAP values with their ensemble distribution rather than as a single point estimate.

17.5 Instance 5: Causal Discovery Within a Markov Equivalence Class

We instantiate Theorem 3.6 for causal structure learning, where the fundamental obstruction is Markov equivalence: distinct directed acyclic graphs (DAGs) can be observationally indistinguishable yet disagree on causal edge directions.

The explanation system. Define the 6-tuple $\mathcal{S}_{\text{causal}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows.

- Θ is the set of DAGs over a fixed vertex set V (the space of possible causal structures, i.e., data-generating processes).
- $\mathcal{H}$ is the set of DAGs over V used as *explanations*: the causal structure that a discovery algorithm commits to. The attribution map $\theta : \mathcal{H} \rightarrow \Theta$ sends each explanation DAG to the underlying data-generating DAG it was fitted to.
- Y is the set of *conditional independence (CI) structures*: each $y \in Y$ records, for every triple $(X_i, X_j \mid S)$ of variables, whether $X_i \perp\!\!\!\perp X_j \mid S$ holds in the associated distribution. Equivalently, Y indexes Markov equivalence classes.
- $\text{obs} : \mathcal{H} \rightarrow Y$ maps each DAG $h \in \mathcal{H}$ to the CI structure it encodes (the set of d -separation statements implied by h via the Markov condition). Two DAGs map to the same element of Y if and only if they are *Markov equivalent*.
- $\text{exp} : \mathcal{H} \rightarrow \mathcal{H}$ is an *orientation rule*: given a DAG $h \in \mathcal{H}$ as fitted by an algorithm, the rule commits to a fully oriented causal graph. In practice, exp is any causal discovery procedure (PC algorithm, GES, FCI, etc.) that returns a single DAG rather than an equivalence class.
- $\perp \subseteq \mathcal{H} \times \mathcal{H}$ is the *edge-reversal incompatibility*: $(h, h') \in \perp$ if there exists a directed edge $X \rightarrow Y$ in h and $Y \rightarrow X$ in h' (or vice versa). Two causal explanations are incompatible when they assert opposite orientations for at least one edge.

The Rashomon property: Markov equivalence. The key structural fact is that Markov equivalence classes are non-trivial for nearly all graph structures of practical interest.

Definition 17.8 (Markov Equivalence). *Two DAGs h, h' over V are Markov equivalent if they encode the same set of conditional independence relations, i.e., $\text{obs}(h) = \text{obs}(h')$. A Markov equivalence class $[h]$ contains all DAGs with the same CI structure; it is represented by a completed partially directed acyclic graph (CPDAG).*

Verma and Pearl [1991] established that two DAGs are Markov equivalent if and only if they share the same skeleton (undirected graph) and the same set of v-structures (immoralities $X \rightarrow Z \leftarrow Y$ with no edge between X and Y). Consequently, any undirected edge in the CPDAG corresponds to a reversible orientation: there exist two Markov-equivalent DAGs that orient it in opposite directions. Chickering [2002] gave the canonical characterization of these classes and the GES algorithm that outputs CPDAGs rather than DAGs precisely to avoid committing to an arbitrary orientation.

Proposition 17.9 (Causal Rashomon Property). *$\mathcal{S}_{\text{causal}}$ has the Rashomon property: for any CPDAG containing at least one undirected edge $\{X, Y\}$, there exist DAGs $h, h' \in \mathcal{H}$ with $\text{obs}(h) = \text{obs}(h')$ (same CI structure) yet $(h, h') \in \perp$ (opposite orientations of $X-Y$).*

Proof. Let h be any DAG in the equivalence class orienting the edge as $X \rightarrow Y$, and let h' be the DAG obtained by reversing that edge to $Y \rightarrow X$ (which exists in the equivalence class precisely because the edge is undirected in the CPDAG). By Markov equivalence, $\text{obs}(h) = \text{obs}(h')$. By definition of $\perp$, $(h, h') \in \perp$. $\square$

This is the formalization captured in `CausalInstance.lean` and `CausalDiscovery.lean`: the `rashomon` field of `causalSystem` is a concrete witness pair (h, h') from Proposition 17.9 (specifically, the 3-node chain vs. fork). *Note:* The Lean formalization uses the broader relation $\text{incompatible}(h_1, h_2) \iff h_1 \neq h_2$ rather than the edge-reversal relation defined above. Since $(h_1 \neq h_2)$ implies edge-reversal incompatibility for Markov-equivalent DAGs with a reversed edge, the impossibility under the broader relation implies the impossibility under the narrower one *a fortiori*. The paper’s edge-reversal definition is the scientifically meaningful one; the Lean code uses the broader relation for simplicity of formalization.

Instance 5: causal discovery impossibility.

Theorem 17.10 (Causal Discovery Impossibility — Instance 5). *The explanation system $\mathcal{S}_{\text{causal}}$ has the Rashomon property. Therefore, by Theorem 3.6, no orientation rule exp can simultaneously be:*

1. **Faithful:** $\text{exp}(h) = h$ for every DAG $h \in \mathcal{H}$ (the returned DAG matches the true data-generating structure);
2. **Stable:** $\text{exp}(h) = \text{exp}(h')$ whenever $\text{obs}(h) = \text{obs}(h')$ (Markov-equivalent structures receive identical causal explanations); and
3. **Decisive:** exp always commits to a single fully oriented DAG (never returns a partial orientation or equivalence class).

Proof. By Proposition 17.9, $\mathcal{S}_{\text{causal}}$ has the Rashomon property. Apply Theorem 3.6. $\square$

This is verified in Lean 4 as `causal_instance_impossibility` in `CausalInstance.lean` (three axioms), and as the self-contained `causal_discovery_impossibility` in `CausalDiscovery.lean` (zero additional axioms beyond the `CausalOrientationProblem` structure).

Empirical illustration. Three causal discovery algorithms (PC with $\alpha = 0.05$, PC with $\alpha = 0.01$, GES) on 1,000 samples from the Asia network show overall edge agreement of only 0.17 and *orientation* agreement of 0.33 (Mann-Whitney $p = 4.4 \times 10^{-38}$ across 100 random seeds per condition). At $N = 100,000$, orientation agreement reaches 1.0 (all algorithms agree on every directed edge), though overall agreement is 0.625 because the three undirected edges in the CPDAG are counted as non-agreeing by the overall metric. This distinction is important: the algorithms converge to the *same CPDAG*, confirming that finite-sample ambiguity (not algorithmic limitation) drives the small- N disagreement.

Connection to the classical Markov equivalence problem. Theorem 17.10 is the classical Markov equivalence underdetermination recast in the universal framework. The three properties above correspond exactly to the three goals that any “complete” causal discovery method would need to achieve: return the true DAG (faithfulness), be invariant to observationally equivalent parameterizations (stability), and produce a unique answer (decisiveness). The Markov equivalence theorem of Verma and Pearl [1991] has long been understood to block this combination; our contribution is to show it is a *special case* of the single four-line proof of Theorem 3.6.

Resolution: CPDAGs as G -invariant explanations. The uniform resolution of Section 5 specializes to causal discovery as follows. Let $G = [h]$ denote the Markov equivalence class (the symmetry group acting by edge reversals on the set of reversible edges). The G -invariant explanation map aggregates over all member DAGs:

$$\text{exp}^*(h) = \text{CPDAG}([h]),$$

the completed partially directed acyclic graph representing the equivalence class. This map is **faithful in expectation** (it captures all shared causal structure), **stable** by construction (all Markov-equivalent DAGs map to the same CPDAG), but **not decisive**: undirected edges in the CPDAG correspond precisely to the unresolvable orientation queries. Reporting the CPDAG is therefore the principled response to the impossibility—trading decisiveness for faithfulness and stability, exactly as predicted by Theorem 3.6. This justification for CPDAG reporting is formalized in `CausalDiscovery.lean` as `cpdag_is_stable`.

17.6 Instance 6: Model Selection Under the Rashomon Effect

Model selection—choosing a single best model from a trained set—is perhaps the most commonplace decision in applied machine learning. Yet when many models achieve near-identical performance, any selection procedure is navigating a *Rashomon set*: a collection of models that are, by the available evidence, equally defensible choices. We show that this setting is an instance of Theorem 3.6, and that the resulting instability is not a failure of any particular selection criterion but a mathematical inevitability.

Explanation system. We instantiate the abstract framework as follows.

Definition 17.11 (Model Selection Explanation System). *The model selection explanation system is the tuple $\mathcal{S}_{\text{sel}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ where:*

- $\Theta = \Theta$ is the space of model configurations (hyperparameter settings, architecture choices, training seeds, or any parameterization of the model class under comparison);
- $\mathcal{H}$ is the hypothesis space of trained models; each $h \in \mathcal{H}$ corresponds to a configuration $\theta(h) \in \Theta$ via the attribution map $\theta : \mathcal{H} \rightarrow \Theta$;
- Y is the space of performance metrics (validation loss, accuracy, AUC, or any scalar summary of model quality on a held-out evaluation set); formally, Y is equipped with a strict total order $>$ representing “performs better than”;
- $\text{obs} : \mathcal{H} \rightarrow Y$ maps each hypothesis to its observed performance metric on the evaluation set;
- $\text{exp} : \mathcal{H} \rightarrow Y$ is the selection rule—the criterion used to rank or choose among models; and
- $\perp \subseteq Y \times Y$ is defined by $(y, y') \in \perp$ whenever $y \neq y'$ and the two metrics imply opposite selection decisions (i.e., the model recommended under y is strictly disfavored under y').

The Rashomon property for model selection. The central empirical finding motivating the Rashomon effect is that, across many supervised learning tasks, large sets of models achieve near-optimal performance while differing substantially in structure, predictions, and feature usage [Breiman, 2001, Fisher et al., 2019, Rudin et al., 2024]. Formally, define the ε -Rashomon set as

$$\mathcal{R}_\varepsilon = \{h \in \mathcal{H} : \text{obs}(h) \geq \text{obs}(h^*) - \varepsilon\},$$

where h^* is the empirical risk minimizer and $\varepsilon > 0$ is a tolerance. The Rashomon property holds for $\mathcal{S}_{\text{sel}}$ whenever $\mathcal{R}_\varepsilon$ contains pairs of models that are (i) indistinguishable in performance ($\theta(h) = \theta(h')$ in the sense that they share the same configuration class) yet (ii) produce incompatible selection metrics on different evaluation samples.

Proposition 17.12 (Rashomon Property for Model Selection). *If the ε -Rashomon set $\mathcal{R}_\varepsilon$ contains two models $h, h' \in \mathcal{H}$ with $\theta(h) = \theta(h')$ such that $\text{obs}(h) \neq \text{obs}(h')$, then $\mathcal{S}_{\text{sel}}$ satisfies the Rashomon property (Definition 3.5).*

This is immediate from the definition: $\text{obs}(h)$ and $\text{obs}(h')$ are incompatible (one favors h , the other h'), and both are witnessed by models sharing the same underlying configuration. The condition is generically satisfied: Fisher et al. [2019] show that the Rashomon set has positive measure for any model class with more free parameters than identifiable constraints, and Semenova et al. [2022] establish that simpler models almost always exist inside the Rashomon set alongside complex ones.

Lean formalization. The model selection impossibility is mechanically verified in two Lean 4 files. `ModelSelection.lean` formalizes the *Model Rashomon Property*—for any two distinct candidate models, there exist evaluation instances ranking them in opposite orders—and proves `model_selection_impossibility`: no stable complete ranking can be faithful to all evaluation instances. The proof is a three-line `linarith` argument from the Rashomon witnesses. `ModelSelectionInstance.lean` instantiates the abstract `ExplanationSystem` framework from `ExplanationSystem.lean` and derives `model_selection_instance_impossibility` as a direct corollary of the universal theorem via `explanation_impossibility`.

Instance 6 (Model Selection Impossibility).

Theorem 17.13 (Model Selection Impossibility). *The model selection explanation system $\mathcal{S}_{\text{sel}}$ (Definition 17.11) satisfies the Rashomon property whenever the ε -Rashomon set has positive measure. Consequently, by Theorem 3.6, no selection rule $\text{exp} : \mathcal{H} \rightarrow Y$ can simultaneously be:*

1. **Faithful:** $\text{exp}(h) = \text{obs}(h)$ for all h (the selected model is the one the evaluation metric actually endorses);
2. **Stable:** $\text{exp}(h) = \text{exp}(h')$ whenever $\theta(h) = \theta(h')$ (the same configuration always yields the same selection); and
3. **Decisive:** $\text{exp}(h)$ is a strict recommendation for every h (no ties or abstentions are reported).

Design space. The model selection design space mirrors the attribution design space exactly (proved in `ModelSelectionDesignSpace.lean`): every selection rule falls into one of two families.

- **Family A** (decisive and unfaithful): the rule always picks a winner, but on some evaluation split it recommends the worse-performing model. Standard single-criterion selection rules (e.g., argmax validation accuracy) fall here whenever the Rashomon set is non-trivial.
- **Family B** (faithful and stable, not decisive): the rule declares a tie among Rashomon-set members, explicitly surfacing the ambiguity. This corresponds to reporting the Rashomon set $\mathcal{R}_\varepsilon$ rather than a single selected model.

Implications. Theorem 17.13 formalizes and sharpens the empirical observation of Fisher et al. [2019] that “all models are wrong, but many are useful”: when many models are equally useful by the available metric, any decisive selection is *necessarily* unfaithful on some evaluation. The instability practitioners observe when re-running AutoML pipelines, changing random seeds, or switching validation folds is not a software bug—it is the Rashomon property in action.

The resolution is the same as for feature attribution (Section 17.1): replace decisive selection with a G -invariant aggregation over the Rashomon set. Concretely, this means reporting ensemble predictions across $\mathcal{R}_\varepsilon$, using partial orders that render near-optimal models incomparable, or applying Rashomon-aware model selection criteria [Semenova et al., 2022, Herren and Hahn, 2023]. Each of these trades decisiveness for the pair (faithful, stable), which Theorem 3.6 shows is the only achievable combination under underspecification.

17.7 Instance 7: Saliency Maps (GradCAM) Under Weight-Space Multiplicity

Gradient-weighted Class Activation Mapping (GradCAM) [Selvaraju et al., 2017] is the dominant saliency method for convolutional neural networks: it highlights spatial regions that drive a classification decision by weighting the final convolutional feature maps by their class-specific gradients. GradCAM saliency maps are widely used for model debugging, clinical imaging, and regulatory documentation. We show that no saliency map explanation can be faithful, stable, and decisive under weight-space underspecification.

The explanation system. Let $H \times W$ denote the spatial resolution of the final convolutional feature maps. We instantiate the abstract framework (Definition 3.1) as the 6-tuple $\mathcal{S}_{\text{sal}} = (\Theta_{\text{sal}}, \mathcal{H}_{\text{sal}}, Y_{\text{sal}}, \text{obs}_{\text{sal}}, \text{exp}_{\text{sal}}, \perp_{\text{sal}})$ where:

- Θ_{sal} is the space of CNN weight configurations θ (convolutional filters, batch-norm parameters, classifier head);
- $\mathcal{H}_{\text{sal}} = [0, 1]^{H \times W}$ is the space of ReLU-activated GradCAM heatmaps, normalized to $[0, 1]$;
- Y_{sal} is the space of input-to-output prediction functions $f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$;
- $\text{obs}_{\text{sal}}(\theta) = f_\theta$ records the observable prediction behavior;
- $\text{exp}_{\text{sal}}(\theta)$ is the GradCAM heatmap produced by θ on a fixed input image x_0 ; and
- $(h_1, h_2) \in \perp_{\text{sal}}$ when $\arg\max_{(i,j)} h_1 \neq \arg\max_{(i,j)} h_2$ —the two heatmaps place peak activation at different spatial locations.

The Rashomon property. Sanity checks for saliency methods [Adebayo et al., 2018] demonstrate that saliency maps from networks that differ only in their random initialization can diverge substantially while maintaining identical test accuracy. Hooker et al. [2019] further show that saliency-guided feature removal affects models inconsistently across retraining runs. These findings establish the Rashomon property for $\mathcal{S}_{\text{sal}}$: there exist weight configurations θ_1, θ_2 with $f_{\theta_1} \approx f_{\theta_2}$ yet peak GradCAM activation at different spatial locations.

Proposition 17.14 (Instance 7: Saliency Map Impossibility). *If $\mathcal{S}_{\text{sal}}$ has the Rashomon property, then no saliency map explanation can simultaneously be faithful, stable, and decisive.*

Proof. Direct application of Theorem 3.6 to $\mathcal{S}_{\text{sal}}$. □

Empirical illustration. We validate Proposition 17.14 by creating 10 perturbed ResNet-18 models on CIFAR-10 ($\sigma = 0.0005$ Gaussian weight noise). With **78.8% prediction agreement** (confirming approximate functional equivalence; prediction agreement is 78.8%, reflecting sensitivity of ImageNet-pretrained ResNet-18 to perturbation on CIFAR-10), the **peak flip rate** is **9.6%** [95% CI: 8.8–10.5%]: nearly one in ten model pairs places peak GradCAM activation at a different spatial location despite the same prediction. The mean pairwise spatial IoU (top-20% threshold) is **95.4%** [CI: 95.3–95.5%], confirming high but imperfect overlap. The negative control (identical weights) yields 100% IoU and 0% flip rate. The resolution test (averaged GradCAM heatmaps) yields a mean IoU of **96.9%**—a **+1.5%** improvement over pairwise IoU, confirming that orbit averaging reduces instability. Full results appear in Table 17.

Table 17: GradCAM Saliency Map Instability (ResNet-18, CIFAR-10). Ten functionally equivalent models (pretrained ResNet-18 + Gaussian perturbations, $\sigma = 0.0005$). Spatial IoU uses top-20% threshold; peak flip rate counts pairs where argmax differs by >2 pixels. Negative control uses identical weights (model 0 for all 10 slots). Resolution test averages heatmaps; mean IoU vs. average exceeds pairwise IoU. 95% bootstrap CIs in parentheses.

Metric	Positive	Control	Resolution
Prediction agreement	0.788 (0.752, 0.822)	1.000	—
Mean pairwise IoU	0.954 (0.953, 0.955)	1.000 (1.000, 1.000)	0.969 (0.968, 0.971)
Peak flip rate	0.096 (0.088, 0.105)	0.000 (0.000, 0.000)	—

Moderate instability compared to attention. The 9.6% peak flip rate for GradCAM is substantially lower than the 19.9% full-retraining argmax flip rate for attention maps (Section 17.2). This disparity reflects the structural constraint imposed by spatial locality: GradCAM operates on 7×7 or 8×8 feature maps where neighboring activations are highly correlated, whereas token-level attention distributes mass across an unstructured sequence. Spatial smoothness constrains the degree to which perturbations can rearrange saliency, producing a tighter Rashomon set for the peak-location incompatibility relation. Nevertheless, the impossibility still holds: even moderate instability suffices for the Rashomon property, and the 9.6% flip rate is statistically significant with a bootstrap CI excluding zero.

17.8 Instance 8: Token Citation Explanations Under Model Multiplicity

Large language models are increasingly asked to explain their own predictions: “Classify this review as positive or negative and explain which words drove your decision.” We formalize such self-explanations as citation-based summaries—template explanations citing the top- k most-attended tokens—and show that they inherit the impossibility of Theorem 3.6. Recent work on unfaithful chain-of-thought reasoning [Turpin et al., 2024] and measuring faithfulness of LLM explanations [Lanham et al., 2023] motivates the concern that LLM self-explanations are unreliable even for models that agree on predictions.

The explanation system. We instantiate the abstract framework (Definition 3.1) as the 6-tuple $\mathcal{S}_{\text{llm}} = (\Theta_{\text{llm}}, \mathcal{H}_{\text{llm}}, Y_{\text{llm}}, \text{obs}_{\text{llm}}, \text{exp}_{\text{llm}}, \perp_{\text{llm}})$ where:

- Θ_{llm} is the space of transformer weight configurations θ (attention heads, feed-forward layers, embeddings);
- $\mathcal{H}_{\text{llm}} = \binom{[T]}{k}$ is the set of size- k subsets of token positions, representing the k tokens cited in the explanation (we use $k = 3$);

- Y_{llm} is the prediction space (e.g., binary sentiment labels);
- $\text{obs}_{\text{llm}}(\theta)$ is the predicted label for input sentence x ;
- $\text{exp}_{\text{llm}}(\theta)$ extracts the top- k content tokens by attention rollout importance as the “cited evidence”; and
- $(S_1, S_2) \in \perp_{\text{llm}}$ when the highest-ranked token differs: $\max(S_1) \neq \max(S_2)$ under the attention ordering—i.e., the explanations disagree on the primary evidence token.

The Rashomon property. The attention rollout instability documented in Section 17.2 directly implies the Rashomon property for $\mathcal{S}_{\text{llm}}$: models with identical predictions but different attention distributions cite different tokens as primary evidence. Ye and Durrett [2022] demonstrate empirically that LLM-generated explanations are unreliable, changing substantially under minor prompt or model variations while predictions remain stable. Turpin et al. [2024] show that chain-of-thought explanations can be systematically unfaithful to the model’s actual reasoning process.

Proposition 17.15 (Instance 8: Token Citation Impossibility). *If $\mathcal{S}_{\text{llm}}$ has the Rashomon property, then no token citation explanation method can simultaneously be faithful, stable, and decisive. Token citations are extracted via attention rollout as a proxy for language model self-explanation.*

Proof. Direct application of Theorem 3.6 to $\mathcal{S}_{\text{llm}}$. □

Empirical illustration. We validate Proposition 17.15 using 10 perturbed DistilBERT models ($\sigma = 0.01$, matching the attention experiment) on 100 sentiment sentences. Despite **100% prediction agreement**, the **explanation flip rate** (fraction of model pairs citing a different primary token) is **34.5%** [95% CI: 33.2–35.9%], and the mean pairwise **citation overlap** (Jaccard similarity of top-3 token sets) is only **0.798** [CI: 0.791–0.806]. The **negative control** (20 single-content-token sentences, each tokenizing to [CLS] word [SEP], forcing a single citation target) achieves perfect citation overlap (Jaccard = 1.000), confirming that instability is concentrated in genuinely ambiguous sentences. The **resolution test** (consensus top-3 from averaged attention rollout) yields Jaccard overlap of **0.857** [CI: 0.842–0.871], a **+5.8%** improvement over pairwise overlap, confirming that orbit averaging reduces citation instability. Full results appear in Table 18.

Implications for regulatory requirements on LLM explanations. The EU AI Act [EU Parliament, 2024] requires “meaningful explanations” for high-risk AI systems, and practitioners increasingly rely on LLM self-explanations to satisfy this requirement. Proposition 17.15 and the 34.5% flip rate demonstrate that such explanations are structurally unreliable: retraining (or fine-tuning) the same architecture on the same data can change which tokens the model cites as evidence, even when the prediction itself is unchanged. This is not a failure of any particular prompting strategy or explanation method; it is a mathematical consequence of underspecification. Regulators should require that LLM explanations be accompanied by stability audits—e.g., reporting citation overlap across multiple model instances—rather than treating a single model’s self-explanation as ground truth.

17.9 Instance: Mechanistic Interpretability (Circuit Discovery)

Mechanistic interpretability (MI) has emerged as a central paradigm for AI safety research, aiming to reverse-engineer neural networks into human-understandable circuits [Elhage et al., 2022, Conmy et al., 2023, Anthropic, 2025]. We instantiate Theorem 3.6 for circuit discovery,

Table 18: LLM Explanation Instability (Attention-Based Fallback). 10 perturbed DistilBERT models ($\sigma = 0.01$), 100 sentences. “Citation overlap” = Jaccard similarity of top-3 cited tokens. “Flip rate” = fraction of model pairs where the #1 cited token differs. 95% CIs via bootstrap.

Metric	Value	95% CI
Prediction agreement	100.0%	—
<i>Positive test (all 100 sentences)</i>		
Citation overlap (Jaccard)	0.798	[0.791, 0.806]
Explanation flip rate (#1 token)	0.345	[0.332, 0.359]
<i>Negative control (20 single-content-token sentences: forced top-1 citation)</i>		
Citation overlap (Jaccard)	1.000	[1.000, 1.000]
Flip rate	0.000	[0.000, 0.000]
<i>Resolution test (consensus top-3 vs individual top-3)</i>		
Consensus overlap (Jaccard)	0.857	[0.843, 0.871]

where the fundamental obstruction is that functionally equivalent networks admit multiple valid circuit decompositions.

The explanation system. Define the 6-tuple $\mathcal{S}_{\text{circuit}} = (\Theta, \mathcal{H}, Y, \text{obs}, \text{exp}, \perp)$ as follows.

- Θ is the set of neural network weight configurations (architecture plus trained parameters).
- $\mathcal{H}$ is the set of *circuit decompositions*: subgraphs of the computational graph together with an interpretation of which components implement which sub-computations (e.g., “heads 1.3 and 2.1 form an induction circuit”).
- Y is the set of input–output functions: each $y \in Y$ is the function $f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$ computed by the network on the domain of interest.
- $\text{obs} : \Theta \rightarrow Y$ maps each weight configuration θ to the function it computes.
- $\text{exp} : \Theta \rightarrow \mathcal{H}$ maps each network to its circuit decomposition—the internal computational mechanism identified by a circuit discovery method.
- $\perp \subseteq \mathcal{H} \times \mathcal{H}$ is the *component-attribution incompatibility*: $(h, h') \in \perp$ if h and h' attribute the same computation to disjoint or contradictory subsets of network components (e.g., h says the computation is implemented by heads in layer 2, while h' says it is implemented by heads in layer 4).

The Rashomon property: circuit multiplicity. [Méloux et al. \[2025\]](#) provide striking empirical evidence that the Rashomon property holds for circuit discovery. For a simple XOR task, they identified **85 distinct valid circuits** with *zero circuit error* (each circuit perfectly reproduces the network’s behavior on the task). Moreover, each circuit admitted an average of **535.8 valid interpretations**—different assignments of computational roles to components, all equally faithful to the network’s input–output behavior.

Complementary evidence comes from [Paulo and Belrose \[2026\]](#), who show that sparse autoencoders trained on the *same* data learn *different* features depending on the random seed, demonstrating that even the feature-level decomposition underlying circuit identification is not unique.

Proposition 17.16 (Circuit Rashomon Property). $\mathcal{S}_{\text{circuit}}$ has the Rashomon property: there exist networks $\theta_1, \theta_2 \in \Theta$ with $\text{obs}(\theta_1) = \text{obs}(\theta_2)$ (identical I/O behavior) yet $(\text{exp}(\theta_1), \text{exp}(\theta_2)) \in \perp$ (incompatible circuit decompositions).

Proof. Consider any pair of distinct valid circuits from the 85 found by [Méloux et al. \[2025\]](#): by construction, both achieve zero circuit error (same I/O behavior on the task), yet they attribute the computation to different subsets of network components. $\square$

Instance: mechanistic interpretability impossibility.

Theorem 17.17 (Mechanistic Interpretability Impossibility). *The explanation system $\mathcal{S}_{\text{circuit}}$ has the Rashomon property. Therefore, by Theorem 3.6, no circuit explanation method exp can simultaneously be:*

1. **Faithful:** $\text{exp}(\theta)$ reflects the actual internal circuitry of network θ ;
2. **Stable:** $\text{exp}(\theta_1) = \text{exp}(\theta_2)$ whenever $\text{obs}(\theta_1) = \text{obs}(\theta_2)$ (functionally equivalent networks receive the same circuit explanation); and
3. **Decisive:** exp commits to a single circuit decomposition (rather than reporting the space of valid circuits).

Proof. By Proposition 17.16, $\mathcal{S}_{\text{circuit}}$ has the Rashomon property. Apply Theorem 3.6. $\square$

This is verified in Lean 4 as `mech_interp_impossibility` in `MechInterpInstance.lean` (three type axioms plus the `circuitSystem` axiom encoding the Rashomon witness).

Critical framing: what this does and does not say. This result does *not* undermine mechanistic interpretability—it clarifies what MI can and cannot achieve. Circuit discovery methods such as activation patching [[Conmy et al., 2023](#)], sparse probing [[Elhage et al., 2022](#)], and circuit tracing [[Anthropic, 2025](#)] remain powerful tools for understanding neural network internals. What the impossibility establishes is that the *search for THE circuit*—a single, definitive decomposition—is provably constrained: any method that commits to a unique answer must sacrifice either faithfulness or stability.

The constructive path forward is to *characterize the space of valid circuits*, not to search for a single canonical one. This is precisely the direction that leading MI work is already moving toward: [Anthropic \[2025\]](#) report confidence levels rather than single circuits, and [Conmy et al. \[2023\]](#) frame automated circuit discovery as search over a structured space.

Resolution: circuit equivalence classes. The uniform resolution of Section 5 specializes to mechanistic interpretability as follows. Let $G = [C]$ denote the equivalence class of circuits that reproduce the same I/O behavior. The G -invariant explanation aggregates over all valid circuits:

$$\text{exp}^*(\theta) = \{C \in \mathcal{H} : C \text{ is a valid circuit for } f_\theta\},$$

the set of all circuit decompositions consistent with the network’s behavior. This map is **faithful in expectation** (it captures all structure shared across valid circuits), **stable** by construction (functionally equivalent networks have the same set of valid circuits), but **not decisive**: it returns a set rather than a single circuit. The unresolvable queries—which specific components “really” implement a computation when multiple valid assignments exist—correspond precisely to the undirected edges in the causal discovery analogy (Section 17.5).

Reporting the full equivalence class is the principled response to the impossibility, trading decisiveness for faithfulness and stability, exactly as predicted by Theorem 3.6. This reframes the MI research agenda: the goal is not to find the one true circuit, but to characterize which circuit features are *invariant* across all valid decompositions (the “skeleton” of the circuit space) and which are decomposition-dependent (the “reversible edges”).

Table 19: Empirical illustration of the limits of explanation across four explanation types. These experiments demonstrate that the Rashomon property produces measurable instability in practice, not that the theorem depends on experimental evidence. All instability measures have 95% bootstrap confidence intervals excluding zero. Prediction equivalence confirms models are functionally identical.

Instance	Instability metric	Value	Pred. equiv.
Attribution (GBDT)	SHAP flip rate	33%	>99%
Attention (DistilBERT)	Argmax flip rate	19.9%	100%
Counterfactual (XGBoost)	Direction flip rate	23.5%	$\Delta\text{AUC} < 0.03$
Concept probe (MLP)	$1 - \cos $	0.90	97.2%
Model selection (XGBoost)	Best-model flip rate	80.0%	$\Delta\text{AUC} < 0.02$
Causal discovery	<i>Classical (Verma & Pearl, 1991); Rashomon Derived</i>		
GradCAM (ResNet-18)	Spatial IoU	9.6% peak flip	>78.8%
Token citation (DistilBERT)	Citation flip rate	34.5%	100%
Mech. interp.	Circuit importance rho	0.518 (full) $\rightarrow$ 0.929 (G -inv.)	10 grokked transformers

^{||}Weight perturbation yields 60.0%; we report the conservative full-retraining figure (20 models, last 2 transformer layers retrained).

Empirical confirmation: modular addition circuits. We test this prediction by training 10 two-layer transformers from *independent random initialisation* on modular addition ($a + b \bmod 113$, following Nanda et al. [2023]). All 10 models reach 100% test accuracy. Activation patching on 10 components (8 attention heads + 2 MLPs) yields mean Spearman $\rho = 0.518$ across $\binom{10}{2} = 45$ model pairs (permutation $p < 10^{-6}$), confirming partial but significant circuit agreement.

The architecture determines the symmetry group $G = S_4 \times S_4$ (within-layer head permutations): the 4 heads per layer are architecturally identical, distinguished only by random initialisation. The G -invariant subspace V^G is 4-dimensional (mean layer-0 head importance, mean layer-1 head importance, MLP0, MLP1), giving $\eta = \dim(V^G)/\dim(V) = 4/10 = 0.4$. Projecting onto V^G recovers near-perfect agreement: **mean Spearman** $\rho = 0.929$ (Pearson $r = 0.967$, min $\rho = 0.800$). The stability is not trivially driven by the dominant MLP1 component: the invariant excluding MLP1 still achieves $\rho = 0.822$ ($p < 10^{-6}$).

Within-layer head pairs flip at mean rate 0.500 (exactly coin-flip, $\sigma = 0.071$), confirming the Noether counting prediction (Section 11.5). Head-vs-MLP1 pairs never flip. Fourier analysis reveals that all 10 models learn *completely different* frequency representations (top-5 Jaccard = 0.022), confirming genuine Rashomon: same function, different algorithms, partially shared circuit skeleton.

This directly demonstrates the resolution in practice: the G -invariant projection—averaging head importance within each layer—recovers the “skeleton” of the circuit space ($\rho = 0.929$), while the per-head rankings are the “reversible edges” that depend on the specific Fourier algorithm each model discovers.

Adversarial vs. structural instability. Adversarial attacks on SHAP and LIME Ghorbani et al. [2019], Slack et al. [2020] demonstrate a complementary vulnerability: explanations can be manipulated by adversarial model design. Our impossibility is distinct: instability arises from the *legitimate* Rashomon property of well-specified models, not from adversarial manipulation. The trilemma holds even for honestly trained models with no adversarial intent.

Instance overlap and independence. We note that instances 2 (attention), 7 (saliency), and 8 (token citation) share a common Rashomon mechanism—weight-space multiplicity affecting internal representations. They differ in modality and metric but are not structurally inde-

Statistics (Fisher/Rao–Blackwell)	Framework
Sufficient statistic T	Observation map obs
Non-injectivity of T	Rashomon property
Rao–Blackwell estimator $\mathbb{E}[\delta \mid T]$	Orbit average (G -invariant resolution)
Rao–Blackwell variance reduction	Pareto-optimality

Table 20: Dictionary between sufficient statistics and the abstract framework. Fisher’s sufficiency (1922), the Rao–Blackwell theorem (1945/47), and the explanation impossibility are the same result in different notation.

pendent. A conservative count identifies six genuinely distinct Rashomon mechanisms across the nine instances.

18 Bridge Theorems: Connections to Established Mathematics [Tier A: Hardened]

The cross-domain transfer experiments of the previous section show that the framework *generates* novel predictions. This section establishes a complementary result: the framework *subsumes* classical theorems from statistics, physics, and information theory. We present five “bridge theorems” that identify existing results as special cases of the abstract impossibility structure, and derive one new quantitative prediction for AI interpretability.

18.1 Sufficient Statistics as G -Invariants [Tier A: Hardened]

Fisher’s concept of sufficiency [Fisher, 1922] can be recovered exactly within the impossibility framework. Map the components as follows: let configurations be the data space $\mathcal{X}$, let the observation map obs be a statistic $T: \mathcal{X} \rightarrow \mathcal{Y}$, and let the explanation map be the identity $\text{explain} = \text{id}$ (reporting the full data). Under this identification, the Fisher–Neyman factorization theorem [Neyman, 1935] states that T is sufficient for θ if and only if the conditional distribution $P(X \mid T(X))$ is independent of θ —precisely the stability condition of the abstract framework (the explanation factors through the observation).

The Rashomon property holds whenever T is non-injective: there exist $X_1 \neq X_2$ with $T(X_1) = T(X_2)$, so the full data cannot be faithfully, stably, and decisively reported. This is the typical case—sufficiency is useful precisely because T compresses.

The identification extends to the resolution. The Rao–Blackwell theorem [Rao, 1945, Blackwell, 1947] states that for any estimator $\delta(X)$, the conditional estimator $\delta^*(T) = \mathbb{E}[\delta(X) \mid T(X)]$ has variance no greater than that of δ . This conditioning is *exactly* orbit averaging: the estimator is averaged over the fiber $\{X : T(X) = t\}$, weighted by the conditional distribution. The Rao–Blackwell variance reduction is therefore the framework’s Pareto-optimality theorem—the G -invariant resolution dominates all alternatives—expressed in the language of classical estimation.

These are not analogies. Fisher’s sufficiency criterion (1922), the Rao–Blackwell theorem (1945/47), and the explanation impossibility are the same theorem in different notation (Table 20). Under a uniform prior on the fiber $\{\theta : \text{obs}(\theta) = y\}$, the orbit average $\bar{E}(\theta) = (1/|G|) \sum_{g \in G} E(g \cdot \theta)$ is literally the conditional expectation $\mathbb{E}[\text{explain}(\theta) \mid \text{obs}(\theta) = y]$. The identification is exact when the prior is uniform and the group acts transitively on fibers; it is approximate otherwise.

EM Algorithm	Framework
Latent variable Z	Configuration θ
Observed data X	Observable $\text{obs}(\theta)$
Posterior $P(Z X, \theta)$	Resolution weight on fiber
E-step: $\mathbb{E}[f(Z) X]$	Orbit average (G -invariant resolution)
M-step: update θ	Update resolution parameters
Rao–Blackwell optimality	Pareto-optimality of orbit average

Table 21: Dictionary between the EM algorithm and the abstract framework. The E-step IS the G -invariant resolution with posterior weights; the general EM generalizes from uniform to non-uniform orbit averaging.

18.2 EM as Generalized Orbit Averaging [Tier B: Empirical]

The Expectation-Maximization algorithm [Dempster et al., 1977] is the most widely used computational method for latent-variable models (mixture models, HMMs, factor analysis, topic models). Its E-step can be identified with the framework’s G -invariant resolution.

Map the EM components: configurations are latent variables Z , observations are the observed data X (with $\text{obs}(Z) = X$ via marginalization $P(X) = \sum_Z P(X, Z)$), and explanations are the latent variables themselves. The Rashomon set at X is $\{Z : P(X | Z) > 0\}$ —all latent configurations consistent with the data. The E-step computes

$$\mathbb{E}_{Z|X,\theta}[f(Z)] = \sum_Z f(Z) \cdot P(Z | X, \theta),$$

a weighted average over the fiber $\{Z : \text{obs}(Z) = X\}$. Under a uniform prior ($P(Z)$ constant) with equal likelihoods for all consistent Z (the microcanonical case), this reduces to $(1/|\text{fiber}(X)|) \sum_{Z:\text{obs}(Z)=X} f(Z)$ —exactly the orbit average of the framework’s G -invariant resolution.

The general case—non-uniform priors—extends the orbit average to a posterior-weighted average, with the prior $P(Z | \theta)$ introducing additional structure beyond the bare group action (Table 21). This provides the most practically impactful bridge: the EM algorithm, gauge-invariant averaging in physics, and codon usage tables in biology are all instances of the same resolution strategy.

Part IV

Part IV — Structural Analogies and Speculative Extensions

Status disclaimer. *The material in this Part consists of structural parallels and conditional or illustrative formalizations—not domain results. Where a Lean object appears (the $\text{GL}(n)$ trace biletta, the adelic model, the Gödel conditional, Navier–Stokes), the genuine theorem is stated in Part I or III and only its interpretation lives here. Claims connecting the framework to the Langlands programme, adelic physics, Gödel’s theorem, or the Riemann Hypothesis are analogies, not proofs.*

19 Discussion: What Explainability Can and Cannot Do [Tier C: Speculative]

19.1 Cross-Domain Transfer: From Gauge Theory to ML [Tier C: Speculative]

The gauge theory resolution (restricting to observables invariant under gauge transformations) and the ML resolution (DASH, averaging SHAP values over an ensemble) share identical mathematical structure: both are the G -invariant projection for their respective symmetry groups (the Reynolds operator—§5.4). The structural differences—continuous vs. discrete symmetry, local vs. global action, dynamical vs. static—mean this identity is at the level of the group action formalism, not a claim that the physical mechanisms coincide.

Table 22: Gauge–attribution correspondence: shared and divergent structure.

Property	Gauge theory	ML attribution
Shared		
Group action on descriptions	✓	✓
Orbit averaging (Reynolds operator)	✓	✓
Pareto-optimality of G -projection	✓	✓
Divergent		
Symmetry type	Local, continuous	Global, discrete
Dimension	Infinite	Finite
Dynamics	Equations of motion	Static computation

Beyond ML. As evidence that the impossibility extends beyond ML, we note that spacetime geometry under Ostrowski’s classification of completions of $\mathbb{Q}$ is a maximally incompatible system: the two completions yield exactly two geometries (smooth and ultrametric), triggering the biletta (Theorem 3.13). The adelic resolution — formulating physics over all completions simultaneously — is the enrichment escape, and it is the *unique* faithful + stable value on the Rashomon fiber. The full development, including the Ostrowski–Design Space correspondence and physical scoping (conditional on the Freund–Witten adelic product formula), is in a companion paper [Caraker et al. \[2026b\]](#).

Value alignment. The framework applies to AI value alignment. When multiple RLHF training runs produce models with identical behavioral performance but different implicit value rankings (the value Rashomon property), no value explanation is simultaneously faithful and stable (the value biletta). The enrichment mechanism — adding a ‘balanced’ option that doesn’t commit to either value priority — corresponds to alignment frameworks that explicitly balance competing values rather than ranking them. This formalizes a structural constraint on

alignment: under value multiplicity, precise value specification is impossible without sacrificing either accuracy (faithfulness) or consistency (stability). We formalize this as a Lean-verified `ExplanationSystem` instance (`ValueAlignment.lean`, 0 sorry). Whether RLHF models empirically exhibit value Rashomon is an open question; the Lean formalization proves the structural consequence *if* the property holds, but does not establish that it does. The full complexity of AI alignment (strategic agents, distributional shift, ontological uncertainty) exceeds the framework’s scope.

19.2 Scope and Applicability Beyond Machine Learning [Tier C: Speculative]

The theorem applies to *any* explanation of *any* underspecified system. The four empirical instances in the Nature paper (genomics, mechanistic interpretability, causal inference, neuroimaging) demonstrate the theorem in specific domains; the 14 cross-domain Lean instances (Arrow, Bell, Gödel, gauge theory, statistical mechanics, genetic code, phase problem, QM interpretation, linguistics, value alignment, view-update, linear systems, and two scientific revolutions) prove the algebraic structure. But the scope extends further. The Rashomon property arises wherever multiple valid configurations of a system produce the same observables, which is the default condition in high-dimensional science.

Table 23: Scope of the impossibility theorem beyond the four main instances. Each row identifies a scientific domain, the source of the Rashomon property (what creates multiple valid configurations), and the form the resolution takes.

Domain	Source of Rashomon	What is unstable	Resolution form
Epidemiology	Correlated risk factors (smoking, diet, alcohol, exercise)	Which factor is “the” cause of disease	Report attributable fractions with stability bounds
Climate science	Multiple valid GCMs fit the same temperature record	Which feedback drives warming (clouds, albedo, ocean)	Multi-model ensemble (CMIP)
Economics	Multiple valid macroeconomic models	Which policy lever explains GDP change	Bayesian model averaging
Ecology	Correlated environmental variables	Which variable drives species distribution	Ensemble species distribution models
Pharmacology	Correlated molecular targets	Which pathway mediates drug effect	Pathway-level (not target-level) analysis
Forensics	Multiple valid statistical models of evidence	Strength of evidence attribution	Likelihood ratio over model space
Social science	Researcher degrees of freedom	Which specification produces significance	Multi-analyst studies
Psychology	Garden of forking paths	Which analysis yields the effect	Pre-registration + multiverse analysis

In each domain, the pattern is identical: (1) correlated variables or underdetermined models create the Rashomon property; (2) the specific explanation (which gene, which circuit, which risk factor, which feedback mechanism) depends on the particular model or analysis pipeline; (3) the resolution involves some form of averaging, ensembling, or abstraction to a coarser level where the instability vanishes. The theorem proves that step (3) is not a methodological preference but the unique Pareto-optimal response to a structural constraint.

Extension to mathematical physics: 3D Navier–Stokes. The tightness classification extends beyond data science to open problems in mathematical physics. The 3D incompressible Navier–Stokes equations admit three desiderata: existence of solutions (P1), smoothness/regularity of solutions (P2), and uniqueness (P3). Whether all three hold simultaneously is one of the seven Clay Millennium Problems. We apply the framework computationally: for the Taylor–Green vortex ($A = 1$, $T_{\text{final}} = 10$, capturing peak dissipation at $t \approx 9$ per Brachet 1983), 54 pseudospectral DNS runs across five resolutions (32^3 , 48^3 , 64^3 , 96^3 , 128^3) with Reynolds numbers $\text{Re} \in [50, 10000]$ establish a complete phase diagram.

Phase diagram (dealiasing ON): At low Re , all three properties hold (FULL tightness) at every resolution. Above a critical Re_{crit} , smoothness (P2) fails while existence (P1) and uniqueness (P3) survive — SMOOTH-BLOCKED tightness. The critical Reynolds number increases with resolution:

$$\text{Re}_{\text{crit}} = 2.91 \times N^{1.56}, \quad R^2 = 0.94,$$

where N is the grid size. The exponent 1.56 lies between the Kolmogorov prediction $4/3$ and 2, consistent with the inertial-range scaling (slopes -1.8 to -1.87 at 96^3 and 128^3 , approaching the $k^{-5/3}$ Kolmogorov spectrum).

Four confirmed claims:

1. P2 (smoothness) fails first at ALL five resolutions tested (dealiasing ON).
2. Re_{crit} increases monotonically with N : $707 \rightarrow 1414 \rightarrow 1414 \rightarrow 3162 \rightarrow 7071$.
3. Dealiasing protects P1: with standard 2/3-rule dealiasing, P1 never fails. Without it, P1 always fails after P2 (cascade within 5 ms at 128^3).
4. The inertial range develops at high resolution with near-Kolmogorov slopes.

The 2D control: In 2D, where global regularity is proved (Ladyzhenskaya 1969), the framework correctly reports FULL tightness at all Reynolds numbers tested. The 2D/3D comparison is a clean falsification test: the framework predicts FULL where regularity is known and SMOOTH-BLOCKED where it is unknown, and both predictions are confirmed.

The dealiasing finding: The standard 2/3-rule dealiasing mask, present in all production DNS codes, artificially decouples smoothness from existence. Without dealiasing, the P2→P1 cascade gap shrinks with resolution (0.4–0.6 s at 32^3 ; 0.005–0.01 s at 128^3 ; effectively simultaneous at the continuum limit). This is a methodological finding relevant to the computational PDE community: the dealiasing mask is a Reynolds projection for aliasing symmetry, with capacity $(2N/3)^3$ modes.

What this does not claim: The numerical experiments classify the *computational* tightness type (which property fails first in DNS at finite resolution). They do not prove or disprove global regularity for the analytical 3D Navier–Stokes equations, which remains an open Millennium Problem. The resolution-convergent trend ($\text{Re}_{\text{crit}} \rightarrow \infty$ as $N \rightarrow \infty$) is consistent with either eventual global regularity (FULL at the continuum limit) or a finite-Re breakdown that DNS cannot yet resolve.

Divergence check: $\max |\nabla \cdot \mathbf{u}| < 10^{-8}$ across all 54 runs, confirming the pseudospectral solver preserves incompressibility.

The Explanation Capacity Audit. The explanation capacity theorem was validated at scale across 149 datasets from 53 scientific domains (genomics, neuroscience, particle physics, manufacturing, veterinary medicine, and 48 others). For each dataset, 50 XGBoost models with stochastic regularisation (subsample = 0.8, colsample_bytree = 0.5) were trained with different random seeds, and TreeSHAP importance rankings were computed. Correlation groups were identified via SAGE (hierarchical clustering on $|\text{Spearman } \rho|$ at threshold $\rho^* = 0.70$), giving the explanation capacity $C = g$ (number of groups). Note: the Lean theorem `mi_sufficient_`

`for_impossibility` in `MutualInformation.lean` proves the *sufficiency* direction—mutual information $I(X_j; X_k) > 0$ (not correlation) yields the attribution impossibility—in the abstract framework *conditional* on the bridge hypothesis `hdep_implies_diff` (statistical dependence produces differing attributions). That bridge cannot be discharged abstractly (a degenerate constant attribution falsifies it), but it is *not* merely assumed: `MIWitness.lean` exhibits an explicit two-feature model in which the mutual information is *computed* from a joint distribution (two perfectly-correlated fair bits, $I = \log 2 > 0$ via the Shannon formula) and the bridge is a *theorem*—the attribution difference is derived from the feature-swap Rashomon pair that the same correlation forces—so `concrete_mi_impossibility` holds unconditionally, with no custom axioms (`#print axioms` shows only Lean’s core). The MI-to-attribution converse is thus machine-verified on a concrete witness, though not in full abstract generality. The *escape* direction (MI = 0 implies a stable ranking exists) is resolved in `MIConverse.lean`: it is *false* unconditionally—`mi_converse_fails_without_bridge` formalizes a counterexample with provably independent features (MI computed = 0 from a product distribution) whose feature-swap Rashomon pair admits no faithful stable ranking, exactly the target-symmetry instability noted elsewhere in this paper—and it holds under precisely one named mirror bridge, `hindep_implies_agree`, proven minimal, yielding the conditional two-bridge biconditional `mi_exact_boundary_biconditional`. For continuous features, Spearman correlation is an adequate proxy: an MI-vs-correlation comparison across 131 PMLB benchmark datasets confirms that the two grouping methods produce identical groups 77% of the time (ARI = 1.000) and nearly identical groups 74% of the time (ARI > 0.80; mean ARI = 0.84). However, for binary or categorical features, correlation can fail completely—the drug discovery experiment (Section 11.10) demonstrates this: Pearson clustering predicts 0% instability on binary fingerprints while MI-based clustering recovers the observed 23% rate. The 75% exceedance rate below is therefore exact for the continuous-feature audit but a *lower bound* for domains with binary, categorical, or nonlinear feature dependencies.

Exceedance: 111 of 149 datasets (75%) exceed their explanation capacity — the feature rankings reported in typical SHAP analyses contain more pairwise claims than the correlation structure can stably support. A TreeSHAP replication on 70 PMLB datasets with $M = 50$ confirms: 84% exceed capacity (conservative relative to gain-based 87% on the same datasets, because TreeSHAP is 34% more stable). Exceedance is robust to threshold: 83% at $\rho^* = 0.50$, 81% at 0.60, 75% at 0.70, 69% at 0.80, 50% at 0.90. The directional prediction (within > between flip rate) is confirmed at all five thresholds (Wilcoxon $p \leq 0.021$ at each), with the strongest signal at $\rho^* = 0.70$ and 0.90. The slight weakening at $\rho^* = 0.80$ (77% directional vs. 81% at 0.70) reflects group fragmentation: 8 datasets lose testability as within-group pairs are split, and the mean number of within-group pairs drops from 81 to 62.

Directional prediction: The stability theorem predicts within-group pairwise comparisons (beyond capacity) are less stable than between-group comparisons (within capacity). This prediction is confirmed in 72 of 84 testable datasets (86%; 89 of 110 at the broader testability criterion). A cross-dataset Wilcoxon signed-rank test gives $p = 5.1 \times 10^{-11}$; on 59 real-world datasets only (excluding PMLB benchmarks and synthetic controls), $p = 1.7 \times 10^{-6}$ (47/59 positive). 27 per-dataset confirmations are significant at $p < 0.005$; 22 survive Bonferroni correction ($p < 0.05/149 = 3.4 \times 10^{-4}$). One dataset (Dermatology) shows a significant reversal ($p = 3.6 \times 10^{-4}$; see below). The confirmation-to-reversal ratio is 27:1. A family-level analysis (averaging Friedman variants within families to address non-independence) confirms the result: 59/72 families positive, Wilcoxon $p = 2.5 \times 10^{-8}$; block-bootstrap 95% CI for the mean gap is [0.036, 0.069], excluding zero.

Effect sizes: The strongest confirmations span diverse domains: Satellite Image (remote sensing, $d = 1.27$, $p = 2 \times 10^{-53}$), HAR Activity (wearables, $d = 0.91$, $p = 4 \times 10^{-72}$), Horse Colic (veterinary, $d = 0.99$, $p = 4 \times 10^{-40}$), AP Colon Kidney (genomics, $d = 0.90$, $p = 5 \times 10^{-35}$), MiniBooNE (particle physics, $d = 0.92$, $p = 2 \times 10^{-6}$), Vehicle Silhouette (vehicle engineering,

$d = 0.82$, $p = 10^{-8}$), MiceProtein (neuroscience, $d = 0.67$, $p = 6 \times 10^{-8}$). Median Cohen’s d across all testable datasets: 0.40 (medium effect).

capacity theorem degradation: On synthetic data with known groups at varying correlation ($\rho \in [0.50, 0.99]$), the within-group flip rate increases monotonically with ρ ($r = 0.898$, $p = 0.006$). The bimodal gap (within – between) also increases ($r = 0.864$, $p = 0.012$), confirming the capacity theorem’s degradation curve: the prediction is sharpest at exact exchangeability and degrades continuously with approximate symmetry.

Model-class robustness: XGBoost, Random Forest, and Ridge agree on the direction of the within/between flip rate difference in 4 of 5 representative datasets (80%). The prediction is a property of the data’s dependence structure, not of a specific model class.

Regularisation creates the Rashomon property: Increasing stochastic regularisation (lower subsample, smaller trees) *increases* instability (Breast Cancer: flip rate 0.134 at subsample = 1.0 vs. 0.194 at subsample = 0.5), confirming that the Rashomon set is created by the regularisation, not by data noise. Under deterministic training, a single ranking emerges. This has a practical implication: instability is not a price paid for *better* models—it is a consequence of the stochastic training strategies that produce the Rashomon property in the first place.

Explanation-method robustness: TreeSHAP and native XGBoost feature importance agree on direction in all tested datasets. Native importance shows stronger effects in some cases (Diabetes: $d = 1.81$ native vs. 0.94 TreeSHAP).

The over-explanation penalty: Part (iv) of the stability theorem predicts that the MSE of the orbit average for within-group importance differences equals the full squared norm of the target: $\text{MSE}/\|w\|^2 = 1$. On synthetic data across five correlation levels ($\rho = 0.70$ to 0.99), the observed ratio is **1.000** at every level. Note: this is a consistency check rather than a prediction test — the orbit average zeroes within-group differences by construction, so $\text{MSE}/\|w\|^2 = 1$ holds algebraically. The empirical content is that between-group MSE is substantially lower ($\text{MSE}_{\text{between}}/\text{MSE}_{\text{within}} \approx 0.75$), confirming that within-capacity estimation benefits from orbit averaging while beyond-capacity estimation does not.

Convergence rate: The orbit average’s MSE decays as $M^{-\beta}$ with $\beta = 1.28$ ($R^2 = 0.97$; predicted $\beta = 1.00$). The faster-than-predicted convergence likely reflects conservative i.i.d. assumptions; in practice, models with shared training data are correlated, but the empirical rate is still $O(1/M)$.

The Dermatology reversal: The single significant reversal ($P = 34$, $g = 22$, within = 0.039, between = 0.137, $d = -0.80$) persists across all thresholds ($\rho^* = 0.50$ to 0.90). The explanation: dermatology features are *complementary* clinical indicators (scaling + erythema = same disease), not *substitutable* (competing for signal). The framework predicts instability for substitutable features (the common case in high-dimensional continuous data); complementary features violate the exchangeability assumption. This identifies a principled scope condition: correlation-as-proxy-for-exchangeability works when features compete for the same signal but fails when features co-occur as joint indicators.

Null model test: A natural alternative explanation for the directional prediction is that within-group features have smaller importance differences (by construction, since they are correlated), and smaller differences flip more regardless of group structure. We tested this via pooled per-pair analysis across 14 real-world datasets (13,511 feature pairs, 1,309 within-group, 12,202 between-group; 5 `sklearn` + 9 PMLB benchmark datasets spanning clinical, genomics, economics, computer vision, remote sensing, and speech domains).

Parametric: OLS regression (flip_rate $\sim$ |diff| + |diff|² + within + dataset FE) gives a within-group coefficient of 0.037 (dataset-level cluster-bootstrap 95% CI [0.014, 0.045], excluding zero). The quadratic |diff| term substantially improves fit ($R^2 = 0.42$ vs. the linear baseline), confirming the nonlinear relationship between importance magnitude and flip rate.

Nonparametric: The within>between gap is concentrated in the lowest |diff| quintile (gap = +0.061, $p < 10^{-4}$), where features have similar mean importances and are thus genuinely ex-

changeable. In the upper four quintiles—where within-group features have increasingly different importances despite being correlated—the gap attenuates to near zero. This is consistent with the stability theorem’s mechanism: the effect operates where features are *exchangeable* (similar importance + statistical dependence), not merely where they are correlated.

SAGE vs. correlation vs. MI: A cross-validated comparison of three group-discovery methods (correlation-based, MI-based, and SAGE flip-rate-based clustering, each calibrated to the same number of groups) reveals that SAGE produces substantially larger within–between gaps on 6 of 7 datasets tested (mean SAGE gap = 0.25 vs. correlation gap = 0.05). SAGE clusters by empirical exchangeability (cross-validated: first 25 seeds for discovery, last 25 for evaluation), directly capturing the theoretical criterion. MI-based and correlation-based groups are nearly identical for continuous features (ARI = 0.84 across 131 PMLB datasets; identical in 77%), confirming that the MI generalisation matters for binary/categorical features but not for continuous data.

Per-dataset randomisation: on synthetic data with correctly structured groups ($\rho \geq 0.80$), the real partition produces a significantly larger gap than 2,000 random partitions ($p = 0.001$ for all four settings).

The interpretation: group membership carries predictive information *beyond* |diff|, concentrated where features are genuinely exchangeable (low |diff|, high dependence). The quadratic OLS coefficient is robust across sample sizes (14-dataset cluster-bootstrap CI [0.014, 0.045]). The stratified effect is strongest in the lowest quintile, where correlated features have similar importances and thus compete for the same predictive signal. In higher quintiles, correlation groups include non-exchangeable pairs (correlated but with different true importances), attenuating the within–between gap—consistent with the distinction between correlation and exchangeability (Section 11.8). The cross-dataset aggregate test (Wilcoxon $p = 5 \times 10^{-11}$) confirms that group structure matters at the dataset level; the per-pair OLS confirms it at the feature-pair level after nonlinear control for importance magnitude.

Capacity theorem across heterogeneous datasets: The headline $R^2 = 0.957$ for the capacity theorem is from 7 well-characterised groups. On real-world datasets with power-adequate inferred groups ($n_w \geq 10$), the formula partially generalises. The prediction strength depends on the correlation threshold used for group discovery:

ρ^*	n (real, $n_w \geq 10$)	Spearman $\rho(\eta, \text{gap})$	p	Interpretation
0.50	49	+0.55	$< 10^{-4}$	Large groups, high power
0.60	44	+0.55	10^{-4}	Similar
0.70	39	+0.40	0.011	Primary threshold
0.80	33	+0.29	0.10	Group fragmentation
0.90	28	+0.16	0.42	Insufficient within-group pairs

The formula works best at $\rho^* = 0.50\text{--}0.60$ ($\rho \approx 0.55$, $p = 3.8 \times 10^{-5}$ at $\rho^* = 0.50$, surviving Bonferroni correction for the 5 pre-specified thresholds tested) where groups are large enough to provide statistical power, degrades at $\rho^* = 0.70$ ($\rho = 0.40$, $p = 0.011$ —marginal under Bonferroni), and fails at $\rho^* \geq 0.80$ where aggressive thresholds fragment groups below testable size. Power weighting (by $\sqrt{n_w}$) modestly improves the correlation at $\rho^* = 0.70$ ($\rho = 0.09 \rightarrow 0.15$), confirming that underpowered datasets add noise. The “real-world” filter excludes PMLB benchmark and synthetic control datasets, whose engineered correlation structures could inflate the η correlation. The practical recommendation: use $\rho^* = 0.50\text{--}0.60$ for the capacity theorem prediction, not the $\rho^* = 0.70$ threshold used for the directional prediction.

A comparison of three group-discovery methods on 14 datasets with cached importance matrices (split-half cross-validated for SAGE to avoid circularity) shows that SAGE does *not* improve the η prediction ($\rho_{\text{SAGE}} = 0.46$ vs $\rho_{\text{corr}} = 0.54$), even though SAGE produces larger within–between gaps for the directional prediction. This dissociation suggests the capacity

formula is more sensitive to group *granularity* (how many features per group) than to group *accuracy* (whether the grouped features are truly exchangeable).

Group stability: Bootstrap resampling (200 resamples) confirms that correlation-based groups are stable: mean ARI > 0.94 for real-world datasets (Breast Cancer, Diabetes, California Housing), and ARI = 1.000 for synthetic data at $\rho \geq 0.80$. At $\rho = 0.70$ (threshold = correlation), ARI = 0.33—groups are not recovered, consistent with the null-model test failure at this boundary.

Conformal prediction intervals: Split conformal inference (25 calibration + 25 test seeds) produces per-feature importance intervals with 84–88% coverage (target: 90%; the shortfall reflects the small calibration set). Features with interval width $> 2\times$ the median width are flagged as unreliable (1–11 per dataset). With $M \geq 50$ models (25+ per split), conformal intervals provide a model-free uncertainty quantification for SHAP values that requires no distributional assumptions—a practical complement to the Gaussian flip formula for practitioners who need per-feature (rather than per-pair) reliability estimates.

Dimensionality boundary: High-dimensional datasets ($P \geq 50$) systematically show smaller or reversed gaps (Spearman $\rho = -0.35$ between P and gap, $p = 2 \times 10^{-4}$). This is a power artefact: at high P , correlation clustering yields few within-group pairs ($n_w \leq 3$), making the per-dataset test unreliable. The cross-dataset test is unaffected (excluding all $P \geq 50$ datasets: Wilcoxon $p = 4.6 \times 10^{-8}$).

The theorem as a capacity law and uncertainty principle for scientific explanation. The Uncertainty from Symmetry theorem (Section 8) formalises this scope. For any system with symmetry group G , the explanation capacity is $C = \dim(V^G)$ and the explanation loss rate is $\eta = 1 - C/\dim(V)$. These are not specific to ML explanations, SHAP values, or feature importance. They apply to any vector-valued description of any system with a symmetry group—physical, biological, economic, or social. The *explanation uncertainty bound* ($\text{unfaith}(\theta_1) + \text{unfaith}(\theta_2) \geq \Delta - \delta$, the quantitative form of the blemma) provides the complementary bound: reducing unfaithfulness at one configuration forces increased unfaithfulness at another, with a lower bound determined by the system’s structure.

The practical consequence for working scientists in *any* field: if your system has statistically dependent explanatory variables (whether linearly correlated or not—the operative dependence condition is $I(X_j; X_k) > 0$), the specific attribution you report (which variable is most important, which mechanism is primary) is not a stable property of the data. It is a joint property of the data and the particular model or analysis pipeline. The stable content—what survives averaging over valid pipelines—is at a coarser level (variable groups, pathway classes, mechanism families). The theorem identifies exactly what that coarser level is (V^G , the symmetry-invariant subspace) and proves the orbit average is the optimal way to access it.

19.3 The “Isn’t This Obvious?” Defense

A natural reaction is that practitioners already know explanations are unstable—so what is new? The contribution is fivefold. First, we make explicit that eight scientific domains and nine ML explanation types share *one* mechanism—the Rashomon property—rather than seventeen unrelated pathologies. Second, the Lean 4 formalization (Section 10: 130 files, 715 theorems, 0 sorry) eliminates logical gaps that informal arguments inevitably leave. Third, the empirical experiments (Section 17) show quantifiable severity: 33% attribution flip rates, measurable counterfactual and concept instability on standard benchmarks. Fourth, the resolution via G -invariant maps is constructive and actionable, not merely diagnostic. Fifth, Lean-verified tightness theorems (Proposition 3.10) prove the impossibility is non-trivial: each pair of properties *is* simultaneously achievable, so the theorem is not an artifact of overly strong definitions—it identifies the exact boundary between the possible and the impossible. Beyond these five points, three structural arguments establish that the result is non-trivial:

1. **Tightness.** Each pair of the three properties (faithful + stable, faithful + decisive, stable + decisive) is simultaneously achievable, with Lean-verified witnesses (Proposition 3.10). If the definitions were rigged to be trivially unsatisfiable, no pair would work.
2. **Biconditional.** The impossibility holds *if and only if* the Rashomon property is present (Theorem 3.6 and its converse in `Necessity.lean`). This means the axiom set precisely characterizes the phenomenon—it is tight in both directions.
3. **Quantitative corollaries.** The invariance counting theorem (exactly $g(g-1)/2$ stable queries, invariant across all $\rho \in [0.5, 0.99]$; Section 11.5) and the capacity theorem ($R^2 = 0.957$; Section 11.8) are non-obvious predictions that follow from the framework but not from any simpler observation about non-injectivity. A reviewer who says “this is just the pigeonhole principle” must also explain why the pigeonhole principle predicts a 50-percentage-point bimodal gap that is invariant across $\rho \in [0.5, 0.99]$.

Comparison with Arrow’s theorem. Arrow’s impossibility theorem [Arrow, 1951] is more surprising than ours in one respect: its axioms appear independently mild, and the impossibility emerges from their conjunction. Our axioms are more transparently in tension once the Rashomon property is stated. The framework captures Arrow’s *setting*—underdetermination in preference aggregation—but not Arrow’s *theorem*, which proves dictatorship from weaker axioms (unanimity and IIA). Our faithfulness is strictly stronger than Arrow’s unanimity (requiring agreement on all profiles, not just unanimous ones); our stability is strictly stronger than IIA (requiring global invariance, not just pairwise). The framework’s impossibility is therefore complementary to Arrow’s: it applies to a broader class of systems under stronger assumptions. The contribution is not surprise but *precision*: the axiom set is tight (each pair of properties is achievable), equipped with machine-checked necessity converses (the full biconditional for the maximal-incompatibility class; the fully-specified converse in general), and quantitatively productive (the counting theorem and capacity theorem). Arrow’s formalization created the field of social choice theory not because it was surprising but because it replaced intuition with a theorem that could be built upon; we aim for the same role in explainability.

19.4 The “All Instances Are Trivial” Defense

A reviewer might object that each individual instance is already known: attention instability was documented by Jain and Wallace [2019], Markov equivalence is classical [Verma and Pearl, 1991], fairness impossibility was proved by Chouldechova [2017], and SHAP instability was observed by D’Amour et al. [2022] and Krishna et al. [2022]. The contribution is not any single instance but the *meta-theorem*: the same four-line proof (Theorem 3.6) closes all seventeen cases—eight derived scientific instances and nine ML instances—under a single structural condition. Before this work, a practitioner encountering attention instability might switch to SHAP, then to counterfactual explanations, then to concept probes—each time hoping the new method escapes the problem. Our theorem shows that no such escape exists: the impossibility is structural (it follows from the Rashomon property), not method-specific. Recognizing this unifies the practitioner’s response: rather than debugging each method individually, one diagnoses the Rashomon property and applies the G -invariant resolution uniformly.

19.5 Implications Beyond Machine Learning [Tier C: Speculative]

The Explanation Impossibility is not specific to machine learning.

Philosophy of science: formalizing Quine–Duhem. Quine [Quine, 1951] and Duhem [Duhem, 1954] argued that empirical evidence underdetermines theory choice (Duhem on the impossibility of testing hypotheses in isolation; Quine on confirmational holism—sharing the core

implication that observational equivalence leaves explanation underdetermined). Theorem 3.6 proves an impossibility theorem for explanation under observational equivalence—providing a constructive witness of the underdetermination logic articulated by the Duhem-Quine thesis of holist underdetermination [Laudan and Leplin, 1991]; note that Laudan himself was skeptical of the underdetermination thesis, but the citation documents the philosophical context. Prior formal analyses of underdetermination in Bayesian frameworks [Dorling, 1979, Howson and Urbach, 2006] did not yield impossibility results for explanation methods.

Scope of the underdetermination analogy. Duhem–Quine underdetermination concerns *confirmational* holism: no hypothesis can be tested in isolation because auxiliary assumptions are always involved [Stanford 2006]. Our trilemma addresses *explanatory* underdetermination: multiple models explain the same observations equally well but offer incompatible explanations. These are related but distinct: confirmational underdetermination asks “which theory is supported by the evidence?” while explanatory underdetermination asks “which explanation of the prediction is correct?” Our framework formalises the latter. The Rashomon property is a sufficient condition for explanatory underdetermination; it is not a formalisation of Duhem–Quine holism per se. The G -invariant resolution is structurally reminiscent of epistemic structural realism [Worrall, 1989] (Worrall’s argument is diachronic—structural content survives theory change across scientific revolutions; our framework is synchronic—structural content is shared across equivalent models within one paradigm): stable explanatory content is what is invariant across equivalent configurations. The connection is suggestive; a rigorous mapping would require engaging with the ontological commitments of structural realism, which our descriptive framework does not address.

Structural realism. The G -invariant resolution—projecting onto the invariant subspace—is a formalization of the structural realist’s position in philosophy of science: only the invariant structure is stably describable [Worrall, 1989, Ladyman et al., 2007]. The framework provides a quantitative version of this thesis: the fraction of structure that is stably accessible equals $\dim(V^G)/\dim(V)$. The symmetry-to-reality inference — whether invariant structure is genuinely “more real” than variant structure — is itself contested [Dasgupta, 2016]. Our framework is neutral on this ontological question: it claims only that invariant quantities are *stably reportable*, not that they are *real*.

Dropping decisiveness corresponds to van Fraassen’s [van Fraassen, 1980] constructive empiricist stance that science aims for empirical adequacy, not full theoretical truth. Okasha [Okasha, 2011] connected theory choice to Arrow’s impossibility in social choice; our result provides the complementary connection for underdetermination of explanation.

The Shannon parallel. The framework bears the same structural relationship to explanation that Shannon’s channel capacity theorem [Shannon, 1948] bears to communication. Shannon proved: (1) you cannot transmit above the channel capacity (impossibility); (2) you CAN approach capacity with random codes (constructive achievability); (3) the capacity is $C = B \log(1 + S/N)$ (quantitative formula with no free parameters); (4) this applies to any communication channel (universality within domain). The present framework provides the same package for explanation: (1) you cannot have faithful + stable + decisive under Rashomon (impossibility); (2) you CAN achieve faithful + stable via orbit averaging (constructive resolution); (3) the instability rate is $\eta = 1 - \dim(V^G)/\dim(V)$ (quantitative formula, $R^2 = 0.957$ across seven domains); (4) this applies to any explanation system with the Rashomon property (universality across domains). The difference: Shannon’s universality is within one domain (communication); this framework’s universality spans eight scientific domains, explaining a century of convergent evolution.

Precedents and distinctiveness. No prior work combines all of the following: a formally verified impossibility theorem (1178 distinct theorems across three repositories, 0 sorry), a constructive resolution proved Pareto-optimal, a quantitative prediction with zero free parameters ($R^2 = 0.957$), domain-specific empirical consequences (gene expression, AI safety, neuroimaging), an explanation of convergent evolution across eight scientific communities, a classification of 20 impossibility theorems from 12 domains by structural type, and a formalisation of the Duhem–Quine thesis with quantitative predictions. Each piece has precedents—Arrow (impossibility), Shannon (impossibility + achievability + formula), Bilodeau et al. (attribution impossibility), Blau and Wilson (axiom independence)—but the combination is, to our knowledge, new.

Neuroscience: degeneracy as the Rashomon property. Marder and Goaillard [Marder and Goaillard, 2006] documented that crustacean stomatogastric ganglia producing identical motor patterns are built from qualitatively different ionic conductance profiles. Edelman and Gally [Edelman and Gally, 2001] formalized this as *degeneracy*: structurally distinct elements performing the same function. Degeneracy is precisely the Rashomon property in our framework.

Economics: partial identification. Manski [Manski, 2003] developed partial identification theory for settings where observational data do not pin down a unique model. The identified set corresponds to an observe-fiber in our framework; partial identification is the economic instantiation of the Rashomon property.

Medical diagnosis. When multiple conditions produce the same symptoms, diagnosis faces the trilemma. Differential diagnosis is precisely G -invariant resolution: it sacrifices decisiveness for faithfulness and stability.

Arrow’s theorem as structural analogy. The result shares structural features with Arrow’s impossibility theorem [Arrow, 1951]: three individually reasonable desiderata that are pairwise satisfiable but jointly impossible under a natural background condition. Like Arrow’s theorem, the proof is short once the axioms are identified; the contribution lies in identifying the right axiomatic framework.

Spacetime geometry as a maximally incompatible system. As evidence that the impossibility extends beyond ML explanation methods, we note that spacetime geometry under Ostrowski’s classification of completions of $\mathbb{Q}$ is a maximally incompatible system: the two completions yield exactly two geometries (smooth and ultrametric), triggering the bilemma (Theorem 3.13). The adelic resolution—formulating physics over all completions simultaneously—is the enrichment escape. The full development, including the Ostrowski–Design Space correspondence and physical scoping, is in a companion paper.

The pattern extends to many further domains: metamerism in color perception (where the cone response matrix is the observation map, making it an instance of the linear system pattern), molecular chirality (where enantiomers share achiral observables), and others.

Relation to sheaf-theoretic contextuality. Abramsky and Brandenburger [Abramsky and Brandenburger, 2011] characterise quantum contextuality via non-extendability of local sections to global sections. Our bilemma captures a related but distinct obstruction: the impossibility of a faithful, stable explanation when H is maximally incompatible. The sheaf-theoretic approach provides the *mechanism* (topological obstruction); ours provides the *consequence* (property-pair achievability classification). Whether the bilemma can be expressed as a sheaf-theoretic obstruction is an open question.

Compatibility complex characterisation. The tightness classification has a topological interpretation via the *compatibility complex* K : the simplicial complex on H whose simplices are pairwise-compatible subsets. Discrete K (no edges) implies maximal incompatibility and the biletma. Contractible K (cone point = neutral element) implies faithful+stable is achievable. The 2×2 tightness table corresponds to the topology of K . This characterisation identifies the biletma as a *lower-level* obstruction than Abramsky-Brandenburger sheaf-theoretic contextuality: the biletma is an existence failure (empty stalks, $H^0 = 0$) on a discrete cover, not a gluing failure ($H^1 \neq 0$) on an overlapping cover. The approximate Rashomon extension (overlapping ε -fibers) may have genuine sheaf-theoretic content—this remains open.

Relation to structuralist philosophy of science. The trilemma framework sits within the structuralist tradition of Sneed [Sneed \[1971\]](#) and Stegmüller, which formalises scientific theories as set-theoretic predicates. Our explanation system $(\Theta, Y, H, \text{observe}, \text{explain}, \text{incomp})$ is a Sneedian structure; the trilemma is a structural constraint on admissible maps. The enrichment pattern (adding a neutral element to resolve the biletma) parallels Sneed’s treatment of theory specialisation, where new structure is adjoined to an existing core.

— Speculative Connections and Open Questions —

The following subsections document structural connections that are conjectural, historical interpretations, or beyond the current scope of empirical validation. They are included as a research program record, not as validated claims of the framework.

19.6 Scientific Revolutions as Enrichment Events [Tier C: Speculative]

The biletma characterization—enrichment is necessary iff the explanation space is maximally incompatible—maps onto the structure of scientific revolutions. Table 24 shows five instances where a paradigm shift follows the biletma pattern: a pre-revolution explanation space is maximally incompatible (the biletma proves $F+S$ impossible), and the revolution adds a neutral element (enrichment restores $F+S$ at the cost of decisiveness).

Table 24: Five instances of biletma $\rightarrow$ enrichment across science.

Transition	Pre-revolution H	Enrichment	Lost
SHAP instability	{positive, negative}	DASH tie	Feature rank
Feature selection	{selected, not}	Uncertain	Feature relevance
Classical $\rightarrow$ quantum	{definite, superposition}	No fact of the matter	Pre-measurement state
Newton $\rightarrow$ relativity	{simultaneous, not}	Frame-dependent	Absolute simultaneity
Single $\rightarrow$ adelic	{smooth, ultrametric}	Indeterminate	Spacetime geometry

The pattern suggests a conjecture: *paradigm shifts that preserve the configuration space contain an enrichment component in the sense of the biletma characterization.* The pre-revolution system is maximally incompatible. The revolution adds a neutral element. The enrichment is unique (by the enrichment uniqueness theorem) and sacrifices decisiveness (by the biletma). The all-or-nothing theorem explains why paradigm shifts are discontinuous: there is no interpolation between the pre-revolution and post-revolution frameworks.

Not all resolutions of interpretive disagreement are enrichment events. Bohmian mechanics resolves the measurement problem by changing the configuration space (adding particle positions and a guiding equation), not by enriching the explanation space with a neutral element. Many-Worlds resolves it by changing the observable space (accepting all branches as real). These are *system changes* rather than *explanation-space enrichments*, and fall outside the biletma’s scope. The biletma predicts that within a fixed system, enrichment is the only resolution. Resolutions

that change the system itself represent an alternative strategy that the framework characterizes but does not constrain.

We state this as a conjecture, not a theorem. The mathematical structure (bilemma $\rightarrow$ enrichment necessary $\rightarrow$ enrichment unique $\rightarrow$ decisiveness sacrificed) is proved. The claim that scientific revolutions follow this structure is supported by five formalized instances but requires engagement with historians and philosophers of science. The enrichment mechanism is closer to Lakatos’s progressive research programmes (where the new theory subsumes the old plus new content) than to Kuhn’s incommensurable revolutions (where the old framework’s terms cannot be translated). We use ‘paradigm shift’ loosely; the formal claim is about framework enrichment, not incommensurability.

Prospective prediction. *If the bilemma-enrichment pattern applies to scientific revolutions, then a future resolution of a maximally-incompatible domain would add a neutral element to H (restoring faithfulness + stability at the cost of decisiveness). This is consistent with the historical pattern; whether the pattern applies predictively is the open question. A concrete test: if the black hole information paradox is resolved by a framework that introduces a new category of information state (neither preserved nor destroyed, but complementary), this would instantiate the pattern. Conversely, a resolution that modifies Θ (e.g., adding firewall degrees of freedom) rather than enriching H would fall outside it.*

Falsifiability. The enrichment stack is falsifiable in two senses: (i) a physical argument showing that the Rashomon property fails at some level would cap the stack; (ii) a paradigm shift that changes Θ rather than enriching H would fall outside the pattern (e.g., Bohmian mechanics adds hidden variables to Θ rather than enriching H). The pattern predicts that H -enrichment is forced and unique when H is maximally incompatible; observing a non-unique enrichment would falsify the uniqueness claim.

The Lean formalizations use singleton observable spaces (both preparations map to the same observation), making the Rashomon property vacuously satisfied. This is a deliberate simplification: the structural claim is about the bilemma pattern (binary H , no neutral element $\rightarrow$ F+S impossible), not about the observation-matching mechanism. A richer formalization would model the specific measurement statistics that make the preparations observationally equivalent.

The framework does not distinguish between global and local symmetries. In physics, this distinction is fundamental: Noether’s theorem applies to global symmetries (producing conservation laws), while gauge redundancy arises from local symmetries (producing constraints). The framework’s G -invariant resolution applies equally to both, which is a simplification. A richer framework distinguishing global from local invariance would be needed to capture gauge-fixing and BRST cohomology.

The proof immediately yields a quantitative version. For each Rashomon pair (θ_1, θ_2) , the four-step chain forces at least one failure: unfaithfulness at θ_2 , instability on (θ_1, θ_2) , or indecisiveness at θ_1 . Summing over the Rashomon set under a probability measure: $\varepsilon + \delta + (1 - d) \geq \mu(\text{Rashomon})$, where ε is the unfaithfulness rate, δ the instability rate, d the decisiveness rate, and μ the Rashomon set measure. This tradeoff surface—which includes the binary impossibility as the special case $\varepsilon = \delta = 0$ —is developed with domain-specific bounds in forthcoming work.

The per-fiber bound (Proposition 3.10 in the main text) is the qualitative skeleton of this inequality. The full measure-theoretic development, including domain-specific instances and tightness analysis, is the subject of ongoing work.

19.7 Enrichment as Abstraction [Tier C: Speculative]

The enrichment mechanism — adding a neutral element to a maximally incompatible explanation space — has a cognitive interpretation: it is the mathematical structure of *abstraction*.

Consider a child learning to classify animals. The pre-abstraction explanation space is maximally incompatible: $H = \{\text{poodle, labrador, beagle, } \dots\}$, where each breed is incompatible with every other (incompatible $= \neq$). The Rashomon property holds: different exemplars of “the same animal” (same observable behavior) are assigned different breed labels (incompatible explanations). The biletma applies: no breed-level explanation is simultaneously faithful (matches the specific animal) and stable (consistent across behaviorally equivalent exemplars).

The resolution is abstraction: adding the concept “dog” as a neutral element (compatible with all breeds). The enriched explanation space $H' = \{\text{poodle, labrador, } \dots, \text{dog}\}$ restores faithful + stable via $E(\theta) = \text{dog}$, at the cost of decisiveness (can’t specify the breed).

The biletma thus offers a mathematical *interpretation* of why abstraction might exist: it is the unique minimal enrichment that restores stability under perceptual multiplicity. The all-or-nothing theorem explains why abstraction is discrete (“dog” or a specific breed, no interpolation). The enrichment uniqueness theorem explains why abstract concepts are canonical (the neutral element is unique on each equivalence class).

This interpretation is a conjecture requiring empirical validation through cognitive science experiments. It connects the impossibility theorem to prototype theory in cognitive science [Rosch, 1975]: prototypes are the neutral elements of perceptual categories. A full development — formalizing category formation as a sequence of enrichment events and testing against categorization experiments — is an open direction that would require engagement with the cognitive science literature.

Caveat. The analogy between mathematical enrichment and cognitive abstraction is speculative. We are not aware of cognitive science evidence connecting these concepts, and we do not develop this connection further.

19.8 Recursive Impossibility: The Enrichment Stack [Tier C: Speculative]

The enrichment mechanism of Section 3.7 resolves the biletma at one level: adding a neutral element to $\mathcal{H}$ restores $F + S$ at the cost of decisiveness. But the enriched system is itself an ExplanationSystem — and if it has its own Rashomon property at the enriched level, the biletma applies again. This recursion never terminates.

The enrichment stack. Define an *enrichment stack of depth k* as a sequence of ExplanationSystems $S_0, S_1, \dots, S_k$ where each S_{i+1} is the enrichment of S_i . At each level i , the biletma holds: no explanation of S_i is simultaneously faithful and stable (for maximally incompatible $\mathcal{H}_i$). The enrichment adds a neutral element at level i , resolving the biletma there, but the enriched system S_{i+1} may have a new Rashomon property at a finer resolution — and thus a new biletma.

The *cumulative indecisiveness theorem* (proved in Lean 4 in the companion repository): for a stack of depth k , a faithful + stable theory sacrifices decisiveness at *all* k levels. Enrichment at level i does not restore decisiveness at level $j \neq i$.

Infinite depth. The enrichment stack can have genuinely infinite depth. For the concrete system $\Theta = \mathbb{N}$, $\mathcal{H} = \text{Bool}$, $Y = \{*\}$ with $\text{explain}(\theta) = \text{bit}_k(\theta)$ (the k -th bit of θ), the biletma holds at every level k : bits 0 and 2^k agree on the observation but disagree on bit k . The enrichment at level k (adding a neutral value for bit k) does not resolve level $k + 1$ because different bits are independent. The stack has countably infinite depth, proved constructively in Lean 4 (`infiniteStack_biletma`).

The Gödel parallel. The recursive enrichment structure shares a pattern with Gödel’s incompleteness theorems. Both instantiate a common abstraction (a *recursive impossibility system*, formalized in Lean as `RecursiveImpossibility`): impossibility holds at every level, and the resolution at each level creates the conditions for the next.

	Gödel	Enrichment stack
System	Formal arithmetic	Faithful + stable theory
Three properties	Consistent, complete, r.e.	Faithful, stable, decisive
Impossibility	Can’t have all three	Can’t have all three
Escape	Add unprovable as axiom	Enrich $\mathcal{H}$ with neutral
Recursion	New system, new unprovable	New level, new bilemma

This is a shared *pattern* — both produce impossibility at every level with unbounded depth — not a deep unification. The mechanisms differ: Gödel’s arises from self-reference (the diagonal lemma), while the enrichment stack arises from the Rashomon property at finer resolutions. Both are proved as instances of `RecursiveImpossibility` in Lean (the enrichment stack proved constructively; Gödel derived from the diagonal property of formal systems).

Tarski undefinability as alternative parallel. Tarski’s undefinability theorem may provide a tighter parallel than Gödel’s incompleteness for the enrichment stack. Tarski shows that truth for a language cannot be defined within that language, forcing ascent to a metalanguage—which faces the same limitation. This hierarchy (object language $\rightarrow$ metalanguage $\rightarrow$ meta-metalanguage $\rightarrow \dots$) is structurally closer to the enrichment stack ($H_0 \rightarrow H_1 \rightarrow H_2 \rightarrow \dots$) than Gödel’s self-referential diagonal construction. We retain `RecursiveImpossibility` as capturing the shared *consequence* (impossibility at every level with unbounded depth) while acknowledging that the *mechanisms* differ: Gödel uses diagonalisation, Tarski uses definability ascent, and the trilemma uses enrichment.

The Galois analogy. The enrichment stack mirrors the algebraic structure of a tower of field extensions in Galois theory. The correspondence: the explanation space $\mathcal{H}$ plays the role of a field F ; the coverage conflict (no neutral element) is an irreducible polynomial (no root); adding a neutral element is adjoining a root; the enriched space $\mathcal{H}'$ is the extension $F(\alpha)$; new impossibility at the next level corresponds to new irreducible polynomials over $F(\alpha)$. The `infiniteStack_independence` theorem (level k enrichment gives zero information about level $k + 1$) is the analogue of linear disjointness of extensions, and the unbounded depth of the stack corresponds to the algebraic closure being infinite-dimensional. This parallel is structural (the proved independence matches linear disjointness) but not yet formalised as a functor; the full Galois correspondence—degree bounds, automorphism groups, obstruction theory—is a separate direction.

Unification consequence. If the enrichment stack continues to arbitrary depth, any unified framework (one that resolves the bilemma at every level) must sacrifice decisiveness at every level—an infinite hierarchy of indeterminacies. This is a structural analogue of the observation that no finite extension of Peano arithmetic can prove all arithmetic truths. Whether the physical stack extends beyond depth 3 is an empirical question.

Speculative application to fundamental physics. *If* the enrichment stack applies to physics, and *if* the Rashomon property holds at each level, the framework would interpret as:

- **Level 0 (quantum):** “Which measurement outcome?” — enrichment via superposition.
(*Proved; the bilemma for quantum measurement is Lean-verified.*)

- **Level 1 (adelic):** “Which geometry?” — enrichment via the indeterminate geometry. (*Lean-verified for finite fields; conjectural for number fields.*)
- **Level 2 (black hole information):** “What happens to information?” — enrichment via complementarity. (*Speculative application: the framework does not establish that the Rashomon property holds at this level.*)
- **Level 3 (spacetime emergence):** “How does spacetime emerge?” — enrichment via pre-geometric structure. (*Speculative: beyond current formalization.*)

We emphasise that Levels 2–3 represent speculative applications of the framework to open physical questions, not predictions the framework itself makes. Whether the Rashomon property holds at each putative level is precisely the open physics question; the framework says only “if it holds, the biletma applies.”

Confidence tiers for the enrichment stack. **Proved (zero domain-specific axioms):** Levels 0–1 (quantum contextuality, adelic geometry). The biletma, tightness classification, and resolution spectrum at these levels are machine-checked. **Axiomatised (conditional):** Level 2 (black hole information). Depends on the contested claim that the information paradox constitutes a genuine Rashomon pair. **Speculative:** Level 3 (spacetime emergence). Grounded in AdS/CFT and causal set theory but not empirically verifiable with current technology. A companion paper [Caraker et al., 2026b] formalises the physical enrichment stack.

20 Analogy: Arrow’s Impossibility Theorem [Tier C: Speculative]

The structure of our result parallels Arrow’s impossibility theorem in social choice theory [Arrow, 1951]. Arrow showed that no voting rule satisfying unanimity, independence of irrelevant alternatives, and non-dictatorship exists for three or more candidates—a universal impossibility that unified disparate paradoxes in voting theory. Nipkow [2009] later formalized Arrow’s theorem in Isabelle/HOL. Our Explanation Impossibility plays an analogous role for explainability: it unifies nine domain-specific impossibilities under one structural cause (the Rashomon property) and one four-line proof, and the resolution (orbit averaging) is the XAI counterpart of accepting partial orderings in social choice. The framework does not currently express the Gibbard–Satterthwaite theorem, which concerns strategy-proofness (incentive compatibility)—a property involving strategic agents, not truthful explanation. Extending the framework to capture strategic considerations would require adding an agent model, which is outside the current scope.

Comparison with topological impossibility. Chichilnisky’s theorem Chichilnisky [1980] shows that no continuous, anonymous, unanimous aggregation rule exists for preferences on a sphere. Our trilemma differs in three respects: (i) it requires no topological structure on the explanation space, (ii) the impossibility arises from the Rashomon property rather than from the topology of the preference space, and (iii) our resolution (orbit averaging) is constructive while Chichilnisky’s theorem has no finite-dimensional escape. The trilemma applies to discrete, finite systems where Chichilnisky’s theorem is vacuous.

Comparison with Fourier-theoretic generalisations. Kalai Kalai [2002] generalises Arrow’s theorem using Fourier analysis on the Boolean cube, showing that aggregation functions close to dictatorships must be close on most inputs. Our trilemma concerns explanation maps (not aggregation rules), requires the Rashomon property (not preference diversity), and applies to non-social-choice domains. The formal relationship between Kalai’s Fourier-theoretic approach and our representation-theoretic resolution (G -invariant projection) is an open question.

(The full impossibility is formalized as `ubiquity_impossibility` in `Ubiquity.lean`.)

21 Bayesian Interpretation [Tier C: Speculative]

The framework admits a natural Bayesian reading. Faithfulness corresponds to likelihood consistency (the explanation does not contradict the data-generating mechanism). Stability corresponds to posterior invariance (the explanation depends only on the observable evidence, not on which particular configuration produced it). Decisiveness corresponds to point estimation (the explanation commits to a single value rather than a distribution). The enrichment mechanism—adding a neutral element to restore faithful+stable—corresponds to Bayesian model averaging: replacing a point estimate with a posterior predictive distribution that averages over the equivalence class. The G -invariant resolution is then the orbit average $\bar{E}(\theta) = (1/|G|) \sum_{g \in G} E(g \cdot \theta)$. Under a uniform prior over the equivalence class, the orbit average coincides with the conditional expectation $E[\text{explain}(\theta) \mid \text{obs}(\theta)]$. For non-uniform priors, the orbit average and the posterior predictive diverge. The connection to Rao–Blackwell is therefore exact under uniformity and analogical in general. The Pareto-optimality of orbit averaging (Theorem 5.5) is a consequence of the admissibility of Bayes estimators under proper priors [Berger, 1985]. Specifically, under squared-error loss with a proper prior over the equivalence class, Bayes estimators are admissible (Berger, 1985, Theorem 5.2.4). The orbit average is the Bayes estimator under the uniform prior; its Pareto-optimality among stable maps follows. Under non-uniform priors, the Bayes estimator may differ from the orbit average, but the orbit average remains minimax among symmetric estimators.

22 Quantum Error Correction [Tier C: Speculative]

The dimension ratio $\eta = \dim(V^G)/\dim(V)$ has a precise interpretation in quantum error correction: it equals the code rate k/n of a quantum error-correcting code, where V^G is the code subspace invariant under the error group G . We *conjecture* that the Knill–Laflamme conditions correspond to the stability condition of the framework applied to the code subspace. A full formal bridge—analogue to the Fisher/Rao–Blackwell dictionary (Table 20)—remains an open direction. The enrichment mechanism (adding a neutral element to restore faithful+stable) corresponds to adding ancilla qubits: enlarging the physical space to create room for fault-tolerant encoding [Knill and Laflamme, 1997].

23 Classification of Impossibility Theorems by Tightness [Tier C: Speculative]

Impossibility theorems appear across science, but they differ in *tightness*: which pairs of desirable properties survive when the full triple fails. We formalise this as the `AbstractImpossibility` framework (Lean-verified, 481 theorems, 0 sorry across 38 files in a companion repository) and machine-check the *placement* of 21 classical impossibility theorems from 13+ domains into the four tightness types—each with an explicit, kernel-checked witness of which property pairs survive (e.g. CHSH by 16-strategy enumeration, Kochen–Specker by a parity argument). This atlas—not the elementary exhaustiveness theorem—is the contribution: it is, to our knowledge, the first uniform, mechanically verified classification of impossibility results by structural type.

23.1 Tightness Types [Tier C: Speculative]

Four tightness types exist (proved exhaustive and mutually exclusive):

- **Full tightness:** every pair of properties is independently achievable; only the triple conjunction fails. Resolution strategy: “pick two.”
- **Collapsed tightness:** property pairs are blocked, not just the triple. A *universal poison* property (e.g. stability, linearity) is individually incompatible with both remaining desiderata. Resolution strategy: enrich the solution space or weaken the poison property.
- **p12-blocked** and **p23-blocked:** one specific pair is blocked. These are naturally 2-ary impossibilities embedded in a 3-property framework.

23.2 The Classification [Tier C: Speculative]

Table 25: Twenty-one impossibility theorems classified by tightness type. “Structural”: non-trivial Lean proof over unbounded solution space. “Model”: 3-element type that forces its tightness by construction. Bell and Kochen-Specker tightness is decomposition-dependent. $GL(n)$ is proved for all $n \geq 2$ and all primes p over finite fields. [†]ML fairness transitions from full to PPV-blocked under performance constraints ($TPR \geq 0.3$); see text.

#	System	Domain	Tightness	Evidence
1	Bilemma	ML explanation	Collapsed	structural
2	Quantum linearity	Quantum information	Collapsed	model
3	$GL(n)$ representation	Langlands programme	Collapsed	structural
4	Arrow	Social choice	Full	structural
5	Gibbard-Satterthwaite	Social choice	Full	structural
6	Bell (CHSH)	Quantum physics	Full	structural
7	Kochen-Specker	Quantum physics	Full	structural
8	Gödel incompleteness	Mathematical logic	Full	structural
9	ML fairness [†]	ML prediction	Full [†]	model
10	Bias-variance	Statistics	Full	model
11	Sen’s liberal paradox	Social choice	Full	model
12	CAP theorem	Distributed systems	Full	model
13	FLP impossibility	Distributed systems	Full	model
14	Mundell-Fleming	Economics	Full	model
15	Perpetual motion	Thermodynamics	Full	model
16	Second law	Thermodynamics	Full	model
17	Münchhausen trilemma	Epistemology	Full	model
18	Competitive exclusion	Ecology	Full	model
19	Penrose-Hawking	General relativity	Full	model
20	Eastin-Knill	Quantum computing	p12-blocked	model
21	Shannon secrecy	Cryptography	p23-blocked	model

23.3 Key Findings [Tier C: Speculative]

Full tightness is the overwhelming default. Sixteen of 21 impossibilities have full tightness—every pair is achievable, and only the triple conjunction fails. This holds across ML, physics, logic, economics, ecology, epistemology, thermodynamics, and distributed systems. Full tightness is the “normal” state; it propagates under embeddings (*enlargement monotonicity*: if the triple conjunction is impossible and every pair is achievable, embedding into a larger solution space preserves these properties).

Collapsed tightness is rare and mechanistically specific. Only 3/21 impossibilities have collapsed tightness: the bilemma, the quantum linearity trilemma, and the $GL(n)$ representation

impossibility. All three arise from the “universal poison” mechanism: a property (P_2) that forces uniform treatment, simultaneously incompatible with two different desiderata (P_1 and P_3). For the biletma, stability is the poison. For quantum linearity, universality. For $\text{GL}(n)$, conjugation invariance (same output for conjugate group elements $\Rightarrow$ can’t faithfully track full matrix data AND can’t decisively distinguish conjugates).

The structural parallel across all three is precise: stability plays the role of linearity plays the role of conjugation invariance. All are resolved by enrichment—adding neutral elements (biletma), ancillary systems (quantum), or the trace/character ($\text{GL}(n)$).

The Langlands boundary. The $\text{GL}(n)$ instance is proved for all $n \geq 2$ and all primes p over finite fields (`GLnLanglands.lean`, companion Ostrowski repository). The Rashomon witnesses are the upper and lower unipotent matrices $I + E_{01}$ and $I + E_{10}$: same trace, distinct matrices. For $n = 1$, the trace determines the 1×1 matrix uniquely—no Rashomon, full tightness. The boundary $n = 1$ (full) vs $n \geq 2$ (collapsed) is the abelian/non-abelian boundary that organises the Langlands programme [Langlands, 1970].

The trace (= character) is the optimal stable resolution (`reynolds_best_approximation`): it is conjugation-invariant, correctly normalized ($\text{tr}(I_n) = n$), and forced by the biletma. The impossibility framework proves the correspondence’s NECESSITY—the alternative (stable tracking of full representation data) is impossible—without constructing it (which required Harris-Taylor, Wiles, Arthur, et al.).

Three known Langlands instances match the structural predictions:

- $\text{GL}(1)$: class field theory = no biletma (abelian, full tightness).
- $\text{GL}(2)$ over $\mathbb{F}_q$: Deligne-Lusztig theory = character construction = biletma resolution.
- $\text{GL}(2)$ over $\mathbb{Q}$: modularity theorem = modular forms as the stable resolution.

The structural argument (distinct conjugacy class representatives exist in any non-abelian group) extends beyond finite fields, but the Lean formalization covers the finite-field case.

The DASH–Langlands bridge. The DASH ensemble average for feature attribution and the Langlands character for representations are instances of the same algebraic structure: the Reynolds projection onto the trivial representation of the symmetry group. For DASH, the group is S_M (permutations of equivalent models); for $\text{GL}(n)$, the group acts by conjugation. Both are proved Pareto-optimal (`reynolds_best_approximation`).

The *impossibility trace formula* (`impossibility_trace_formula`) proves character orthogonality within the impossibility framework: for simple representations V and W , the inner product of characters is 1 if $V \cong W$ and 0 otherwise. This is the impossibility-theoretic reading of Schur orthogonality: characters separate representations because the biletma leaves no other option—the trace is the unique stable invariant that distinguishes non-isomorphic representations.

The *adjoint connection* (`adjoint_connection`) quantifies the information loss in both settings: for a representation of dimension d , the loss at a group element is $(d^2 - 1 - \chi_{\text{Ad}}(g))/d$, where χ_{Ad} is the adjoint character. This is the explanatory information loss identity $\|v - Rv\|^2 = \|v\|^2 - \|Rv\|^2$ applied to the conjugation action: the stable content ($\|Rv\|^2 = |\chi(g)|^2/d$) captures the character, and the unstable content ($\|v - Rv\|^2$) is the adjoint. The formula is computationally verified for $\text{GL}(2, \mathbb{F}_p)$ at $p = 3, 5, 7$ and proved algebraically for all dimensions.

The characters-separate-representations theorem (`characters_separate_reps`) proves that the character map is injective on isomorphism classes of simple representations: if two simple representations have the same character, they are isomorphic. This means the Reynolds projection (character) is not just optimal but *complete*—it loses the minimum possible information (the gauge-dependent part) while retaining all distinguishing information (the isomorphism class).

The *Reynolds naturality* theorem (`reynolds_naturality` in `UncertaintyFromSymmetry.lean`) proves that equivariant maps commute with Reynolds projections: if $\varphi : V \rightarrow W$ intertwines the group actions ($R_W \circ \varphi = \varphi \circ R_V$), then the stable resolutions are compatible. We deliberately do *not* phrase this as “predicting Langlands functoriality”: it is the elementary statement that a natural transformation commutes with a group average, carries no arithmetic content, and does not bear on the automorphic-to-Galois correspondence the Langlands programme concerns. The formal resemblance is noted only to orient readers from representation theory.

The gauge reading. The three collapsed instances share a gauge-theoretic description. In each case, a “gauge freedom” (observational equivalence between configurations) makes faithful tracking of gauge-dependent data incompatible with stability:

- **ML explanation:** the Rashomon set is the gauge orbit; SHAP values are gauge-dependent; DASH averages over the gauge orbit.
- **Quantum information:** conjugate preparations are gauge-equivalent; the decoherence-free subspace is gauge-invariant.
- **Langlands:** conjugacy classes are gauge orbits in $\mathrm{GL}(n)$; the character (trace) is the gauge-invariant observable.

The gauge principle of physics, the Langlands correspondence, and the interpretability crisis of AI are all instances of the same impossibility: you cannot stably track gauge-dependent data. The orbit average (Reynolds operator / character / DASH) is the unique optimal gauge-invariant projection.

The explanation impossibility is structurally more severe than fairness. Within ML: the blemma (collapsed) blocks property pairs, while ML fairness (full) only blocks the triple. Resolution strategies that work for fairness (“pick two”: choose which fairness metric to sacrifice) do NOT work for explanation (some pairs are blocked; must enrich). This is a precise structural distinction with practical consequences: explanation methods require qualitatively different design approaches than fairness methods.

However, the fairness tightness type is not static—it is **operating-point dependent**. Under performance constraints ($\mathrm{TPR} \geq 0.3$), the Chouldechova–Kleinberg impossibility transitions from full to PPV-blocked, with calibration as the universal poison (Section 24, item 5). At practical performance levels, the standard “pick two” advice is wrong: equalized odds is the *only* achievable pairwise compromise ($\mathrm{gap} < 0.006$), validated on Adult Income and synthetic data. This is arguably the framework’s most novel empirical finding: the tightness classification is not merely a taxonomy but a diagnostic that reveals how impossibility structure changes with operating conditions.

Intermediate types are embedded 2-ary impossibilities. Eastin–Knill (p12-blocked) and Shannon secrecy (p23-blocked) are naturally 2-property impossibilities (transversal + universal; unconditional + short-key) with a third property added. Any 2-ary impossibility can produce an intermediate type this way. The meaningful binary classification is *collapsed vs. full*.

Obstruction taxonomy. The classification determines the resolution strategy:

- *Resolvable obstruction* (collapsed tightness): the obstruction can be resolved by enriching the solution space with a neutral element. The blemma is resolvable (DASH adds “tied”; CPDAG adds “undirected”).

- *Intrinsic obstruction* (full tightness): each property must be genuinely sacrificed; no enrichment eliminates the conflict. Arrow’s theorem is intrinsic (no voting system satisfies all three properties, even with additional options).

These types are proved mutually exclusive. Full tightness implies irreducibility: no property is redundant (each pair is independently achievable, so removing any one property yields a satisfiable conjunction).

23.4 Evidence Tiers and Caveats [Tier C: Speculative]

Six instances have *structural* evidence: non-trivial Lean proofs over unbounded solution spaces with substantive tightness witnesses (bilemma, Arrow, Gibbard-Satterthwaite, Bell, Kochen-Specker, Gödel). Fourteen use *trilemma models*: 3-element types where each constructor satisfies exactly 2 of 3 properties. Trilemma models FORCE full tightness by construction; they confirm the cross-domain pattern but do not independently test the classification.

Bell (CHSH) and Kochen-Specker are naturally 2-ary impossibilities whose 3-property versions involve splitting one property (e.g. separating local realism into locality and determinism). Their full tightness is decomposition-dependent. The quantum linearity trilemma is model-level evidence (3-element type), not a structural proof over the space of all quantum channels; the bilemma has strictly stronger evidence (proved over all explanation functions on an infinite solution space).

Prior work on axiom independence for Arrow’s theorem (Blau 1972, Wilson 1972) established tightness for individual social choice results. The cross-domain classification—surveying 20 impossibilities from 12+ domains and identifying collapsed tightness as rare and mechanistically specific—is new.

Part V

Part V — Falsified Predictions and Boundaries

Predictions that were tested and rejected, and the framework’s known limits. Reported prominently by design: the boundary of a theory is part of its content.

24 Limitations and Open Questions [Tier D: Falsified]

We identify five limitations and directions for future work.

- (1) **Axiomatized Rashomon properties for ML instances.** The Rashomon property for each ML instance is axiomatized—verified to hold under stated assumptions—rather than derived from first principles of the learning algorithm. (The eight non-ML derived instances in Section 16 are the exceptions: each derives the Rashomon property constructively with zero domain-specific axioms.) Deriving the Rashomon property from, e.g., the loss landscape geometry of gradient descent would strengthen the ML results.
- (2) **Attention: retraining vs. perturbation.** The attention argmax flip rate is 19.9% under full retraining (20 independently converged DistilBERT models, the canonical Rashomon setup) and 60% under weight perturbation (dropout masks on a single model, a supplementary proxy). The 19.9% is the primary result; the 60% demonstrates that perturbation overestimates instability.
- (3) **Non-attribution quantitative bounds.** For the attribution instance, the companion paper provides sharp quantitative bounds (divergence rates, ensemble size requirements, DASH optimality). Analogous quantitative bounds for the counterfactual, concept, and fairness instances remain open.
- (4) **Decisiveness as a practical cost.** The G -invariant resolution sacrifices decisiveness: it reports ties where the data are ambiguous. Whether downstream decision-makers can tolerate partial rankings or tied causal structures—and how to present such outputs effectively—is an empirical question we do not address here.
- (5) **Fairness tightness transition.** The fairness audit impossibility (Instance 6, formalized in `FairnessAudit.lean`) has been experimentally validated. The Chouldechova–Kleinberg impossibility (calibration + equal FPR + equal FNR) exhibits a tightness *transition* under performance constraints: at $\tau = 0$ (unconstrained), full tightness holds (all pairs achievable by degenerate classifiers); at $\tau \geq 0.3$ (practical performance: $\text{TPR} \geq 30\%$ for both groups), the tightness transitions to PPV-blocked—calibration becomes a universal poison, individually incompatible with both equal FPR and equal FNR. The only achievable pairwise compromise at practical performance is equalized odds (gap < 0.006). Validated on Adult Income (sex as protected attribute, base-rate gap 0.20) and synthetic data (gap 0.25, 0.40), with bootstrap retraining (20 iterations, 100% stable). This changes the practical recommendation from “choose a fairness criterion” to “equalized odds is the only achievable criterion at practical performance; calibration is structurally incompatible.”
- (6) **Exact vs. approximate Rashomon.** The framework now includes a Lean-verified approximate extension (`ApproximateRashomon.lean`). An `ApproxExplanationSystem` adds a ‘nearness’ relation on observables, with ε -stability requiring E to be constant on approximately-equivalent observations. The approximate trilemma follows: if approximately-equivalent configurations have incompatible explanations (approximate

Rashomon), then faithful + ε -stable + decisive is impossible. The exact framework is the special case where nearness is equality. The Gaussian flip rate formula provides the quantitative bridge: the SNR measures how ‘near’ two observations must be for their explanations to flip.

- (7) **Approximate symmetry: null model comparison.** At intermediate correlations ($\rho < 1$), a simple null model $\Phi(-c\sqrt{1-\rho^2})$ with one fitted parameter ($c = 1.06$, $R^2 = 0.925$) predicts the average within-group flip rate as well as or better than the framework’s Gaussian formula ($R^2 = -6.08$ at the aggregate level). The framework’s distinctive contribution at approximate symmetry is the bimodal *structure* (within-group vs. between-group separation, dip $p < 0.002$), not the aggregate magnitude. At exact symmetry ($\rho = 1$) and cross-domain (capacity theorem across 7 domains), the framework’s predictions are unchallenged. The practical recommendation: use the Gaussian formula for per-pair prediction (Spearman 0.28 on real data), coverage conflict for per-feature screening (Spearman $\rho = 0.92$ – 0.98 per dataset, 0.96 aggregate across 15 datasets; permutation-validated), and the capacity theorem for cross-domain understanding.
- (8) **Elementary core theorem.** The qualitative impossibility follows from elementary logic: if a map is not injective, its fiber contains distinct elements that cannot be simultaneously preserved and collapsed. The contribution is not the proof’s depth but the axiomatic framework—tightness (each pair achievable), necessity (biconditional with Rashomon, machine-checked for inequality-incompatibility; sufficiency plus the fully-specified converse in general), and the quantitative corollaries (invariance counting, capacity theorem) that follow from the framework but not from any simpler observation about non-injectivity.
- (9) **Broken symmetry in trained networks.** The $1/n$ interpretability ceiling (Corollary 12.1) assumes unbroken S_n permutation symmetry. Training breaks this symmetry via gradient descent and weight initialization; the bound is an upper limit on the fraction of neuron-level interpretations that are *necessarily* stable, not a characterization of trained-network interpretability.
- (10) **Synonymous codons are not biologically equivalent.** Different codons for the same amino acid differ in translational efficiency [Ikemura, 1985], mRNA stability, and epigenetic marking (CpG dinucleotides). The gauge analogy applies to the protein-coding function only; richer models incorporating tRNA abundance and codon context would break the S_k symmetry.
- (11) **Cross-domain instances are structural analogies, not domain-specific results.** No individual domain instance was surprising to domain experts—each formalized a known phenomenon (microstate multiplicity, gauge redundancy, Markov equivalence, codon degeneracy). The formalization verifies the structural correspondence but does not produce new domain-specific results. The contribution is the unification: showing that these known phenomena share a common trilemma structure, that the trilemma strengthens to a dilemma under maximal incompatibility, and that the resolution (enrichment via neutral element) is a single mechanism underlying DASH averaging, CPDAG reporting, Bayesian model averaging, and quantum error correction.
- (12) **Group identification bottleneck.** The universal capacity theorem ($\dim(V^G)/\dim(V)$ predicts instability) achieves $R^2 = 0.957$ for seven instances with well-characterized symmetry groups but only holdout $R^2 = 0.24$ for the nine instances with approximate or unknown groups. The nine poorly-fit instances share a common cause: the assigned symmetry group was approximate. This group identification problem is fundamental, not merely technical—determining the correct symmetry group from data remains open. The SAGE algorithm sidesteps it heuristically but without a formal optimality guarantee.

The drug discovery experiment (Section 11.10) demonstrates this bottleneck concretely: Pearson-based clustering predicts 0% instability on binary fingerprints while MI-based clustering recovers the observed 23% rate within 16%. However, the Noether bimodal structure transfers only for domains with discrete group structure, not for diffuse overlap patterns.

- (13) **Coverage conflict is a conceptual naming, not a deep theorem.** The biconditional “coverage conflict $\leftrightarrow$ no neutral element” (proved in `BilemmaCharacterization.lean`) is a definitional equivalence: both express that no element of $\mathcal{H}$ is compatible with all explain-values. The value of naming this “coverage conflict” is conceptual—it identifies the algebraic feature that powers the bilemma and clarifies why binary $\mathcal{H}$ strengthens the impossibility (maximum coverage) while larger $\mathcal{H}$ weakens it (room for neutral elements).
- (14) **SNR as coverage conflict is a modeling assumption.** The connection between coverage conflict (a discrete algebraic property) and SNR (a continuous statistical quantity) rests on the assumption that importance differences are approximately Gaussian. Under this assumption, $\text{SNR} < \tau$ operationalizes “no compatible direction exists for this pair.” The assumption is validated by Shapiro-Wilk tests (60–100% of pairs pass across datasets) but is not a theorem. Ridge regression on California Housing gives $R^2 = -1.96$ due to heavy-tailed coefficient distributions; the Gaussianity diagnostic should be run before applying the formula.
- (15) **Confirmatory vs. exploratory analyses.** The stable fact count prediction, capacity theorem (well-characterized subset), and three falsified extensions were pre-registered (`PRE_REGISTRATION.md`). The coverage conflict bridge (15 datasets), clinical TreeSHAP audit (5 datasets), enrichment statistical tests, and regulatory compliance metrics are exploratory analyses conducted after observing initial results. The Noether permutation test is confirmatory (pre-specified test on a pre-registered prediction). We label each analysis accordingly throughout the text. *Caveat on temporal precedence:* `PRE_REGISTRATION.md` entered version control in the same commit as the knockout results it governs, so its precedence cannot be independently established from the repository history; and its own primary criterion—capacity $R^2 > 0.90$ on the *full* domain set—was not met (the all-domain fit is $R^2 \approx 0.60$). The headline $R^2 = 0.957$ is the pre-specified well-characterised-group subset, defined by an a-priori criterion but a subset nonetheless; we report the full-set figure alongside it and treat the capacity law as Tier B (empirical), not as a consequence of the impossibility theorem.
- (16) **Adversarial audit summary.** A comprehensive adversarial audit (24 simulated reviewers, 6 adversarial stress tests on the SAGE algorithm, 4 stress tests on the MI experiment, circularity analysis of the quantum instance, and a cross-cutting “so what” test) identified the following boundary conditions: (1) the core theorem is elementary—a consequence of definitions, not a deep result; (2) every domain instance relabels a known result; (3) the quantum η connection is correct but circular (the general POVM case relies on an unproven assertion that $\text{span}\{M_i\} = V^G$). The ExplanationSystem instance for the Peres-Mermin magic square captures the *structural analogy* between quantum contextuality and explanation instability: both involve two perspectives on the same system that demand incompatible interpretations. The combinatorial content of the Peres-Mermin theorem (the parity impossibility, verified separately as `peres_mermin_bool` in Lean) is independent of the framework; the framework shows that this content has the same trilemma structure as SHAP instability, Markov equivalence, and codon degeneracy; (4) the S_n prediction for neural networks is falsified (the actual symmetry is continuous, not discrete; see Section 11.7). The strongest surviving contributions are the invariance counting theorem

(bimodal, invariant across ρ , novel) and the SAGE algorithm ($R^2 = 0.81$ on 8 datasets, 6/6 audit tests passed).

- (17) **Frontier-scale training from scratch.** The TinyStories circuit stability experiment (Section 11.16) validates the impossibility at 16M and 45M parameter scales trained from scratch, but the GPT-2 fine-tuning configuration shows no instability because fine-tuning does not create the Rashomon property. Training 10+ models at GPT-2 Small scale (124M parameters) from independent random initialisation on a standard corpus (e.g., OpenWebText) would resolve whether the impossibility holds at production scale. This requires $\sim \$500$ – $1,000$ in compute and is the single highest-priority experiment for the research programme.
- (18) **Algorithmic identification vs. importance ranking.** The impossibility as proved concerns *importance rankings* (which component contributes most to the output). A deeper question for AI safety is whether *algorithmic identification* (what algorithm the model implements, what function each circuit computes) is also subject to the impossibility. If the function computed by “the most important head” is stable across seeds even when the head’s identity is not, the safety implications are less severe than the importance instability suggests. Functional stability analysis—examining attention patterns, OV matrices, or activation probes across seeds—would address this.
- (19) **Isotypic decomposition of attribution variance.** The G -invariant projection captures the trivial-representation component of the attribution vector. Decomposing the full attribution space into isotypic components under G would separate the variance into independent contributions from each irreducible representation, providing per-component uncertainty bounds (not just the aggregate η). For the $S_4 \times S_4$ symmetry of the grokking experiment, this would decompose the 10-dimensional attribution space into the trivial (4-dim), standard (6-dim), and potentially sign representations, each with a predictable variance contribution. The isotypic decomposition would make the Reynolds operator framework predictive at a finer granularity.
- (20) **General two-family decomposition.** The Design Space Theorem shows that G -invariant attribution systems decompose into exactly two achievable families (Family A: faithful + decisive; Family B: faithful + stable). Three instances (feature attribution, model selection, causal discovery) are proved. A general theorem—that G -invariance of ANY symmetric decision problem yields a two-family decomposition—would be a genuine contribution to impossibility theory. The connection to Hunt–Stein theory in statistical decision theory (equivariant estimators decompose into minimax families) suggests a path.

Mechanistic interpretability: confirmed with resolution demonstrated. A definitive experiment trains 10 two-layer transformers from *independent random initialisation* on modular addition ($a+b \bmod 113$), all reaching 100% test accuracy. Activation patching on 10 components (8 attention heads + 2 MLPs) yields mean Spearman $\rho = 0.518$ across 45 model pairs (permutation $p < 10^{-6}$ vs. null). The symmetry group $S_4 \times S_4$ (within-layer head permutations, derived from the architecture) gives $\dim(V^G)/\dim(V) = 4/10$. Projecting onto V^G (layer-mean head importance, MLP0, MLP1) recovers near-perfect agreement: **Spearman** $\rho = 0.929$ (Pearson $r = 0.967$, min $\rho = 0.800$). The stability is not trivially MLP-driven: the invariant excluding MLP1 still achieves $\rho = 0.822$ ($p < 10^{-6}$).

Within-layer head pairs flip at exactly 50% (12 pairs, mean = 0.500, $\sigma = 0.071$), confirming the stable fact count prediction. Head-vs-MLP1 pairs never flip (8 pairs, flip rate = 0.000). Mann-Whitney within > between: $p = 2.4 \times 10^{-5}$.

Corrected controls (matched $n = 2000$): step 50 (random init, 0.2% accuracy) $\rho = 0.300$; step 500 (memorised, 100% train/50% test) $\rho = 0.668$; step 50k (grokked) $\rho = 0.518$. Post-grokking $>$ random: $p = 2.0 \times 10^{-4}$, Cohen’s $d = 0.76$. Fourier analysis shows all 10 models learn *completely different* frequency representations (top-5 Jaccard = 0.022), confirming genuine Rashomon—same function, different algorithms.

Earlier pilot experiments were inconclusive: zero-ablation on GPT-2/SST-2 (25.8% flip) was dominated by classifier noise; LoRA fine-tuning ($\rho = 0.885$) was confounded by 99.3% shared weights. The definitive experiment eliminates both confounds.

Theorem scope vs. pipeline scope. A critical distinction: the *theorem* is universal (zero axioms, applies to any explanation system under Rashomon), but the *diagnostic pipeline* (SAGE, DASH, Gaussian formula, coverage conflict) is validated for **tabular feature-importance explanations** on classification and regression tasks. The factorial validation (30 datasets, 6 model/method configurations, Cochran’s Q confirming model-agnosticity) establishes the pipeline’s scope as: any model type, any importance method, $P \geq 11$ features, tabular data.

For non-tabular modalities—text (token importance), images (pixel/patch saliency), graphs (node/edge attribution), time series (temporal importance)—the *impossibility applies* (proved) but the *pipeline tools require adaptation*:

- **Text:** SAGE clustering on token flip rates is feasible but requires sequence-aware distance metrics (tokens are ordered, not exchangeable columns).
- **Images:** Patch-level SAGE requires spatial grouping (nearby patches may form natural orbits under translation symmetry).
- **Graphs:** Node importance under GNN Rashomon requires graph-isomorphism-aware clustering.

These adaptations are natural extensions of the framework but require domain-specific validation and are left to future collaborative work.

Extension to deep learning and LLMs. The qualitative impossibility (Theorem 3.6) applies to any model class with the Rashomon property, including deep neural networks and large language models. The quantitative bounds ($1/(1-\rho^2)$ for GBDTs) are architecture-specific and do not transfer directly. Key open problems: (i) characterising the Rashomon set geometry for overparameterised neural networks (cf. D’Amour et al. [2022]); (ii) defining the symmetry group for weight-space permutations in deep networks; (iii) computational cost of orbit averaging over weight-space symmetry groups (potentially intractable for large models, requiring approximate methods). For LLM self-explanations specifically, the token citation flip rate (34.5%) suggests the Rashomon property holds empirically, but the symmetry group structure is unknown.

25 Falsification: The Rashomon Boundary [Tier D: Falsified]

A scientific theory must be refutable. The framework makes a sharp negative prediction: the impossibility holds *if and only if* the Rashomon property is present. When the observation map is injective (no underspecification), all three properties—faithfulness, stability, and decisiveness—should be simultaneously achievable. We tested this boundary condition in three domains:

1. **Linear systems.** No Rashomon ($m = n$, square invertible): pairwise RMSD = 9.4×10^{-8} . Rashomon ($m < n$, underdetermined): RMSD = 0.112. Ratio: 1.2×10^6 .
2. **Genetic code.** No Rashomon (injective code, degeneracy = 1): Shannon entropy = 0 bits. Rashomon (standard degenerate code): entropy = 1.40 bits.




Figure 10: Falsification test across three domains. When the Rashomon property is absent (left of each pair), all three desiderata are achieved. When present (right), the impossibility manifests as nonzero RMSD, nonzero entropy, or zero decisiveness. The boundary is absolute.

- Causal discovery.** No Rashomon (fully identified DAG, $|\text{MEC}| = 1$): decisiveness = 1.0.
Rashomon (chain DAG, $|\text{MEC}| > 1$): decisiveness = 0.0.

The boundary is absolute: in every domain tested, removing the Rashomon property eliminates the impossibility entirely. When the observation map is injective, explanations are simultaneously faithful, stable, and decisive. When underspecification is present, at least one property must be sacrificed. The framework is therefore scientifically refutable—a single domain exhibiting the Rashomon property yet achieving all three properties would constitute a counterexample (Figure 10).

26 Boundary Conditions: The Quantitative Boundary and Three Falsified Extensions [Tier D: Falsified]

This section records the boundary of the framework’s *quantitative* reach. It opens with the overarching finding—that the capacity η -law fails as a point-predictor of instability magnitude—and then the three specific conjectured extensions that were pre-registered and tested, all three of which were falsified. (These three, together with the spectral-gap and flip-correlation predictions reported elsewhere in this Part, are the five pre-registered falsified predictions summarised at its close; the η -law below is the umbrella they instantiate, not a sixth separate prediction.)

The quantitative capacity law (η -law), falsified thrice. The strongest *point-predictive* conjecture—that $\eta = 1 - \dim(V^G)/\dim(V)$ predicts the *magnitude* of explanation instability—has now failed three independent ways, and we record the boundary as settled. (i) It failed prospectively (pre-registered; holdout $R^2 = 0.24$, prospective $\rho = 0.31$). (ii) From-data recovery of the integer capacity dimension under approximate symmetry is not reliably possible (five principled estimators). (iii) A structural-floor reframe motivated by the structural/statistical dichotomy—predict the large- n residual floor rather than the finite- n total—was tested and also fails to differentiate: η fits the finite- n total at least as well as the floor, and the from-data $\hat{\eta}$ still does not recover it. The lesson is epistemic, not merely empirical. The capacity theory is a *bounds* theory, not a point-predictive law; reframed as a bound—the number of transfer-confirmed stable facts is at most the capacity—it holds with 100% coverage out of sample (§8). The point-prediction is retired on evidence; the bound stands. This is the same discipline that distinguishes a distribution-free guarantee from a fitted regression throughout the program.

Phase transition at the interpolation threshold. We predicted a sharp transition in explanation instability at the overparameterization ratio $r^* = \dim(\Theta)/\dim(Y) \approx 1$. We varied r from 0.05 to 5.0 (25 log-spaced values) using Ridge regression and XGBoost (100 models per r -value, $N_{\text{obs}} = 200$). A sharp sigmoid transition was confirmed in all model classes (sigmoid $R^2 = 0.92\text{--}0.98$), but the critical point is $r^* \in [0.01, 0.12]$ —an order of magnitude below the predicted $r^* \approx 1$. The design confounds overparameterization with signal dilution: adding features adds zero-coefficient features, which are unstable regardless of the interpolation threshold.

Quantitative uncertainty bound. We conjectured $\alpha + \sigma + \delta \leq 2$, where α (faithfulness), σ (stability), and δ (decisiveness) are continuous scores. The maximum observed sum was 2.86 (linear ensemble: $\alpha = 0.873$, $\sigma = 0.986$, $\delta = 1.0$), falsifying the bound. The binary impossibility—no method achieves $\alpha = \sigma = \delta = 1.0$ —remains confirmed (maximum $\alpha = 0.873$). The trilemma manifests as an exclusion zone around $(1, 1, 1)$, not a linear budget constraint.

Molecular evolution from character theory. We tested whether the S_k representation theory predicts synonymous substitution structure beyond simple biochemistry. The neighbor-count model (number of synonymous 1-mutation variants per codon) achieves $R^2 = 1.0$; the character-theoretic $(k-1)/k$ model achieves $R^2 = 0.65$ per codon but adds zero residual variance after controlling for neighbor count (partial $R^2 = 0.0$, $p = 1.0$). The codon table’s structure is fully explained by local nucleotide degeneracy, not abstract group theory.

Codon null model. The Spearman $\rho = 1.0$ between degeneracy and usage entropy is a calibration check, not an independent validation. Under uniform codon usage within each degeneracy class, entropy $= \log_2 d$, which is monotone in d ; the rank correlation is therefore 1.0 by construction. The empirical $\rho = 1.0$ confirms that observed entropy is monotonically related to degeneracy (consistent with the framework) but does not discriminate between the trilemma prediction and any model predicting monotonicity. The constructive Rashomon witness (zero domain-specific axioms) is the primary validation of the codon instance.

Two additional knockout candidates from the approximate-symmetry regime were tested and eliminated:

Spectral gap $\rightarrow$ ensemble size. The spectral gap of S_k predicts a convergence rate $\sim k^{-M}$ for ensemble flip rate as a function of ensemble size M . Empirically, the predicted rate is 14–100 $\times$ too fast (even at $\rho = 0.999$). The exact-group spectral gap does not predict approximate-group convergence.

Flip correlations from irreducible decomposition. The irreducible decomposition of the S_k representation predicts that features in the same irreducible should flip *together* (correlated flip indicators). Empirically, within-group flip correlation (-0.031) $\approx$ between-group (-0.013): features in the same irreducible do NOT flip together.

Summary. In total, five pre-registered predictions were falsified: phase transition location, quantitative uncertainty bound, molecular evolution from character theory, spectral gap convergence rate, and flip correlations from irreducible decomposition. These falsifications are scientifically valuable: they define the boundary between the framework’s proven domain (the impossibility theorem and its resolutions) and its speculative extensions. The impossibility and resolution work universally; the quantitative predictions work only when the symmetry group is exactly known and features are genuinely exchangeable.

Proportionality axiom sensitivity. The design space theorem assumes exact proportionality ($\text{CV}(c) = 0$). Simulation across $\text{CV} \in [0, 1.0]$ (1,000 trials per value, 11 values) confirms the two-family structure is robust: Family A flip rate ≈ 0.50 at all CV values, Family B tie rate ≥ 0.93 at all CV values, and zero Pareto violations across 11,000 trials. The instability source is collinearity-driven split competition, not the proportionality constant. The core impossibility (Theorem 3.6) is independent of proportionality.

27 Attribution PCA Eigenspectrum: A Negative Result [Tier D: Falsified]

We tested whether PCA on the SHAP value matrix across multiple models recovers the symmetry group structure. On synthetic data with $g = 3$ groups of $k = 4$ features ($\rho = 0.95$), the eigenspectrum gap appears at PC 10 (expected: PC 3), and $\eta_{\text{PCA}} = 0.50$ (expected: 0.75). The gap only matches the group count at $\rho = 0.99$ (near-exact symmetry). On the NARPS dataset (48 teams $\times$ 700 dimensions), the gap appears at PC 1 (expected: PC 7 for 7 networks). On Breast Cancer data, the stable-PC loading vs. flip rate correlation is $\rho = -0.27$ ($p = 0.15$ —not significant).

The eigenspectrum approach fails because PCA captures variance (cardinal), not rankings (ordinal). A feature with high SHAP variance can have a perfectly stable sign; a feature with low variance can flip easily near zero. The flip rate matrix (pairwise ranking instability) is the correct object for capturing ranking stability, not the SHAP covariance matrix. This negative result clarifies that extending the capacity theorem to approximate symmetry requires rank-based methods, not spectral decomposition of cardinal values.

References

- Samson Abramsky and Adam Brandenburger. The sheaf-theoretic structure of non-locality and contextuality. *New Journal of Physics*, 13:113036, 2011.
- Julius Adebayo, Justin Gilmer, Michael Muelly, Ian Goodfellow, Moritz Hardt, and Been Kim. Sanity checks for saliency maps. In *Advances in Neural Information Processing Systems*, volume 31, 2018.
- Samuel K. Ainsworth, Jonathan Hayase, and Siddhartha Srinivasa. Git re-basin: Merging models modulo permutation symmetries. In *International Conference on Learning Representations*, 2023.
- Anthropic. Circuit tracing: Revealing computational mechanisms in language models. Technical report, Anthropic, 2025.
- Kenneth J Arrow. *Social Choice and Individual Values*. Wiley, 1951.
- François Bancilhon and Nicolas Spyratos. Update semantics of relational views. *ACM Transactions on Database Systems*, 6(4):557–575, 1981.
- James O. Berger. *Statistical Decision Theory and Bayesian Analysis*. Springer Series in Statistics. Springer, 2nd edition, 1985.
- Blair Bilodeau, Natasha Shandilya, Blossom Gao, and Umang Bhatt. Impossibility theorems for feature attribution. *Proceedings of the National Academy of Sciences*, 121(2):e2304406120, 2024.
- David Blackwell. Conditional expectation and unbiased sequential estimation. *The Annals of Mathematical Statistics*, 18(1):105–110, 1947.

- Tolga Bolukbasi, Kai-Wei Chang, James Zou, Venkatesh Saligrama, and Adam Kalai. Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. In *Advances in Neural Information Processing Systems*, volume 29, 2016.
- Leo Breiman. Random forests. *Machine Learning*, 45:5–32, 2001.
- Drake Caraker, Bryan Arnold, and David Rhoads. The attribution impossibility: No feature ranking is faithful, stable, and complete under collinearity. *arXiv preprint*, 2026a. Companion paper. Software: <https://github.com/DrakeCaraker/dash-shap>.
- Drake Caraker, Bryan Arnold, and David Rhoads. Impossibility of faithful stable explanation under Ostrowski’s classification: A formally verified theorem applied to spacetime geometry. *Foundations of Physics*, 2026b. Under review.
- Graciela Chichilnisky. Social choice and the topology of spaces of preferences. *Advances in Mathematics*, 37(2):165–176, 1980.
- David Maxwell Chickering. Optimal structure identification with greedy search. *Journal of Machine Learning Research*, 3:507–554, 2002.
- Noam Chomsky. *Syntactic Structures*. Mouton, The Hague, 1957.
- Alexandra Chouldechova. Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2):153–163, 2017.
- Bilal Chughtai, Lawrence Chan, and Neel Nanda. A toy model of universality: Reverse engineering how networks learn group operations. In *International Conference on Machine Learning*, pages 6243–6267. PMLR, 2023.
- Arthur Conmy, Augustine N Mavor-Parker, Aidan Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. In *Advances in Neural Information Processing Systems*, 2023.
- Francis H C Crick. Codon—anticodon pairing: the wobble hypothesis. *Journal of Molecular Biology*, 19(2):548–555, 1966.
- Alexander D’Amour, Katherine Heller, Dan Moldovan, et al. Underspecification presents challenges for credibility in modern machine learning. *Journal of Machine Learning Research*, 23(226):1–61, 2022.
- Shamik Dasgupta. Symmetry as an epistemic notion (twice over). *British Journal for the Philosophy of Science*, 67(3):837–878, 2016.
- Leonardo de Moura and Sebastian Ullrich. The Lean 4 theorem prover and programming language. In *International Conference on Automated Deduction (CADE)*, 2021.
- Thomas Decker, Luisa-Marie Herm, Robert Feldhans, Michael Kamp, and Barbara Hammer. Provably better explanations with optimized aggregation of feature attributions. In *International Conference on Machine Learning*, 2024.
- Arthur P. Dempster, Nan M. Laird, and Donald B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society: Series B (Methodological)*, 39(1):1–38, 1977.
- David L. Donoho and Philip B. Stark. Uncertainty principles and signal recovery. *SIAM Journal on Applied Mathematics*, 49(3):906–931, 1989.

- Jon Dorling. Bayesian personalism, the methodology of scientific research programmes, and Duhem’s problem. *Studies in History and Philosophy of Science*, 10(3):177–187, 1979.
- Pierre Duhem. *The Aim and Structure of Physical Theory*. Princeton University Press, 1954. English translation of the 1906 French original.
- Gerald M Edelman and Joseph A Gally. Degeneracy and complexity in biological systems. *Proceedings of the National Academy of Sciences*, 98(24):13763–13768, 2001.
- Ronen Eldan and Yuanzhi Li. Tinstories: How small can language models be and still speak coherent english? In *International Conference on Learning Representations*, 2023.
- Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, Roger Grosse, Sam McCandlish, Jared Kaplan, Dario Amodei, Martin Wattenberg, and Christopher Olah. Toy models of superposition. *Transformer Circuits Thread*, 2022.
- EU Parliament. Regulation (EU) 2024/1689 (artificial intelligence act). Official Journal of the European Union, OJ L, 12.7.2024, 2024.
- Aaron Fisher, Cynthia Rudin, and Francesca Dominici. All models are wrong, but many are useful: Learning a variable’s importance by studying an entire class of prediction models simultaneously. *Journal of Machine Learning Research*, 20(177):1–81, 2019.
- Ronald A. Fisher. On the mathematical foundations of theoretical statistics. *Philosophical Transactions of the Royal Society of London. Series A*, 222:309–368, 1922.
- Amirata Ghorbani, Abubakar Abid, and James Zou. Interpretation of neural networks is fragile. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 3681–3688, 2019.
- Allan Gibbard. Manipulation of voting schemes: A general result. *Econometrica*, 41(4):587–601, 1973.
- Nicholas Goldowsky-Dill, Chris MacLeod, Lucas Sato, and Aryaman Arora. Localizing model behavior with path patching. In *arXiv preprint arXiv:2304.05969*, 2023.
- Herbert Hauptman and Jerome Karle. *Solution of the Phase Problem. I. The Centrosymmetric Crystal*, volume 3. Polycrystal Book Service, 1953.
- Andrew Herren and P Richard Hahn. Statistical inference for feature rankings under model multiplicity. *Journal of Computational and Graphical Statistics*, 2023.
- Ruth Hershberg and Dmitri A Petrov. Selection on codon bias. *Annual Review of Genetics*, 42: 287–299, 2008.
- Sara Hooker, Dumitru Erhan, Pieter-Jan Kindermans, and Been Kim. A benchmark for interpretability methods in deep neural networks. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Colin Howson and Peter Urbach. *Scientific Reasoning: The Bayesian Approach*. Open Court, 3rd edition, 2006.
- Xuanxiang Huang and Joao Marques-Silva. On the failings of Shapley values for explainability. *International Journal of Approximate Reasoning*, 171:109112, 2024.
- Toshimichi Ikemura. Codon usage and tRNA content in unicellular and multicellular organisms. *Molecular Biology and Evolution*, 2(1):13–34, 1985.

- John David Jackson. *Classical Electrodynamics*. Wiley, New York, 3rd edition, 1999.
- Sarthak Jain and Byron C Wallace. Attention is not explanation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, pages 3543–3556, 2019.
- Edwin T Jaynes. Information theory and statistical mechanics. *Physical Review*, 106(4):620–630, 1957.
- Daniel Jurafsky and James H. Martin. *Speech and Language Processing*. Prentice Hall, 3rd edition, 2024. Draft available at <https://web.stanford.edu/~jurafsky/slp3/>.
- Gil Kalai. A Fourier-theoretic perspective on the Condorcet paradox and Arrow’s theorem. *Advances in Applied Mathematics*, 29(3):412–426, 2002.
- Amir-Hossein Karimi, Gilles Barthe, Borja Balle, and Isabel Valera. Algorithmic recourse under imperfect causal knowledge: a probabilistic approach. In *Advances in Neural Information Processing Systems*, volume 33, pages 265–277, 2020.
- Jerome Karle and Herbert Hauptman. The symbolic addition procedure for phase determination for centrosymmetric and non-centrosymmetric crystals. *Acta Crystallographica*, 9(7):635–651, 1956.
- Been Kim, Martin Wattenberg, Justin Gilmer, Carrie Cai, James Wexler, Fernanda Viegas, and Rory Sayres. Interpretability beyond classification error: Applying testing with concept activation vectors (TCAV). In *International Conference on Machine Learning*, pages 2668–2677, 2018.
- Jon Kleinberg, Sendhil Mullainathan, and Manish Raghavan. Inherent trade-offs in the fair determination of risk scores. In *Innovations in Theoretical Computer Science (ITCS)*, 2017.
- Emanuel Knill and Raymond Laflamme. Theory of quantum error-correcting codes. *Physical Review A*, 55(2):900–911, 1997.
- Satyapriya Krishna, Tessa Han, Alex Ber, Shaun Buzzelli, Michael Kearns, et al. The disagreement problem in explainable machine learning: A practitioner’s perspective. In *NeurIPS 2022 Workshop on Trustworthy and Socially Responsible Machine Learning*, 2022.
- Yury G. Kudryashov. SardMoreira: a Lean 4 formalisation of Moreira’s version of Sard’s theorem. <https://github.com/urkud/SardMoreira>, commit 14bc8a1, 2025. Apache-2.0 licensed; vendored and forward-ported in this work.
- Ivan Kupka. Counterexample to the Morse–Sard theorem in the case of infinite-dimensional manifolds. *Proceedings of the American Mathematical Society*, 16(5):954–957, 1965.
- Gabriel Laberge, Yann Pequignot, Alexandre Mathieu, Foutse Khomh, and Mario Marchand. Partial order in chaos: Consensus on feature attributions in the rashomon set. *Journal of Machine Learning Research*, 24(364):1–50, 2023.
- James Ladyman, Don Ross, David Spurrett, and John Collier. *Every Thing Must Go: Metaphysics Naturalized*. Oxford University Press, 2007.
- Robert P. Langlands. Problems in the theory of automorphic forms. *Lectures in Modern Analysis and Applications III*, pages 18–61, 1970.

- Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, Kamilé Lukićsić, Karina Nguyen, Nicholas Schiefer, Catherine Olsson, Tom Henighan, Hume Jones, Arthur Jermyn, Chris Olah, Dario Amodei, Ethan Perez, Jared Kaplan, and Jack Clark. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023.
- Larry Laudan and Jarrett Leplin. Empirical equivalence and underdetermination. *The Journal of Philosophy*, 88(9):449–472, 1991.
- Erich L Lehmann and Joseph P Romano. *Testing Statistical Hypotheses*. Springer, 3rd edition, 2005.
- Scott M Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Charles F Manski. *Partial Identification of Probability Distributions*. Springer, 2003.
- Eve Marder and Jean-Marc Goaillard. Variability, compensation and homeostasis in neuron and network function. *Nature Reviews Neuroscience*, 7(7):563–574, 2006.
- Charles Marx, Flavio Calmon, and Berk Ustun. But are you sure? an uncertainty-aware perspective on explainability. In *Advances in Neural Information Processing Systems*, 2024.
- Nicolai Meinshausen and Peter Bühlmann. Stability selection. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 72(4):417–473, 2010.
- Maxime Méloux, Silviu Maniu, François Portet, and Maxime Peyrard. Everything, everywhere, all at once: Is mechanistic interpretability identifiable? In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2502.20914.
- Christoph Molnar, Gunnar König, Julia Herbringer, Timo Freiesleben, Giuseppe Casalicchio, Marvin N Wright, and Bernd Bischl. General pitfalls of model-agnostic interpretation methods for machine learning models. *International Journal of Data Science and Analytics*, 2022.
- Carlos Gustavo T. de A. Moreira. Hausdorff measures and the Morse–Sard theorem. *Publicacions Matemàtiques*, 45(1):149–162, 2001.
- Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. *arXiv preprint arXiv:2301.05217*, 2023.
- Jerzy Neyman. Su un teorema concernente le cosiddette statistiche sufficienti. *Giornale dell’Istituto Italiano degli Attuari*, 6:320–334, 1935.
- Tobias Nipkow. Social choice theory in HOL: Arrow and Gibbard-Satterthwaite. *Journal of Automated Reasoning*, 43:289–304, 2009.
- Emmy Noether. Invariante Variationsprobleme. *Nachrichten von der Gesellschaft der Wissenschaften zu Göttingen, Mathematisch-Physikalische Klasse*, 1918:235–257, 1918.
- Miquel Noguer i Alonso. Mathematical foundations of explainability. *SSRN Working Paper*, 2025.
- Office of the Comptroller of the Currency. Supervisory guidance on model risk management (SR 11-7/OCC 2011-12). OCC Bulletin 2011-12, 2011.
- Samir Okasha. Theory choice and social choice: Kuhn versus Arrow. *Mind*, 120(477):83–115, 2011.

- Gonçalo Paulo and Nora Belrose. Sparse autoencoders trained on the same data learn different features. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2501.16615.
- Willard Van Orman Quine. Two dogmas of empiricism. *The Philosophical Review*, 60(1):20–43, 1951.
- C. Radhakrishna Rao. Information and the accuracy attainable in the estimation of statistical parameters. *Bulletin of the Calcutta Mathematical Society*, 37:81–91, 1945.
- Shrisha Rao. The limits of AI explainability: An algorithmic information theory approach. *arXiv preprint arXiv:2504.20676*, 2025.
- Ieva Rauluseviciute, Rafael Riudavets-Puig, Romain Blanc-Mathieu, et al. JASPAR 2024: 20th anniversary of the open-access database of transcription factor binding profiles. *Nucleic Acids Research*, 52(D1):D142–D147, 2024.
- Shauli Ravfogel, Yanai Elazar, Hila Gonen, Michael Twiton, and Yoav Goldberg. Null it out: Guarding protected attributes by iterative nullspace projection. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 7237–7256, 2020.
- Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. “why should i trust you?”: Explaining the predictions of any classifier. In *ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016.
- Eleanor Rosch. Cognitive representations of semantic categories. *Journal of Experimental Psychology: General*, 104(3):192–233, 1975.
- Cynthia Rudin et al. Position: Amazing things come from having many good models. In *International Conference on Machine Learning (ICML)*, 2024.
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- Mark Allen Satterthwaite. Strategy-proofness and Arrow’s conditions: Existence and correspondence theorems for voting procedures and social welfare functions. *Journal of Economic Theory*, 10(2):187–217, 1975.
- Andrew D Selbst and Solon Barocas. The intuitive appeal of explainable machines. *Fordham Law Review*, 87:1085, 2018.
- Ramprasaath R Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh, and Dhruv Batra. Grad-CAM: Visual explanations from deep networks via gradient-based localization. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 618–626, 2017.
- Lesia Semenova, Cynthia Rudin, and Ronald Parr. On the existence of simpler machine learning models. *ACM/IMS Journal of Data Science*, 2022.
- Claude E. Shannon. A mathematical theory of communication. *The Bell System Technical Journal*, 27(3):379–423, 1948.
- Paul M Sharp and Wen-Hsiung Li. The codon adaptation index—a measure of directional synonymous codon usage bias, and its potential applications. *Nucleic Acids Research*, 15(3):1281–1295, 1987.

- Dylan Slack, Sophie Hilgard, Emily Jia, Sameer Singh, and Himanshu Lakkaraju. Fooling LIME and SHAP: Adversarial attacks on post hoc explanation methods. In *AAAI/ACM Conference on AI, Ethics, and Society*, 2020.
- Stephen Smale. An infinite dimensional version of Sard’s theorem. *American Journal of Mathematics*, 87(4):861–866, 1965.
- Joseph D Sneed. *The Logical Structure of Mathematical Physics*. Reidel, 1971.
- P Kyle Stanford. *Exceeding Our Grasp: Science, History, and the Problem of Unconceived Alternatives*. Oxford University Press, 2006.
- Gilbert Strang. *Introduction to Linear Algebra*. Wellesley-Cambridge Press, 5th edition, 2016.
- Bernd Sturmfels. *Algorithms in Invariant Theory*. Texts and Monographs in Symbolic Computation. Springer, 2nd edition, 2008.
- Mukund Sundararajan and Amir Najmi. The many Shapley values for model explanation. In *International Conference on Machine Learning*, 2020.
- The mathlib Community. The Lean mathematical library. In *Certified Programs and Proofs (CPP)*, 2020.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. *Advances in Neural Information Processing Systems*, 36, 2024.
- Berk Ustun, Alexander Spangher, and Yang Liu. Actionable recourse in linear classification. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAccT)*, pages 10–19, 2019.
- Bas C van Fraassen. *The Scientific Image*. Oxford University Press, 1980.
- Thomas Verma and Judea Pearl. Equivalence and synthesis of causal models. In *Conference on Uncertainty in Artificial Intelligence*, pages 255–270, 1991.
- Sarah Wiegrefe and Yuval Pinter. Attention is not not explanation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 11–20, 2019.
- John Worrall. Structural realism: The best of both worlds? *Dialectica*, 43(1–2):99–124, 1989.
- Xi Ye and Greg Durrett. The unreliability of explanations in few-shot in-context learning. *Advances in Neural Information Processing Systems*, 35, 2022.
- Yuanhe Zhang, Jason D Lee, and Fanghui Liu. Statistical learning theory in Lean 4: Empirical processes from scratch. *arXiv preprint arXiv:2602.02285*, 2026. URL <https://arxiv.org/abs/2602.02285>.

A Arrow’s Impossibility Theorem: A Structural Comparison

Two impossibility theorems constrain aggregation under conflicting desiderata: Arrow’s impossibility theorem in social choice [Arrow, 1951] and the Explanation Impossibility developed in this work. Both assert that a small set of individually reasonable properties cannot be jointly satisfied. This section maps the axioms of each theorem onto the other, identifies the precise point where the correspondence breaks, and characterizes the complementary regimes in which the two results operate.

A.1 Formal setup

Arrow’s theorem. Let A be a finite set of alternatives with $|A| \geq 3$, and let $N = \{1, \dots, n\}$ with $n \geq 2$ be a set of voters. Write $\mathcal{L}(A)$ for the set of strict linear orders on A (complete, transitive, antisymmetric, irreflexive). A *preference profile* is a tuple $\pi = (\pi_1, \dots, \pi_n) \in \mathcal{L}(A)^N$. A *social welfare function* is a map $F : \mathcal{L}(A)^N \rightarrow \mathcal{L}(A)$. Arrow’s theorem states that no F satisfies all four of the following axioms:

- A1. Unrestricted Domain (UD).** F is defined on all of $\mathcal{L}(A)^N$.
- A2. Weak Pareto (WP).** If $a \succ_{\pi_i} b$ for every voter i , then $a \succ_{F(\pi)} b$.
- A3. Independence of Irrelevant Alternatives (IIA).** For all $a, b \in A$: if profiles π and π' agree on every voter’s ranking of a versus b , then $F(\pi)$ and $F(\pi')$ agree on a versus b .
- A4. Non-Dictatorship (ND).** There is no voter i such that $F(\pi) = \pi_i$ for all π .

The explanation impossibility. An EXPLANATIONSYSTEM $(\Theta, H, Y, \text{obs}, \text{exp}, \text{incomp})$ consists of a configuration space Θ , an explanation space H , an observable space Y , an observation map $\text{obs} : \Theta \rightarrow Y$, a native explanation map $\text{exp} : \Theta \rightarrow H$, and an irreflexive incompatibility relation on H . We say the system has the *Rashomon property* if there exist $\theta_1, \theta_2 \in \Theta$ with $\text{obs}(\theta_1) = \text{obs}(\theta_2)$ and $\text{incomp}(\text{exp}(\theta_1), \text{exp}(\theta_2))$. For any such system, no explanation $E : \Theta \rightarrow H$ is simultaneously:

- E1. Faithful.** $\neg \text{incomp}(E(\theta), \text{explain}(\theta))$ for all θ .
- E2. Stable.** $\text{observe}(\theta_1) = \text{observe}(\theta_2) \Rightarrow E(\theta_1) = E(\theta_2)$.
- E3. Decisive.** $\text{incomp}(\text{explain}(\theta), h) \Rightarrow \text{incomp}(E(\theta), h)$ for all θ, h .

This is `explanation_impossibility` in the Lean formalization, proved from the Rashomon property alone with zero auxiliary axioms.

A.2 The axiom correspondence

We now map each Arrow axiom to the explanation framework. The mapping uses the *query-relative* formulation (Lean: `QueryRelative.lean`), in which incompatibility is parameterized by a query space Q , with each query $q \in Q$ inducing its own incompatibility relation incomp_q on H .

IIA $\longleftrightarrow$ query-relative stability (exact). Fix a pair $q = (a, b) \in A \times A$ as a query. Define the query-relative observation map $\text{observe}_q(\pi) = (\mathbf{1}[a \succ_{\pi_i} b])_{i \in N}$, the vector of individual pairwise preferences restricted to (a, b) . IIA states: if $\text{observe}_q(\pi) = \text{observe}_q(\pi')$, then the social ranking of a versus b agrees under $F(\pi)$ and $F(\pi')$. This is precisely query-relative stability (Lean: `q_faithful`, `q_decisive` in `QueryRelative.lean`) applied to the pairwise query q : E ’s answer to “does a rank above b ?” depends only on the pairwise observations. The correspondence is exact; no content is lost or added.

Arrow axiom	Explanation property	Quality
IIA	Query-relative stability	Exact
Weak Pareto	Faithfulness (restricted)	Partial
Non-Dictatorship	(no analog)	None
Unrestricted Domain	(implicit totality)	Structural

Table 26: Axiom correspondence between Arrow’s theorem and the explanation impossibility. Only IIA admits an exact translation.

Weak Pareto $\longleftrightarrow$ faithfulness (partial). WP says: if every voter ranks $a \succ b$, then F ranks $a \succ b$. Faithfulness says: $E(\theta)$ never contradicts $\text{explain}(\theta)$. The mapping is indirect. If we set $\text{explain}(\pi) = F(\pi)$, then faithfulness requires that E not contradict F . If F satisfies WP, faithfulness of E to F automatically respects Pareto—but only transitively. Pareto constrains F directly (unanimous preferences are respected); faithfulness constrains E relative to F . This is a genuine structural difference: Pareto is a first-order constraint on the aggregation map, while faithfulness is a second-order constraint on an *explanation* of the aggregation map.

Non-Dictatorship (no analog). ND constrains the internal structure of F : no single voter determines the output. The explanation framework has no comparable axiom. Our three properties—faithful, stable, decisive—constrain an external explanation E relative to a given system S ; they say nothing about the internal structure of the system’s own maps. In Arrow’s framework, F is *both* the system and the object being evaluated. In our framework, explain is the system’s native output, and E is a separate map that interprets it. The E/F split has no counterpart in social choice.

Unrestricted Domain (structural). Both frameworks implicitly assume totality: F is defined on all profiles, and E is defined on all configurations. UD plays the same structural role in each—it ensures the impossibility is not circumvented by restricting the domain.

A.3 Why the embedding fails

The axiom correspondence of Section A.2 suggests a natural question: is Arrow’s theorem an instance of the explanation impossibility? The answer is no, and the reason is precise.

Key insight. *IIA prevents the Rashomon property from holding.*

To apply the explanation impossibility, we need the Rashomon property: two configurations θ_1, θ_2 with the same observation but incompatible explanations. In the social choice mapping with query $q = (a, b)$, the observation is $\text{observe}_q(\pi) = (\mathbf{1}[a \succ_{\pi_i} b])_{i \in N}$ and the explanation is the social ranking of a versus b under $F(\pi)$. Rashomon at query q would require two profiles π, π' with $\text{observe}_q(\pi) = \text{observe}_q(\pi')$ but $F(\pi)$ and $F(\pi')$ disagreeing on (a, b) .

But this is *exactly what IIA forbids*. If F satisfies IIA, then $\text{observe}_q(\pi) = \text{observe}_q(\pi')$ implies that $F(\pi)$ and $F(\pi')$ agree on (a, b) . The social ranking of any pair is a deterministic function of the individual pairwise preferences. There is no Rashomon ambiguity for any query: the observation uniquely determines the explanation.

The two theorems therefore operate in *complementary regimes*:

Regime	Impossibility
Rashomon present, IIA not assumed	No E is faithful + stable + decisive
IIA enforced, no pairwise Rashomon	Only dictatorships satisfy WP + ND

The explanation impossibility says: when the map from observations to explanations is many-to-one (Rashomon), you cannot explain faithfully, stably, and decisively. Arrow’s theorem says: when the map from observations to the social ranking is *required* to be one-to-one per query (IIA), the only aggregation rules that also respect unanimity are dictatorial. Different failure modes of aggregation, activated by opposite structural assumptions.

This complementarity is not a deficiency of either theorem. It reflects a genuine difference in what each result constrains. The explanation impossibility is about *underspecification*: multiple models fit the data, and no single explanation can faithfully represent them all. Arrow’s theorem is about *preference diversity*: multiple voters have conflicting orderings, and no aggregation rule can satisfy a short list of fairness criteria. Underspecification and preference diversity are distinct sources of impossibility, and neither subsumes the other.

A.4 What the theorems share

Despite the embedding failure, the theorems share substantial structure:

1. **Both are impossibility results about aggregation.** Arrow aggregates individual preferences into a social ordering; the explanation impossibility aggregates a Rashomon set of models into a single explanation.
2. **Both have three independent desiderata.** Once unrestricted domain (or totality of E) is factored out as a structural assumption, Arrow has WP, IIA, and ND; the explanation impossibility has faithfulness, stability, and decisiveness.
3. **Both are tight.** Dropping any single axiom makes the remaining set satisfiable. For Arrow: dropping ND yields dictatorship; dropping IIA yields Borda count; dropping WP yields anti-dictatorships. For the explanation impossibility: dropping decisiveness yields G -invariant projections (e.g., CPDAG, DASH); dropping stability yields the native explanation; dropping faithfulness yields an arbitrary constant.
4. **Both have the same proof shape.** Short diagonal arguments. Arrow’s proof (in the Geanakoplos version) constructs a pivotal voter by toggling an alternative between the top and bottom of one voter’s ranking. The explanation impossibility constructs a contradiction by chaining decisiveness, stability, and faithfulness through two Rashomon-equivalent configurations. Both are four-to-five-step proofs.
5. **Both have been formalized in proof assistants.** Nipkow [2009] formalized Arrow’s theorem and the Gibbard–Satterthwaite theorem [Gibbard, 1973, Satterthwaite, 1975] in Isabelle/HOL. The explanation impossibility is formalized in Lean 4 (this work), with zero axiom dependencies.
6. **Both involve a map from a product space to a codomain, where natural conditions are jointly unsatisfiable.** Arrow: $F : \mathcal{L}(A)^N \rightarrow \mathcal{L}(A)$, conditions WP + IIA + ND. Explanation: $E : \Theta \rightarrow H$, conditions faithful + stable + decisive. The abstract shape is the same: a function from a structured input to a structured output, subject to three constraints that form an inconsistent triad.

The shared structure is categorical in nature. Both theorems say: given a map from a product of structured objects to a codomain, imposing compatibility with each factor (faithfulness/-Pareto), locality (stability/IIA), and non-degeneracy (decisiveness/non-dictatorship) yields a

contradiction. The difference is in what prevents joint satisfaction: *underspecification* (multiple models fit the same data) versus *preference diversity* (multiple voters hold conflicting orderings).

A.5 Implications

Independence, not generalization. The explanation impossibility is not a generalization of Arrow’s theorem. The two results are logically independent: neither implies the other. Their shared abstract structure reflects a common pattern—the impossibility of aggregation under conflicting locality and non-degeneracy constraints—not a subsumption relationship. Any claim that one “contains” the other would require either fabricating a Rashomon property where IIA prevents one from existing, or fabricating a non-dictatorship axiom where the explanation framework provides none.

The query-relative bridge. The identification $\text{IIA} = \text{query-relative stability}$ could not have been stated without the query parameterization of the explanation impossibility (Section A.1; Lean: `QueryRelative.lean`). The standard (non-query-relative) framework has a single global incompatibility relation, which does not capture IIA’s pairwise structure. The query-relative generalization is therefore not merely a technical refinement: it is the precise bridge between the two frameworks.

Toward a unified categorical framework. A unified framework subsuming both theorems would need to abstract over the *source of impossibility*: Rashomon (many-to-one observation map) versus preference diversity (many voters with conflicting orderings). In both cases, a structured input space admits conflicting “local views” that no global map can reconcile. Formalizing this common pattern—perhaps as a sheaf-theoretic obstruction or a failure of a limit in a suitable category—is a direction for future work.

Relation to Okasha (2011). Okasha [2011] applied Arrow’s theorem to theory choice, arguing that Kuhnian criteria (accuracy, simplicity, scope, fruitfulness, consistency) play the role of voters, and scientific theories play the role of alternatives. This is a different exercise from ours. Okasha asks: does Arrow’s theorem constrain theory choice? We ask: how does Arrow’s theorem *compare structurally* with the explanation impossibility? Okasha’s analysis uses Arrow as a tool; ours uses Arrow as a point of comparison. The findings are complementary: Okasha shows that preference aggregation is relevant to philosophy of science; we show that the explanation impossibility, while structurally parallel, operates in a complementary regime where IIA is absent and Rashomon is present.

B Complete Lean Theorem Statements

This appendix reproduces the key Lean 4 theorem statements from the formalization. The complete codebase comprises 130 files in `UniversalImpossibility/`; we present the three most important files in full and excerpt key theorems from the instance files.

B.1 ExplanationSystem.lean — Core Framework and Impossibility

This file defines the abstract `ExplanationSystem` structure, the three desiderata (faithfulness, stability, decisiveness), and proves the universal impossibility with zero custom axiom dependencies.

```
1 /-!
2 # Abstract Explanation System
3
```

```

4 The universal impossibility framework. An explanation system  $S : \text{Theta} \rightarrow Y$ 
5 is explained by  $E : \text{Theta} \rightarrow H$ . The Rashomon property says observationally
6 equivalent configurations can produce incompatible explanations.
7
8 ## The Trilemma
9
10 No explanation  $E$  can simultaneously be:
11 - **Faithful**:  $E(\text{theta})$  never contradicts the system's native explanation
12 - **Stable**:  $E(\text{theta}_1) = E(\text{theta}_2)$  whenever  $\text{observe}(\text{theta}_1) = \text{observe}(\text{theta}_2)$ 
13 - **Decisive**:  $E(\text{theta})$  commits to every distinction the native explanation makes
14
15 Each pair is achievable:
16 - Faithful + stable: abstain on controversial distinctions (ties)
17 - Faithful + decisive: use the native explanation (unstable under retraining)
18 - Stable + decisive: pick an arbitrary fixed explanation (unfaithful)
19 -/
20
21 set_option autoImplicit false
22
23 /-- An explanation system with configuration space  $\text{Theta}$ , explanation space  $H$ ,
24 and observable space  $Y$ . -/
25 structure ExplanationSystem (Theta : Type) (H : Type) (Y : Type) where
26 /-- The observation map: what the system produces. -/
27 observe : Theta -> Y
28 /-- The explanation map: how we interpret a configuration. -/
29 explain : Theta -> H
30 /-- Incompatibility relation on explanations. -/
31 incompatible : H -> H -> Prop
32 /-- Incompatibility is irreflexive: no explanation contradicts itself. -/
33 incompatible_irrefl : forall (h : H), not (incompatible h h)
34 /-- The Rashomon property: equivalent configurations produce incompatible
35 explanations. -/
36 rashomon : exists theta_1 theta_2 : Theta,
37   observe theta_1 = observe theta_2
38   ^ incompatible (explain theta_1) (explain theta_2)
39
40 variable {Theta : Type} {H : Type} {Y : Type}
41
42 /-- Faithful:  $E$  never contradicts the native explanation.
43  $E(\text{theta})$  may be less specific than  $\text{explain}(\text{theta})$  but never disagrees
44 with it. -/
45 def faithful (S : ExplanationSystem Theta H Y) (E : Theta -> H) : Prop :=
46   forall (theta : Theta), not (S.incompatible (E theta) (S.explain theta))
47
48 /-- Stability: the explanation factors through the observable map.
49 Equivalent configurations get the same explanation. -/
50 def stable (S : ExplanationSystem Theta H Y) (E : Theta -> H) : Prop :=
51   forall (theta_1 theta_2 : Theta),
52     S.observe theta_1 = S.observe theta_2
53     -> E theta_1 = E theta_2
54
55 /-- Decisive:  $E$  inherits every incompatibility of the native explanation.
56 Whatever the true explanation rules out,  $E$  also rules out.
57  $E$  is at least as committal as  $\text{explain}$ . -/
58 def decisive (S : ExplanationSystem Theta H Y) (E : Theta -> H) : Prop :=
59   forall (theta : Theta) (h : H),
60     S.incompatible (S.explain theta) h -> S.incompatible (E theta) h
61

```

```

62 /-- The Explanation Impossibility.
63
64 No explanation of a system with the Rashomon property can be simultaneously
65 faithful, stable, and decisive.
66
67 Proof structure:
68 1. Rashomon gives theta_1, theta_2 with same output but incompatible explanations
69 2. Decisiveness at theta_1: incompatible(explain theta_1, explain theta_2)
70    -> incompatible(E theta_1, explain theta_2)
71 3. Stability: observe theta_1 = observe theta_2 -> E theta_1 = E theta_2,
72    so incompatible(E theta_2, explain theta_2)
73 4. Faithfulness at theta_2: not incompatible(E theta_2, explain theta_2)
74 5. Contradiction between 3 and 4 -/
75 theorem explanation_impossibility (S : ExplanationSystem Theta H Y)
76   (E : Theta -> H) (hf : faithful S E) (hs : stable S E)
77   (hd : decisive S E) :
78   False := by
79   obtain <theta_1, theta_2, hobs, hinc> := S.rashomon
80   have h1 : S.incompatible (E theta_1) (S.explain theta_2) :=
81     hd theta_1 (S.explain theta_2) hinc
82   have h2 : E theta_1 = E theta_2 := hs theta_1 theta_2 hobs
83   rw [h2] at h1
84   exact hf theta_2 h1
85
86 -- =====
87 -- Tightness: each pair of properties is achievable
88 -- =====
89
90 /-- **Tightness 1: Faithful + Decisive (dropping Stable).**
91   E = explain is always faithful and decisive. -/
92 theorem tightness_faithful_decisive (S : ExplanationSystem Theta H Y) :
93   faithful S S.explain ^decisive S S.explain := by
94   constructor
95   . intro theta
96     exact S.incompatible_irrefl (S.explain theta)
97   . intro theta h hinc
98     exact hinc
99
100 /-- **Tightness 2: Faithful + Stable (dropping Decisive).**
101   The constant function E(theta) = c is faithful and stable
102   when c is compatible with every native explanation. -/
103 theorem tightness_faithful_stable (S : ExplanationSystem Theta H Y)
104   (c : H) (hc : forall (theta : Theta),
105     not (S.incompatible c (S.explain theta))) :
106   faithful S (fun _ => c) ^stable S (fun _ => c) := by
107   constructor
108   . intro theta
109     exact hc theta
110   . intro _ _ _
111     rfl
112
113 /-- **Tightness 3: Stable + Decisive (dropping Faithful).**
114   The constant function E(theta) = c is stable and decisive
115   when c inherits every incompatibility. -/
116 theorem tightness_stable_decisive (S : ExplanationSystem Theta H Y)
117   (c : H) (hc : forall (theta : Theta) (h : H),
118     S.incompatible (S.explain theta) h -> S.incompatible c h) :
119   stable S (fun _ => c) ^decisive S (fun _ => c) := by

```

```

120 constructor
121 . intro _ _ _
122   rfl
123 . intro theta h hinc
124   exact hc theta h hinc

```

Listing 1: ExplanationSystem.lean — complete file

B.2 UniversalResolution.lean — G -Invariant Resolution

This file formalizes the resolution framework: if a group G acts transitively on observe-fibers and the resolution is G -invariant, then it is automatically stable.

```

1 import UniversalImpossibility.ExplanationSystem
2 import Mathlib.GroupTheory.GroupAction.Defs
3
4 /-!
5 # G-Invariant Resolution Framework
6 -/
7
8 set_option autoImplicit false
9
10 variable {Theta : Type} {H : Type} {Y : Type} {G : Type}
11
12 /-- A symmetry structure on an explanation system. -/
13 structure HasSymmetry (S : ExplanationSystem Theta H Y)
14   (G : Type) [Group G] [MulAction G Theta] where
15   observe_invariant :
16     forall (g : G) (theta : Theta),
17       S.observe (g * theta) = S.observe theta
18   fiber_transitive :
19     forall (theta_1 theta_2 : Theta),
20       S.observe theta_1 = S.observe theta_2
21       -> exists (g : G), g * theta_1 = theta_2
22
23 /-- A resolution R is G-invariant if constant on G-orbits. -/
24 def gInvariant (R : Theta -> H) (G : Type)
25   [Group G] [MulAction G Theta] : Prop :=
26   forall (g : G) (theta : Theta), R (g * theta) = R theta
27
28 /-- G-invariance + fiber transitivity implies stability. -/
29 theorem gInvariant_stable
30   (S : ExplanationSystem Theta H Y)
31   (R : Theta -> H)
32   [Group G] [MulAction G Theta]
33   (sym : HasSymmetry S G)
34   (hInv : gInvariant R G) :
35   stable S R := by
36   intro theta_1 theta_2 hobs
37   obtain <g, hg> := sym.fiber_transitive theta_1 theta_2 hobs
38   have h := hInv g theta_1
39   rw [hg] at h
40   exact h.symm

```

Listing 2: UniversalResolution.lean — complete file

B.3 Ubiquity.lean — Structural Ubiquity

```

1  import UniversalImpossibility.ExplanationSystem
2
3  /-!
4  # Structural Ubiquity of Explanation Impossibility
5
6  The impossibility is not a pathological corner case -- it is generic whenever
7  the configuration space has higher dimension than the observable space.
8
9  ## Main results
10
11 - 'generic_underspecification': if  $\dim(\text{Theta}) > \dim(Y)$ , the fiber has
12   positive dimension.
13 - 'fiber_nondegeneracy_implies_impossibility': if the explanation map is
14   non-constant on a fiber, then no explanation can be faithful, stable,
15   and decisive.
16 - 'ubiquity_impossibility': constructs an ExplanationSystem from witness
17   data and derives the impossibility via 'explanation_impossibility'.
18 -/
19
20 set_option autoImplicit false
21
22 -- =====
23 -- Dimensional argument
24 -- =====
25
26 /-- Generic underspecification: if the configuration space has higher
27   dimension than the observable space, the system is generically
28   underspecified. -/
29 theorem generic_underspecification {n m : Nat} (h : m < n) :
30   n - m > 0 := by
31   omega
32
33 -- =====
34 -- Non-degeneracy -> impossibility bridge
35 -- =====
36
37 /-- If two configurations have the same observable output but incompatible
38   explanations, and the incompatibility relation is irreflexive, then
39   no explanation can be faithful, stable, and decisive. -/
40 theorem fiber_nondegeneracy_implies_impossibility
41   {Theta H Y : Type}
42   (observe : Theta -> Y) (explain : Theta -> H)
43   (incomp : H -> H -> Prop)
44   (theta_1 theta_2 : Theta)
45   (hobs : observe theta_1 = observe theta_2)
46   (hinc : incomp (explain theta_1) (explain theta_2))
47   (E : Theta -> H)
48   (hf : forall (theta : Theta), not (incomp (E theta) (explain theta)))
49   (hs : forall a b : Theta,
50     observe a = observe b -> E a = E b)
51   (hd : forall (theta : Theta) (h : H),
52     incomp (explain theta) h -> incomp (E theta) h) :
53   False := by
54   have h1 : incomp (E theta_1) (explain theta_2) :=
55     hd theta_1 (explain theta_2) hinc

```

```

56 have h2 : E theta_1 = E theta_2 := hs theta_1 theta_2 hobs
57 rw [h2] at h1
58 exact hf theta_2 h1
59
60 -- =====
61 -- Connection to ExplanationSystem
62 -- =====
63
64 /-- Given concrete witness data, construct an ExplanationSystem and derive
65 the impossibility. -/
66 theorem ubiquity_impossibility
67   {Theta H Y : Type}
68   (observe : Theta -> Y) (explain : Theta -> H)
69   (incomp : H -> H -> Prop)
70   (incomp_irrefl : forall (h : H), not (incomp h h))
71   (theta_1 theta_2 : Theta)
72   (hobs : observe theta_1 = observe theta_2)
73   (hinc : incomp (explain theta_1) (explain theta_2)) :
74   let S : ExplanationSystem Theta H Y :=
75     { observe, explain, incompatible := incomp,
76       incompatible_irrefl := incomp_irrefl,
77       rashomon := <theta_1, theta_2, hobs, hinc> }
78   forall (E : Theta -> H),
79     faithful S E -> stable S E -> decisive S E
80     -> False := by
81   intro S E hf hs hd
82   exact explanation_impossibility S E hf hs hd

```

Listing 3: Ubiquity.lean — complete file

B.4 Instance Theorems

Each of the nine instances follows the same pattern: construct explicit finite witness configurations, bundle them into an `ExplanationSystem` whose Rashomon witness is checked by `decide/rfl` (zero axioms), and apply `explanation_impossibility`. We list the theorem statements as they appear in the repository (`*InstanceConstructive.lean`).

```

1 -- Attribution (AttributionInstanceConstructive.lean)
2 theorem attribution_impossibility_constructive :
3   forall (E : AttrConfigC -> AttrRanking),
4     faithful attrSystemC E -> stable attrSystemC E ->
5     decisive attrSystemC E -> False :=
6   explanation_impossibility attrSystemC
7
8 -- Attention (AttentionInstanceConstructive.lean)
9 theorem attention_impossibility_constructive
10   (E : AttnConfig -> AttnArgmax)
11   (hf : faithful attentionSystemConstructive E)
12   (hs : stable attentionSystemConstructive E)
13   (hd : decisive attentionSystemConstructive E) : False :=
14   explanation_impossibility attentionSystemConstructive E hf hs hd
15
16 -- Counterfactual (CounterfactualInstanceConstructive.lean)
17 theorem counterfactual_impossibility_constructive
18   (E : CFConfig -> CFDirection)
19   (hf : faithful cfSystemConstructive E)
20   (hs : stable cfSystemConstructive E)

```

```

21 (hd : decisive cfSystemConstructive E) : False :=
22   explanation_impossibility cfSystemConstructive E hf hs hd
23
24 -- Concept Probe (ConceptInstanceConstructive.lean)
25 theorem concept_impossibility_constructive
26   (E : ConceptCfg -> ConceptDirection)
27   (hf : faithful conceptSystemConstructive E)
28   (hs : stable conceptSystemConstructive E)
29   (hd : decisive conceptSystemConstructive E) : False :=
30   explanation_impossibility conceptSystemConstructive E hf hs hd
31
32 -- Causal Discovery (CausalInstanceConstructive.lean)
33 theorem causal_impossibility_constructive :
34   forall (E : CausalConfigC -> EdgeOrientation),
35     faithful causalSystemC E -> stable causalSystemC E ->
36     decisive causalSystemC E -> False :=
37   explanation_impossibility causalSystemC
38
39 -- Model Selection (ModelSelectionInstanceConstructive.lean)
40 theorem model_selection_impossibility_constructive
41   (E : MSCfg -> MSModelId)
42   (hf : faithful msSystemConstructive E)
43   (hs : stable msSystemConstructive E)
44   (hd : decisive msSystemConstructive E) : False :=
45   explanation_impossibility msSystemConstructive E hf hs hd
46
47 -- GradCAM (SaliencyInstanceConstructive.lean)
48 theorem saliency_impossibility_constructive
49   (E : SaliencyCfg -> SaliencyRegion)
50   (hf : faithful saliencySystemConstructive E)
51   (hs : stable saliencySystemConstructive E)
52   (hd : decisive saliencySystemConstructive E) : False :=
53   explanation_impossibility saliencySystemConstructive E hf hs hd
54
55 -- LLM Self-Explanation (LLMExplanationInstanceConstructive.lean)
56 theorem llm_explanation_impossibility_constructive
57   (E : LLMCfg -> LLMCitation)
58   (hf : faithful llmSystemConstructive E)
59   (hs : stable llmSystemConstructive E)
60   (hd : decisive llmSystemConstructive E) : False :=
61   explanation_impossibility llmSystemConstructive E hf hs hd
62
63 -- Mechanistic Interpretability (MechInterpInstanceConstructive.lean)
64 theorem mech_interp_impossibility_constructive
65   (E : MechInterpCfg -> CircuitDecomp)
66   (hf : faithful mechInterpSystemConstructive E)
67   (hs : stable mechInterpSystemConstructive E)
68   (hd : decisive mechInterpSystemConstructive E) : False :=
69   explanation_impossibility mechInterpSystemConstructive E hf hs hd

```

Listing 4: Nine constructive instance impossibility theorems

C Detailed Experiment Methodology

This appendix provides full details of the four empirical experiments validating Instances 2–5 of the universal impossibility. All experiments are reproducible via:

`python paper/scripts/run_all_universal_experiments.py`

C.1 Experiment 1: Attention Map Instability (Instance 2)

Research question. Do functionally equivalent transformer models assign peak attention to different tokens?

Model. DistilBERT-base-uncased [Sanh et al., 2019], a 6-layer, 12-head transformer with 66M parameters.

Perturbation protocol. We create 10 model variants by adding Gaussian noise $\mathcal{N}(0, \sigma)$ to all weight matrices of a pretrained DistilBERT checkpoint. Five models use $\sigma = 0.01$ and five use $\sigma = 0.02$. This simulates the Rashomon set of models reachable by different training trajectories that arrive at the same loss basin.

Dataset. 200 synthetic sentiment sentences (100 positive, 100 negative), constructed from 10 positive and 10 negative templates with 10 variations each.

Attention extraction. For each sentence and each model, we extract attention weights from all 6 layers using attention rollout (the product of attention matrices across layers), producing a distribution over tokens.

Metrics.

- **Prediction agreement:** fraction of sentences where all 10 models produce the same classification.
- **Argmax flip rate:** fraction of model pairs (i, j) and sentences where $\arg \max_t \alpha_t^{(i)} \neq \arg \max_t \alpha_t^{(j)}$.
- **Mean Kendall τ :** average rank correlation between attention distributions across model pairs.

Bootstrap. 95% confidence intervals are computed via the percentile method with 2000 bootstrap resamples over sentence-model-pair combinations.

Results. Under full retraining (20 models, last 2 transformer layers retrained from independent seeds), prediction agreement is 94.5%, the argmax flip rate is **19.9%** [95% CI: 19.1–20.6%], and mean Kendall $\tau = 0.63$ [CI: 0.62–0.63]. Under weight perturbation ($\sigma \in \{0.01, 0.02\}$, 10 models), prediction agreement is 100% and the argmax flip rate rises to 60.0% [CI: 59–61%]; however, weight perturbation overstates instability because perturbing weights is not equivalent to retraining from a new seed (perturbation preserves the loss basin, while retraining explores genuinely different trajectories). We report the full-retraining figure (19.9%) as the primary result (Cohen’s $d = 0.50$).

C.2 Experiment 2: Counterfactual Instability (Instance 3)

Research question. Do equivalent-accuracy models produce contradictory counterfactual explanations?

Dataset. German Credit (UCI, OpenML `credit-g`), 1000 instances, 20 features after one-hot encoding of categorical variables. Numerical features are standardized.

Models. 20 XGBoost classifiers with identical hyperparameters (default `xgb.XGBClassifier` settings), differing only in random seed. Train/test split: 80/20 with a fixed random state for the split.

Counterfactual method. Greedy feature-importance-ordered perturbation toward the positive-class centroid. For each query point x_0 classified as “bad credit,” features are perturbed in order of their importance (measured by XGBoost feature importance) toward the mean of the “good credit” class, until the prediction flips.

Query points. 45–50 test-set instances classified as “bad credit” by at least 18 of the 20 models (to ensure the counterfactual direction is well-defined for most models).

Metrics.

- **Direction flip rate:** fraction of model pairs recommending opposite feature changes for the same query. Two counterfactuals “flip” if any feature is perturbed in opposite directions.
- **Cross-model recourse validity:** fraction of counterfactuals generated for model i that also flip the prediction on model j .
- **Distance instability (CV):** coefficient of variation of ℓ_2 counterfactual distances across models.

Bootstrap. 500 bootstrap resamples over query points for 95% percentile CIs.

Results. Mean AUC: 0.797 [CI: 0.783–0.810]. Direction flip rate: 23.5% [CI: 20.6–26.6%]. Cross-model validity: 61.3% [CI: 55.8–66.2%]. Distance CV: 0.328 [CI: 0.281–0.372].

C.3 Experiment 3: Concept Probe Instability (Instance 4)

Research question. Do functionally equivalent neural networks encode the same concept in incompatible directions in representation space?

Dataset. `sklearn.datasets.load_digits`: 1797 handwritten digit images (8×8 pixels, 10 classes).

Models. 15 `MLPClassifier` models with architecture (128, 64), ReLU activation, trained with different random seeds (`random_state = 42, 43, \dots, 56`). All models achieve $> 95\%$ test accuracy on a 70/30 split.

Concept definitions.

- **Curved digits:** $\{0, 2, 3, 5, 6, 8, 9\}$ vs. angular $\{1, 4, 7\}$.
- **Symmetric digits:** $\{0, 1, 8\}$ vs. rest.

Concept direction extraction. For each model, extract penultimate-layer activations (64-dimensional) on the test set. Fit a `LinearSVC` to separate concept-positive from concept-negative examples. The resulting weight vector is the concept activation vector (CAV).

Metrics.

- **Mean $|\cos(\mathbf{v}_i, \mathbf{v}_j)|$:** pairwise cosine similarity of CAVs (off-diagonal entries of the 15×15 matrix).
- **Concept direction instability:** $1 - |\cos|$.
- **TCAV score std:** standard deviation of TCAV sensitivity scores across models.
- **Prediction agreement:** fraction of test points classified identically by all 15 models.

Bootstrap. 2000 bootstrap resamples for 95% CIs.

Results. Mean $|\cos|$: 0.10 [CI: 0.086–0.114]. Direction instability: 0.90. TCAV std (curved): 0.157. TCAV std (symmetric): 0.123. Prediction agreement: 97.2% [CI: 95.7–98.5%].

C.4 Experiment 4: Model Selection Instability (Instance 6)

Research question. Does the identity of the “best” model change when evaluated on different data splits?

Dataset. German Credit (credit-g, OpenML), same as Experiment 2.

Models. 50 XGBoost classifiers, all trained on a single fixed 80/20 split, differing only in random seed (`random_state` = 0, 1, ..., 49).

Evaluation protocol. Each of 20 independent random 80/20 evaluation splits is used to compute test-set AUC for all 50 models. For each split, the “best” model is the one with the highest AUC.

Metrics.

- **Best-model flip rate:** fraction of consecutive evaluation splits where the identity of the best model changes.
- **Unique winners:** number of distinct models that are “best” on at least one evaluation split.
- **AUC spread:** mean range of AUC values across the 50 models within a single evaluation split.
- **Overall mean AUC:** mean $\pm$ std of AUC across all models and splits.

Results. Best-model flip rate: 80% (Cohen’s $d = 2.0$). Consecutive flip rate: 89%. Unique winners: 11 out of 50. Mean AUC spread per split: 0.0197. Overall AUC: 0.948 ± 0.0023 .

D Proof Sketches for All Nine Instances

Each instance follows the same four-step proof structure: (1) establish the Rashomon property, (2) faithfulness forces $E = \text{obs}$, (3) decisiveness forces $E(\theta_1) \neq E(\theta_2)$, (4) stability forces $E(\theta_1) = E(\theta_2)$, yielding a contradiction. Below we spell out the domain-specific content for all nine instances.

D.1 Instance 1: Feature Attribution

Setup. Θ = data-generating distributions with collinearity structure. $\mathcal{H}$ = trained models (GBDT, Lasso, neural net). $Y = (\text{Fin } P \rightarrow \mathbb{R})$: feature attribution vectors. $\text{obs}(f) = \varphi(f)$: the model’s own SHAP values. exp : the attribution procedure under evaluation. $\perp$: strict reversal of a collinear pair ($\varphi_j > \varphi_k$ vs. $\varphi_k > \varphi_j$).

Rashomon property. Collinear features within a group ℓ are symmetric under the DGP. Surjectivity of the first-mover map guarantees that for any collinear pair j, k , there exist models f, f' with the same DGP but opposite rankings. This is the `RashomonProperty` of `Trilemma.lean`.

Proof.

1. Rashomon: $\exists f, f'$ with $\theta(f) = \theta(f')$ and $\varphi_j(f) > \varphi_k(f), \varphi_k(f') > \varphi_j(f')$.
2. Faithful E reports $\varphi_j(f) > \varphi_k(f)$ for f and $\varphi_k(f') > \varphi_j(f')$ for f' .
3. Decisive: $E(f) \neq E(f')$ (rankings differ).
4. Stable: $E(f) = E(f')$ (same DGP).
5. Contradiction.

Lean. `attribution_impossibility` in `Trilemma.lean` (0 axioms); `attribution_impossibility_constructive` in `AttributionInstanceConstructive.lean` (constructive witness; 0 axioms).

D.2 Instance 2: Attention Maps

Setup. Θ = neural network weight configurations. Y = input→output functions (the behavior the network computes). $\mathcal{H} = \Delta^T$: attention distributions over T tokens. $\text{obs}(\theta) = f_\theta$: the function the network computes. $\text{exp}(\theta) = \alpha_\theta$: the attention distribution. $\perp$: $\arg \max$ differs.

Rashomon property. D’Amour et al. [2022]: $\exists \theta_1, \theta_2$ with $f_{\theta_1} \approx f_{\theta_2}$ yet $\arg \max_t (\alpha_{\theta_1})_t \neq \arg \max_t (\alpha_{\theta_2})_t$.

Proof.

1. Rashomon: equivalent models with different max-attended tokens.
2. Faithful E reports α_{θ_1} for θ_1 and α_{θ_2} for θ_2 .
3. Decisive: $E(\theta_1) \neq E(\theta_2)$.
4. Stable: $E(\theta_1) = E(\theta_2)$.
5. Contradiction.

Lean. `attention_impossibility_constructive` in `AttentionInstanceConstructive.lean`.

D.3 Instance 3: Counterfactual Explanations

Setup. Θ = model parameters (decision boundaries). Y = predictions on a test set. $\mathcal{H} = \mathcal{X}$: nearest counterfactual examples. $\text{obs}(\theta) = f_\theta|_{\text{test}}$. $\text{exp}(\theta) = \arg \min_{x'} \|x' - x_0\|$ s.t. $f_\theta(x') \neq f_\theta(x_0)$. $\perp$: counterfactuals point in different directions.

Rashomon property. Models with identical test predictions place their decision boundaries in different locations near the query point x_0 , yielding counterfactuals in opposite directions [Semenova et al., 2022, Fisher et al., 2019].

Proof. Same four-step structure. Lean: `counterfactual_impossibility_constructive` in `CounterfactualInstanceConstructive.lean`.

D.4 Instance 4: Concept Probes

Setup. Θ = model weights (defines internal representations). Y = predictions (input-output function). $\mathcal{H} = \mathbb{R}^c$: concept activation vectors. $\text{obs}(\theta) = f_\theta$. $\text{exp}(\theta) = v_\theta$: concept direction from linear probe. $\perp$: directions differ (not related by rotation in concept subspace).

Rashomon property. Equivalent models encode concepts in different directions [Bolukbasi et al., 2016, Kim et al., 2018, Ravfogel et al., 2020]. The overparameterized zero-loss manifold has dimension $\geq d - n > 0$, so infinitely many weight configurations yield identical predictions but geometrically distinct representations.

Proof. Same four-step structure. Lean: `concept_impossibility_constructive` in `ConceptInstanceConstructive.lean`.

D.5 Instance 5: Causal Discovery

Setup. Θ = DAGs over a fixed vertex set V . Y = conditional independence structures (Markov equivalence classes). $\mathcal{H}$ = DAGs (causal explanations). $\text{obs}(h)$ = the CI structure encoded by h (d -separation). $\text{exp}(h)$ = the fully oriented DAG returned by a discovery algorithm. $\perp$: opposite edge orientations ($X \rightarrow Y$ vs. $Y \rightarrow X$).

Rashomon property. Markov equivalence [Verma and Pearl, 1991]: any CPDAG with at least one undirected edge contains two DAGs in the same equivalence class with that edge oriented in opposite directions.

Proof. Same four-step structure.

Lean: `causal_impossibility_constructive` in `CausalInstanceConstructive.lean`;
`causal_discovery_impossibility` in `CausalDiscovery.lean`.

D.6 Instance 6: Model Selection

Setup. Θ = model configurations. $\mathcal{H}$ = trained models. Y = performance metrics (with strict total order). $\text{obs}(h)$ = observed performance on evaluation set. $\text{exp}(h)$ = the selection rule’s recommendation. $\perp$: opposite selection decisions.

Rashomon property. The ε -Rashomon set contains models with indistinguishable performance that are ranked differently on different evaluation splits [Fisher et al., 2019, Rudin et al., 2024, Semenova et al., 2022].

Proof. Same four-step structure.

Lean: `model_selection_impossibility_constructive` in
`ModelSelectionInstanceConstructive.lean`; `model_selection_impossibility` in
`ModelSelection.lean`.

E Companion Results: Number Theory and Physics

This section documents results from the companion repository `ostrowski-impossibility` (481 theorems; 0 axiom declarations, with ~ 11 domain hypotheses—including the Selmer 1951 result and the contested/speculative black-hole and spacetime-emergence properties—carried as section variables that gate their downstream results; 38 files; 0 `sorry`). These establish the impossibility framework’s connections to number theory, mathematical physics, and computability theory.

E.1 Ostrowski’s Classification and the Mathlib Bridge

Ostrowski’s theorem classifies all nontrivial absolute values on $\mathbb{Q}$: each is equivalent to the real absolute value or a p -adic absolute value. The companion repository formalises this via Mathlib’s `Rat.AbsoluteValue.equiv_real_or_padic`, converting the completion family (archimedean vs p -adic) from axioms to concrete definitions. This provides the zero-axiom foundation for the physics biletta instantiation. The bridge `ostrowski_classify` (a `noncomputable def` in `OstrowskiFramework.lean`, built on Mathlib’s `Rat.AbsoluteValue.equiv_real_or_padic`) connects the abstract impossibility to the concrete number-theoretic structure.

E.2 Freund–Witten Three-Fold Zeta Symmetry

The Freund–Witten identity $\Lambda(a)\Lambda(b)\Lambda(1-a-b) = \Lambda(1-a)\Lambda(1-b)\Lambda(a+b)$ is proved using three applications of Mathlib’s `completedRiemannZeta_one_sub` (`FreundWitten.lean`). This witnesses the Rashomon property for the physics biletta: archimedean and non-archimedean contributions are related by the functional equation, creating the maximally incompatible explanation system.

E.3 Adelic Resolution

The “adelic resolution” here is a placeholder model, not a formalization of the adèle ring: the underlying type is a finite enumeration of completions with an identity embedding, and the resolution witnessing `adelic_resolves` is literally the constant predicate `fun _ => True` (`AdelicResolution.lean`). It records the *intended* parallel—commit to all completions simultaneously, as DASH does for attributions and CPDAG for causal discovery—but, unlike those two instances, nothing nontrivial is proved: the real adèle ring and its symmetry action are not formalised. We flag this as the weakest companion result [Tier C].

E.4 Circuit Complexity: AC^0 Parity

The Furst–Saxe–Sipser/Håstad AC^0 parity lower bound is formalised as an `AbstractImpossibility` with properties `{polynomial-size, constant-depth, computes-parity}` (`CircuitComplexity.lean`). Exhaustive verification confirms: (a) no depth-1 AND/OR gate computes XOR_3 (54 gates, kernel `decide`); (b) no depth-2 circuit with 2 gates computes XOR_3 (52,488 circuits, the sole `native_decide`). Full tightness is proved with three witnesses (AND_3 , XOR tree, exponential DNF).

A novel *conditional barrier analysis* shows that WITH one-way functions, proof barriers have full tightness; WITHOUT, the triple is achievable. This is one of only two instances *in this classification* where impossibility existence depends on an unresolved conjecture (the other being Navier–Stokes regularity); we do not claim it is exhaustive over all impossibility results.

E.5 DPRM Trilemma and the Selmer Curve

Hilbert’s 10th problem is modelled as a three-property impossibility `{correctness, totality, universality}` with full tightness (`DiophantineImpossibility.lean`). The degree hierarchy forms an enrichment stack: Hasse–Minkowski resolves degree ≤ 2 but not degree 3. The Selmer curve $3x^3 + 4y^3 + 5z^3 = 0$ is locally solvable mod 2, 3, 5, 7 (verified by kernel `decide`) but globally insoluble (assumed as the explicit hypothesis `selmer_no_nontrivial_solution`—Selmer’s 1951 theorem—rather than proved; it is a hypothesis parameter, not an `axiom` declaration). This parallels the local-to-global obstruction of the Tate–Shafarevich group.

E.6 Navier–Stokes Conditional Tightness

The Navier–Stokes trilemma is modelled as a three-strategy tightness *schema*—a finite case analysis over the three classical compromise strategies (Leray weak, local classical, viscosity), with no PDEs formalised: without regularity the schema has full tightness; with regularity all three properties are achievable (`NavierStokesImpossibility.lean`). The Reynolds dichotomy (`ns_reynolds_dichotomy`) encodes the same case split parametrically—below a critical Reynolds number all three properties {smooth, energy-conserving, global} are jointly achievable; above, the schema becomes an `AbstractImpossibility`—as bookkeeping over assumed regime classifications, not as fluid dynamics. The 54 numerical experiments (Section 19.2) probe this classification experimentally.

E.7 The N -Property Generalisation and Myerson–Satterthwaite

The `AbstractImpossibilityN` structure generalises beyond 3 properties to arbitrary N (`GeneralTheory.lean`). For the Myerson–Satterthwaite impossibility with $N = 4$ properties (individual rationality, incentive compatibility, efficiency, budget balance), all $\binom{4}{2} = 6$ pairwise combinations are achievable while the quadruple conjunction fails—full tightness, verified by kernel `decide` on all $\text{Fin } 4 \times \text{Fin } 4$ pairs (`ms_full_tightness_n4`).

E.8 The Enrichment–Riemann Hypothesis Program

A five-phase investigation connects enrichment stack structure to analytic properties of the Riemann zeta function. The results are primarily negative, identifying precise obstacles:

DFT eigenvalue constraint (local RH). The discrete Fourier transform on $\mathbb{Z}/p\mathbb{Z}$ satisfies $F^4 = p^2 \cdot \text{id}$, forcing eigenvalues to have absolute value $\sqrt{p}$ —an eigenvalue-modulus fact sometimes called the “local Riemann Hypothesis” in spectral form, unrelated to the analytic Riemann Hypothesis (`EnrichmentFunctor.lean`).

Prime enrichment stack. The enrichment tower indexed by primes with $H_p = \text{Fin } p$ has unfaithfulness $1 - p^{-1}$ at level p , matching the Euler obstruction term. Cumulative unfaithfulness diverges by Euler’s theorem (1737) (`PrimeEnrichmentStack.lean`).

Enrichment Sha (cohomological obstruction). The “enrichment Sha” is the derived limit ($\varprojlim^1$) obstruction to lifting local bilemma resolutions globally (`EnrichmentSha.lean`). The bilemma is locally resolvable at each prime but globally unresolvable. *Negative result:* Sha = CRT = Euler multiplicativity—the obstruction has no arithmetic content beyond the Euler product.

Forced resolution (negative). Collapsed tightness forces the unique resolution to be the identity at each prime, selecting the trivial Dirichlet character. This is tautological: ANY explain function gives collapsed tightness (`EnrichmentForcedResolution.lean`).

Quantitative bounds (negative). Approximate bilemma bounds on $\text{Fin } p$ depend only on p (the group size), with no prime-specific arithmetic content. Summing over primes only reproduces classical bounds (Mertens for $\sum 1/p$) (`EnrichmentQuantitative.lean`).

These three negative results (Sha = CRT, forced resolution = identity, bounds depend only on p) identify the precise gap between the enrichment framework and the Riemann Hypothesis: the local structure is too weak to encode global cancellation.

E.9 Langlands Extensions

$\text{GL}(2, \mathbb{F}_2) = S_3$. The smallest non-abelian group demonstrates the Langlands boundary concretely (`GL2Impossibility.lean`). For 1-dimensional representations, values are determined by conjugacy class (no Rashomon, full F+S+D). For 2-dimensional representations, conjugate elements have similar but unequal matrices (Rashomon, bilemma).

$\text{SL}(n, \mathbb{F}_p)$. The bilemma extends from $\text{GL}(n)$ to $\text{SL}(n)$ because the unipotent Rashomon pair has determinant 1 (`ClassicalGroups.lean`). The boundary theorem generalises: $n = 1$ full, $n \geq 2$ collapsed.

Trace compatibility / functoriality. The trace (stable resolution of the $\text{GL}(n)$ bilemma) is compatible with natural maps between matrix algebras (`LanglandsFunctoriality.lean`). These are elementary trace-compatibility facts; we do *not* claim they constitute Langlands functoriality, which concerns automorphic representations that are absent from this formalization.

Adelic Tate theory. The chain from local Fourier analysis at each prime to global Poisson summation, Jacobi theta modular transformation, and the functional equation $\xi(s) = \xi(1-s)$ is formalised (`AdelicTate.lean`).

Automorphic gap (`HomExplanationSystem`). When the explain function is required to be a group homomorphism, the enrichment Sha becomes strictly stronger than CRT (`HomExplanationSystem.lean`). For S_3 : CRT allows 6 character combinations but group relations force $\chi(r) = 1$, leaving only 2. The gap—4 out of 6 CRT-allowed combinations are blocked by the group relations—is a structural analogy to the local-global tension in the Langlands programme, not a formalization of the Langlands correspondence.

Galois connection. The enrichment tower is formalised as a field extension tower (`GaloisConnection.lean`). Coverage conflict = irreducible polynomial; neutral element = root; enrichment = field extension. Level independence follows from CRT and connects to Euler totient multiplicativity.

F Companion Results: Attribution-Specific Formalization

This section documents results from the companion repository `dash-impossibility-lean` (368 theorems, 2 axioms, 59 files, 0 sorry). These provide the attribution-specific depth underlying the universal framework.

F.1 Constructive Bilemma Instances

Three binary ML explanation types are formalised with dedicated inductive types and bilemma proofs (`Bilemma.lean`):

- **SHAP sign** (`SHAPSign` type, `shap_sign_bilemma`): No method for determining a feature’s SHAP sign (positive/negative) can be simultaneously faithful and stable.
- **Feature selection** (`FeatureStatus` type, `feature_selection_bilemma`): No method for determining whether a feature is selected can be simultaneously faithful and stable.
- **Counterfactual direction** (`CounterfactualDir` type, `counterfactual_bilemma`): No method for determining counterfactual direction can be simultaneously faithful and stable.

All three are fully constructive (zero axioms, zero sorry).

F.2 BeyondBinary: Enrichment Restores Achievability

The `BeyondBinary.lean` file closes the enrichment loop constructively:

- `concreteMLSystem` / `concreteML_bilemma`: A minimal Bool-based witness that the bilemma is non-vacuous.
- `GradedAttribution` type (positive / weakPositive / negative): Enriching binary to 3-element space restores F+S (`graded_fs_achievable`) because `weakPositive` serves as a neutral element.
- `tightness_dichotomy`: F+S achievability depends entirely on whether $\mathcal{H}$ has a neutral element.
- `concreteML_coverageConflict` / `graded_no_coverageConflict`: Coverage conflict is present in binary (bilemma applies) and absent in graded (bilemma does not apply).

F.3 Mechanistic Interpretability Bilemma

For circuit decompositions, the bilemma strengthens the trilemma (`MechInterp.lean`): F+S alone is impossible because the circuit space is maximally incompatible (`mech_interp_bilemma`). Any stable circuit explanation is unfaithful at ≥ 1 of every 2 functionally equivalent networks (`mech_interp_unfaithfulness`). The only resolution is equivalence classes—there is no “DASH-like” neutral element for binary circuit decompositions.

F.4 Additional Empirical Results

Ranking lottery. 50 random seeds produce 24 distinct top-3 SHAP rankings on Breast Cancer Wisconsin (4.2% agreement). At 100 seeds: 35 distinct. Top-5: 46 distinct. Cross-implementation: XGBoost 24, LightGBM 29, Random Forest 40 distinct rankings. Subsample sensitivity: 17 distinct at subsample= 0.95, 1 at subsample= 1.0.

NN attribution instability. 87% of feature pairs are unstable across independently trained neural networks. Model-vs-SHAP noise ratio: 8:1 (the model contributes $8\times$ more variance than the SHAP estimation procedure).

SNR calibration. $\Phi(-\text{SNR})$ achieves $R^2 = 0.94$ across 1,325 feature pairs from multiple datasets.

Minority fraction diagnostic. A 7-line nonparametric diagnostic achieves Spearman $\rho = 0.92\text{--}0.98$ with empirical flip rates, outperforming the Gaussian $\Phi(-\text{SNR})$ formula by $2\times$ on non-Gaussian data ($\rho = 0.955$ vs 0.463 on California Housing).

Bimodality dip test. Hartigan’s dip test confirms bimodality of the flip-rate distribution: $p < 0.002$ for $\rho \geq 0.5$; control $p = 0.373$ at $\rho = 0$.

G Companion Results: DASH Method and Empirical Validation

This section documents results from the companion repository `dash-shap` (Python package, PyPI: `dash-shap`). These provide the engineering depth underlying the DASH resolution.

G.1 DASH Benchmark Results

All benchmarks use 50 replicates, $M = 200$ models, $K = 30$ diverse models (PAPER_CONFIG).

Linear correlation sweep ($\rho = 0.9$). DASH stability: 0.977 vs Single Best: 0.958 (+0.019, $p < 0.001$); vs LSM: 0.938 (+0.039, $p < 0.001$). DASH top- k_5 : 0.863 vs SB: 0.546 (+0.317, $p < 0.001$). DASH equity (within-group CV): 0.175 vs SB: 0.232 ($p < 0.001$).

Real-world datasets. Breast Cancer: DASH = 0.925 vs SB = 0.376 (+0.549, largest improvement). Superconductor: DASH = 0.964 vs SB = 0.840 (+0.124). California Housing: DASH = 0.978 vs SB = 0.969 (+0.009, $p = 0.063$, not significant—low collinearity).

G.2 Crossed ANOVA: Variance Decomposition

A 7×7 factorial design (7 data seeds $\times$ 7 model seeds) decomposes attribution variance. Single Best: 37.6% data / 40.6% model / 21.8% residual. DASH: 73.6% data / 16.2% model / 10.2% residual. DASH reduces model-selection variance by 60%, shifting the dominant variance component from model choice to data.

G.3 Colsample Ablation

Low `colsample_bytree` (0.1–0.5) is the operative diversity mechanism. At $\rho = 0.9$: low `colsample` = 0.976 vs high `colsample` = 0.953. The effect disappears at $\rho = 0$ ($p = 0.54$), confirming the mechanism is specific to collinearity. Under nonlinearity: low = 0.885 vs high = 0.801 ($\Delta = +0.084$).

G.4 Stochastic Retrain Equivalence

DASH and Stochastic Retrain are statistically equivalent on stability (TOST equivalence confirmed: $\Delta = -0.0009$, $p = 0.40$). This proves the operative mechanism is model independence, not pipeline engineering.

G.5 Var[SHAP] = DASH MSE Identity

The variance of SHAP values across the Rashomon set equals the minimum MSE any stable method can achieve. Validated with 0 violations in 12 tests (identity holds exactly to machine precision).

G.6 Feature Stability Index (FSI)

Cross-model disagreement metric: $\text{FSI}_j = \text{CV}(\phi_j)$ across models. Thresholds: $\text{FSI} < 0.3$ stable; $\text{FSI} > 0.7$ likely collinear; $\text{FSI} > 1.0$ disagreement exceeds signal. The Importance–Stability (IS) plot provides a 4-quadrant classification of features by importance and reliability.

G.7 The F5→F1→DASH Pipeline

A progressive diagnostic pipeline: (1) `screen(model, X)` — single-model screen ($< 1s$) identifying correlated groups and flagged pairs; (2) `validate(models, X)` — multi-model validation computing Z -statistics and flip rates per pair; (3) `consensus(models, X)` — DASH resolution (element-wise mean, tie detection); (4) `report()` — human-readable instability disclosure. The 3-line `check()` API wraps all stages into a single call.

G.8 Retracted Results

Seven results from earlier versions were retracted with documented explanations: (1) entropy bimodality (100% permutation artefact); (2) pairwise “audit pairs” (marginal rates suffice); (3) $\eta = 1/g$ from correlation thresholds; (4) data-only instability prediction (max Spearman 0.26); (5) phase transition in stability curve (gradual, not sharp); (6) “model-class dependence” (confounded comparison); (7) $p/2$ average unfaithfulness bound (corrected to $p \cdot \text{mean_minority_fraction}$). These are documented in the companion repository as a transparency measure.

G.9 Extension Modules

Twelve extension modules provide practitioner tools beyond the core pipeline: BCa bootstrap confidence intervals (`confidence_intervals`), pairwise ordering confidence (`partial_order`), SHAP-substitutability clustering (`feature_groups`), composite importance+stability ranking (`stable_feature_selection`), per-observation uncertainty (`local_uncertainty`), worst-case top- k guarantee (`robust_certification`), SNR/flip-rate/M-recommendation bridge (`theory_bridge`), robust/collinear/fragile/unimportant labels (`causal_flags`), structured audit report (`audit_report`), cosine-distance drift monitoring (`DriftMonitor`), RMSE-stability Pareto selection (`ParetoSelector`), and cross-site federated consensus (`federated_consensus`).

Appendix: Lean 4 Cross-Reference and Theorem Inventory

G.10 Cross-Reference: Monograph $\leftrightarrow$ Lean

Key theorems with their Lean identifiers and source files. All files are in the `UniversalImpossibility/` directory.

Monograph result	Lean name	File
Core impossibility (Thm 3.6)	<code>explanation_impossibility</code>	<code>ExplanationSystem.lean</code>
Bilemma (§3.6)	<code>bilemma</code>	<code>MaximalIncompatibility.lean</code>
Tightness: $\text{neutral} \Rightarrow \text{F+S}$	<code>neutral_implies_faithful_stable</code>	<code>BilemmaCharacterization.lean</code>
Tightness: $\text{bilemma} \Rightarrow \text{no neutral}$	<code>bilemma_implies_no_neutral</code>	<code>BilemmaCharacterization.lean</code>
Enrichment recovers F+S	<code>recovery_faithful_stable</code>	<code>MaximalIncompatibility.lean</code>
G -invariant $\Rightarrow$ stable	<code>gInvariant_stable</code>	<code>UniversalResolution.lean</code>
Pythagorean decomposition	<code>uncertainty_from_symmetry</code>	<code>UncertaintyFromSymmetry.lean</code>
Best approximation	<code>best_approximation</code>	<code>UncertaintyFromSymmetry.lean</code>
Over-explanation penalty	<code>beyond_capacity_penalty</code>	<code>UncertaintyFromSymmetry.lean</code>
DASH within-group dominance	<code>dash_within_group_dominance</code>	<code>ParetoOptimality.lean</code>
Pareto frontier dichotomy	<code>pareto_frontier_dichotomy</code>	<code>ParetoOptimality.lean</code>
Design space dichotomy	<code>universal_design_space_dichotomy</code>	<code>UniversalDesignSpace.lean</code>
MI sufficiency (conditional)	<code>mi_sufficient_for_impossibility</code>	<code>MutualInformation.lean</code>
MI quantitative unfaithfulness	<code>mi_quantitative_unfaithfulness</code>	<code>MIQuantitativeBridge.lean</code>
GBDT impossibility	<code>attribution_impossibility</code>	<code>Trilemma.lean</code>
Ratio divergence	<code>ratio_tendsto_atTop</code>	<code>Ratio.lean</code>
DASH equity	<code>consensus_equity</code>	<code>Corollary.lean</code>
Necessity biconditional (ineq. case)	<code>rashomon_biconditional_neq</code>	<code>NecessityBiconditional.lean</code>
Generic underspecification	<code>generic_underspecification</code>	<code>Ubiquity.lean</code>
Ubiquity impossibility	<code>ubiquity_impossibility</code>	<code>Ubiquity.lean</code>
Recursive enrichment	<code>recursive_impossibility_persists</code>	<code>RecursiveImpossibility.lean</code>
Levels independent	<code>infiniteStack_independence</code>	<code>RecursiveImpossibility.lean</code>

G.11 Complete Theorem Index

All 715 theorem and lemma names from the Lean 4 formalization, grouped by source file. This index enables verification of completeness and traceability from any monograph claim to its machine-checked proof in the public repository.

AlphaFaithful.lean (3)

`stable_ranking_half_unfaithful`
`alpha_faithful_bound`
`faithful_exactly_one_witness`

ApproximateEquity.lean (3)

`global_implies_bounded_one`
`bounded_proportionality_constant_implies_global`

rashomon_from_bounded_proportionality

ApproximateRashomon.lean (4)

approxStable_implies_stable
approx_impossibility
approx_trilemma
exact_implies_approx_rashomon

ArrowInstance.lean (3)

profiles_same_pairwise
profiles_different_social
social_choice_rashomon

AttentionInstanceConstructive.lean (3)

attn_same_output
attn_different_argmax
attention_impossibility_constructive

AttributionInstanceConstructive.lean (2)

attrIncompC_irrefl
attribution_impossibility_constructive

AxiomSubstitution.lean (14)

impossibility_complete
impossibility_complete_three
tightness_faithful_complete
tightness_faithful_stable_complete
no_tightness_stable_complete
resolving_iff_complete
impossibility_resolving
impossibility_resolving_three
tightness_faithful_resolving
tightness_faithful_stable_resolving
no_tightness_stable_resolving
surjective_consistent
decisive_implies_complete
axiom_substitution_summary

BayesOptimalTie.lean (4)

committed_ranking_error_half
tie_disagreement_zero
tie_dominates_commitment
design_space_step3_substantive

BilemmaCharacterization.lean (3)

neutral_implies_faithful_stable
bilemma_implies_no_neutral
maxIncompat_implies_no_neutral

BinaryQuantizer.lean (7)

binary_quantizer_variance
binary_quantizer_fraction
two_over_pi_in_unit_interval
two_over_pi_le_two_thirds
two_over_pi_gt_half
quantizer_signal_loss
quantizer_better_than_half

CPDAGResolution.lean (2)

cpdagResolution_gInvariant
cpdag_gInvariant_implies_stable

CausalDiscovery.lean (3)

causal_discovery_impossibility
cpdag_is_stable
causal_discovery_from_sbd

CausalExplanationSystem.lean (1)

causal_explanation_impossibility

CausalInstanceConstructive.lean (2)

causalIncompC_irrefl
causal_impossibility_constructive

ConceptInstanceConstructive.lean (3)

concept_same_output
concept_different_directions
concept_impossibility_constructive

ConditionalImpossibility.lean (3)

conditional_rashomon_from_symmetry
conditional_attribution_impossibility
conditional_escape

Consistency.lean (7)

consistent_firstMover_surjective
consistent_splitCount_firstMover
consistent_splitCount_nonFirstMover
consistent_proportionality_global
consistent_splitCount_crossGroup_symmetric
consistent_splitCount_crossGroup_stable
axiom_system_consistent

Corollary.lean (5)

consensus_equity
consensus_difference_zero
consensus_variance_rate
consensus_variance_halves
consensus_variance_nonneg

CounterfactualInstanceConstructive.lean (3)

cf_same_prediction
cf_different_directions
counterfactual_impossibility_constructive

DASHResolution.lean (2)

dashResolution_gInvariant
dash_gInvariant_implies_stable

Defs.lean (11)

numTrees_pos
firstMover_surjective
splitCount_firstMover
splitCount_nonFirstMover
proportionality_global
attribution_proportional
attribution_variance_nonneg
consensus_variance_bound
splitCount_crossGroup_symmetric
splitCount_crossGroup_stable
splitCount_eq_of_not_firstMover_j_or_k

DesignSpace.lean (6)

strongly_faithful_impossible
balanced_flip_symmetry
infeasible_faithful_stable_complete
dash_produces_ties
dash_variance_decreases
design_space_theorem

DesignSpaceFull.lean (3)

design_space_exhaustiveness
no_complete_faithful_ranking
family_a_or_family_b

DuhemQuine.lean (6)

dq_same_outcome
dq_blame_different
duhem_quine_impossibility
dq4_same_outcome
dq4_blame_different
duhem_quine_impossibility_4

Efficiency.lean (10)

sum_constraint_swap
swapAt_sum_eq
amplification_factor_m2
amplification_factor_m3
amplification_approaches_one
amplification_gt_one
amplification_bounds
efficiency_sum_varies_across_models
across_model_no_constraint
efficiency_instability_summary

EnrichmentStack.lean (18)

enriched_neutral_faithful_stable
enriched_neutral_not_decisive
cumulative_indecisiveness
incompleteness_depth
enrichment_preserves_deeper_rashomon
neutral_at_n_not_neutral_at_n1
threeLevel_rashomon_0
threeLevel_rashomon_1
threeLevel_rashomon_2
threeLevel_independence
twoLevel_has_multiLevel
InformationFate
informationFate_maxIncompat
bh_bilemma
physicalStack_depth_3
SpacetimeStatus
spacetimeStatus_maxIncompat
qg_bilemma

EnsembleBound.lean (9)

```

consensus_var_eq_over_M
variance_decreases_with_M
double_ensemble_halves_variance
single_model_variance
variance_at_25
ensemble_bound_formula
z_squared_for_5pct
dash_achieves_minimum_variance
sum_squares_ge_inv_M

```

ExplanationLandscape.lean (8)

```

landscape_bottom_no_fs
landscape_top_fs_achievable
landscape_top_achievable
coverage_conflict_iff_no_neutral
enrichment_expands_landscape
coverage_conflict_monotone_restriction
maxIncompat_landscape_chain
local_coverage_conflict_at_pair

```

ExplanationSystem.lean (4)

```

explanation_impossibility
tightness_faithful_decisive
tightness_faithful_stable
tightness_stable_decisive

```

FIMImpossibility.lean (6)

```

excess_loss_nonneg
gaussian_rashomon_witnesses
hessian_eigenvalue_minus
semi_axis_diverges
gaussian_fim_impossibility
gaussian_fim_impossibility_weak

```

FairnessAudit.lean (4)

```

fairness_audit_impossibility
fairness_audit_coin_flip
fairness_audit_coin_flip_reverse
dash_audit_resolution

```

FlipRate.lean (13)

```

non_firstmover_tie
non_firstmover_attribution_tie
firstmover_dominates
balanced_firstmover_count_eq
model_trichotomy

```

firstMoverSets_disjoint
 tieSet_disjoint_left
 tieSet_disjoint_right
 fliprate_partition
 binary_group_firstmover_is_j_or_k
 binary_group_no_ties
 binary_group_flip_rate
 binary_flip_pair_count

GaugeTheory.lean (5)

same_holonomy
 different_configs
 gauge_preserves_holonomy
 gauge_related
 gauge_impossibility

GaussianFlipRate.lean (13)

Phi_mono
 Phi_nonneg
 Phi_le_one
 Phi_tendsto_atBot
 Phi_tendsto_atTop
 gaussianReal_standard_symm
 Phi_neg
 Phi_zero
 gaussianFlipRate_le_half
 gaussianFlipRate_nonneg
 gaussianFlipRate_zero
 gaussianFlipRate_decreasing
 coin_flip_implies_impossibility

General.lean (6)

splitCount_firstMover_gt
 attribution_firstMover_gt
 attribution_reversal
 no_stable_ranking
 gbdt_rashomon
 gbdt_impossibility

GeneralizedBilemma.lean (13)

generalized_bilemma
 maxIncompat_implies_no_common_compatible
 no_bilemma_of_neutral_on_fiber
 fs_achievable_of_fiber_compatible
 incompEquiv_refl
 incompEquiv_symm
 fd_from_incomp_equiv

maxIncompat_incompEquiv_singleton
coverageConflict_implies_no_neutral
maxIncompat_implies_coverageConflict
neutral_destroys_coverageConflict
explanation_collapsed_of_maxIncompat
explanation_full_tightness_of_neutral_committal

GeneticCode.lean (3)

codon_degeneracy
codons_different
genetic_code_impossibility

GoedelIncompleteness.lean (5)

goedel_first_incompleteness
goedel_sentence_true
goedel_recursion
goedel_infinite_incompleteness
proved_shared_structure

Impossibility.lean (4)

attribution_ratio_ge
not_equitable
not_stable
impossibility

IntersectionalFairness.lean (13)

single_audit_wrong
each_audit_has_counterexample
intersectional_audit_impossibility
intersectional_all_audits_wrong
compounding_ge_half
compounding_strict_mono
compounding_nonneg
compounding_le_one
compounding_K2
compounding_K3
compounding_K5
compounding_K10
dash_intersectional_resolution

Iterative.lean (2)

iterative_rashomon
iterative_impossibility

LLMExplanationInstanceConstructive.lean (3)

llm_same_output
llm_different_citations
llm_explanation_impossibility_constructive

LanglandsCorrespondence.lean (6)

trace_mul_comm_applied
dim1_trace_determines
dim2_rashomon_witness
langlands_boundary_statement
scalar_embedding_trace
trace_identity_n

Lasso.lean (2)

lasso_impossibility
lasso_ratio_infinite

LinearSystem.lean (3)

same_output
different_solutions
linear_system_impossibility

LocalGlobal.lean (2)

local_attribution_impossibility
impossibility_is_attribution_agnostic

LocalSufficiency.lean (6)

attribution_ratio_local
ratio_diverges_local
attribution_sum_from_local
equity_violation_local
local_proportionality_suffices
axiom_stratification_documented

LossPreservation.lean (3)

rashomon_from_swap_with_loss
rashomon_from_swap_loss
impossibility_with_quality

MIQuantitativeBridge.lean (5)

mi_implies_positive_gap
gap_implies_unfaithfulness
total_unfaithfulness_bound

pointwise_unfaithfulness_bound
mi_quantitative_unfaithfulness

MarkovEquivalence.lean (4)

chain_fork_same_ci
chain_fork_different_edges
causal_rashomon_derived
causal_impossibility_derived

MaximalIncompatibility.lean (7)

faithful_eq_explain_of_maxIncompat
decisive_eq_explain_of_maxIncompat
bilemma
no_stable_decisive_of_maxIncompat
maxIncompat_tightness_fd
explain_not_stable
recovery_faithful_stable

MechInterpInstanceConstructive.lean (3)

mechInterp_same_output
mechInterp_different_circuits
mech_interp_impossibility_constructive

ModelSelection.lean (2)

model_selection_impossibility
model_selection_impossibility_weak

ModelSelectionDesignSpace.lean (3)

model_stable_complete_unfaithful
model_unfaithfulness_free_implies_tie
model_family_a_or_family_b

ModelSelectionInstanceConstructive.lean (3)

ms_same_performance
ms_different_models
model_selection_impossibility_constructive

MultiAnalystInstance.lean (5)

scientificConclusion_irrefl
multiAnalyst_rashomon
multiAnalyst_impossibility
multiAnalyst_maxIncompat
multiAnalyst_bilemma

MutualInformation.lean (11)

- mi_dependent_iff_not_independent
- mi_independent_symm
- mi_dependent_symm
- gaussian_mi_positive_of_corr_nonzero
- gaussian_mi_zero_of_corr_zero
- rashomon_from_mi_dependence
- impossibility_from_mi
- no_rashomon_from_agreement
- stable_ranking_from_agreement
- mi_sufficient_for_impossibility
- correlation_implies_mi_impossibility

Necessity.lean (3)

- impossibility_from_rashomon
- fully_specified_possibility
- no_rashomon_from_all_three

NecessityBiconditional.lean (3)

- possibility_from_factoring
- not_rashomon_iff_factoring_neq
- rashomon_biconditional_neq

NeuralNet.lean (1)

- nn_impossibility

ParetoOptimality.lean (7)

- no_method_beats_half
- committed_ranking_has_positive_error
- dash_pareto_dominance_within_group
- dash_pareto_gap_exact
- universal_within_group_half
- pareto_frontier_dichotomy
- dash_within_group_dominance

PathConvergence.lean (3)

- impossibility_as_design_space_corollary
- relaxation_paths_converge
- complete_implies_unfaithful

PeresMermin.lean (3)

- peres_mermin_bool
- peres_mermin_parity
- peres_mermin_impossibility

PhaseProblem.lean (4)

same_energy
different_signals
phase_rashomon_derived
phase_impossibility

PredictiveConsequences.lean (5)

stable_implies_unfaithful
rashomon_unfaithfulness
faithful_is_explain
faithful_unique
all_or_nothing

ProportionalityLocal.lean (4)

proportionality_global_implies_local
attribution_firstMover_gt_local
gbdt_rashomon_local
gbdt_impossibility_local

QMInterpretation.lean (3)

qm_same_outcome
qm_ontology_different
qm_interpretation_impossibility

Qualitative.lean (4)

rashomon_from_qualitative
impossibility_qualitative
quantitative_implies_dominates
axiom_implies_surjective

QuantitativeBound.lean (3)

decisive_fails_on_rashomon_pair
not_decisive_at_rashomon_pair
quantitative_impossibility

QuantumMeasurementRevolution.lean (2)

quantum_ontology_maxIncompat
quantum_measurement_revolution

QueryComplexity.lean (16)

testing_constant_pos
le_cam_lower_bound
query_bound_nonneg

```

query_bound_pos
query_complexity_lower_bound
query_bound_scales_with_sigma
query_bound_scales_with_gap
query_bound_mono_sigma
query_bound_antimono_gap
z_test_query_count_nonneg
query_gap_ratio
z_test_optimal_up_to_constant
snr_at_0_15
snr_denominator
snr_multiplier
breast_cancer_query_bound

```

QueryComplexityDerived.lean (12)

```

chebyshev_query_bound
chebyshev_query_lower_bound
chebyshev_bound_simplify
query_complexity_lower_bound_derived
derived_query_bound_nonneg
derived_query_bound_pos
derived_bound_scales_with_sigma
derived_bound_scales_with_gap
derived_bound_mono_sigma
derived_bound_antimono_gap
derived_snr_at_0_15
derived_snr_value

```

QueryComplexityParametric.lean (10)

```

TestingBound
TestingBound
query_complexity_parametric
queryBound_scales_sigma
queryBound_scales_gap
queryBound_mono_sigma
queryBound_antimono_gap
zTest_optimal
leCam_breast_cancer
axiom_free_check

```

QueryRelative.lean (2)

```

query_impossibility
query_possibility

```

RashomonInevitability.lean (5)

```

nondegen_implies_strict_one_direction
rashomon_from_nondegen_and_symmetry

```

rashomon_inevitability
impossibility_inevitable
escape_unanimous_agreement

RashomonUniversality.lean (4)

rashomon_witnesses_from_swap
rashomon_pair_from_swap
rashomon_from_symmetry
impossibility_from_symmetry

Ratio.lean (3)

splitCount_ratio
attribution_ratio
ratio_tendsto_atTop

RecursiveImpossibility.lean (7)

recursive_impossibility_persists
unbounded_undecidability
kthBit_rashomon
infiniteStack_bilemma
infiniteStack_independence
self_referential_undecidability
shared_recursive_structure

RobustnessLipschitz.lean (11)

C_Phi_pos
flipOfSNR_eq
isLipschitzWith_id
isLipschitzWith_neg
isLipschitzWith_comp
isLipschitzWith_const
flipOfSNR_lipschitz
flipOfRho_lipschitz
flip_rate_robust
flip_rate_continuous_at
flip_perturbation_bound_monotone

SBDInstances.lean (5)

attribution_is_sbd_instance
attribution_sbd_unfaithful
model_selection_is_sbd_instance
abstract_aggregation_impossibility
abstract_aggregation_impossibility_weak

SaliencyInstanceConstructive.lean (3)

saliency_same_output
saliency_different_regions
saliency_impossibility_constructive

Setup.lean (1)

attribution_impossibility_bundled

SimultaneityRevolution.lean (2)

simultaneity_maxIncompat
simultaneity_revolution

SpearmanDef.lean (20)

countEqual_pos
countBelow_of_gt
midrank_strict_mono
sumSqRankDiff_nonneg
sumSqRankDiff_pos_of_reversal
spearmanCorr_lt_one_of_reversal
spearmanCorr_bound
attribution_eq_outside_group
attribution_eq_non_firstMover
attribution_swap_structure
attribution_non_fm_stable
countEqual_ge_groupSize_minus_one
midrank_gap_ge_half_groupSize
sumSqRankDiff_ge_sq_groupSize
spearmanCorr_bound_groupSize
spearman_instability_bound
midrank_gap_ge_half_groupSize_tight
sumSqRankDiff_ge_sq_groupSize_tight
spearmanCorr_bound_groupSize_tight
spearman_instability_bound_tight

SplitGap.lean (7)

rho_sq_lt_one
one_minus_rho_sq_pos
denom_pos
denom_ne_zero
split_gap_exact
split_gap_ge_half
split_gap_lower_bound

StatisticalMechanics.lean (5)

same_macrostate
different_microstates

swap_related
swap_preserves_macrostate
stat_mech_impossibility

StumpProportionality.lean (6)

stump_proportionality
stump_attribution_ratio
stump_proportionality_positive
stump_proportionality_exists
stump_proportionality_unique
stump_cv_zero

SymmetricBayes.lean (16)

sbd_unfaithful_witness
sbd_infeasible
tie_is_stable
sbd_faithful_at_most_one
sbd_faithful_unique
sbd_binary_orbit_half_unfaithful
sbd_ternary_orbit_unfaithful
attribution_from_sbd
sbd_faithful_at_most_one_in_orbit
sbd_trichotomy
sbd_family_a_or_b
orbitEquiv_refl
orbitEquiv_symm
orbitEquiv_trans
family_b_output_not_optimal
symmetric_bayes_dichotomy

SymmetryDerive.lean (1)

attribution_sum_symmetric

SyntacticAmbiguity.lean (3)

same_yield
different_bracketings
syntactic_ambiguity_impossibility

Trilemma.lean (2)

attribution_impossibility
attribution_impossibility_weak

Ubiquity.lean (3)

generic_underspecification
fiber_nondegeneracy_implies_impossibility

ubiquity_impossibility

UncertaintyFromSymmetry.lean (17)

inner_Rv_v_eq_inner_Rv_Rv
inner_v_Rv_eq_inner_Rv_Rv
reynolds_orthogonal
uncertainty_from_symmetry
information_loss
stable_bounded
unstable_bounded
reynolds_residual_zero
residual_orthogonal_to_fixed
best_approximation
reynolds_naturality
fixed_perp_residual
reynolds_best_approximation
reynolds_image_in_fixed
fixed_point_in_invariant_subspace
beyond_capacity_penalty
beyond_capacity_optimal

UnfaithfulBound.lean (3)

stable_complete_unfaithful
tie_is_never_unfaithful
unfaithfulness_free_implies_tie

UnfaithfulQuantitative.lean (2)

attribution_prob_half
stable_ranking_unfaithfulness_half

UniversalDesignSpace.lean (1)

universal_design_space_dichotomy

UniversalResolution.lean (1)

gInvariant_stable

ValueAlignment.lean (2)

value_maxIncompat
value_alignment_impossibility

VarianceDerivation.lean (6)

variance_of_mean_from_sum_variance
consensus_variance_from_independence

consensus_variance_nonneg_from_independence
double_M_halves_variance_derived
variance_ratio_from_independence
weighted_variance_ge_consensus_variance

ViewUpdate.lean (3)

same_view
different_rows
view_update_impossibility

MIWitness.lean (9)

pX_eq
pY_eq
MI_eq_log_two
MI_pos
hsym_j
hsym_k
hdep_implies_diff
impossibility_from_mi_general
concrete_mi_impossibility